



Microsoft

# Azure Machine Learning

Level – 300.

Srinag



## Data & Artificial Intelligence

Looking for creating value for your **BUSINESS** that lies at the intersection of **HUMANS** and **INTELLIGENT MACHINES**?

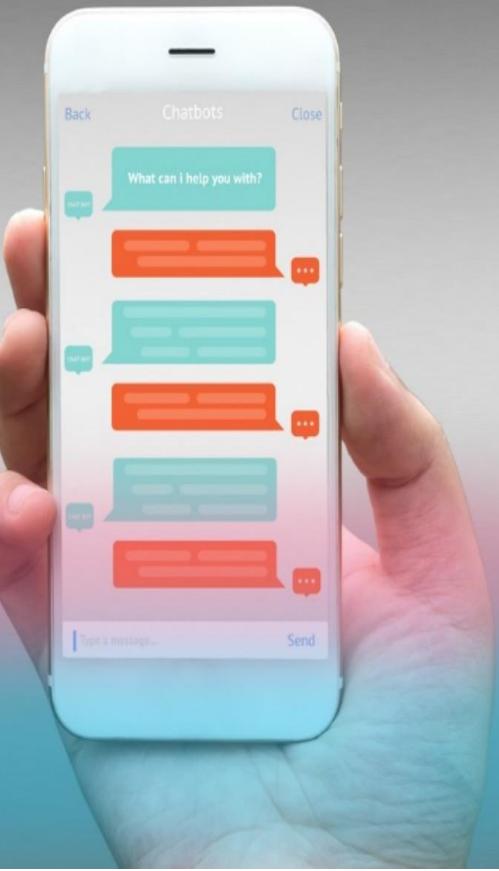
### INTERNET OF THINGS

Not just **CUSTOMERS** but every " **THING** " around you is talking.  
IOT enables innumerable applications ranging from the micro to macro and from the trivial to critical.



### CHATBOTS

Bye-bye **E-Commerce!**  
Welcome Conversational Commerce-A new ecosystem for Artificial Intelligence



A **Boutique Company** that provides Training and Consulting services on **Data + AI + IoT Workloads**

# About Me

Advanced Analytics Specialist.

*Certified from MIT on 'Advanced Analytics and Big Data Challenges'. Technology Enthusiast and Data Specialist.*

## Certifications

- Microsoft Certified Power BI Developer.
- Microsoft Certified Azure Machine Learning Developer from Microsoft Virtual Academy
- Certified on Practical Data Analytics with Microsoft Cortana Intelligence Suite.

## Role

Currently working as Azure Data Specialist and Enterprise Trainer on Azure Data Analytics , Cortana Intelligence Suite and Cognitive Sciences



Sri Nag Sashank

## Experience

7.5 years of total experience as Data Specialist and Data Engineer.

## Training

Expert facilitator and corporate behavioral trainer with 5 years of experience in designing programs, content development , Data Analytics and Machine Learning.

- Hosted many webinars and corporate workshops on Azure Data Analytics & Machine Learning.

# Agenda

- Microsoft Cortana Intelligence Suite
- Introduction to Machine Learning
- Data Preparation in Azure ML
- Classification
- Confusion Matrix
- Ensemble Learning Introduction
- Support Vector Machines
- Hyper Parameter Tuning



# Cortana Intelligence Suite

Microsoft Vision on Data

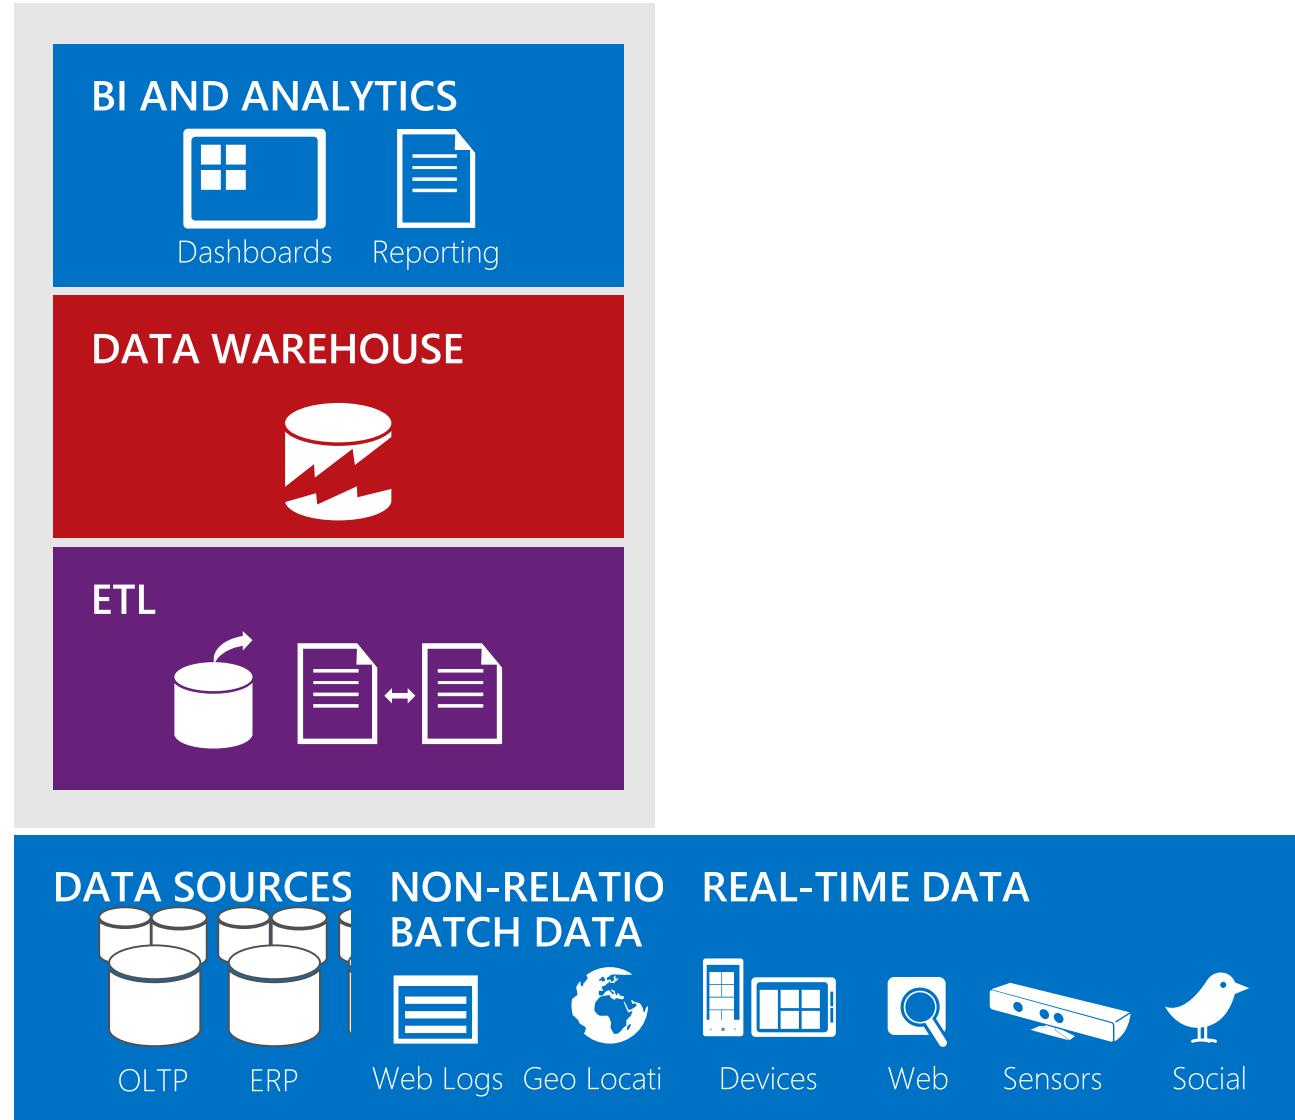


1 Increasing data volumes

2 New data sources and types

3 Real-time data

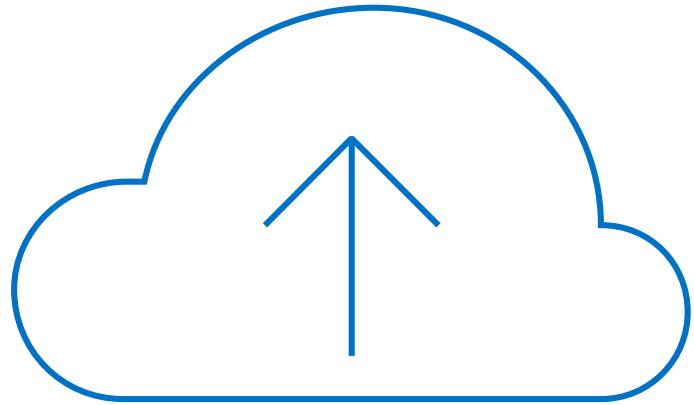
4 Cloud-born data and  
hybrid IT infrastructure



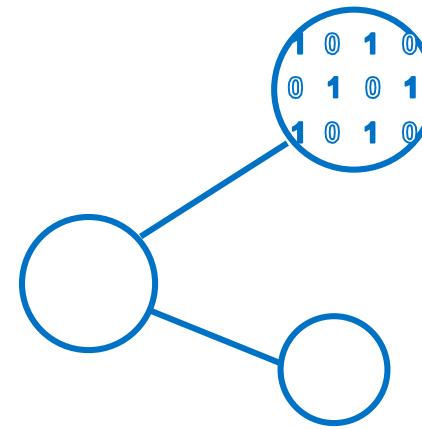
# Insights is a journey



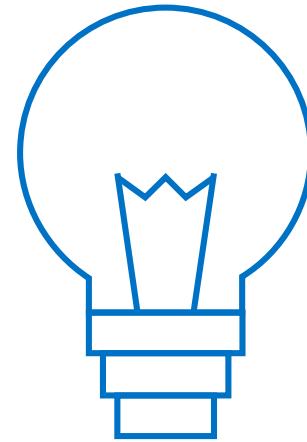
# Convergence accelerates digital transformation



Cloud



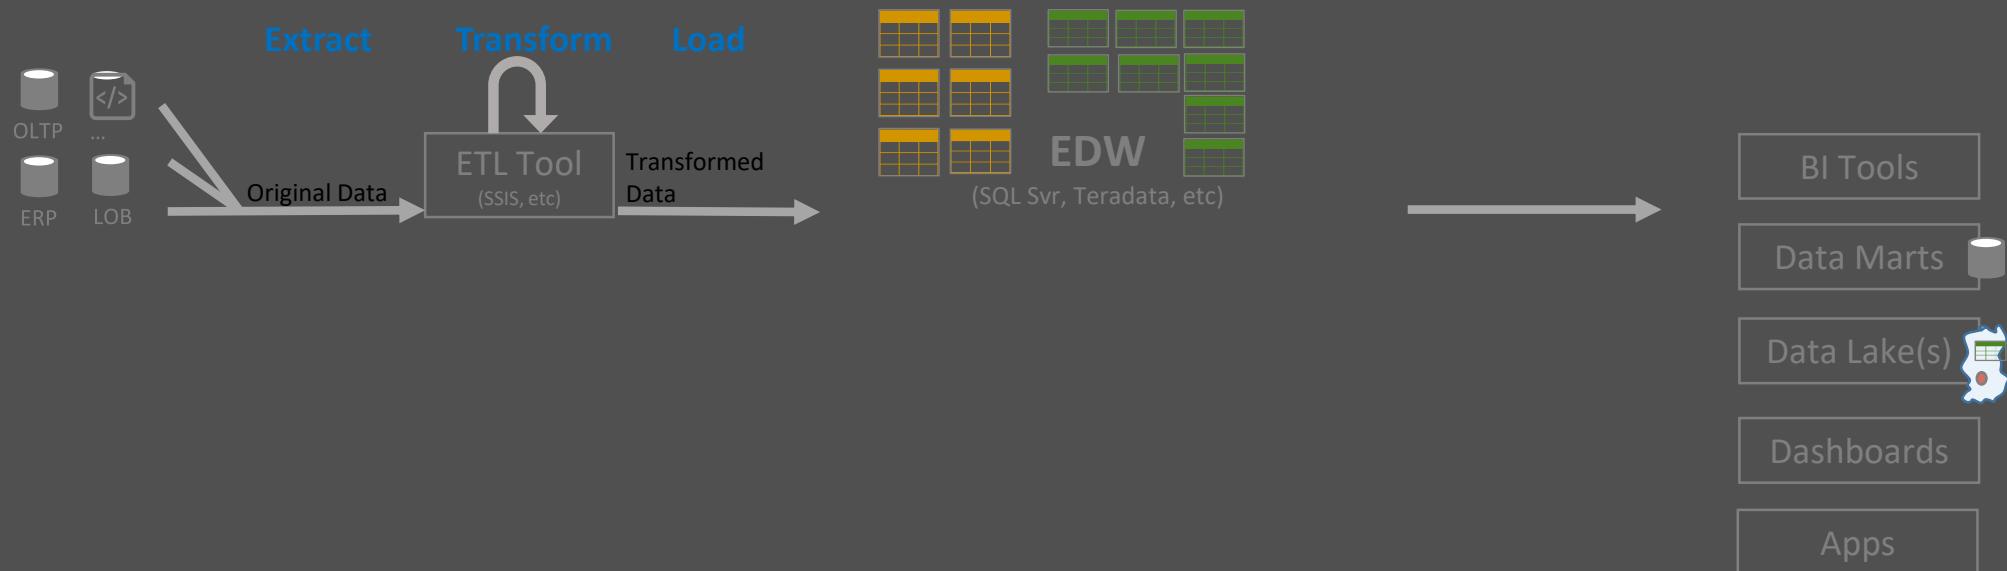
Data



Intelligence

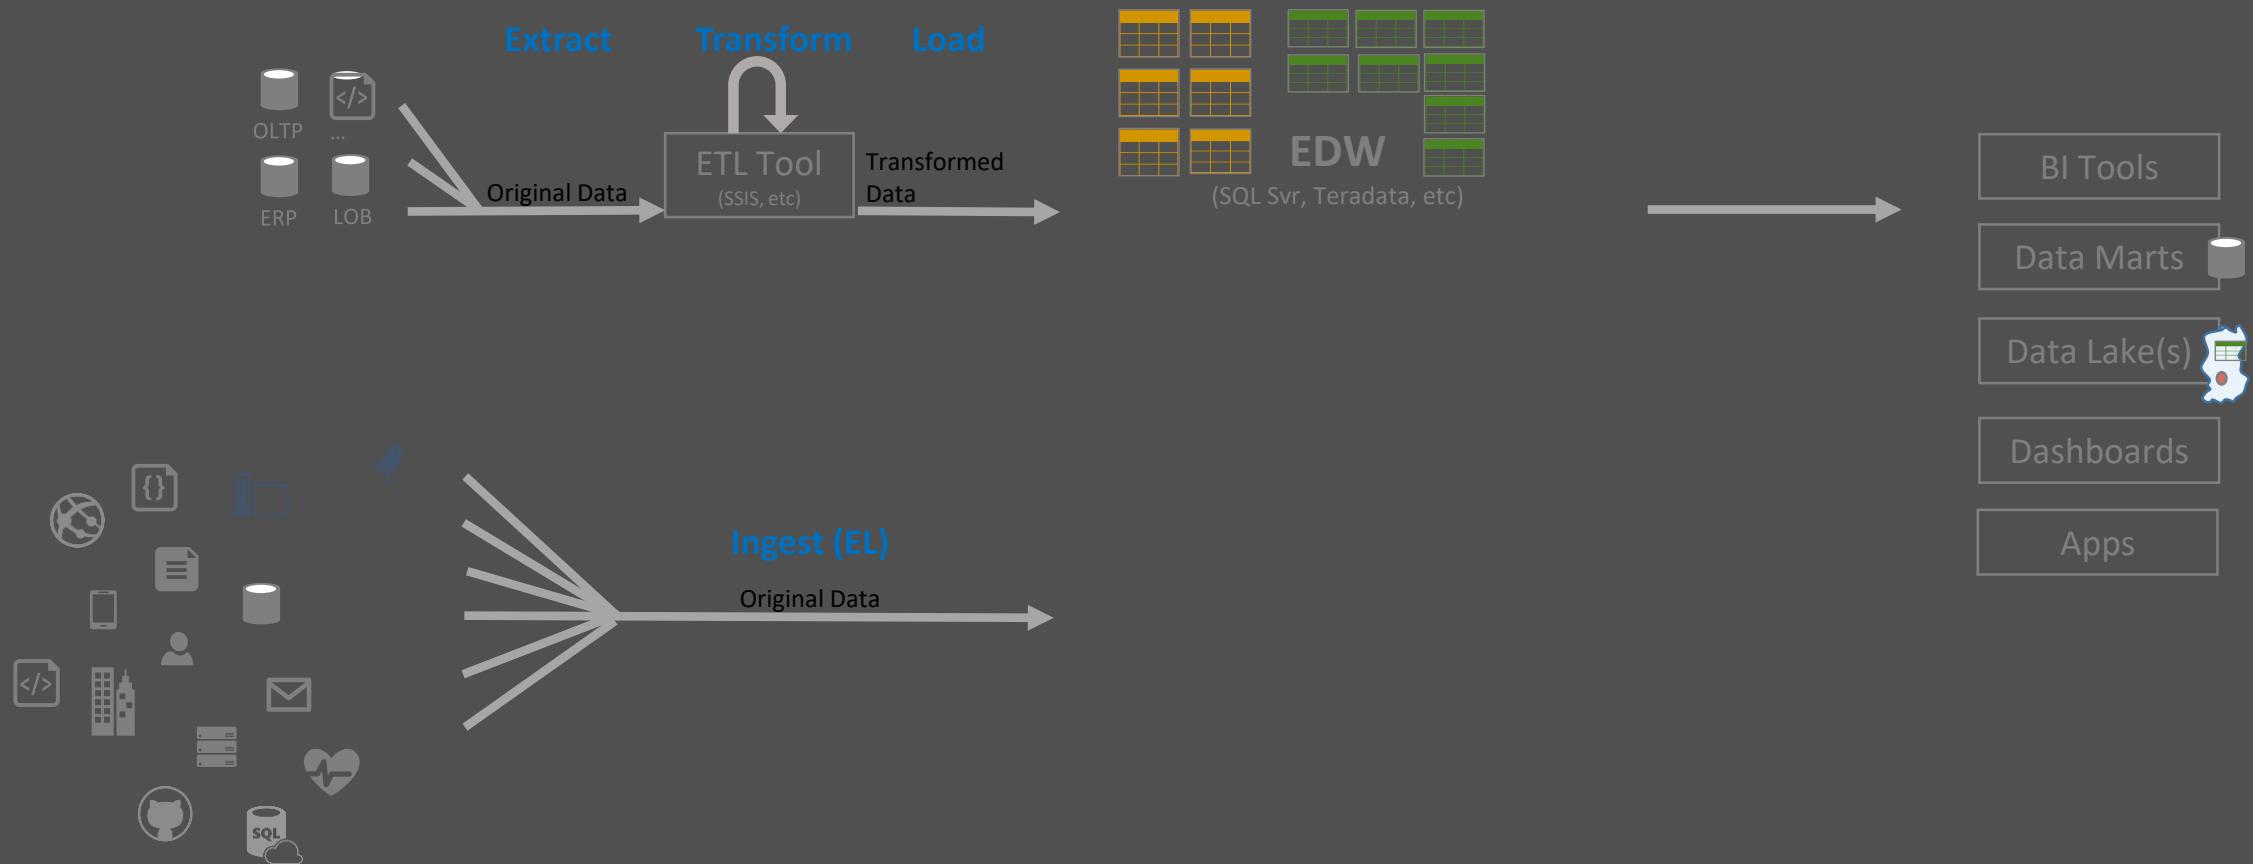
# Evolving Approaches to Analytics

## EDWs to Data Lakes



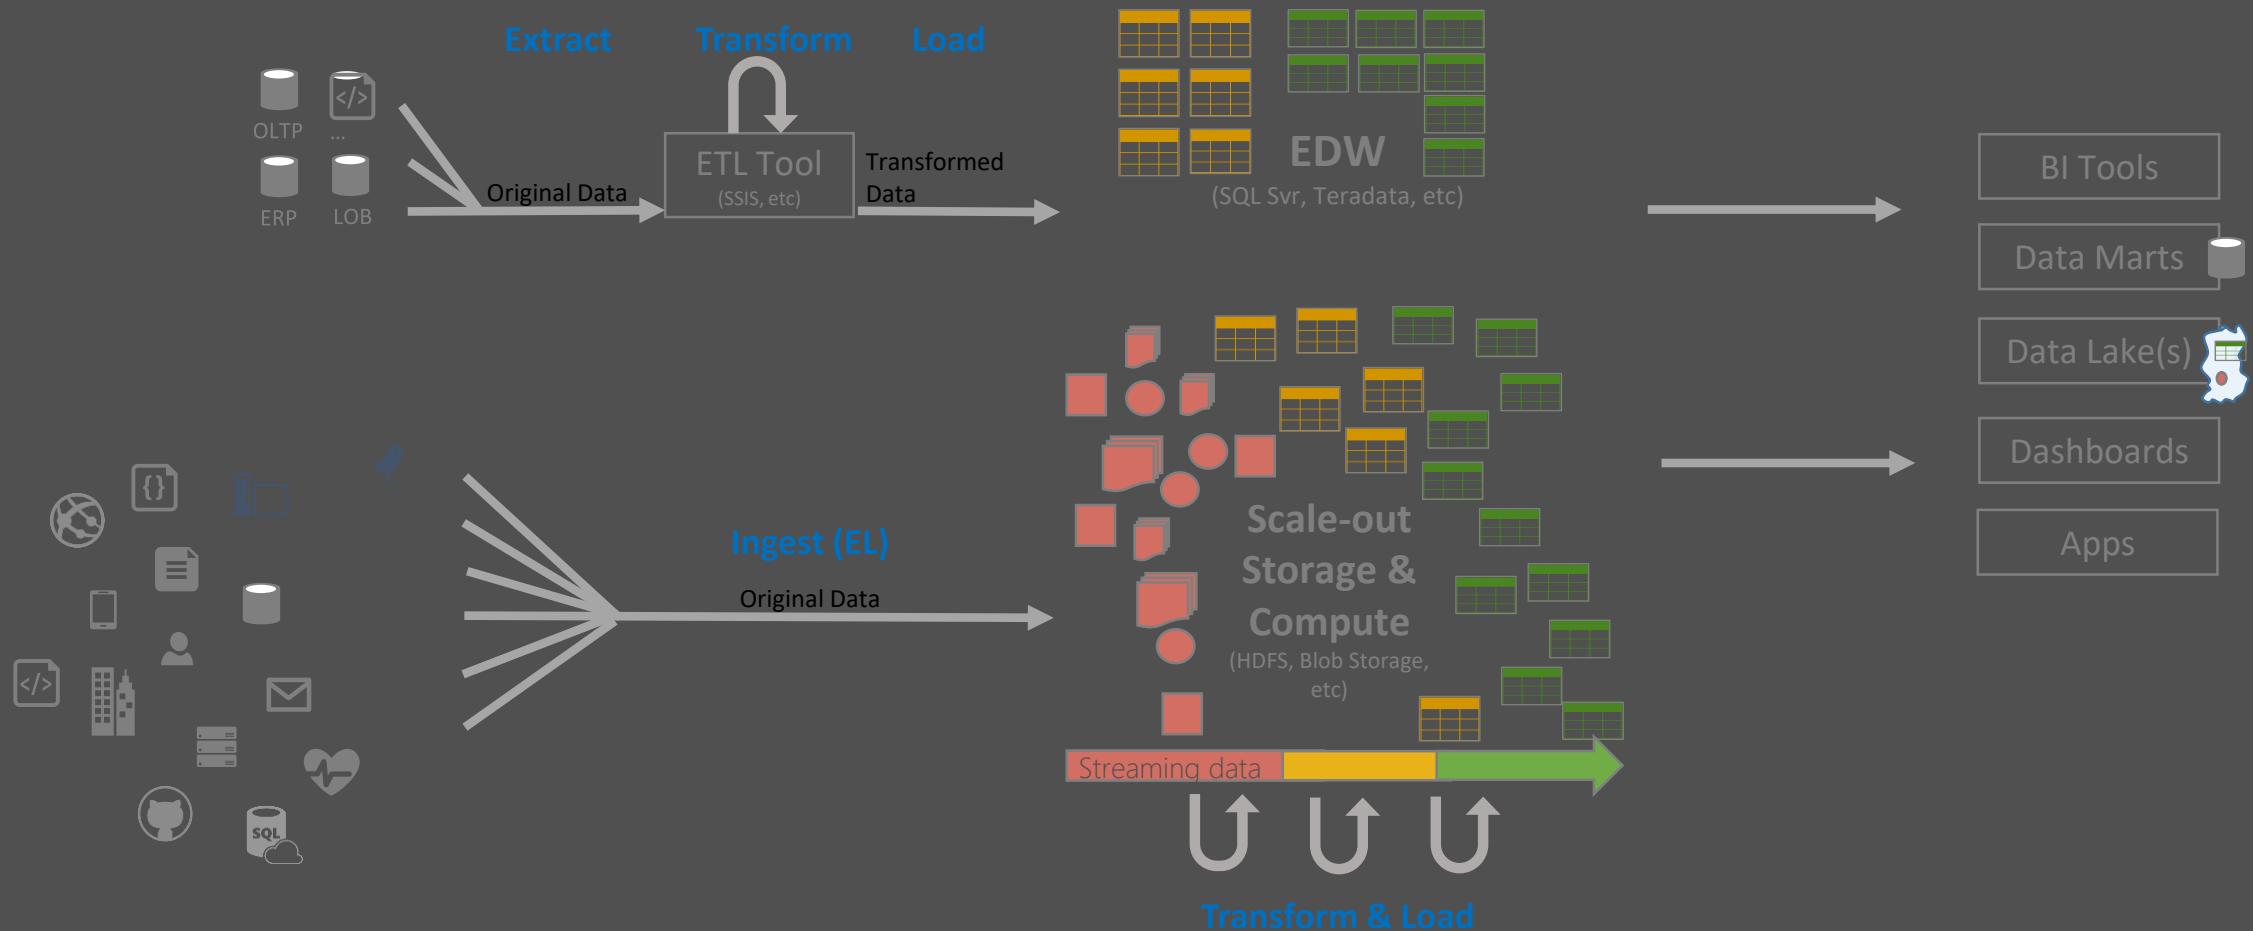
# Evolving Approaches to Analytics

EDWs to Data Lakes



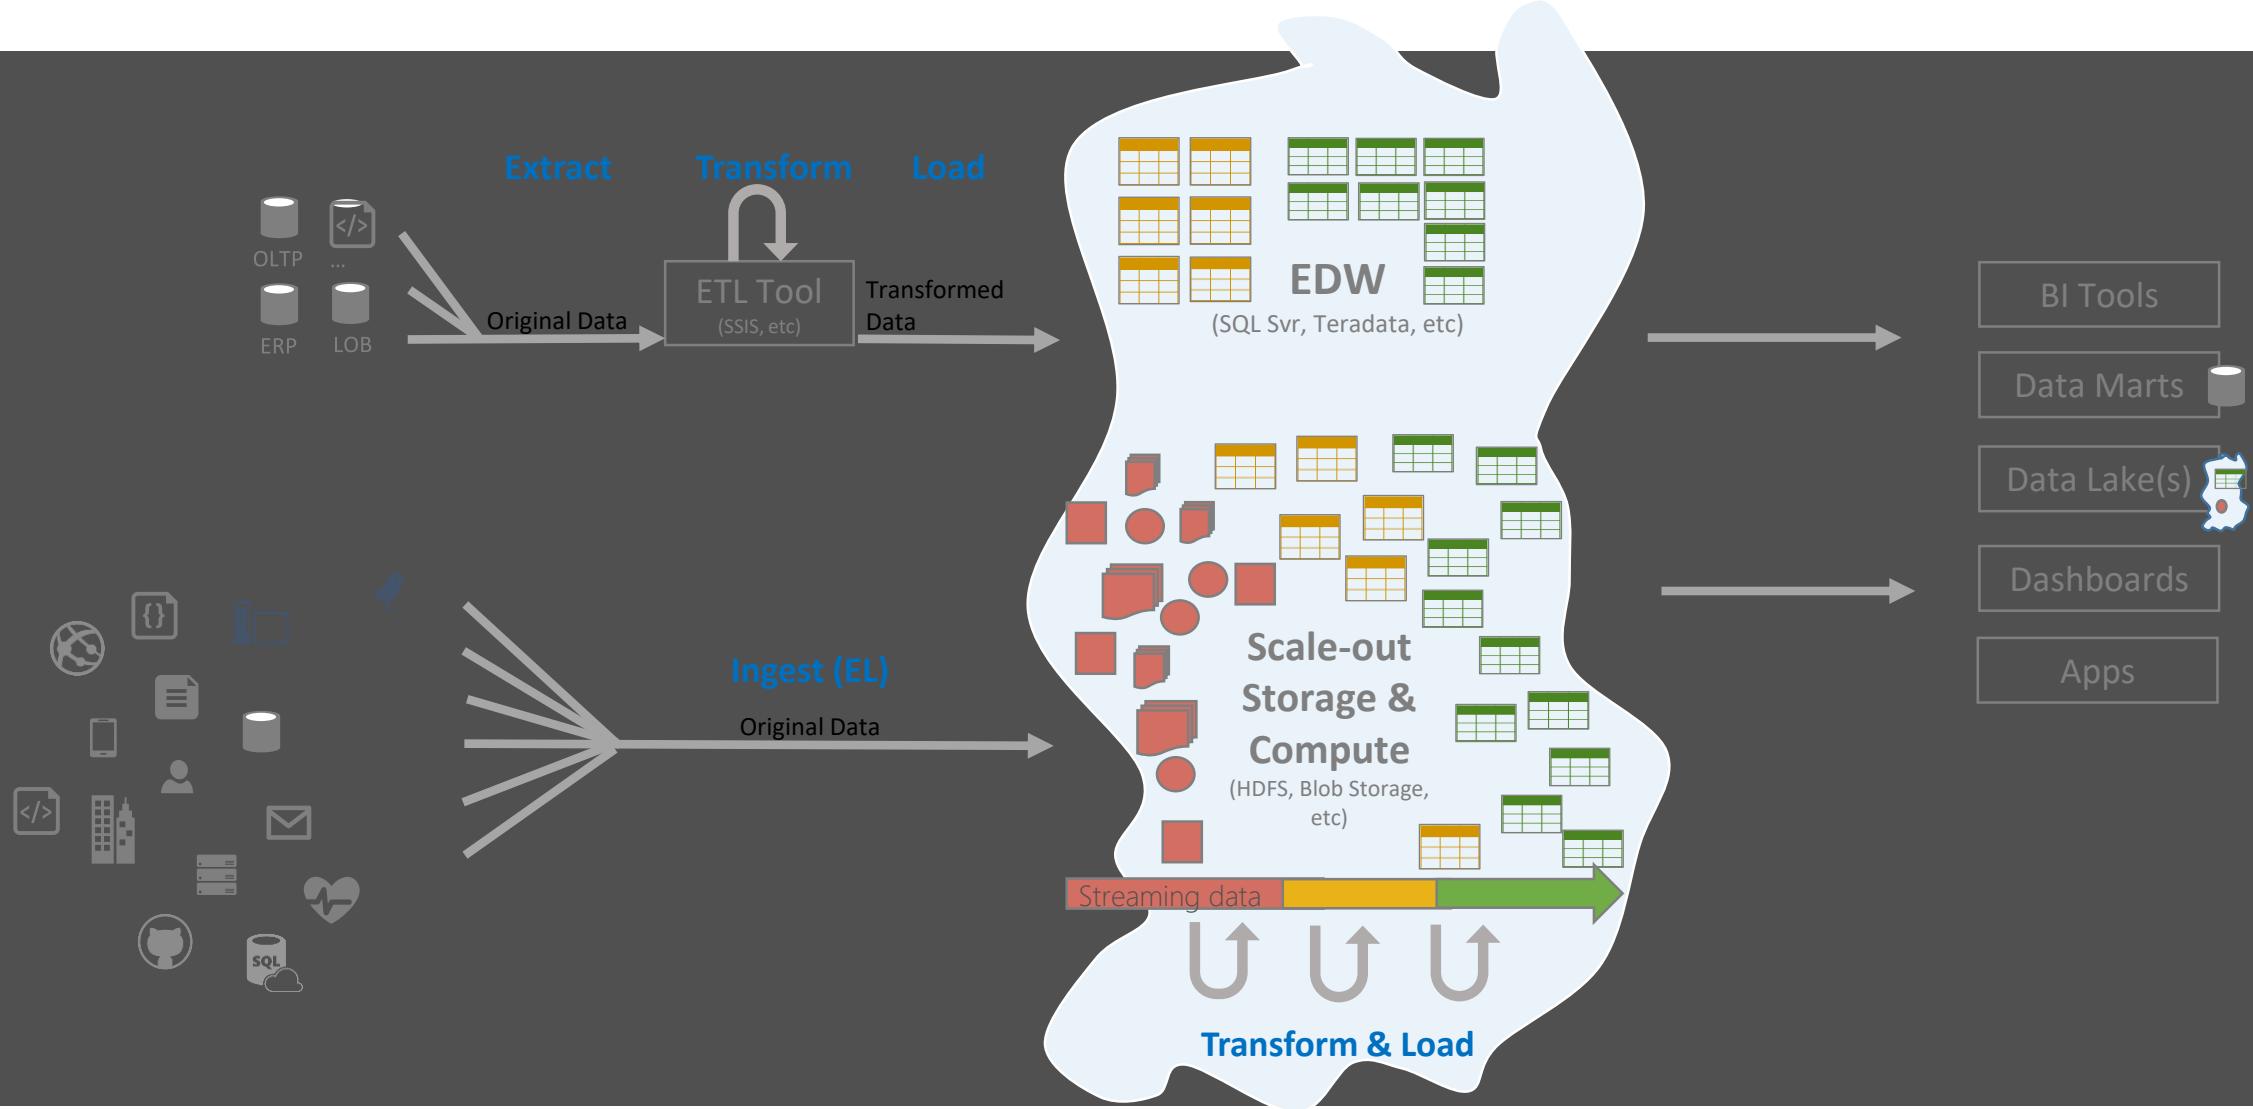
# Evolving Approaches to Analytics

EDWs to Data Lakes

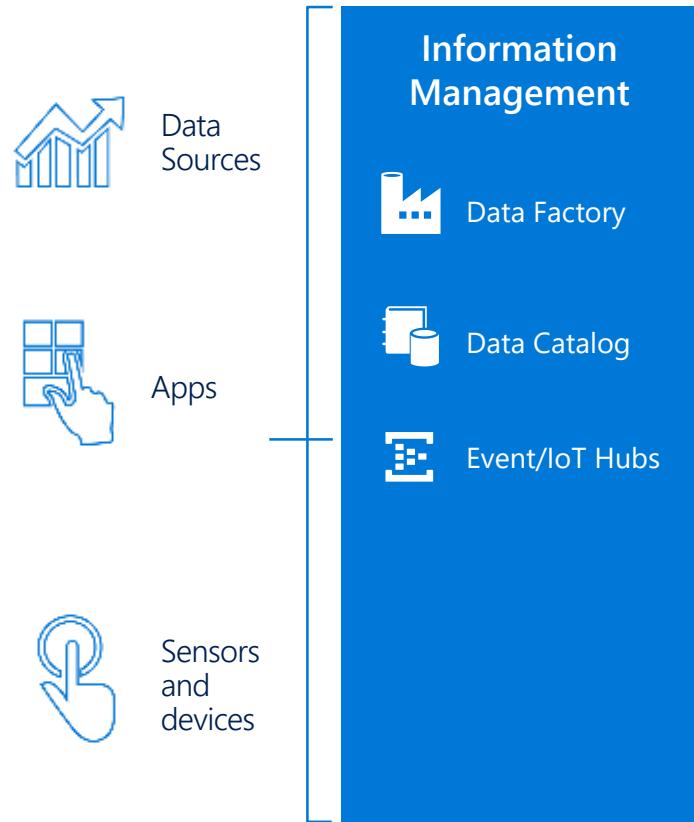


# Evolving Approaches to Analytics

EDWs to Data Lakes

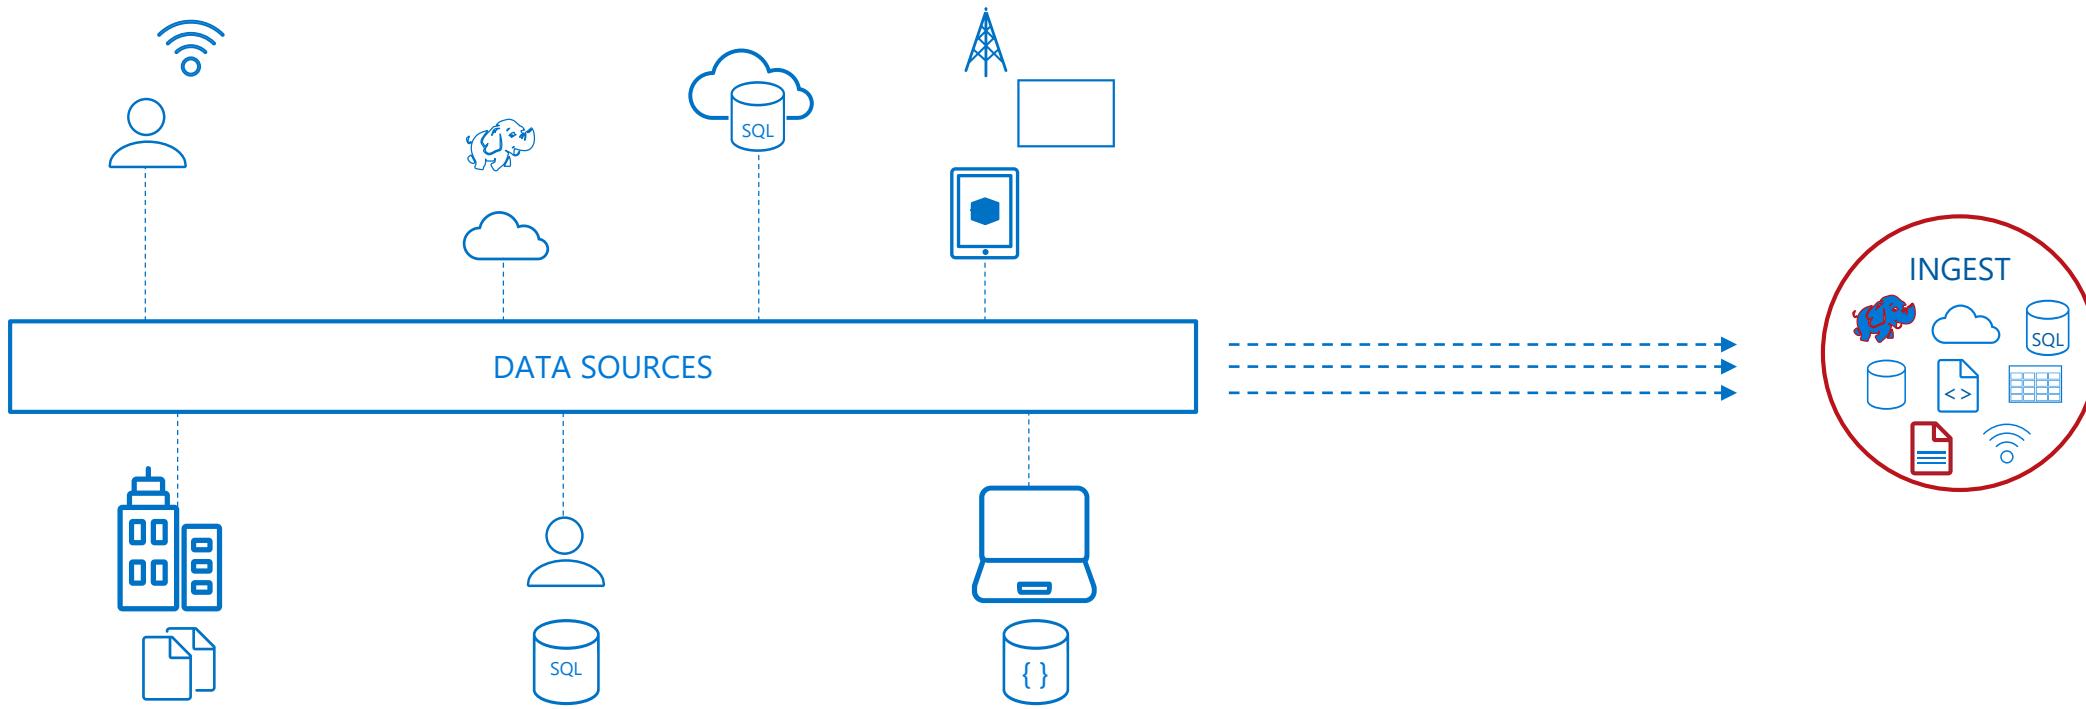
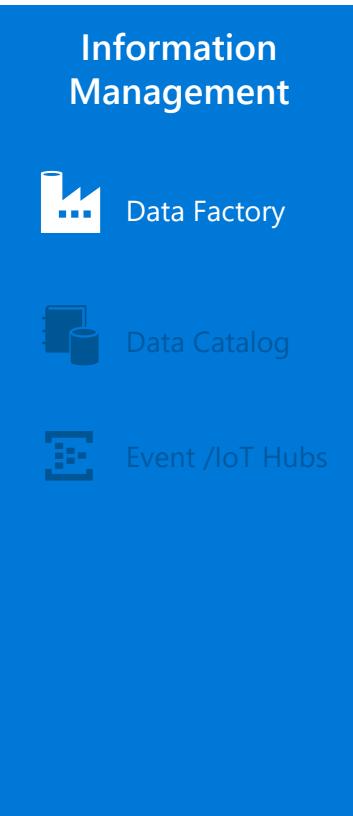


# Information Management



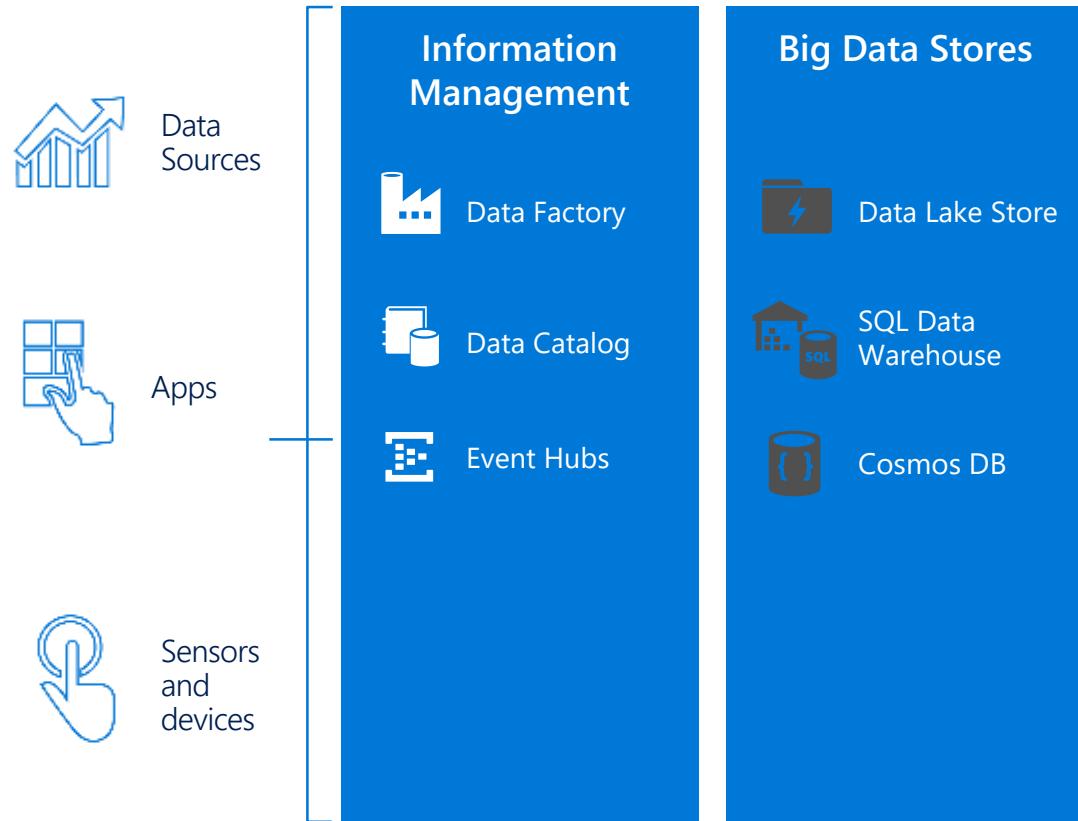
Data

# Compose and orchestrate data services at scale



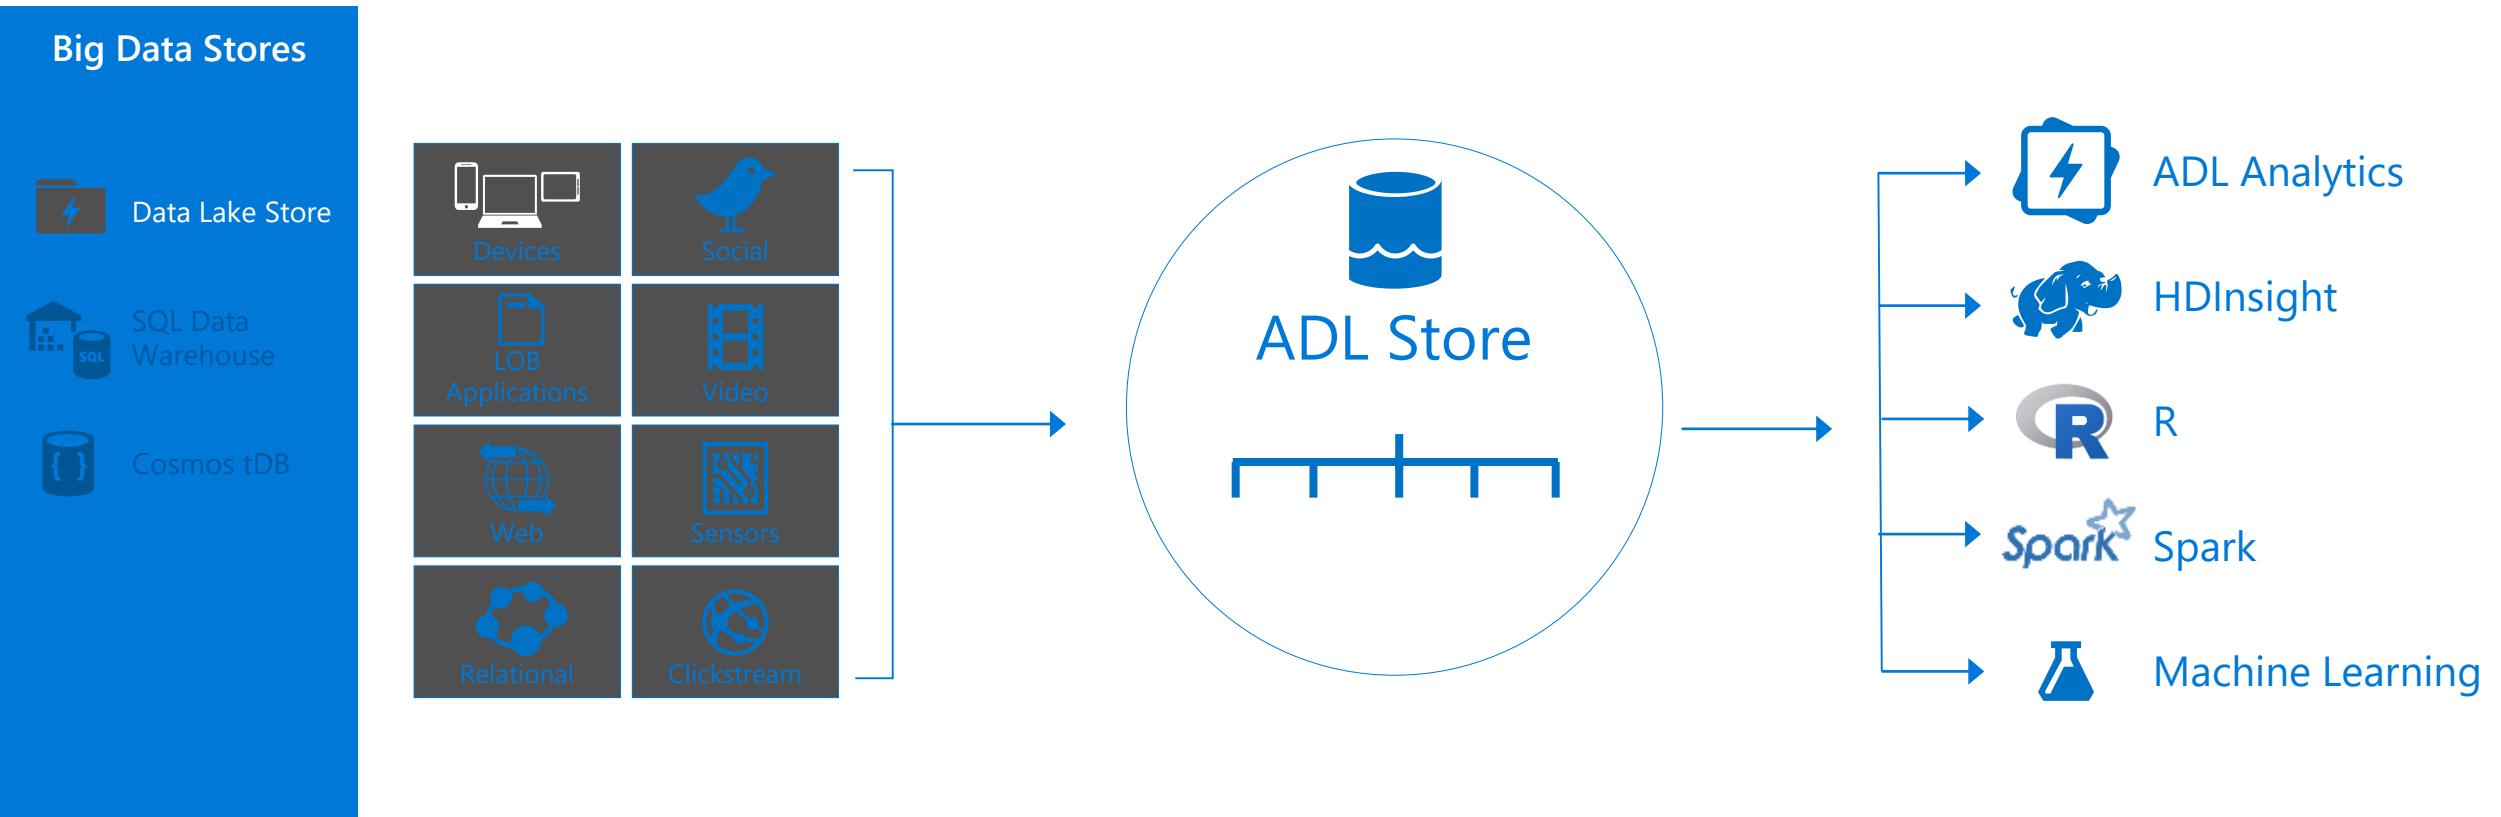
- Create, schedule, orchestrate, and manage data pipelines
- Visualize data lineage
- Connect to on-premises and cloud data sources
- Monitor data pipeline health
- Automate cloud resource management
- Move relational data for Hadoop processing
- Transform with Hive, pig, or custom code

# Big Data Stores



Data

# A hyper-scale repository for big data analytics workloads



- A Hadoop Distributed File System for the cloud
- No fixed limits on file size
- No fixed limits on account size
- Unstructured and structured data in their native format
- Massive throughput to increase analytic performance
- High durability, availability, and reliability
- Azure Active Directory access control

# Elastic data warehouse as a service with enterprise-class features

## Big Data Stores



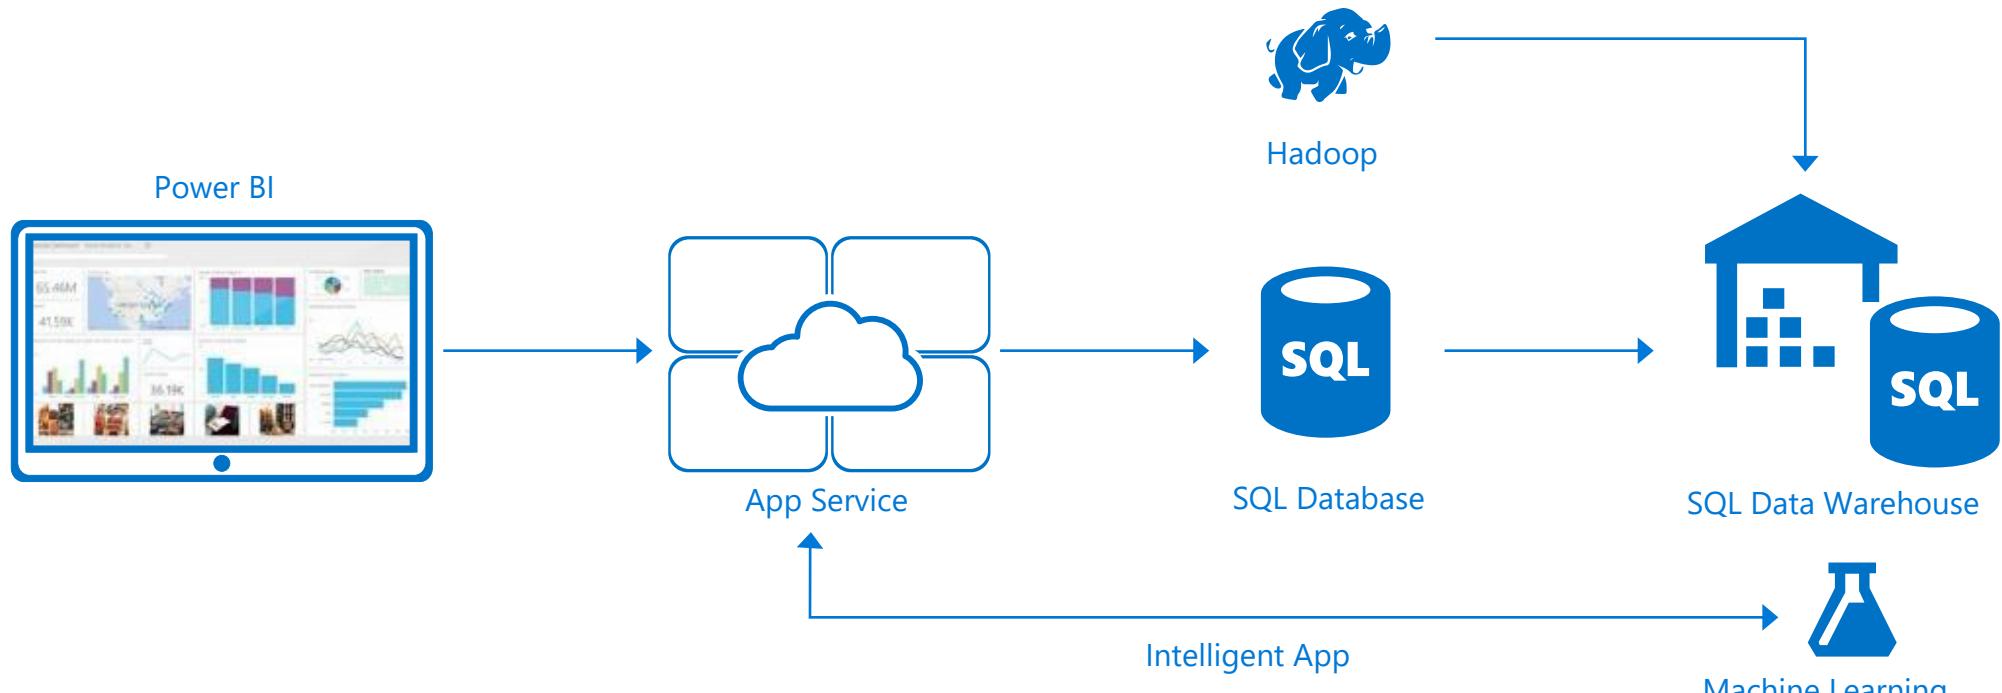
Data Lake Store



SQL Data Warehouse



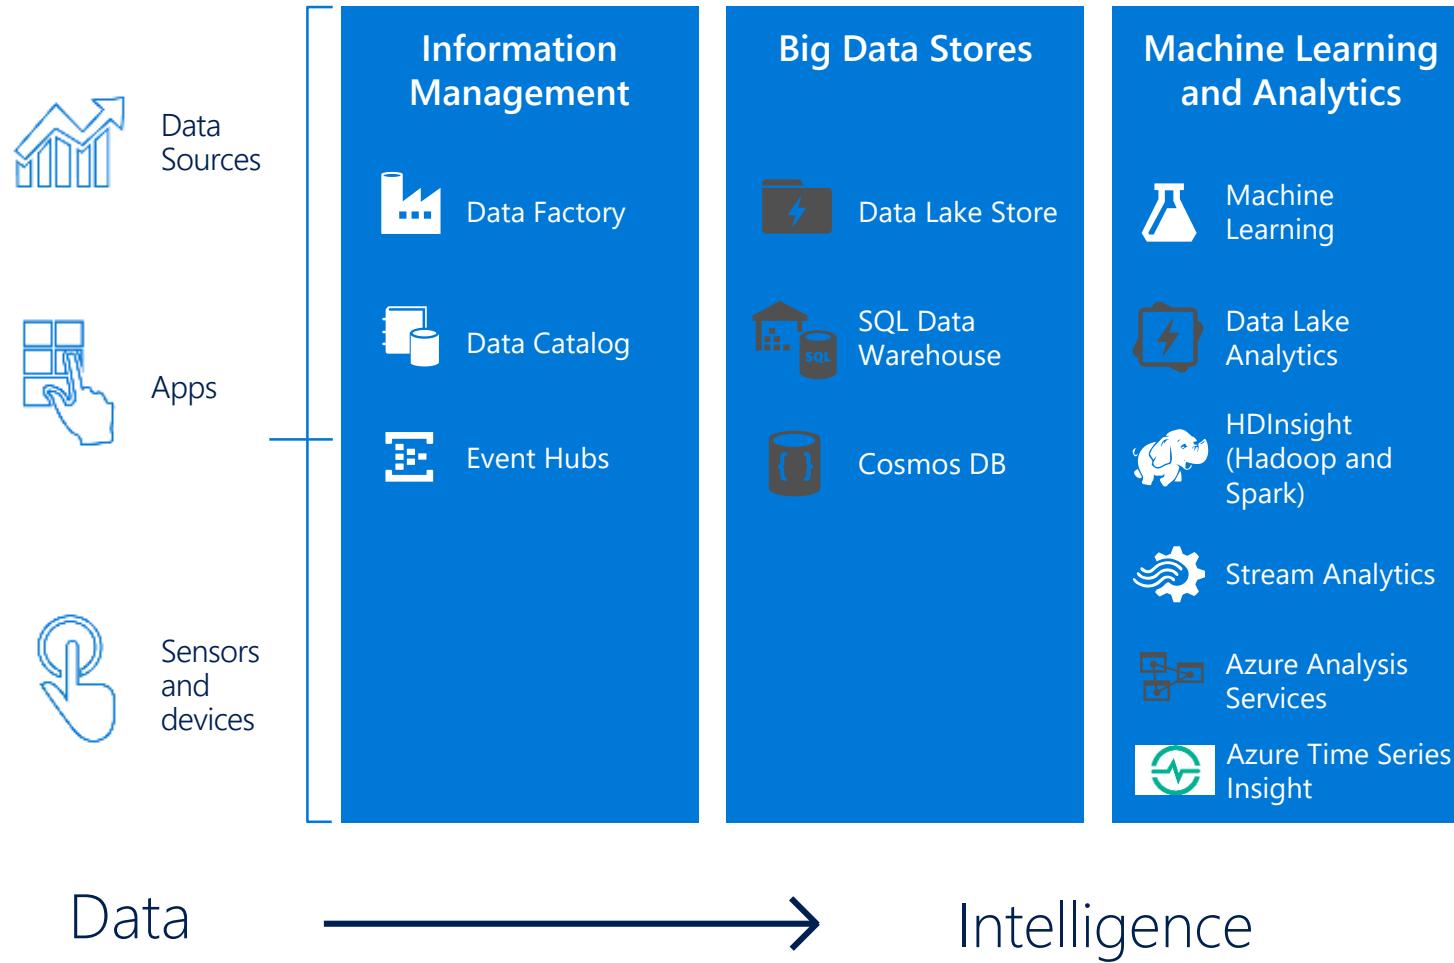
Cosmos DB



- Petabyte scale with massively parallel processing
- Independent scaling of compute and storage—in seconds
- Transact-SQL queries across relational and non-relational data

- Full enterprise-class SQL Server experience
- Works seamlessly with Power BI, Machine Learning, HDInsight, and Data Factory

# Machine Learning and Analytics



# Easily build, deploy, and share predictive analytics solutions

## Machine Learning and Analytics

Machine Learning

Data Lake Analytics

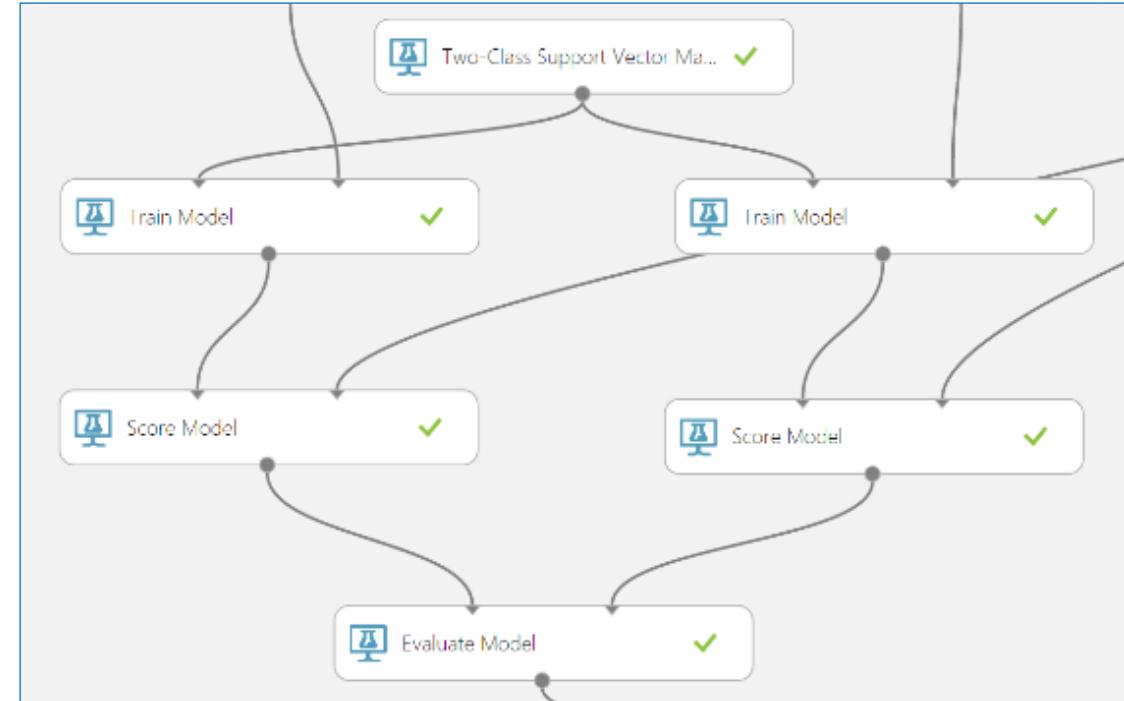
HDInsight  
(Hadoop and Spark)

Stream Analytics

Azure Analysis Services

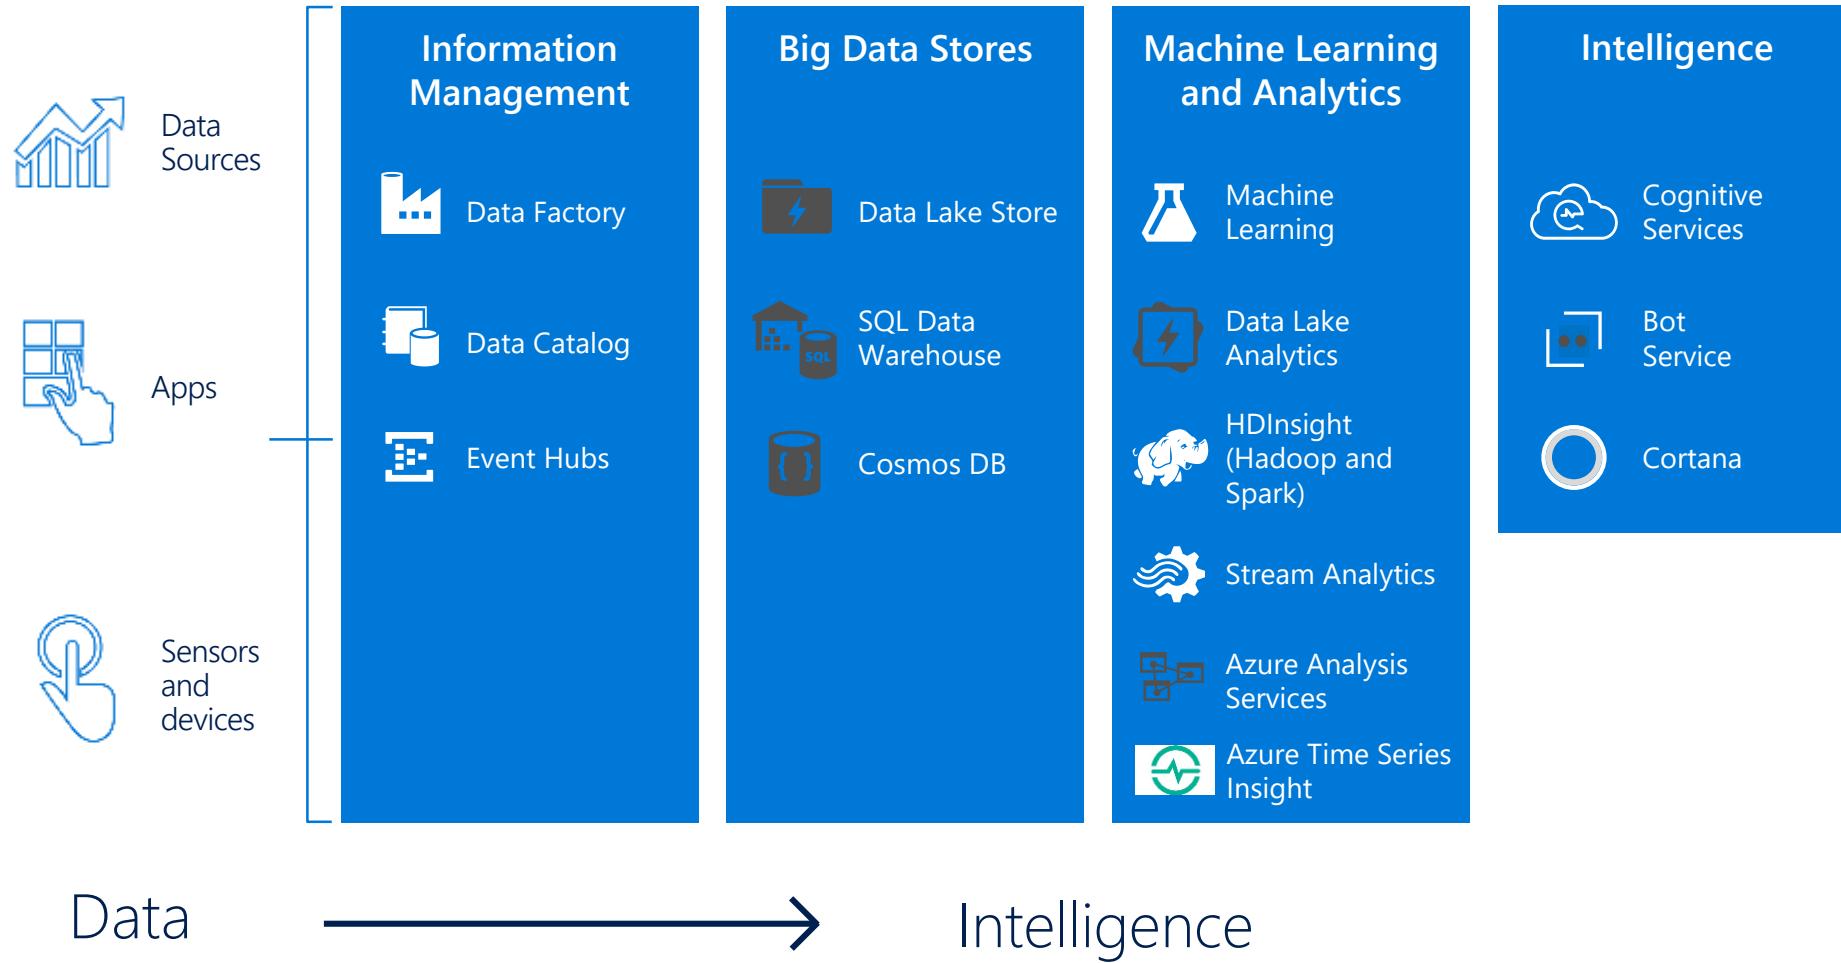
Azure Time Series Insight

The screenshot shows the Cortana Analytics Gallery interface. At the top, there are navigation links: 'Browse all' (highlighted in green), 'Solution Templates', and 'Experiments'. Below this is a 'Refine by' section with dropdown menus for 'CATEGORIES' and 'TAGS'. Under 'CATEGORIES', options like 'Solution Template', 'Experiment', and 'Machine Learning API' are listed. Under 'TAGS', options like 'R', 'Classification', and 'DAT203x' are listed. The main area is titled 'Results' and shows two items: 'Face APIs' and 'Text Analytics'. Each item has a thumbnail image, a title, a brief description, and a timestamp ('10/7/16 7:20 PM' and '20 days ago').



- Simple, scalable, cutting edge. A fully managed cloud service that enables you to easily build, deploy, and share predictive analytics solutions.
- Deploy in minutes. Azure Machine Learning means business. You can deploy your model into production as a web service that can be called from any device, anywhere and that can use any data source.
- Publish, share, monetize. Share your solution with the world in the Gallery or on the Azure Marketplace.

# Intelligence



# Build applications that understand people

Intelligence	Vision	Speech	Language	Knowledge	Search
Cognitive Services	Computer Vision	Speaker Recognition	Text Analytics	Academic	Bing Web Search API
	Content Moderator	Bing Speech	Bing Spell Check	Entity Linking	Bing Image Search API
	Face	Custom Speech Service	Web Language Model	Knowledge Exploration	Bing Video Search API
	Emotion		Linguistic Analysis	Recommendations	Bing News Search API
	Video		Language Understanding	QnA Maker	Bing Auto Suggest API
			Translator		

- Faces, images, emotion recognition and video intelligence
- Spoken language processing, speaker recognition, custom speech recognition
- Natural language processing, sentiment and topics analysis, spelling errors
- Complex tasks processing, knowledge exploration, intelligent recommendations
- Bing engine capabilities for Web, Autosuggest, Image, Video and News

# Get things done in more helpful, proactive and natural ways

Intelligence

Cognitive Services

Bot Service

Cortana



Here are some of the things I can help you with...

Answers

Predictions

Monitoring & Alerts

Task Completion

Cortana for Consumers (today)

Public reference data answers – *"How far is it from Los Angeles to San Francisco?"*

Event predictions – *"Who do you think is going to win the Germany Italy game?"*

Flight status, traffic conditions, changes in weather, ...

Setting reminders, scheduling meetings, getting directions, ...

With the Cortana Intelligence

Answers from organizational data in Power BI  
*"What were our biggest deals that closed last month?"*

Integration with prediction solutions  
*"Which of our customers are most likely to churn in the next quarter?"*

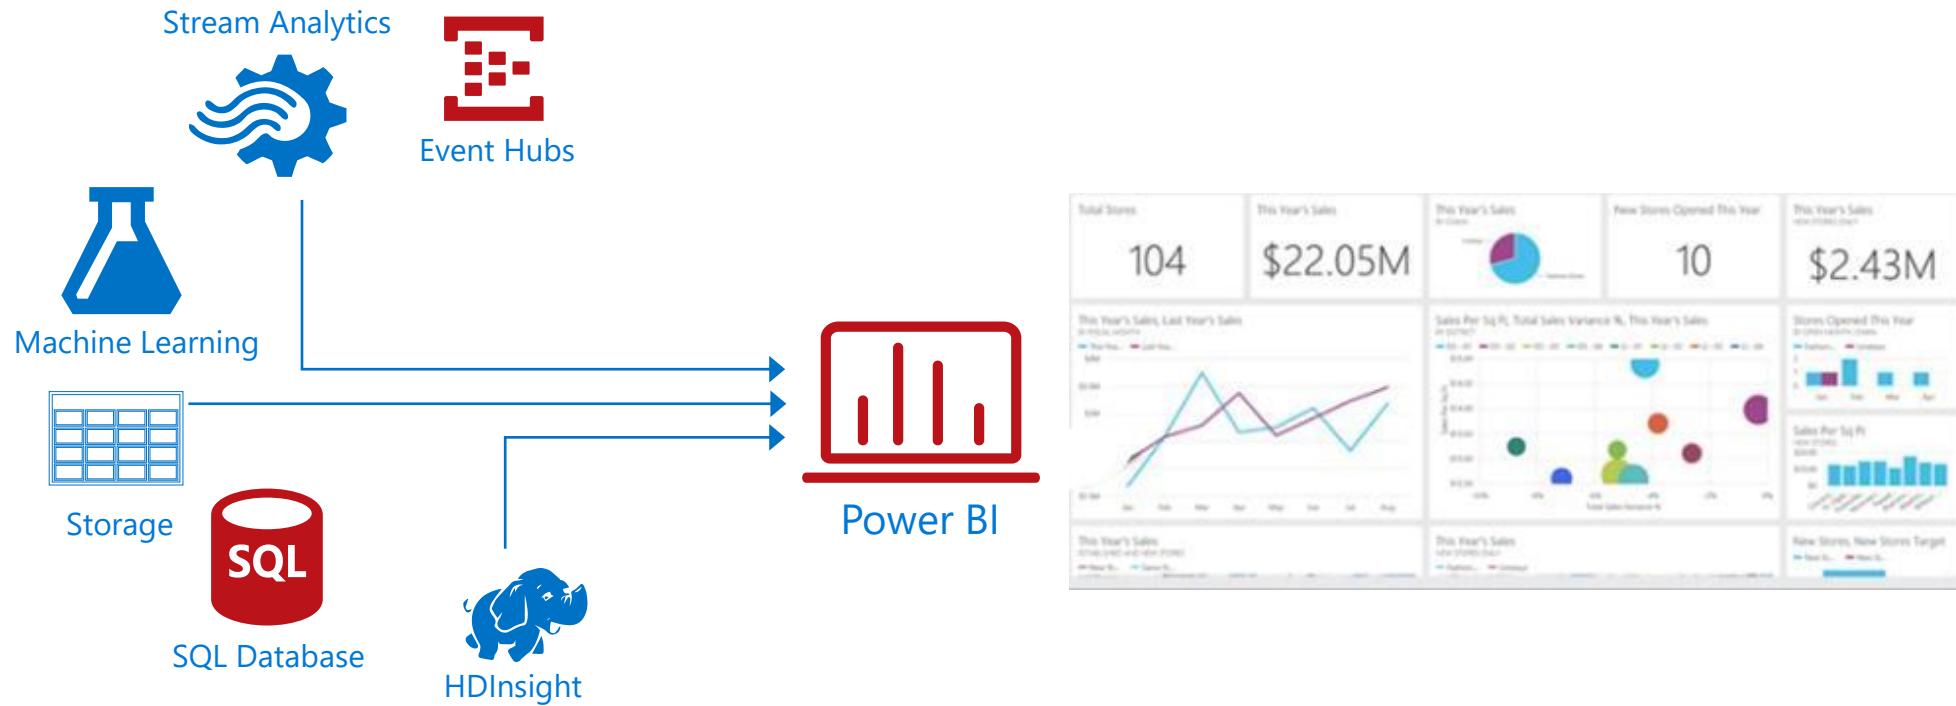
Monitoring KPIs and preemptive alerting  
*"Alert me if this customer ever has a 90% chance of churn in the next 30 days"*

Line of business process integration  
*Assistance with expense report submission on-time within policy*

## Dashboards & Visualizations

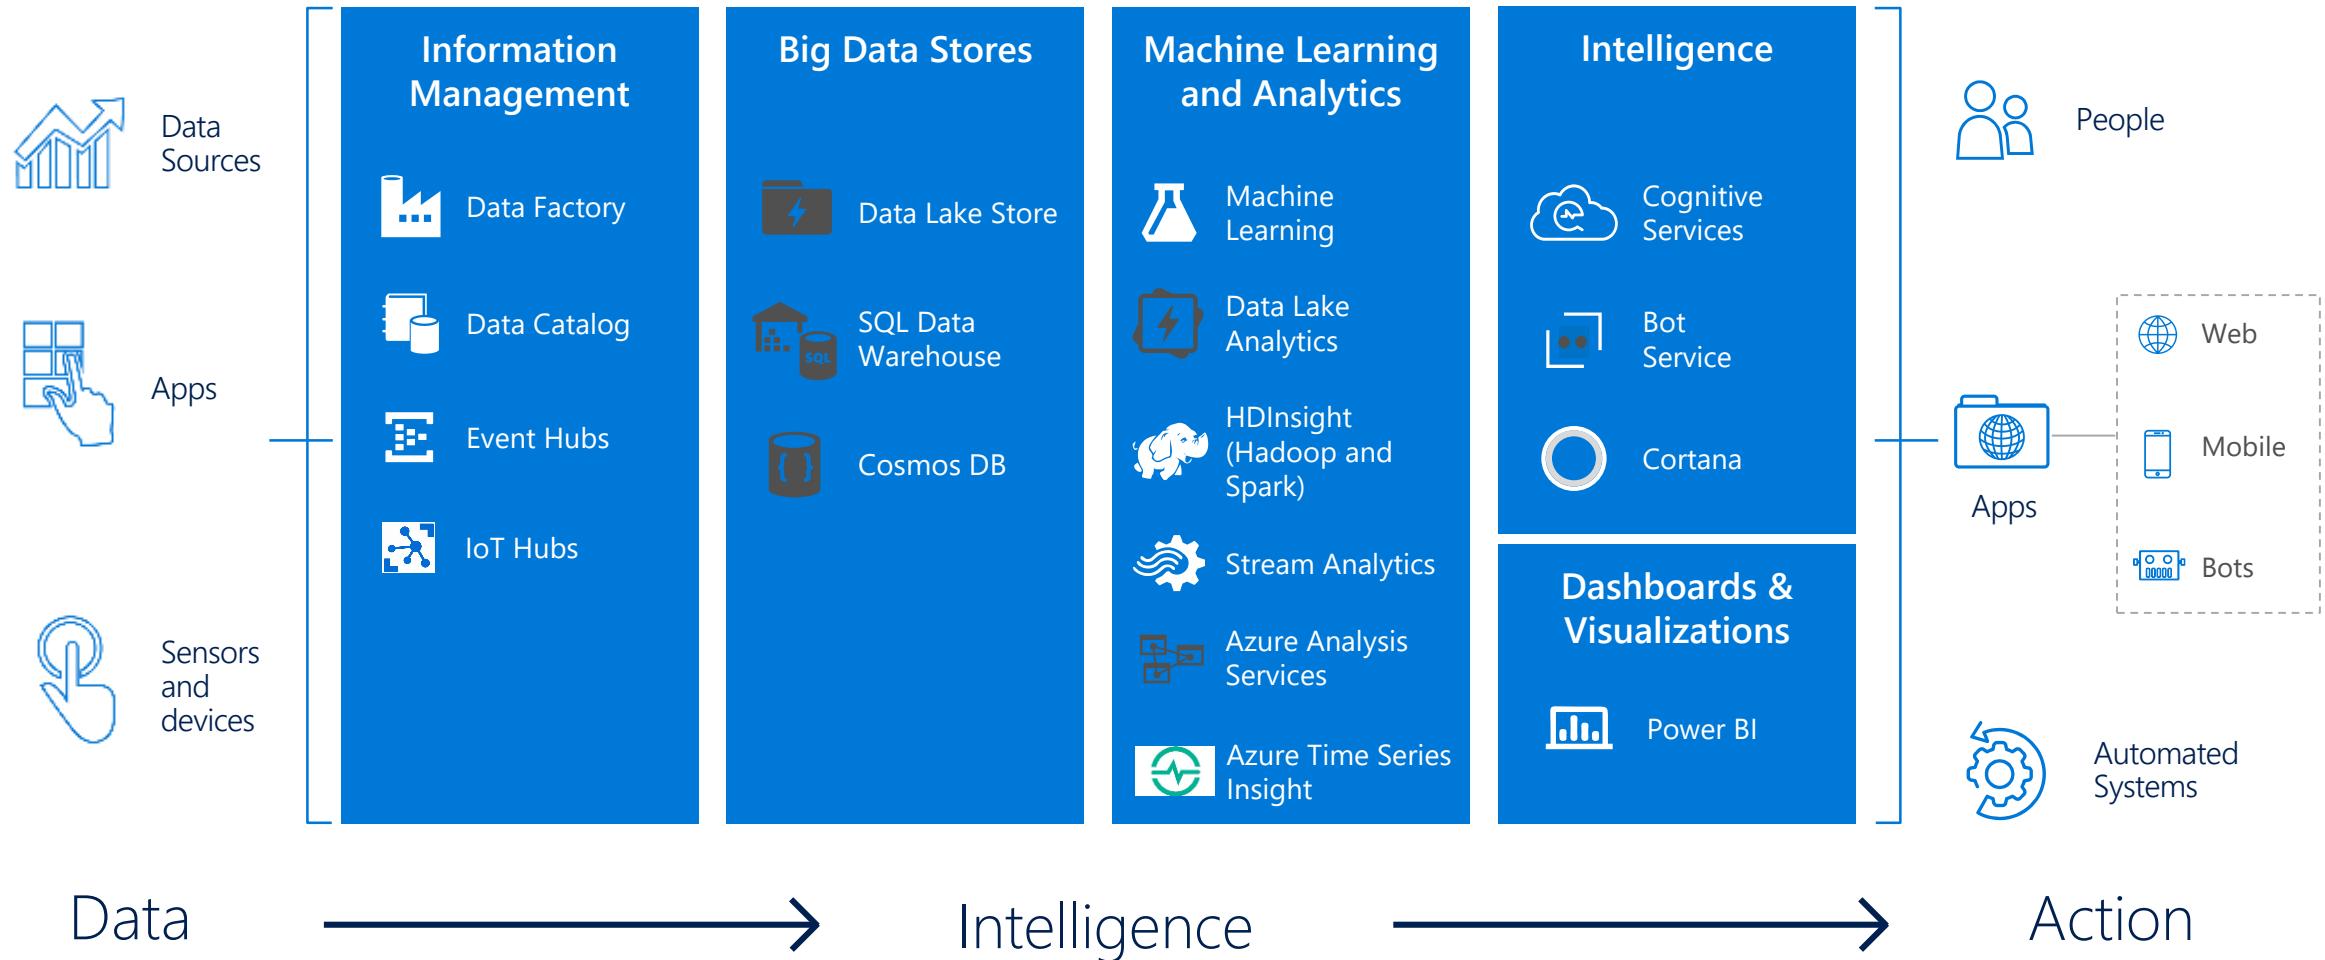


Power BI



- Analytics for everyone, even non-data experts
- Your whole business on one dashboard
- Create stunning, interactive reports
- Drive consistent analysis across your organization
- Embed visuals in your applications
- Get real-time alerts when things change

# Transform data into intelligent action





# Introduction to Machine Learning

Azure Machine Learning



# Why Machine Learning is the Future?

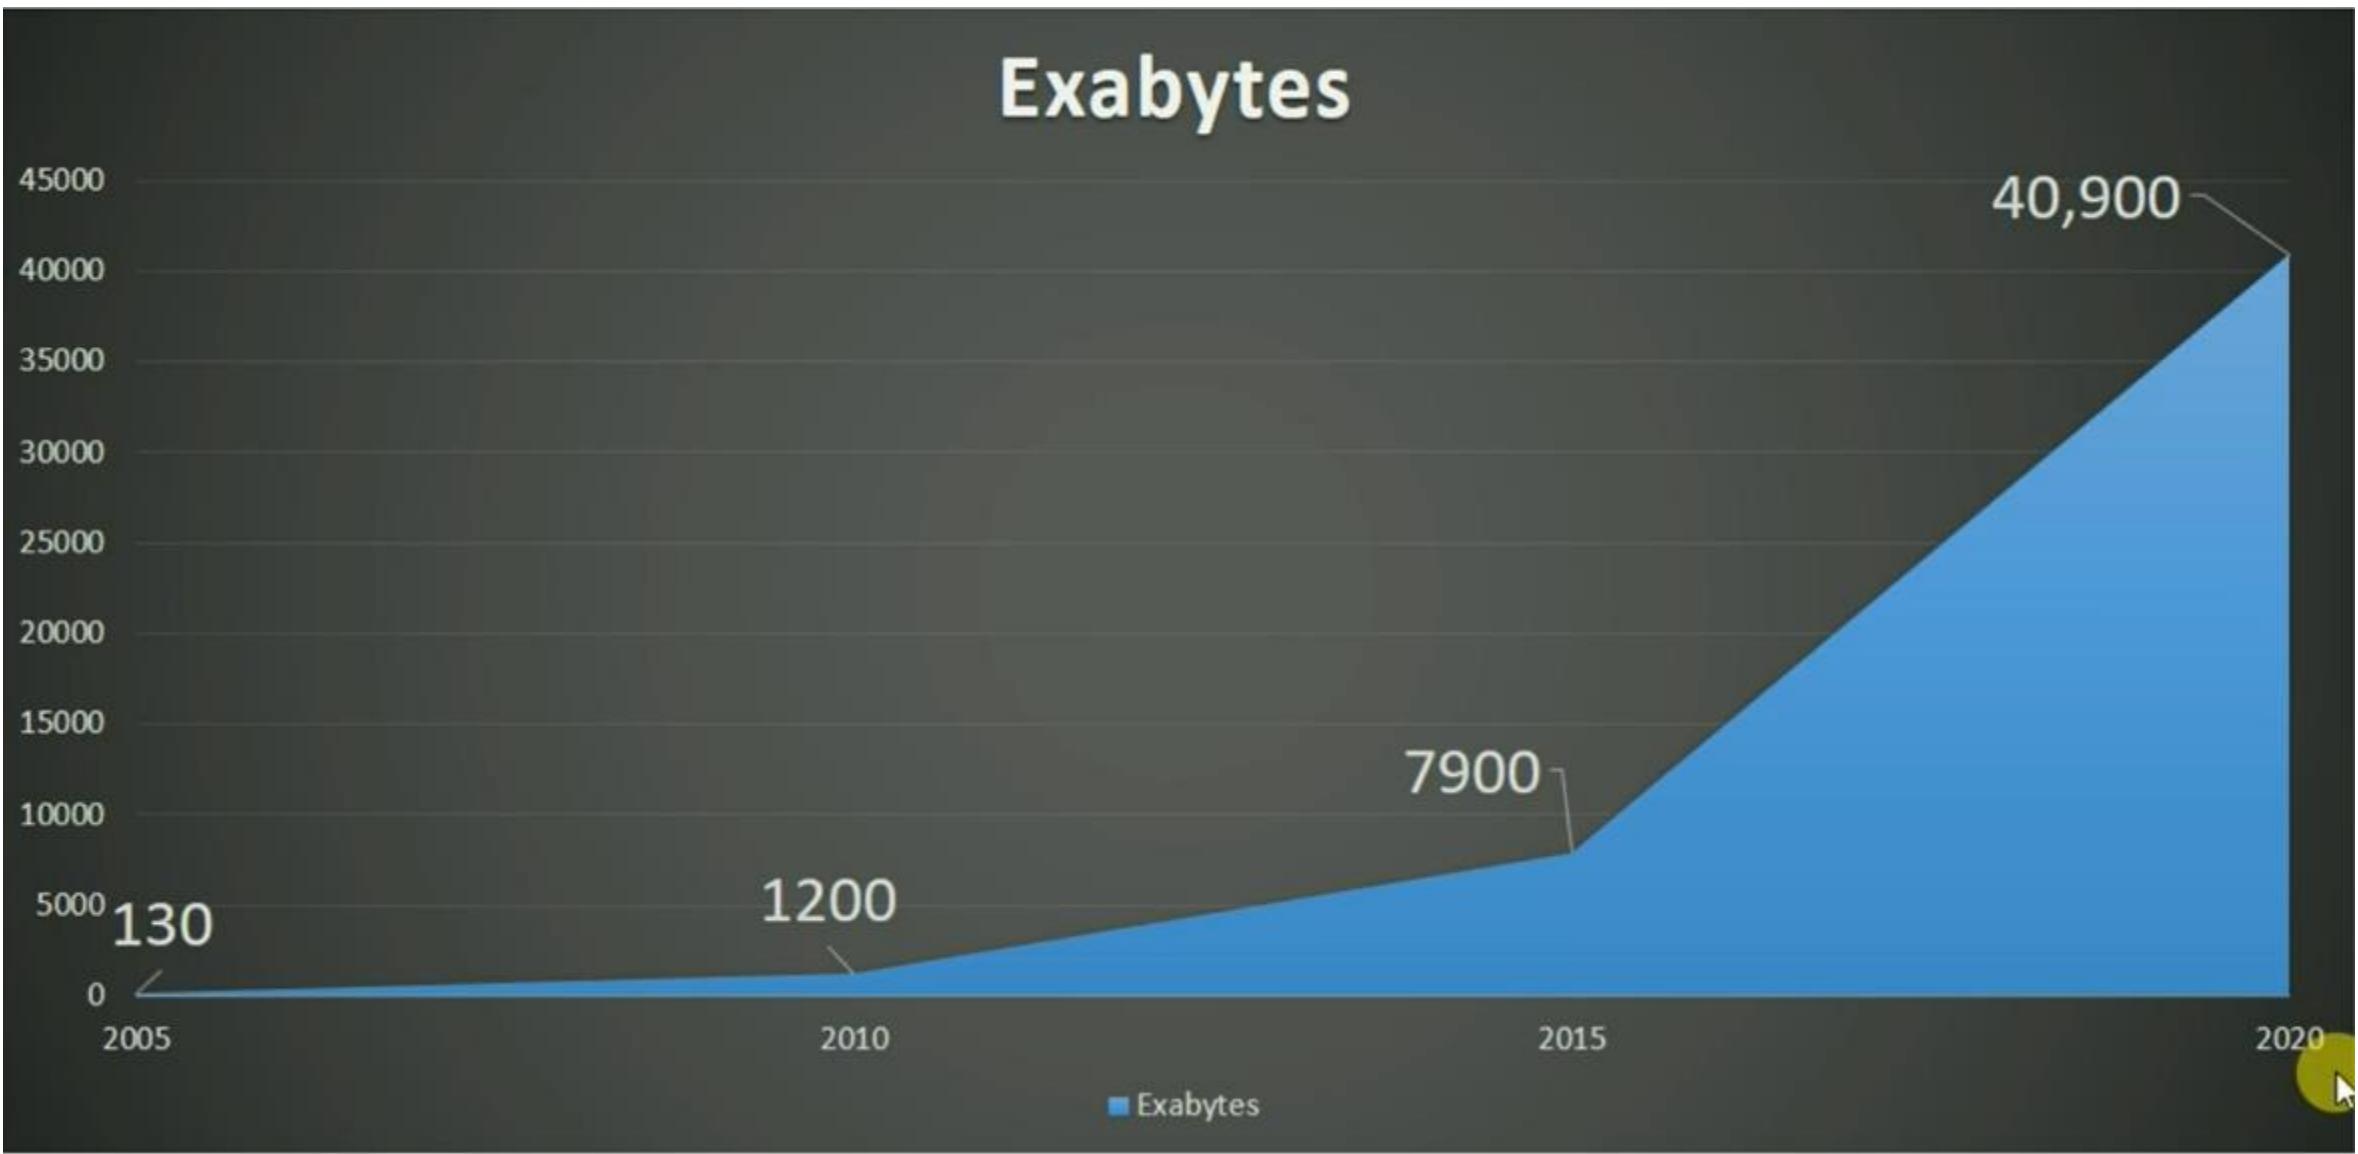
# Measurement of Data

1 TB = 1000 GB

1 PB = 1000 TB

1 EB = 1000 PB

# Growth of Data



# Statistics on Data Growth



90% of today's data has been created in last two years alone

# Statistics on Data Growth



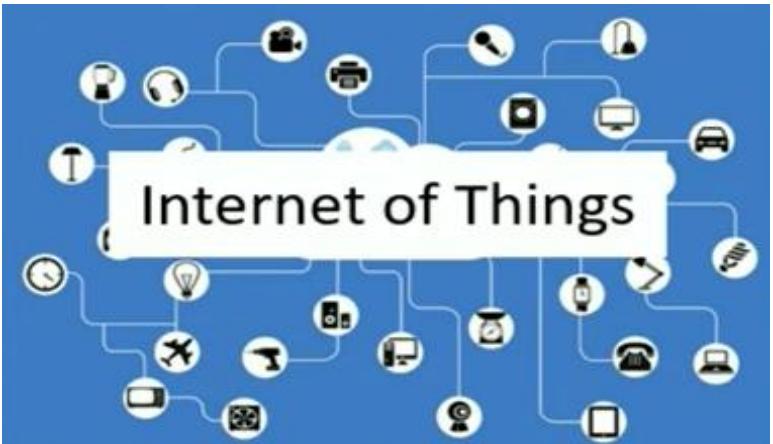
Biotech



Healthcare



Automotive



Internet of Things



Manufacturing



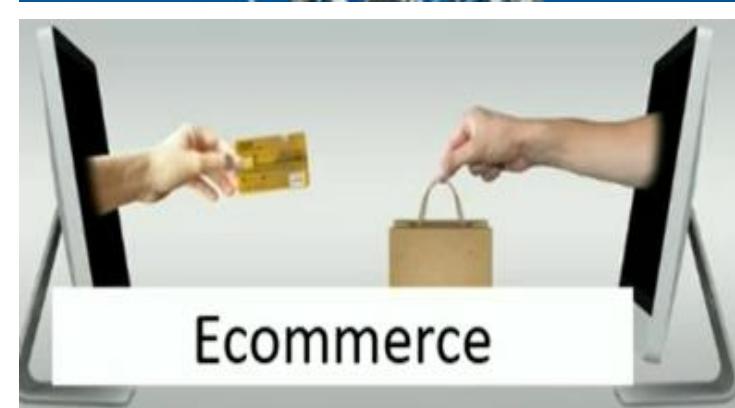
Telecom



Banking and Finance



Social Media



Ecommerce

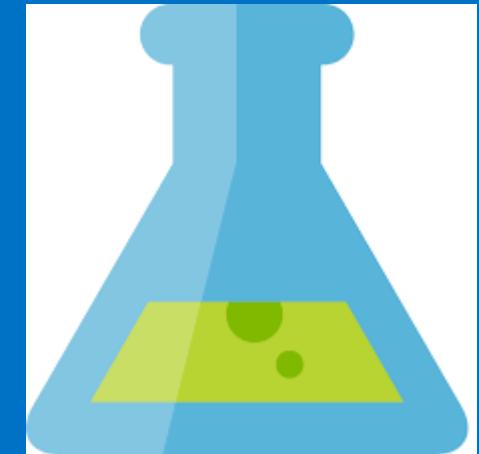
# Heard on the Streets

- IDC Futurescape - Two-thirds of Global 2000 Enterprises CEOs will centre their corporate strategy on digital transformation including machine learning (ML) solutions.
- Harvard Business Review - Data Scientist: The Sexiest Job of the 21st Century
- McKinsey Report - 45 percent of work activities could potentially be automated by currently demonstrated technologies; machine learning can be an enabling technology for the automation of 80 percent of those activities.
- Microsoft CEO Satya Nadella - called out machine learning and the big data that powers it as a key development in his memo to Microsoft last July.



# Agenda – Day1

1. Introduction to Machine Learning
2. How can we leverage Azure Machine Learning
3. Data Preparation in AML Studio
4. Classification
5. Regression
6. Hyper tuning Parameters in AML Studio
7. Deployment of Models using AML Studio.



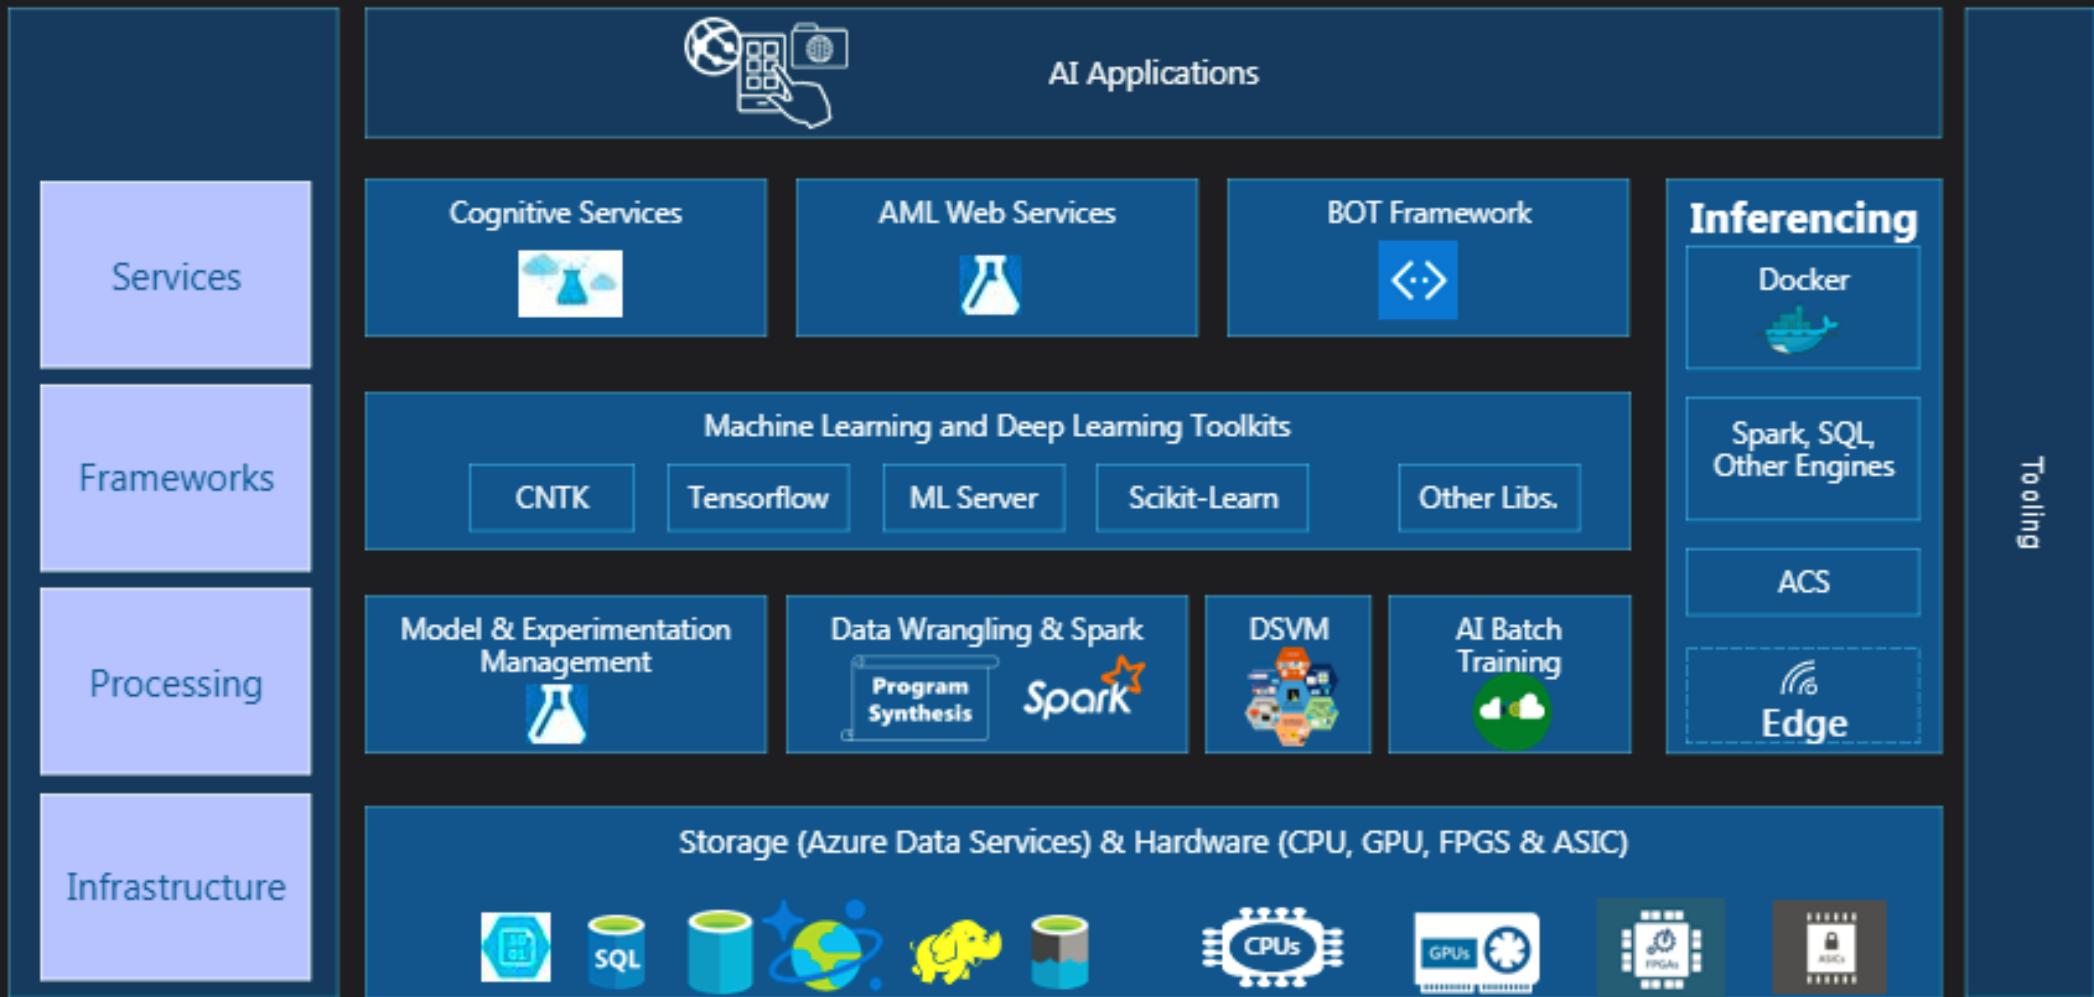
# Agenda – Day2

1. Clustering
2. Recommendation Engine in AML Studio
3. Azure Machine Learning Services
4. Deep Learning & Image Classification
5. AML Work Bench Advanced Data Preparation
6. Building Models in AML WB
7. Deploying Locally & in Dockers.
8. Creating Webservices in Model Management.
9. Demo on Image Classification in AML Workbench.





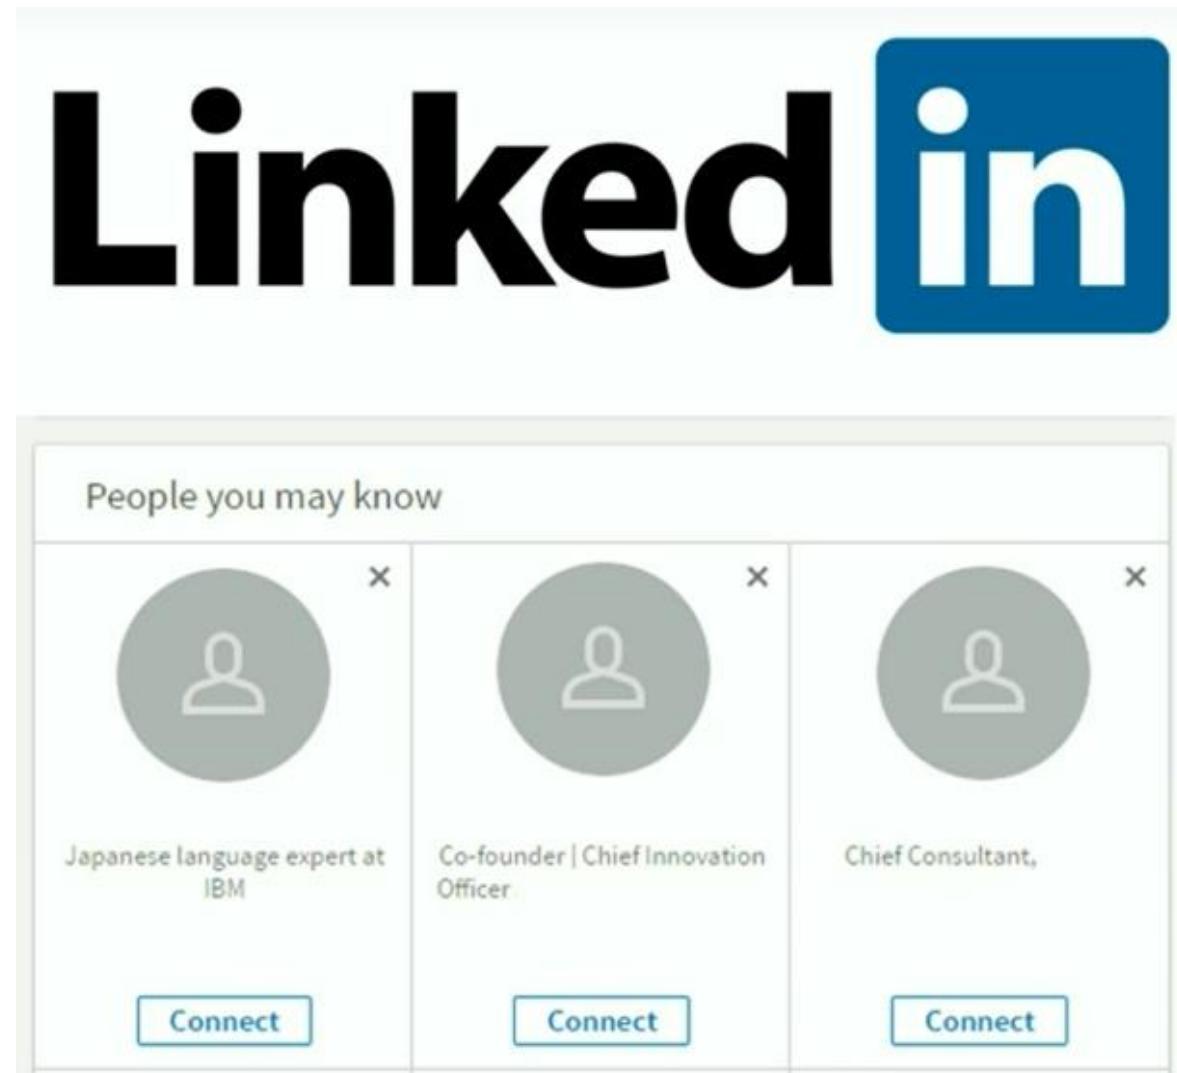
# Cloud AI Stack



# Machine Learning in Real World

# Linkedin

- "People You May Know" ads achieved a click-through rate 30% higher
- They generated millions of new page views.
- Thanks to this one feature, LinkedIn's growth trajectory shifted significantly upward.



# Few More Machine Learning Examples



Customers who



Weber One Touch Gold  
Premium Charcoal  
Grill-57cm

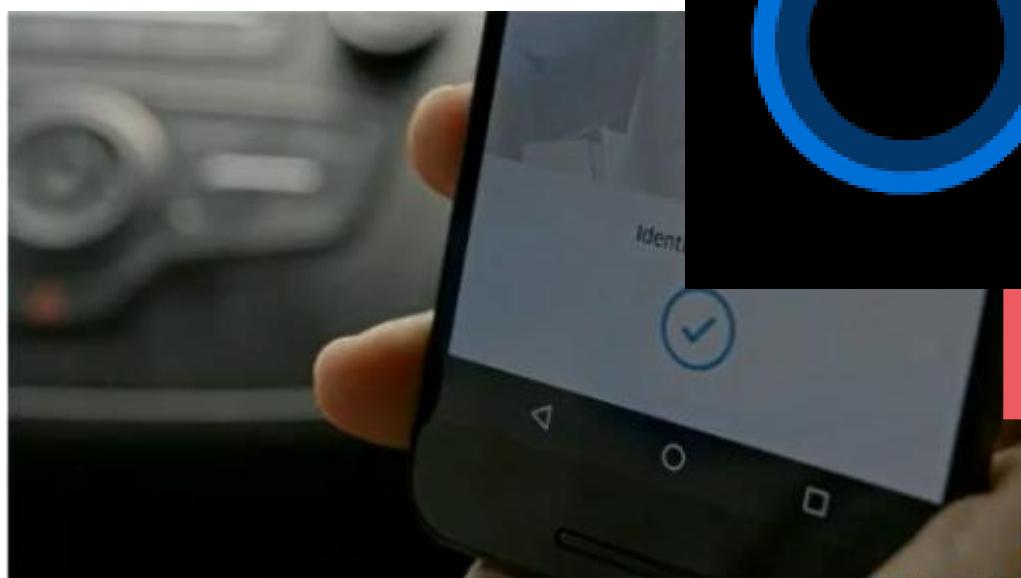
\$225

Add to cart

NoMU Salt Pepper and  
Spice Grinders

\$3

View options



Hi. I'm Cortana.  
Ask me a question!

NETFLIX

Movie Recommendations

# Why Azure ML?

- Drag and Drop interface and no Programming required
- Large variety of algorithm as modules
- From experiment to production API in minutes
- Supports R and Python to bring in your existing code
- Flexibility of data storage; supports variety of data storage options
- Large number of pre-built APIs available as a service

TATA  
MOTORS



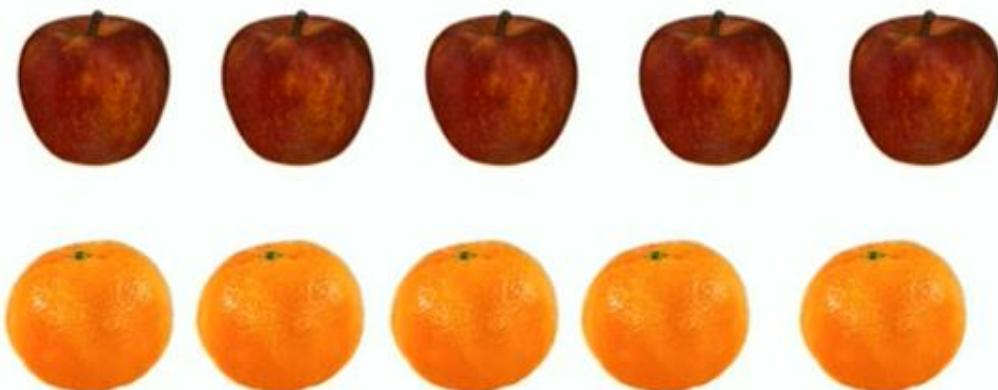
Rolls-Royce



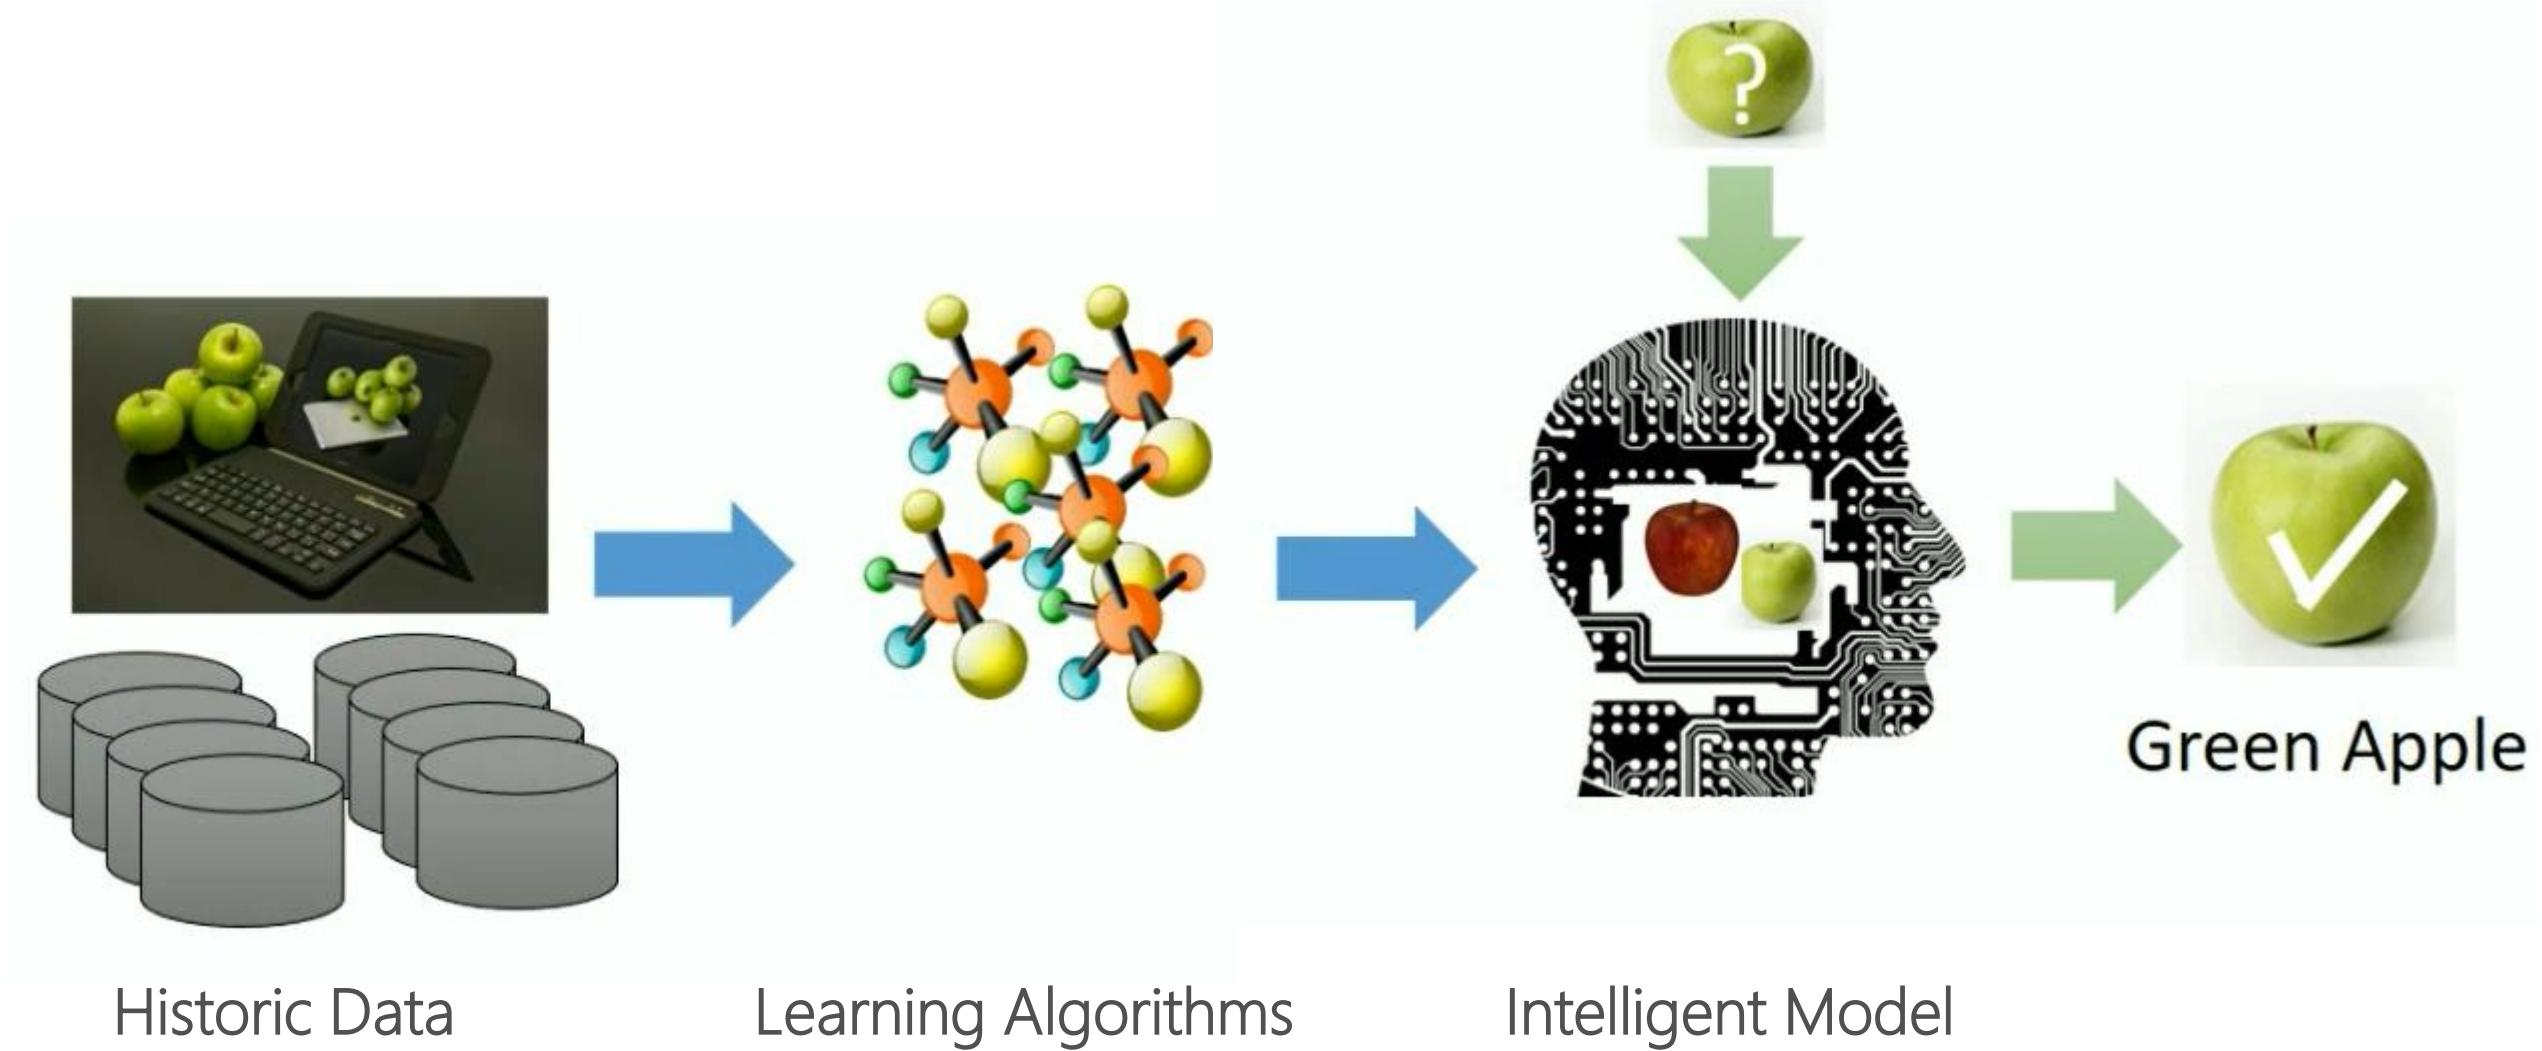
# What is Machine Learning?

# Machine Learning

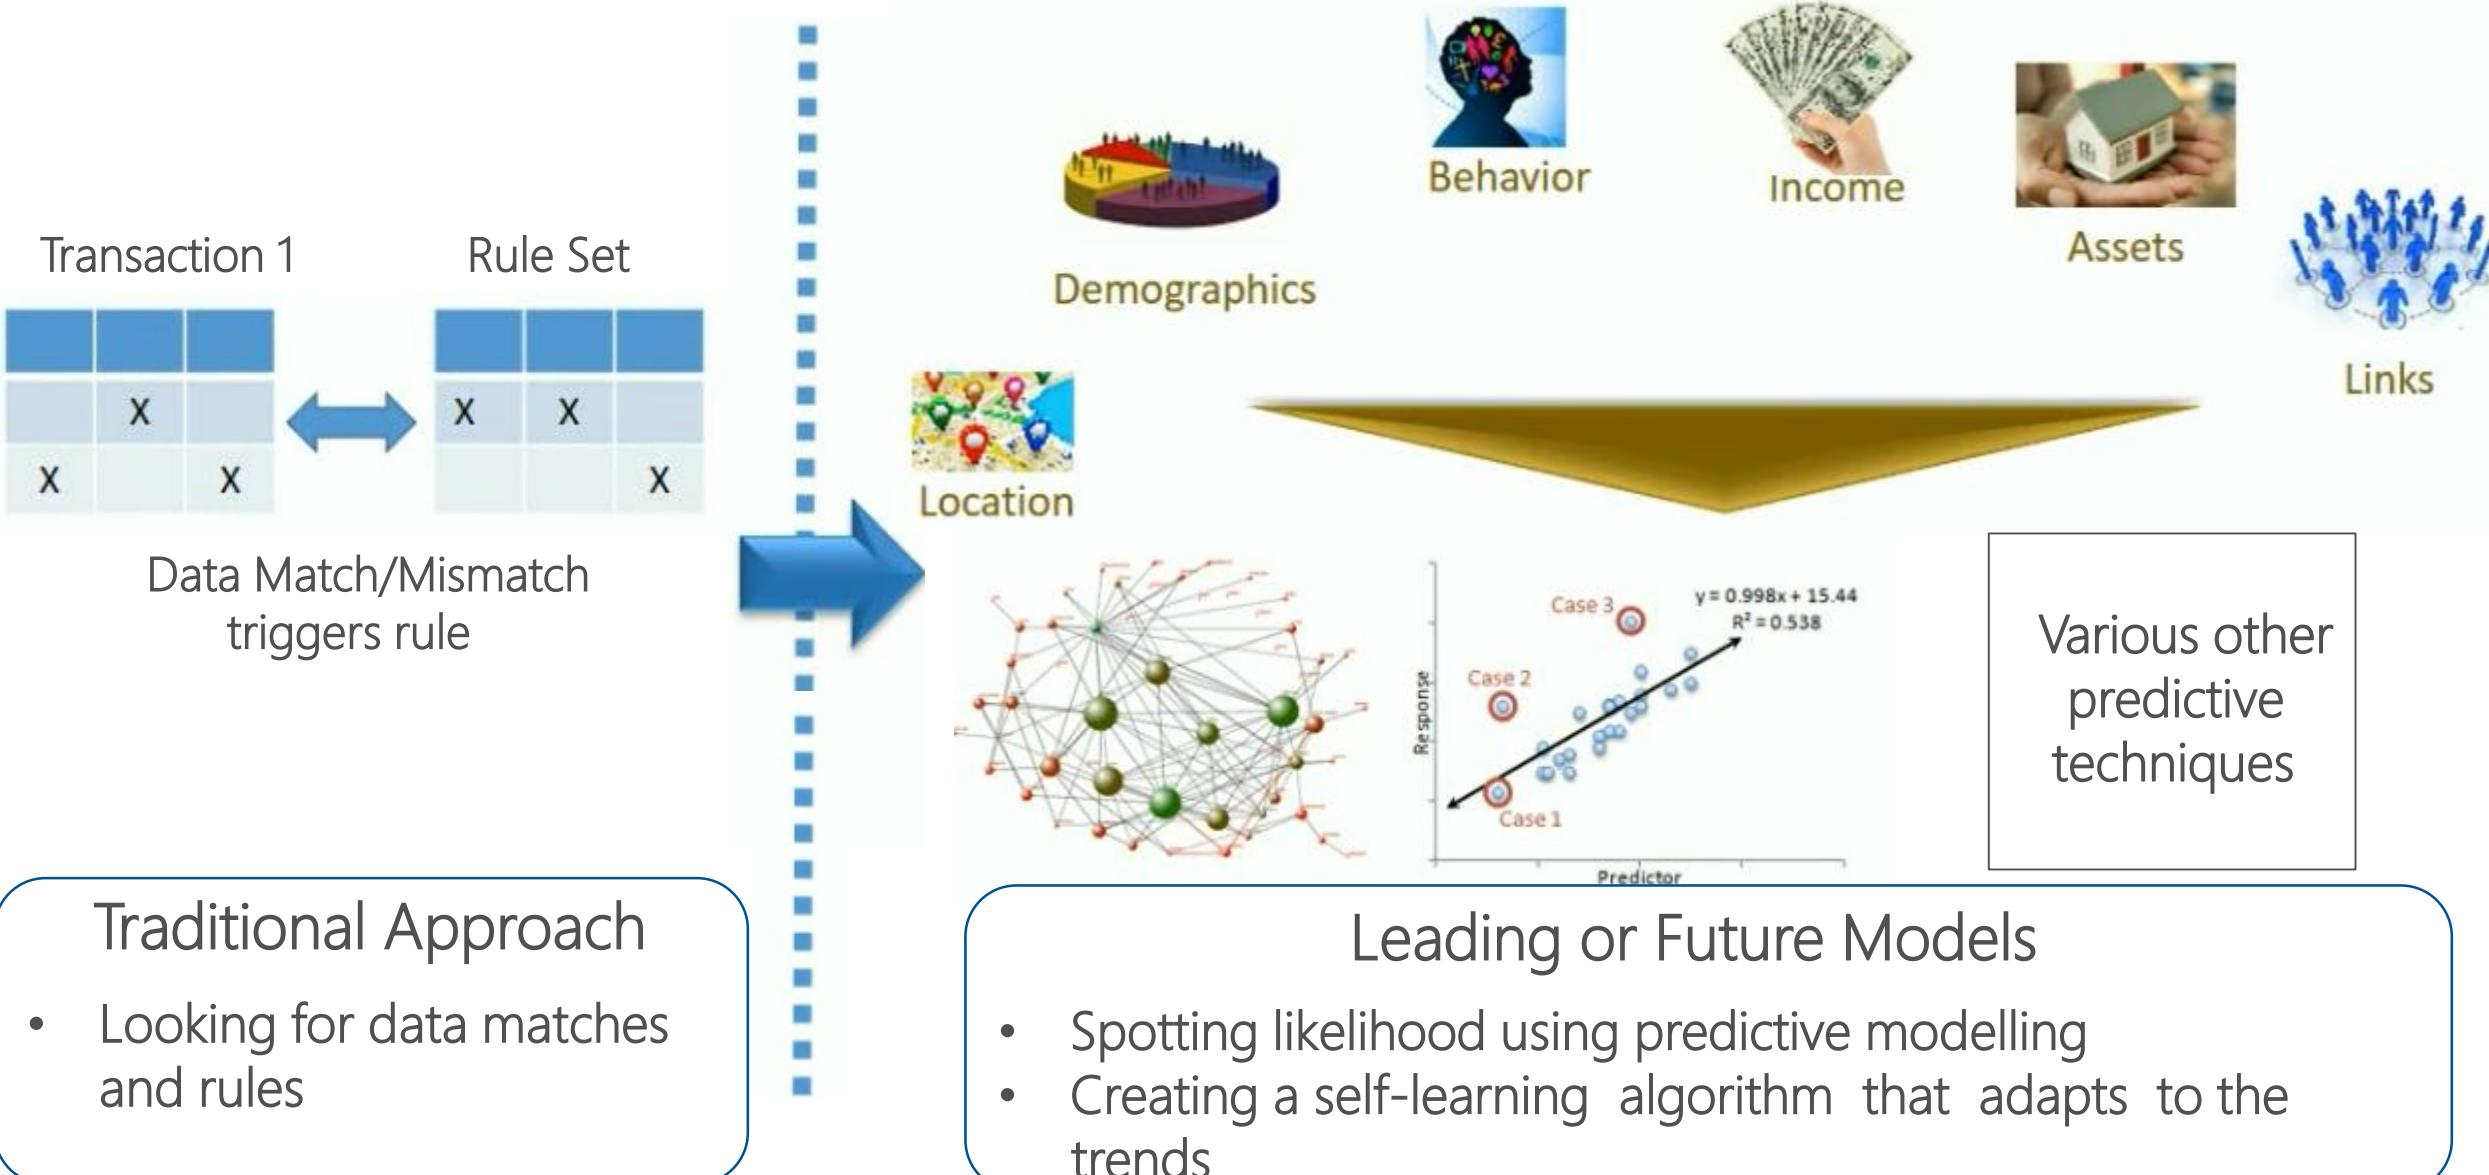
- Machine learning is the subfield of computer science that gives “computers the ability to learn without being explicitly programmed”.



# How Machines Learn?



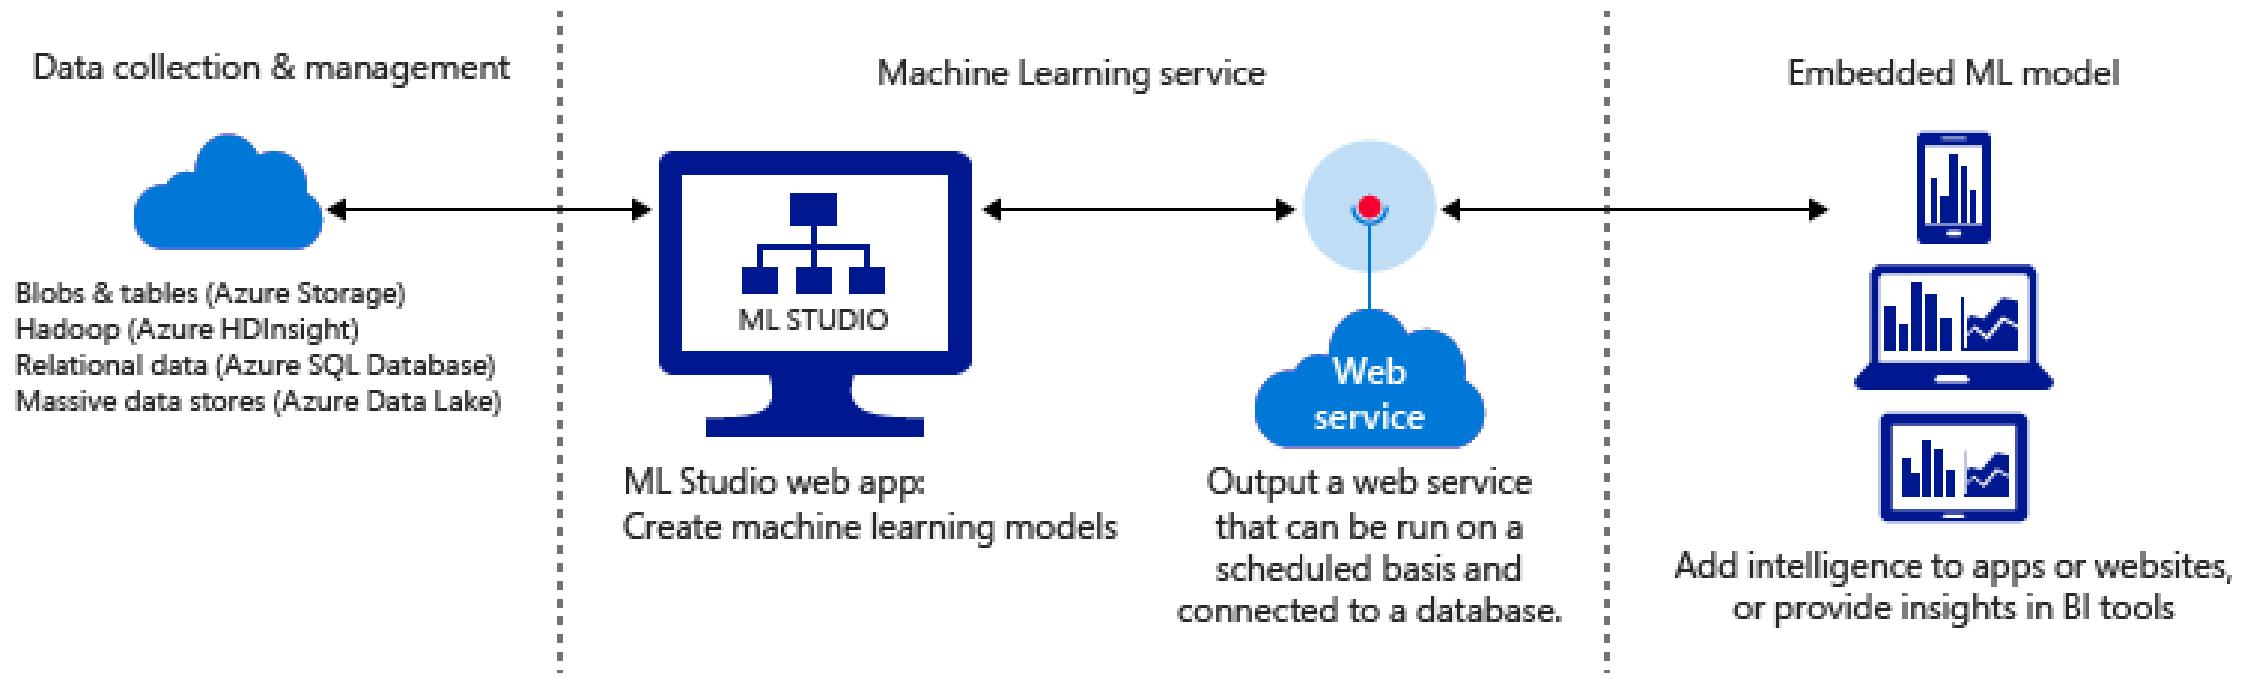
# Futuristic practices and models



# What is Machine Learning in the Microsoft Azure cloud?

## Azure Machine Learning: Basic workflow

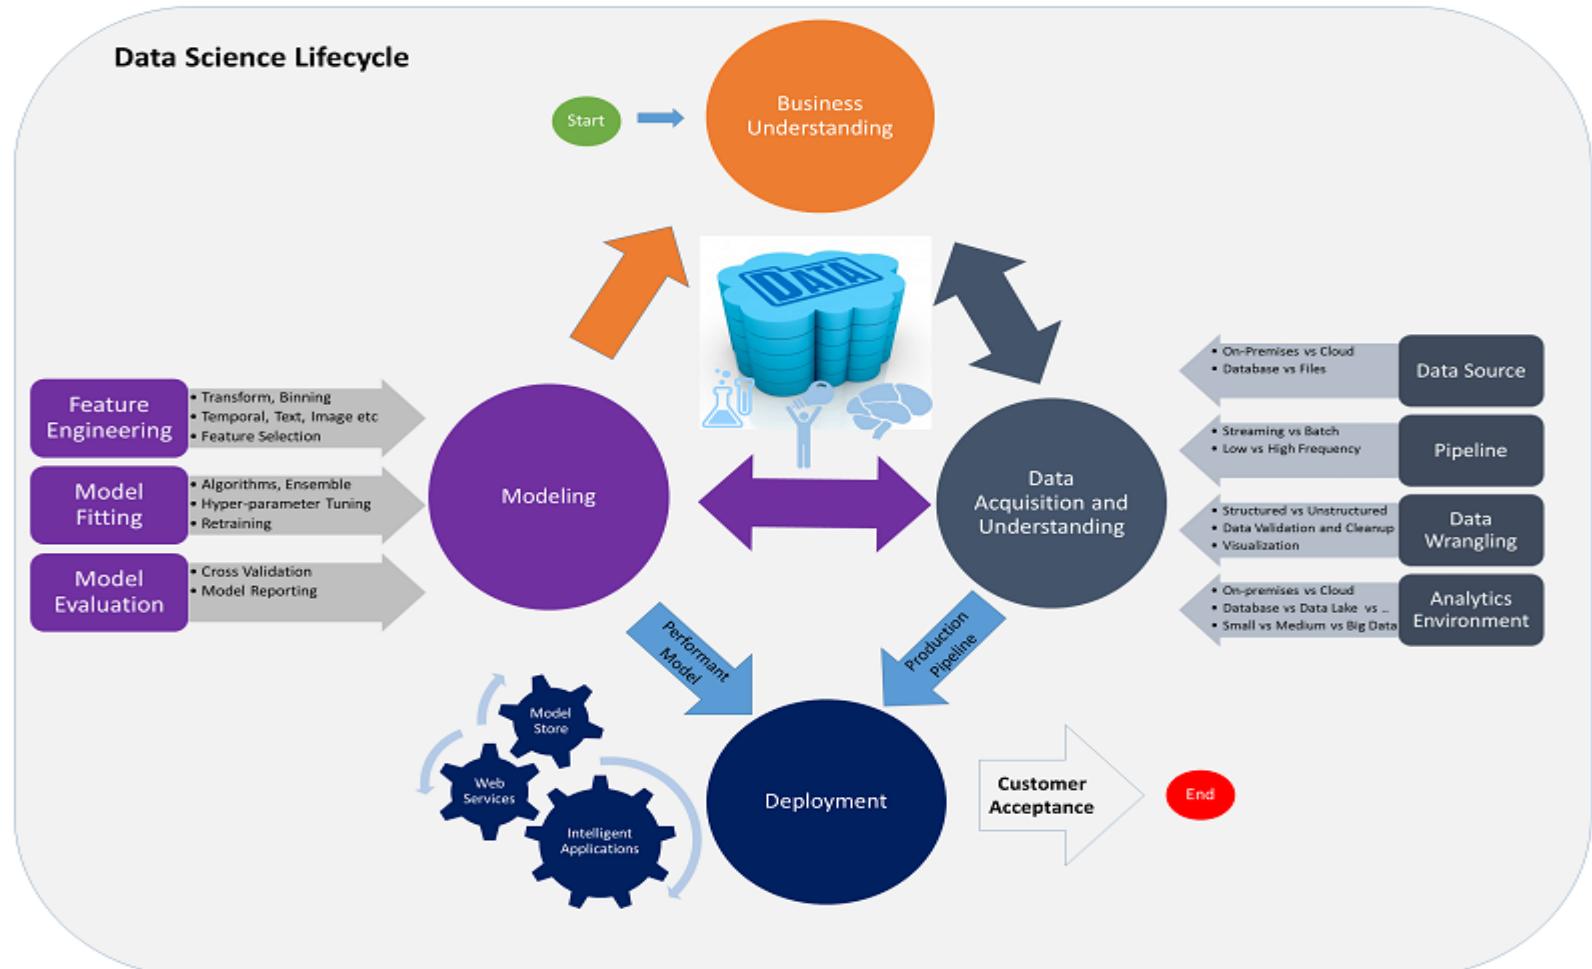
**Build models from data and operationalize a machine learning solution**



# Data Science Process

The TDSP lifecycle is composed of five major stages that are executed iteratively. These include:

- **1. Business Understanding**
- **2. Data Acquisition and Understanding**
- **3. Modeling**
- **4. Deployment**
- **5. Customer Acceptance**



# The 5 questions data science answers

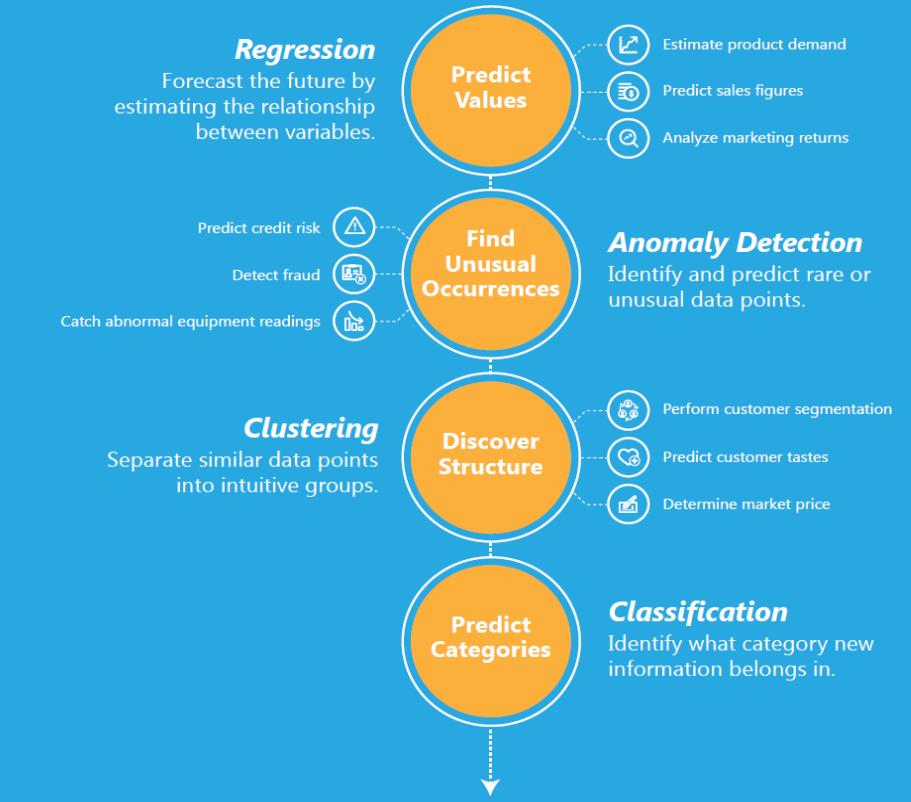
It might surprise you, but *there are only five questions that data science answers:*

- **Is this A or B?**
- **Is this weird?**
- **How much – or – How many?**
- **How is this organized?**
- **What should I do next?**

Each one of these questions is answered by a separate family of machine learning methods, called algorithms.

So, what do you want to find out?

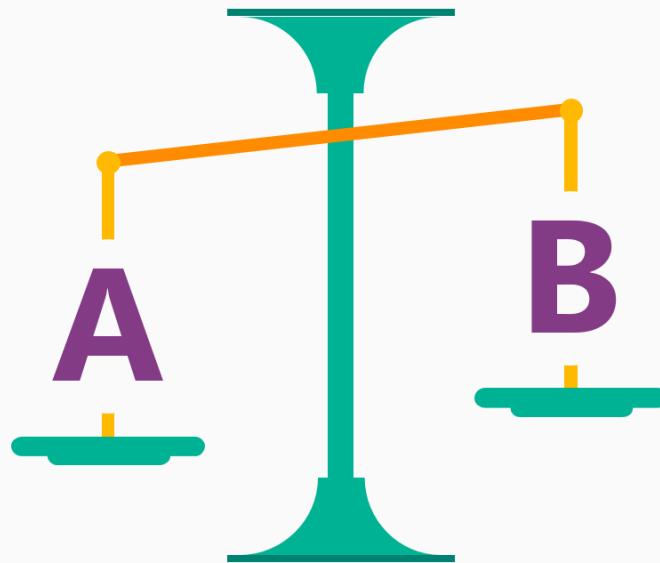
I WANT TO:



# Question 1: Is this A or B? uses classification algorithms

Is this A or B?

Classification algorithms



# Question 2: Is this weird? uses anomaly detection algorithms

Is this weird?

Anomaly detection algorithms



# Question 3: How much? or How many? uses regression algorithms

How much? How many?

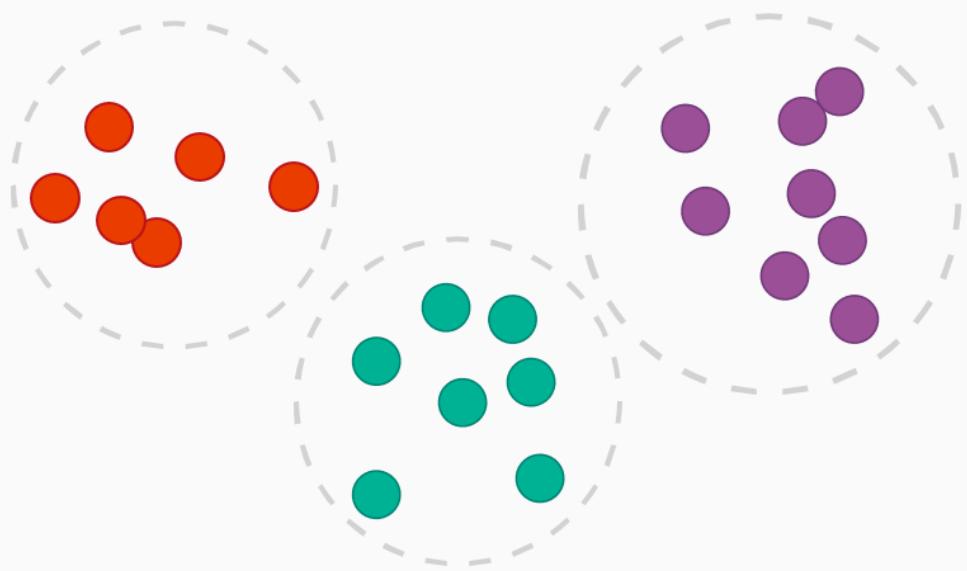
Regression algorithms



# Question 4: How is this organized? uses clustering algorithms

How is this organized?

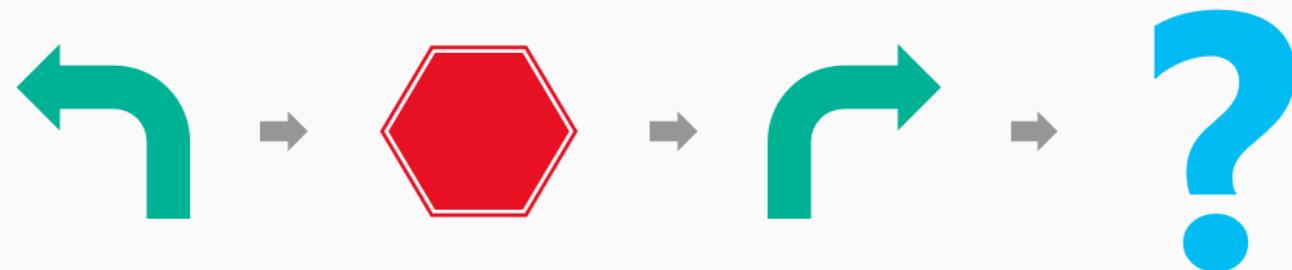
**Clustering Algorithms**



# Question 5: What should I do now? uses reinforcement learning algorithms

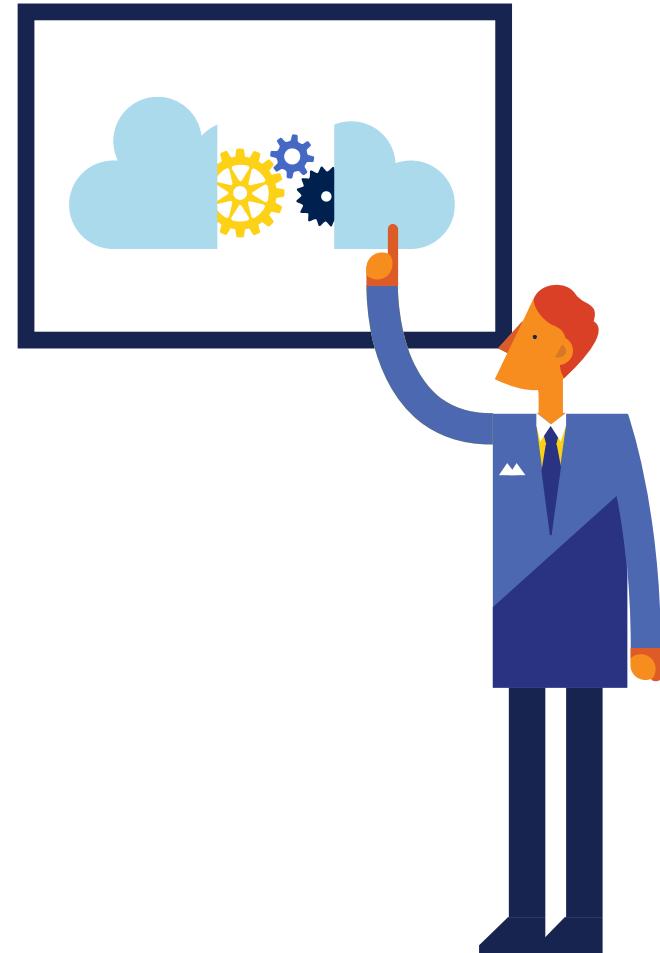
What should I do now?

**Reinforcement Learning Algorithms**



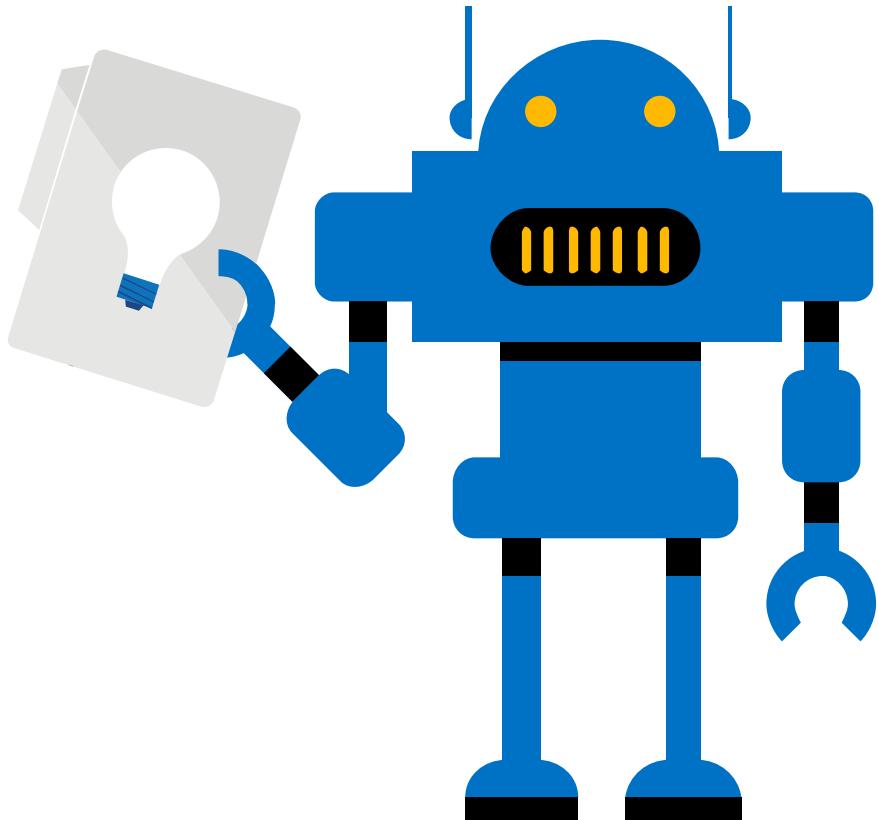
# Supervised learning

1. Purpose
2. Process: Data set and test set
3. Components
4. Computation and recommendations
5. Verification
6. Selection



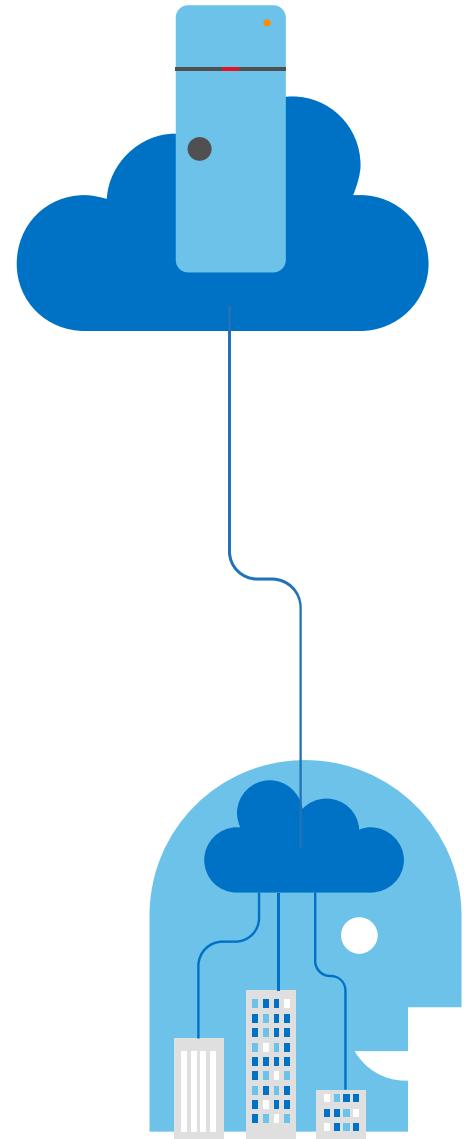
# Unsupervised learning

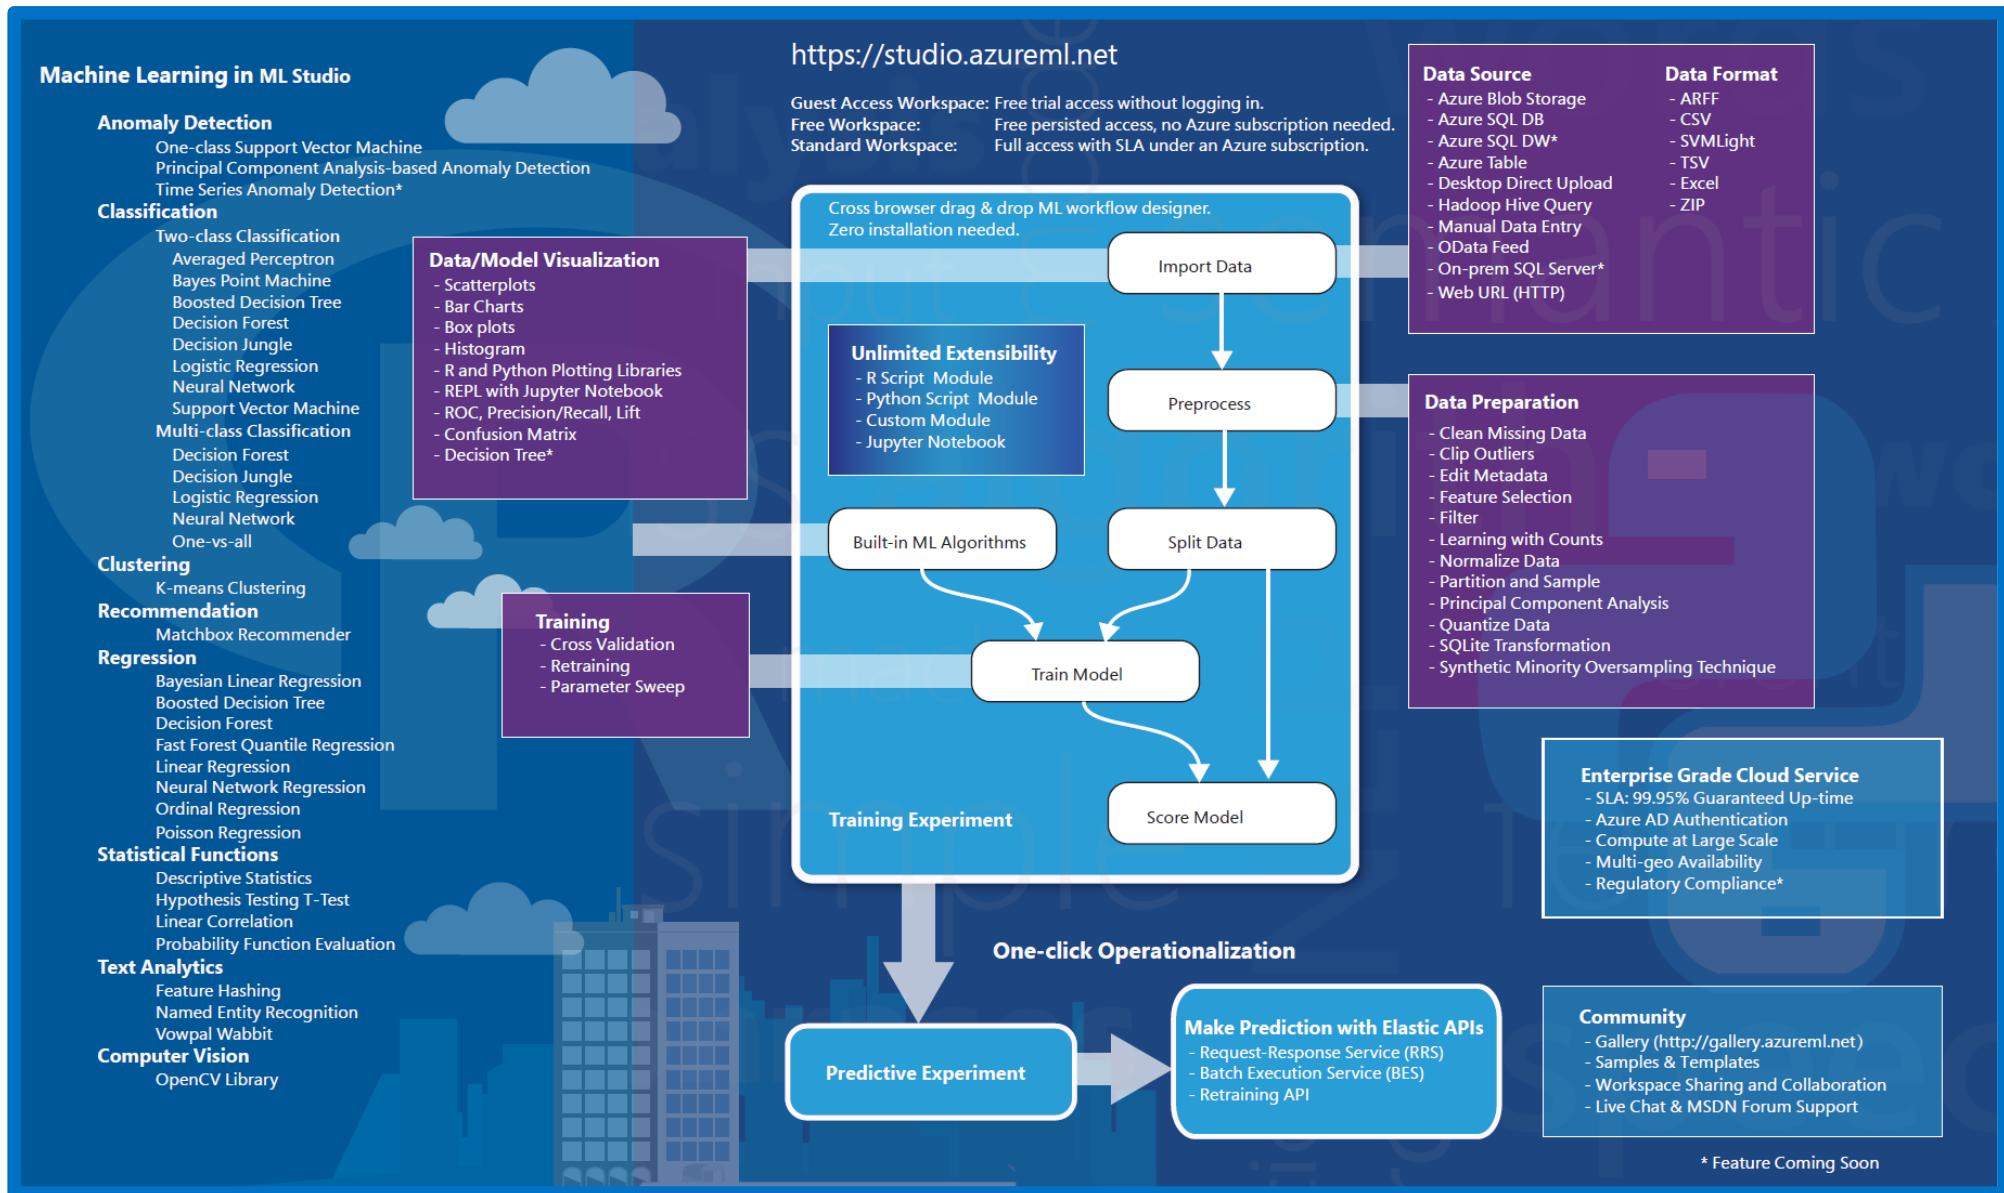
1. Purpose
2. Process
3. Components
4. Computation and recommendation



# Azure Machine Learning

- Why Azure Machine Learning?
- Getting started
  - Setting up a Microsoft Azure account
  - Setting up a storage account
  - Setting up an Azure Machine Learning workspace
- Advantages





# The algorithms clubbed

- Classification: Predict what class case belongs to

Scenarios: Churn analysis, fraud detection, speech recognition

Algorithms: Boosted Decision Tree, Decision Forest, Logistic Regression,...

- Regression: Predict numerical outcomes

Scenarios: Stock prices prediction, sales forecasts, quality control

Algorithms: Bayesian Linear, Linear Regression, Decision Forest

- Clustering: Discover natural groupings of cases

Scenarios: Customer segmentation, pattern recognition

Algorithms: K-means

- Anomaly Detection: Explore unusual patterns

Scenarios: Network intrusion, fraud detection

Algorithms: One class support vector machine ,PCA based anomaly detection

# Understanding Data, Variables/Features

# Understanding The Variables Using a Dataset

Loan_ID	Gender	Married	Dependents	Self_Employed	Income	LoanAmt	Term	CreditHistory	Property_Area	Status
LP001002	Male	No	0	No	\$5,849.00		60	1	Urban	Y
LP001003	Male	Yes	1	No	\$4,583.00	\$128.00	120	1	Rural	N
LP001005	Male	Yes	0	Yes	\$3,000.00	\$66.00	60	1	Urban	Y
LP001006	Male	Yes	2	No	\$2,583.00	\$120.00	60	1	Urban	Y

## Types of Variables

- Predictor/ Independent
  - Gender
  - Married
  - Dependents
  - Self\_Employed
  - Income
  - Loan Amt
  - Term
  - Credit History
  - Property Area
- Target/Dependent
  - Status

## Data Type

- Character/String
  - Gender
  - Married
  - Self-Employed
  - Property Area
  - Status
- Numeric
  - Dependents
  - Income
  - Loan Amt
  - Term
  - Credit History

## Category

- Categorical
  - Gender
  - Married
  - Credit History
  - Self-Employed
  - Property Area
  - Status
- Continuous
  - Dependents
  - Income
  - Loan Amt
  - Term

# Types of Models

# Types of Models

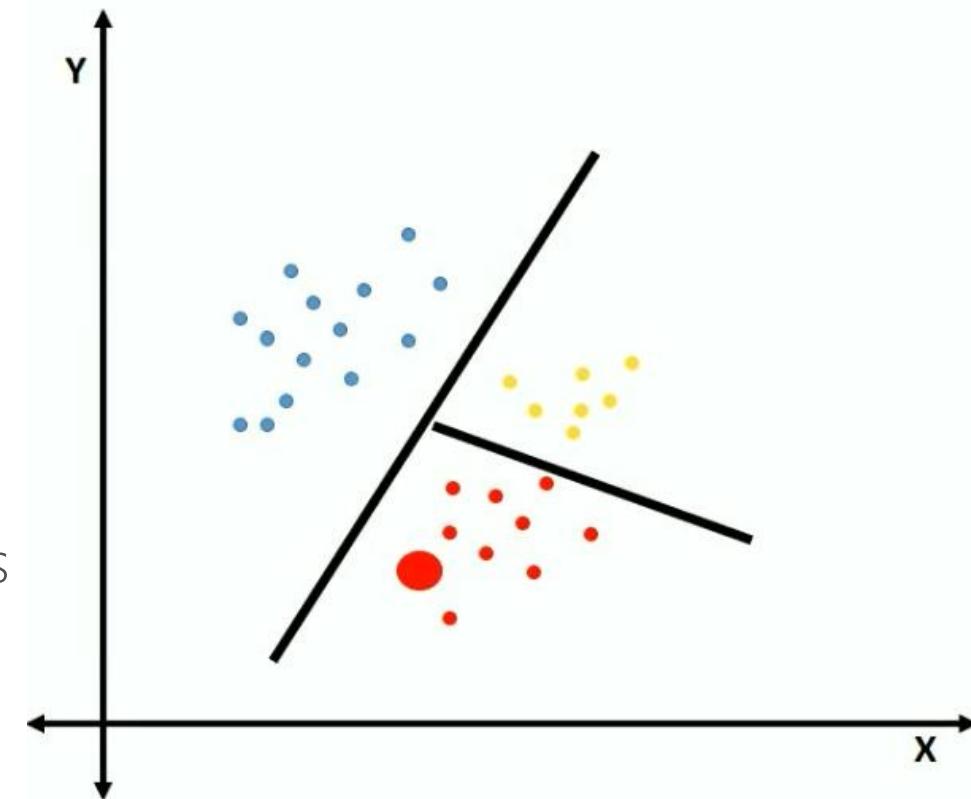
- Classification
- Regression
- Clustering
- Anomaly Detection

# Classification

- Classification is identifying to which set of categories a new observation belongs, on the basis of a training set of data containing observations whose category membership is known.
- Binary/ Two-Class Classification — Either/Or, Yes or No type
- Multi-Class Classification — One of the many alternatives

Some examples could be

- Assigning a given email into "spam" or "non-spam" classes Or Primary, Social or Promotional emails
- Will this customer default on loan repayment?
- Will this customer buy my product?



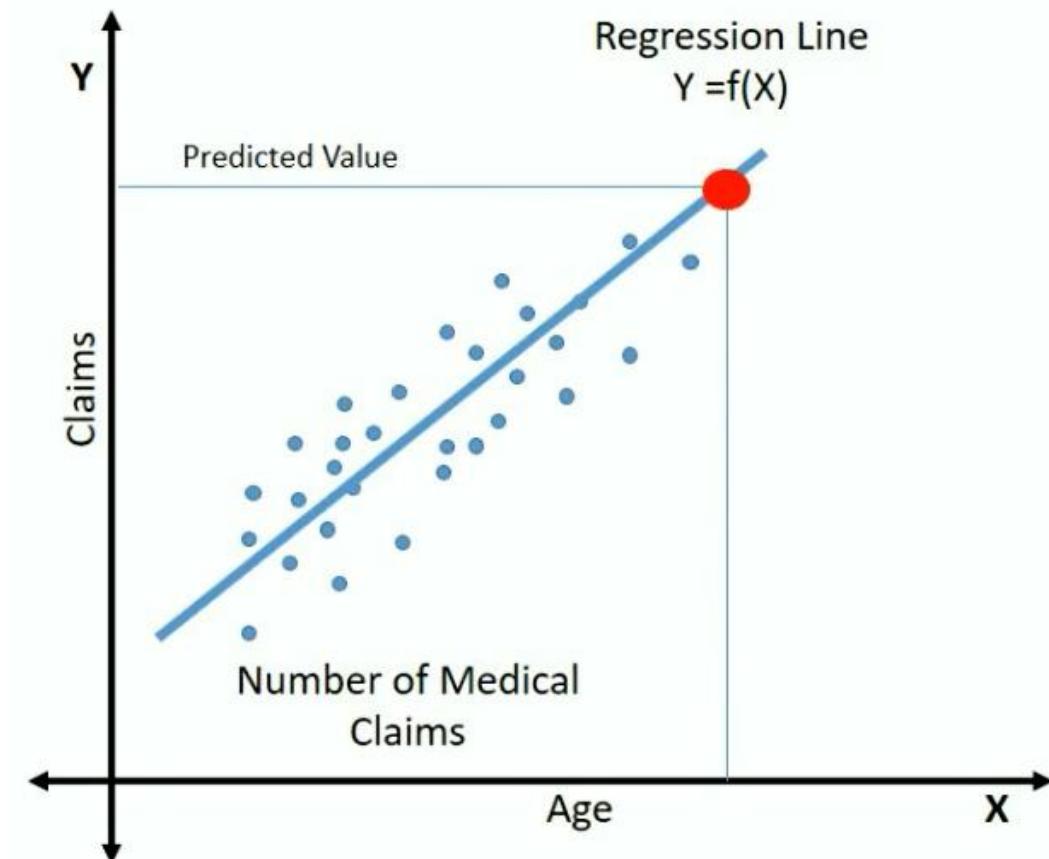
Predicting the value for categorical variable

# Regression Analysis

- Regression Analysis is a statistical process for estimating the relationships among variables where the predictor is a continuous variable
- The focus is on the relationship between a dependent variable and one or more independent variables (or 'predictors')
- One of the most common methods used in Machine Learning
- In certain circumstances, it can also be used to infer causal relationships between dependent and independent variables.

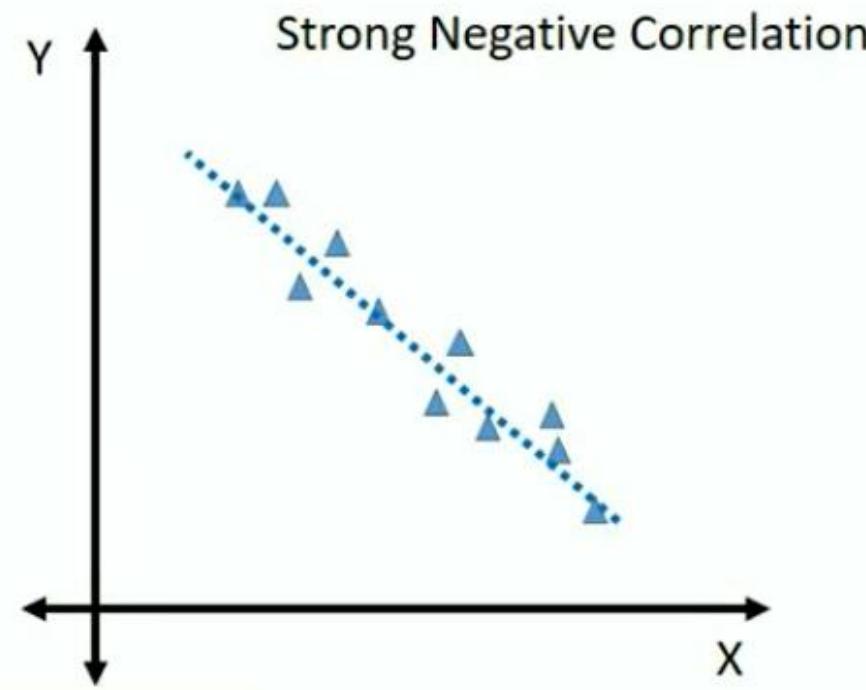
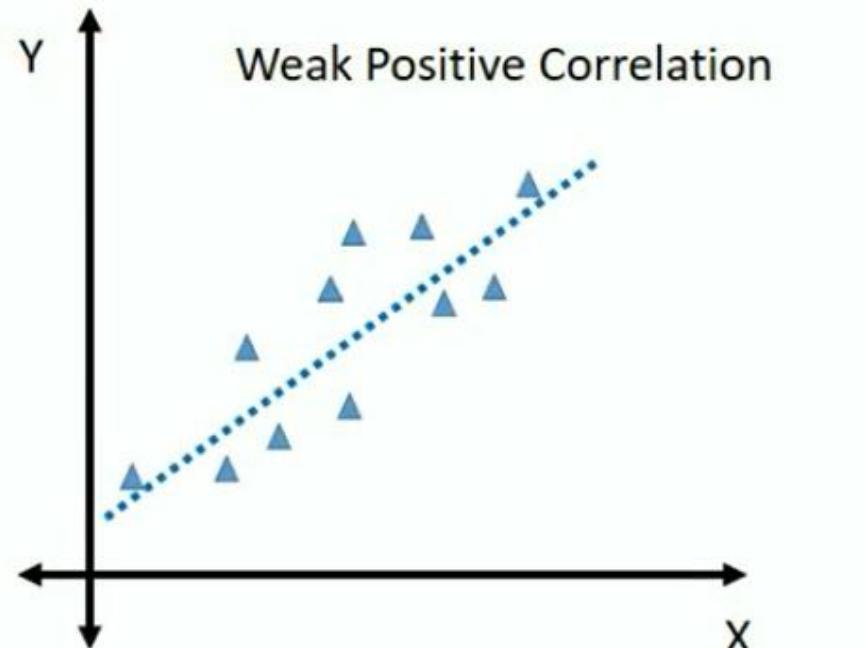
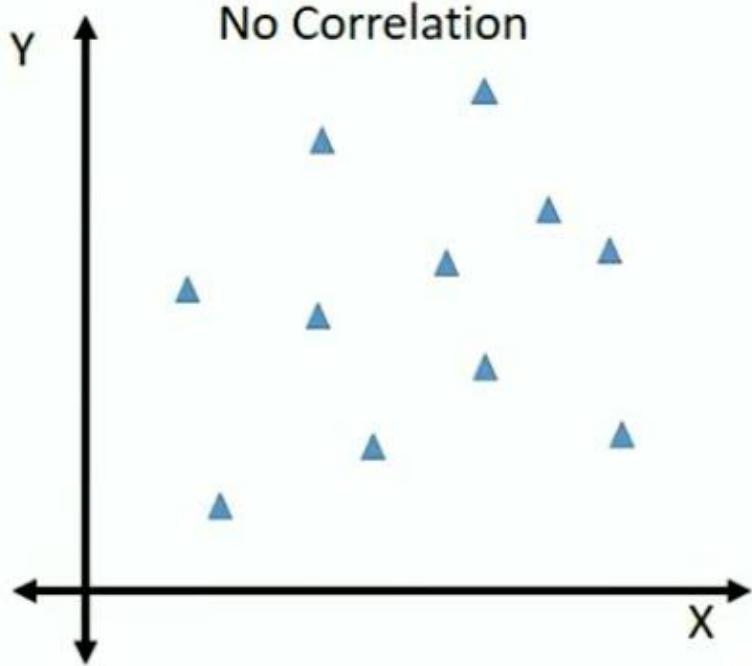
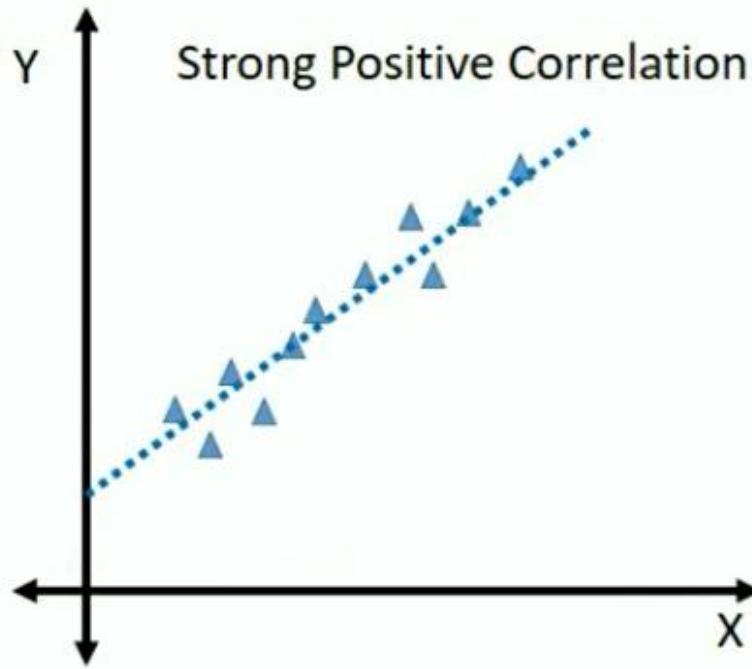
Examples

- Predicting the future sale of products
- Computing fair price of the product or service



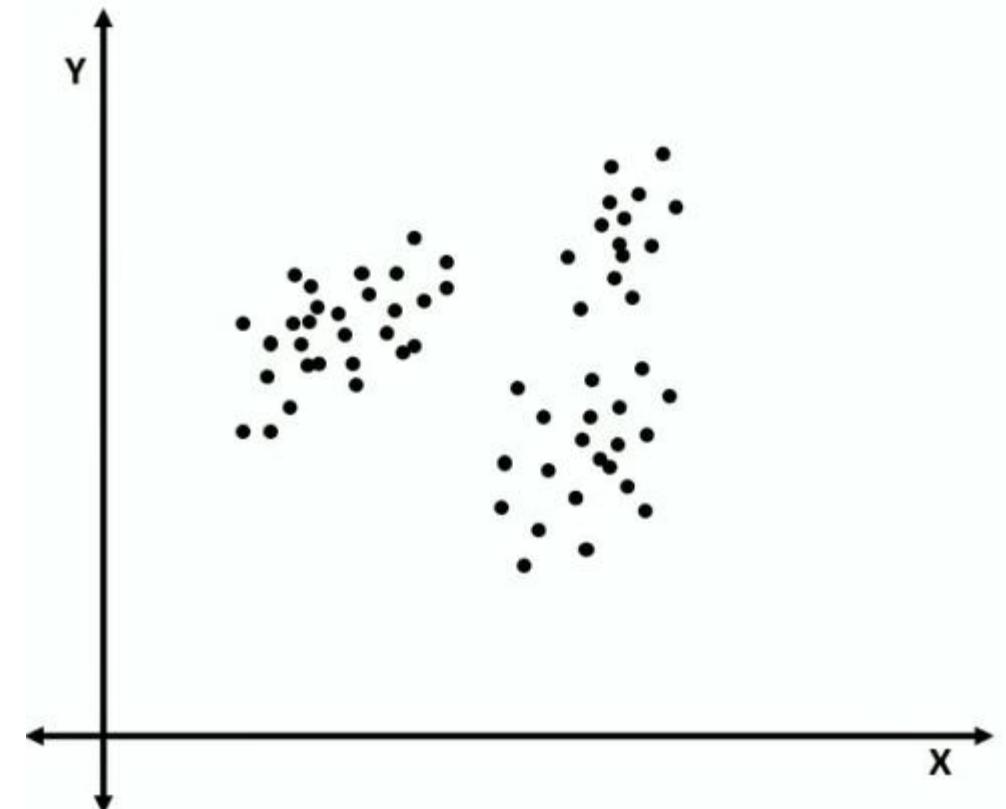
# Casual Relationship?



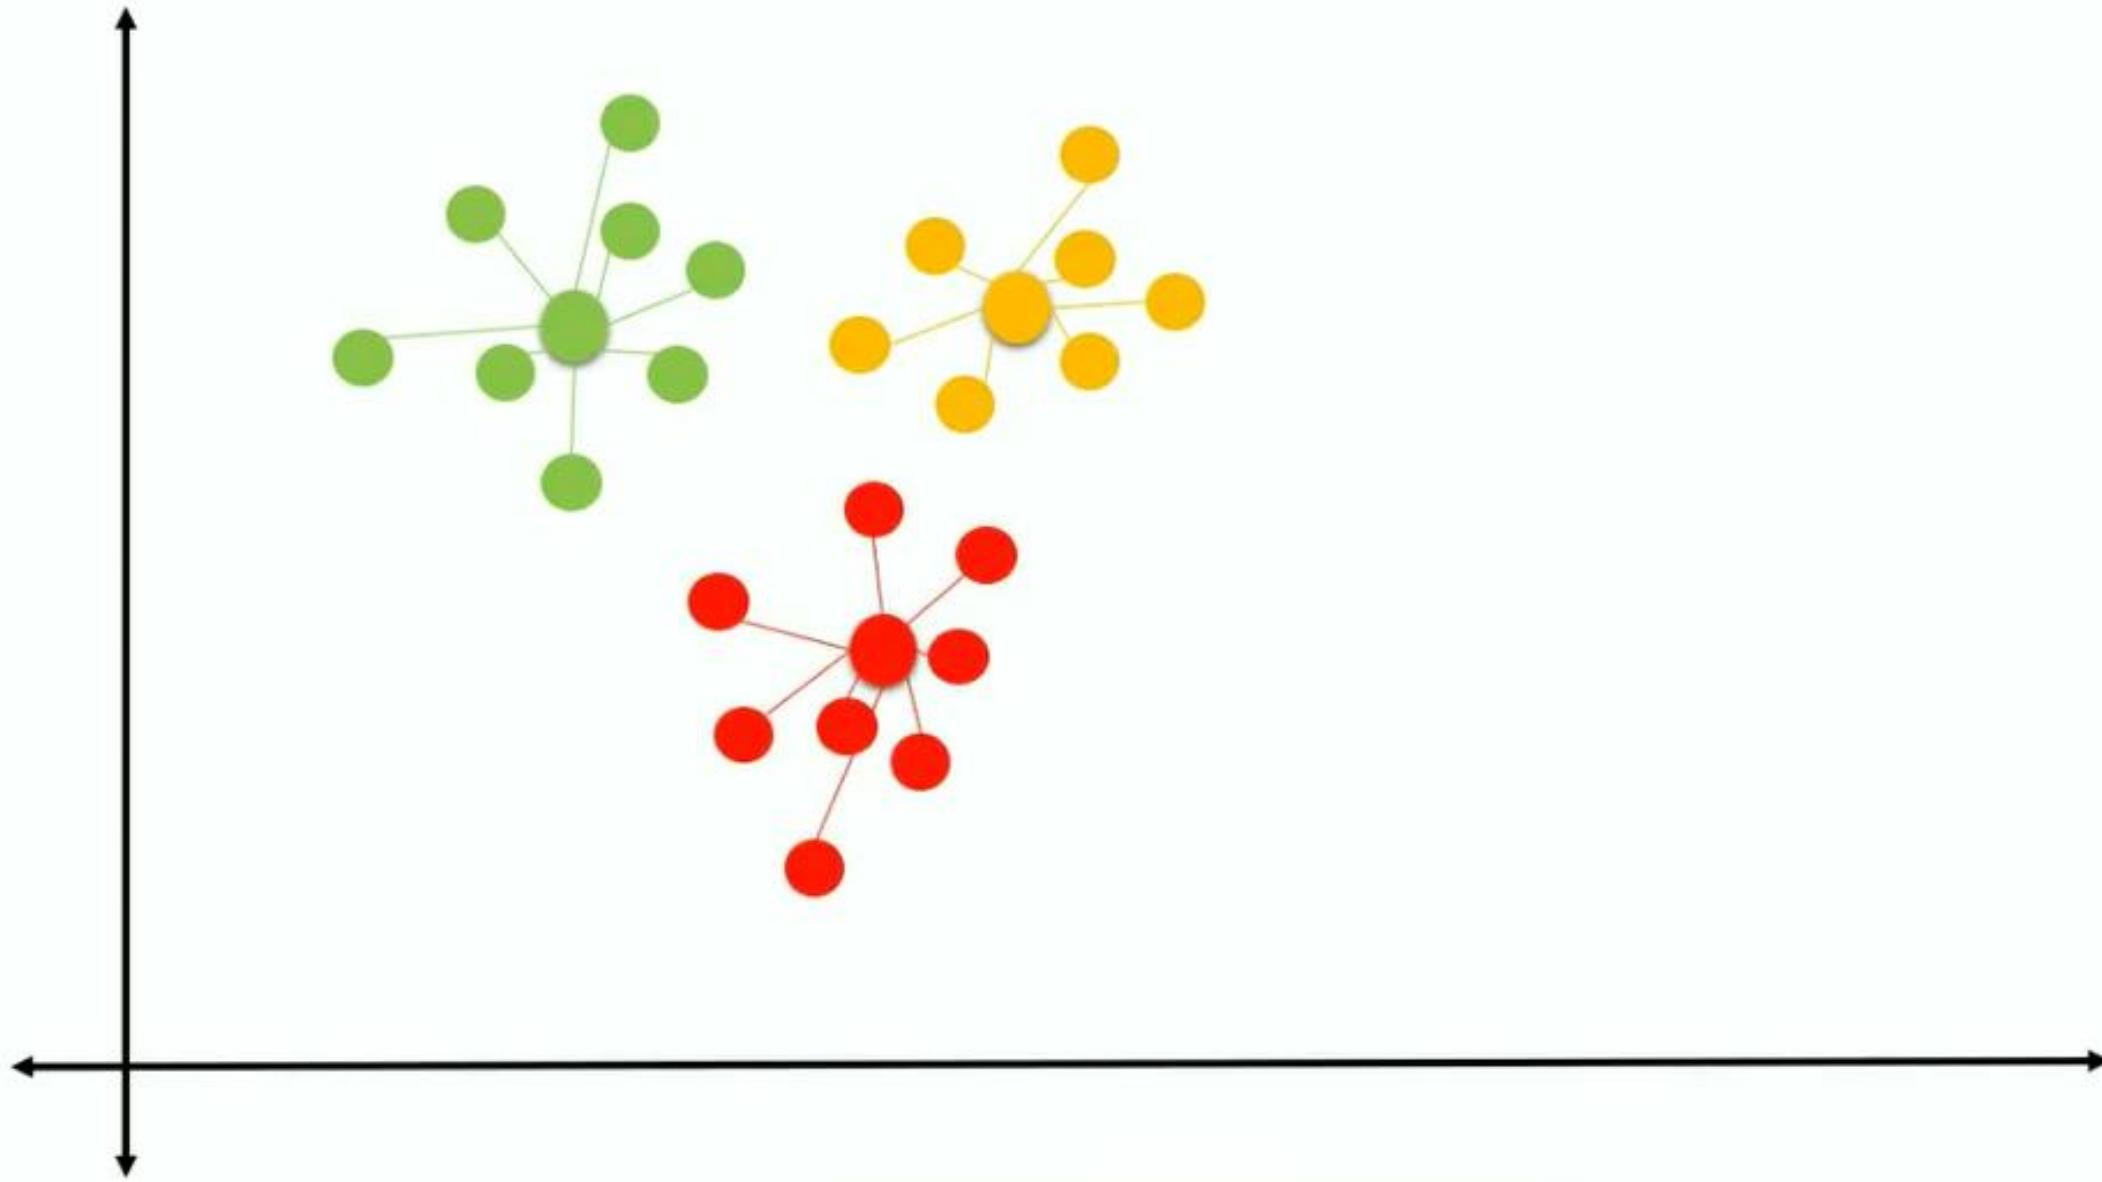


# Clustering or Cluster Analysis

- Clustering is the task of grouping a set of objects in such a way that objects in the same group (called a cluster) are more similar (in some sense or another) to each other than to those in other groups (clusters).
- Unsupervised Learning model
- Customers who make lot of long-distance calls and don't have a job. Who are they?

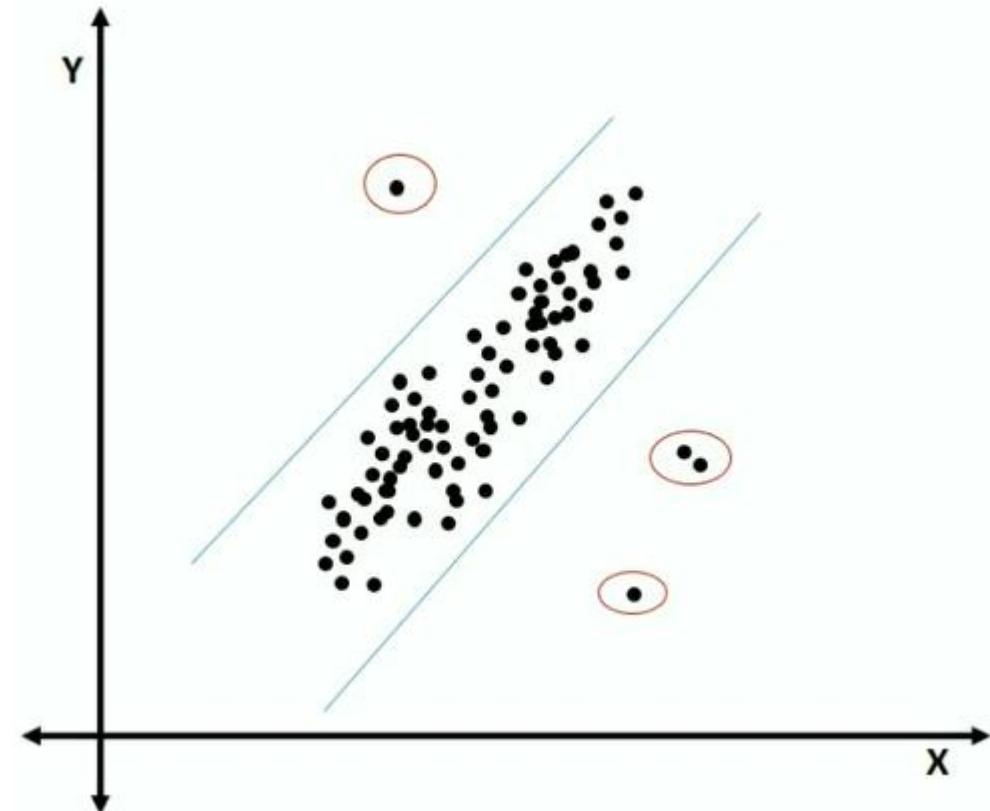


# Cluster Formation



# Anomaly Detection

- Anomaly detection (also outlier detection) is the identification of items, events or observations which do not conform to an expected pattern or other items in a dataset.
- Typically the anomalous items will translate to some kind of problem such as
  - Bank fraud
  - Credit Card Fraud
  - Structural defect
  - Medical problems
- Anomalies are also referred to as outliers, novelties, noise, deviations and exceptions.





# Getting Started with Azure ML

## Walk Through



# Covered– Basics of Machine Learning

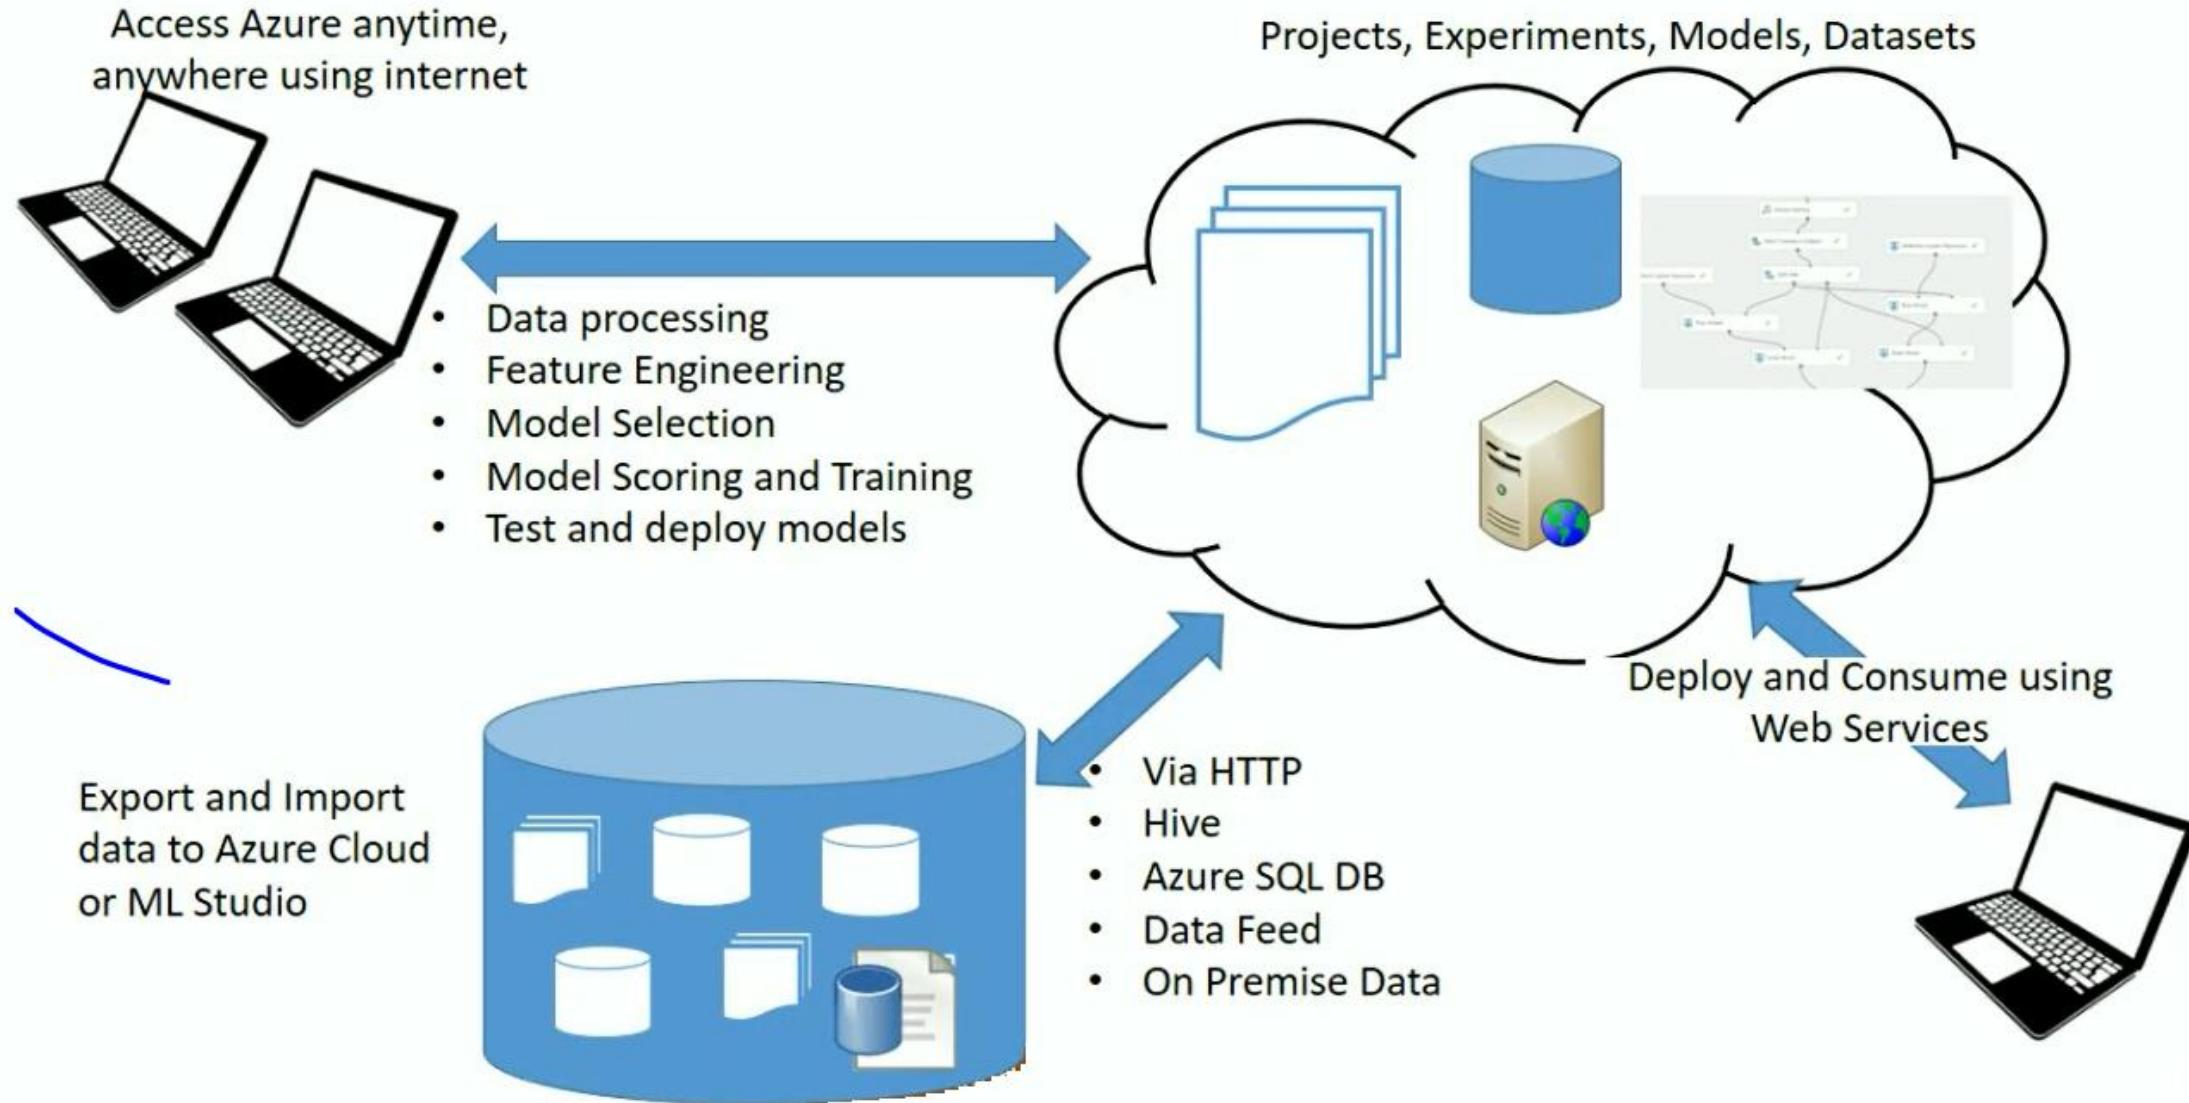
- Why Machine Learning is the Future?
- What is Machine Learning?
- Understanding the data
- Common Machine Learning Terms
- Types of Machine Learning Models

# What is Microsoft Azure ML?

# Azure Machine Learning

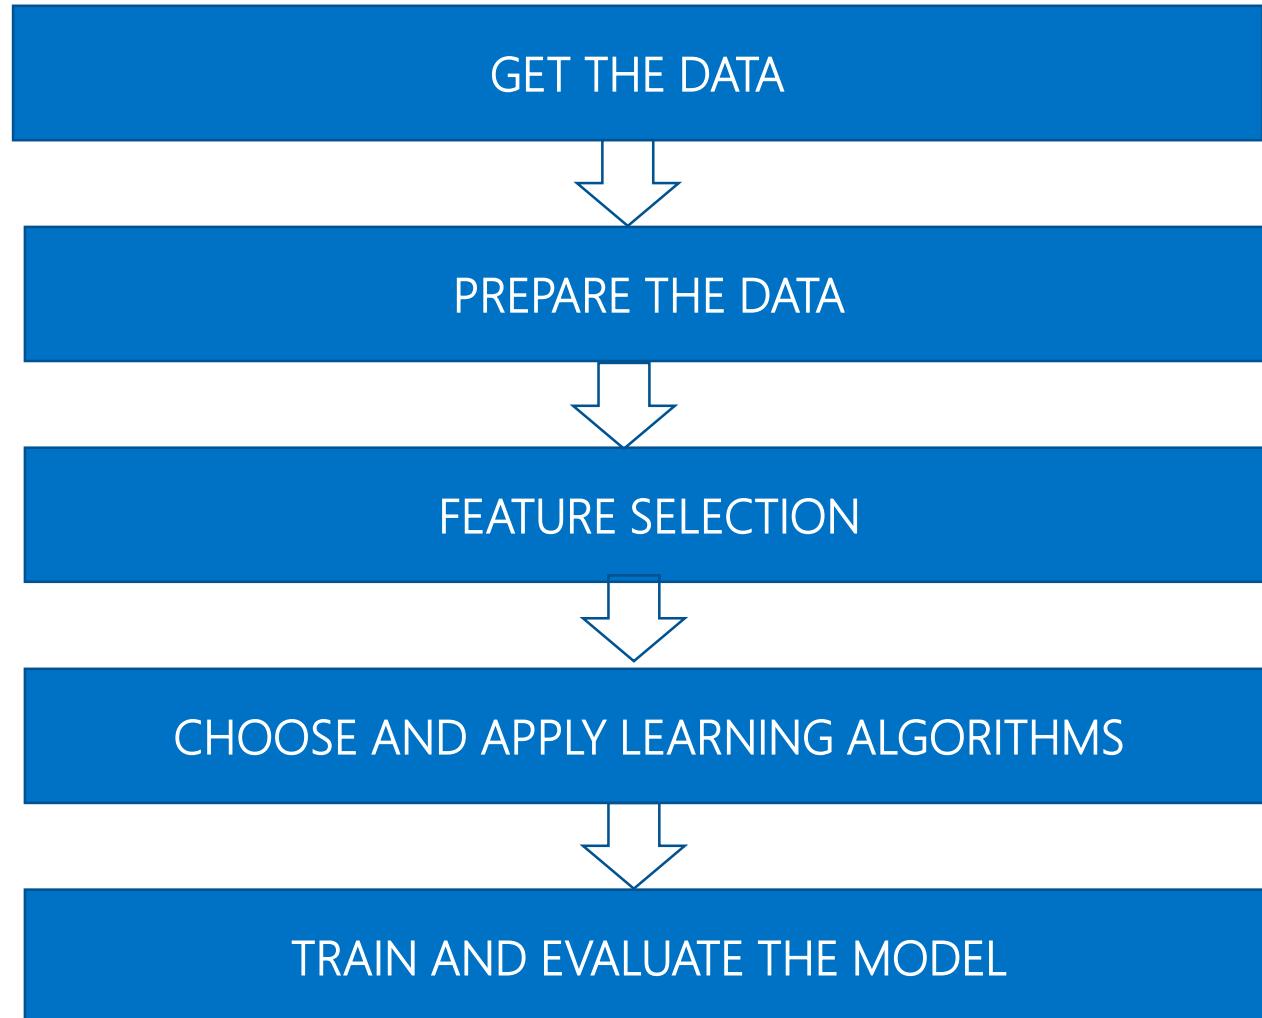
- Cloud based predictive analytics service
- Provides tools to create complete machine learning solution in the cloud
- Quick model creation and deployment using Azure ML Studio
- Allows Models to be deployed as web services
- Provides a large library of Pre-Built Machine learning algorithms and Modules
- Allows for extending your models with custom built R and Python code

# Azure ML Studio Overview



# Creating Azure ML Account

# Workflow of Azure ML



# Workflow of Azure ML

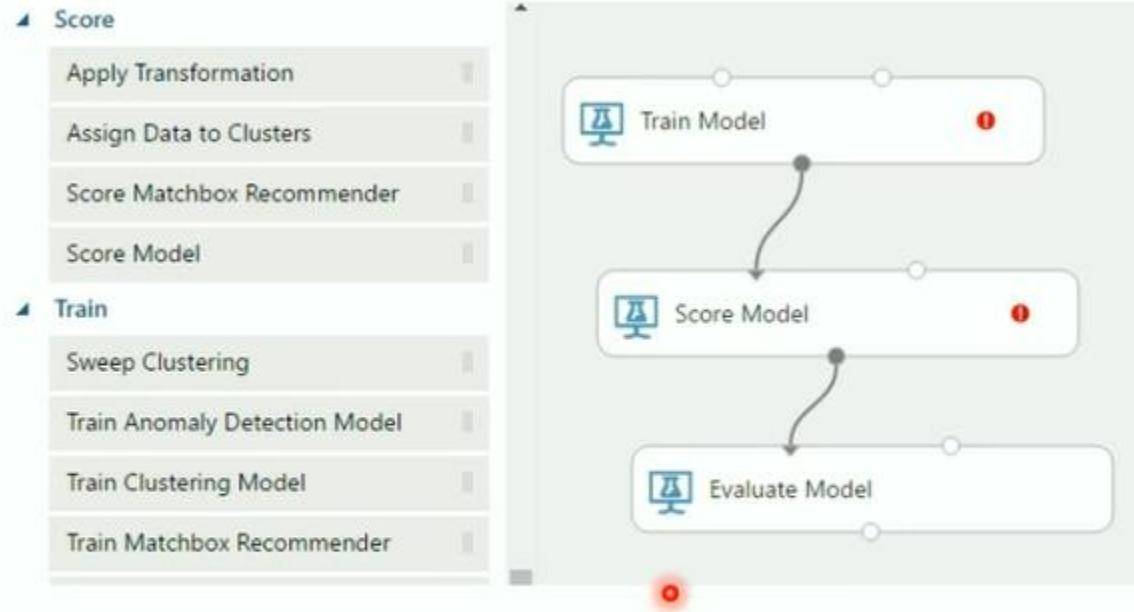
GET THE DATA

PREPARE THE DATA

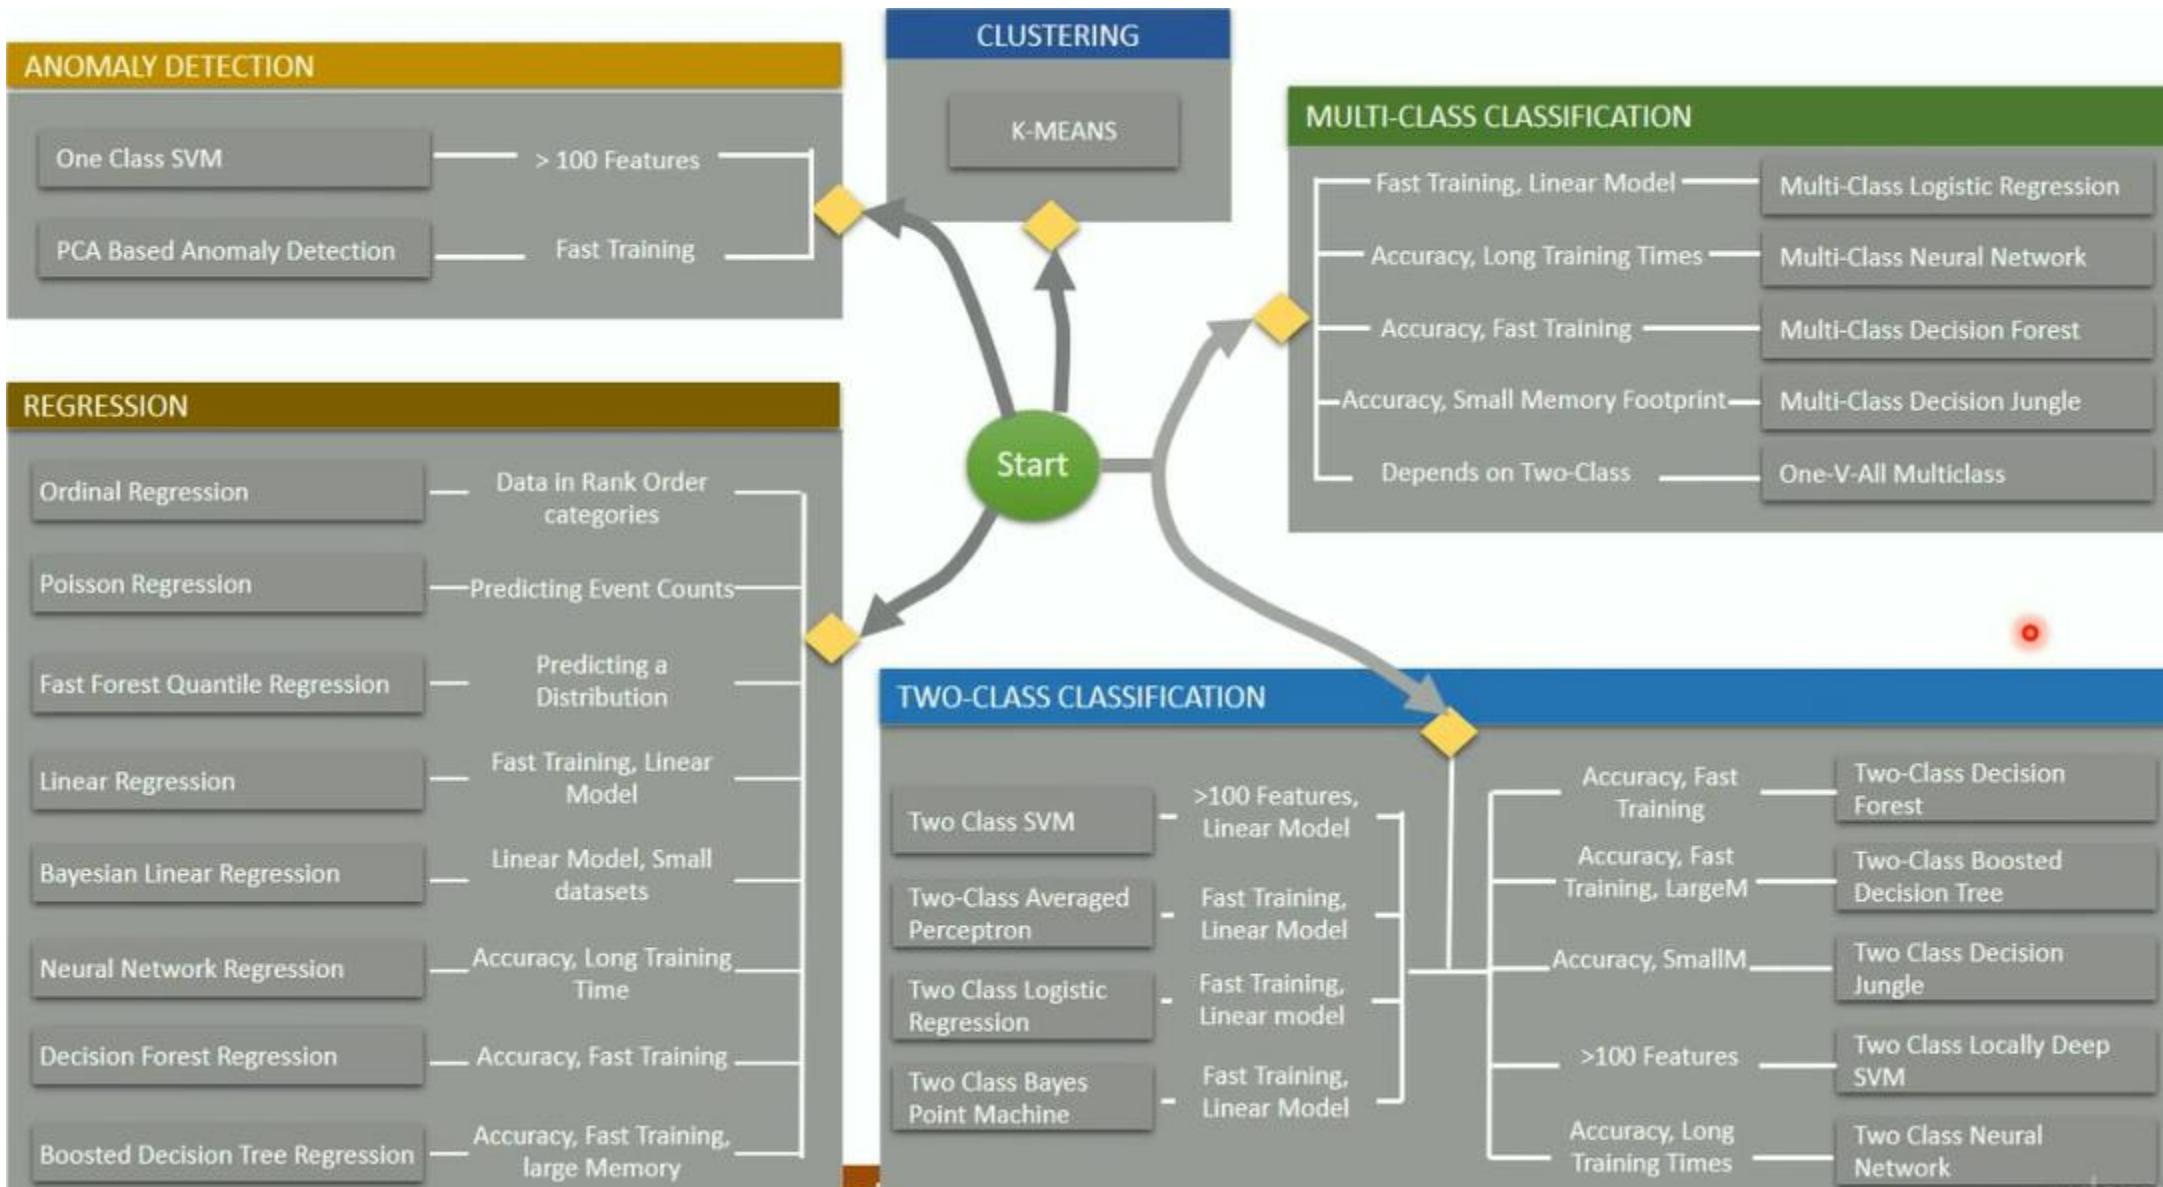
Feature Selection

CHOOSE AND APPLY LEARNING ALGORITHMS

TRAIN AND EVALUATE THE MODEL



# Azure ML Algorithm CheatSheet



## Predicting Categories

Outcome has multiple possibilities? Customer categories etc?

### MULTI-CLASS CLASSIFICATION

Fast Training, Linear Model	Multi-Class Logistic Regression
Accuracy, Long Training Times	Multi-Class Neural Network
Accuracy, Fast Training	Multi-Class Decision Forest
Accuracy, Small Memory Footprint	Multi-Class Decision Jungle
Depends on Two-Class	One-V-All Multiclass

Start

### TWO-CLASS CLASSIFICATION





Microsoft

# Data Processing with Azure ML



# Data Input Output to Azure Workspace

Enter Data Manually

Upload a Dataset

Data Format Conversion

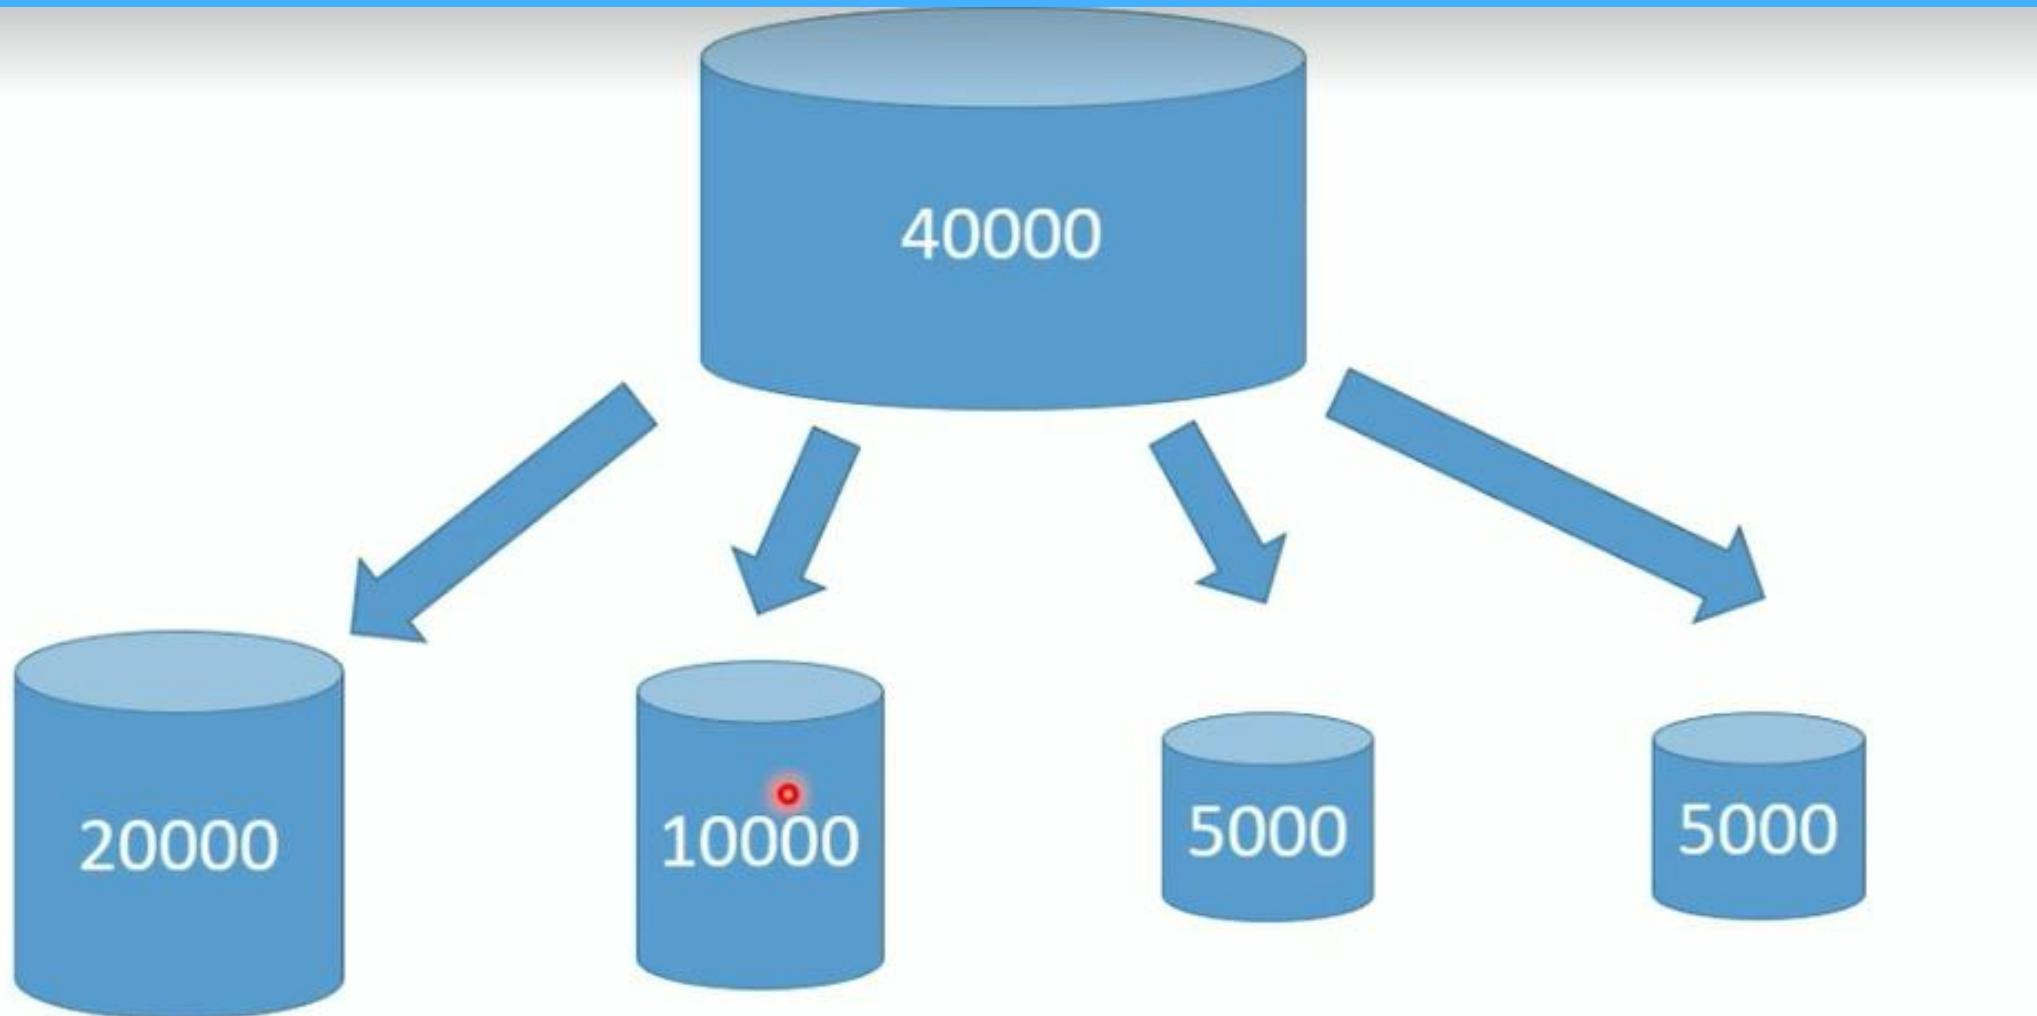
Import Data

Unpacked Zipped Dataset

Lets make the hands dirty

Demo -1. Data Preparation in Azure ML

# Sample and Split Data



Dividing your data into multiple subsections of different sizes.

# Why to Partition Data

- The input dataset is huge and have too many observations
- Partition the data to create train and test datasets
- Create different bins of data for cross validation of your results
- Create a random sample of observations or a more balanced dataset for a particular column

# Demo on Partitioning

# Without Stratification

Rating	Number of Observations
Rating 3	20
Rating 5	5
Total	25

Sampling rate 0.4 or 40%



No Stratification

Rating	Number of Observations
Rating 3	10
Rating 5	0
Total	10

Rating	Number of Observations
Rating 3	5
Rating 5	5
Total	10

Rating	Number of Observations
Rating 3	7
Rating 5	3
Total	10

# With Stratification

Rating	Number of Observations
Rating 3	$20 \times 0.4$ 
Rating 5	$5 \times 0.4$
Total	25

Sampling rate 0.4 or 40%



With Stratification

Rating	Number of Observations
Rating 3	8
Rating 5	2
Total	10



Microsoft

# Classification in Azure ML

## Logistic Regression



# What is Logistics Regression?

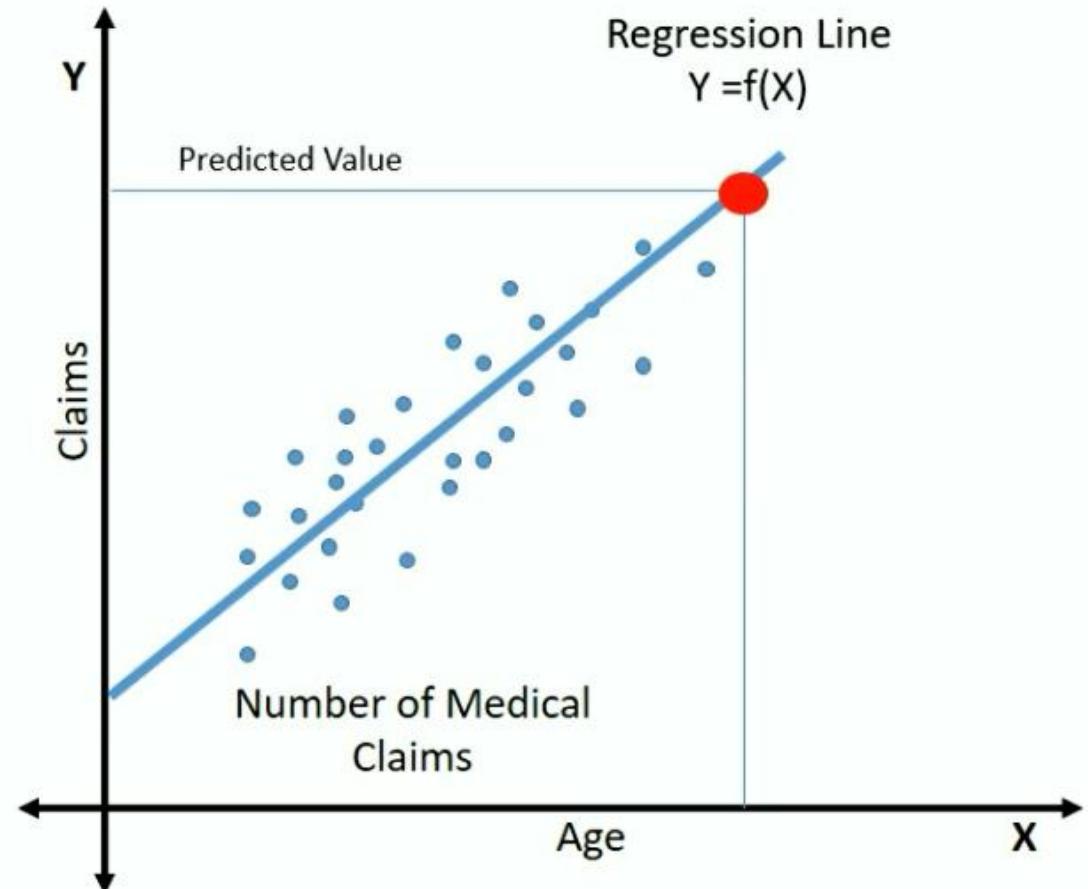
- Used to predict the probability of an outcome
- Can be binary - Yes/No or Multiple
- Supervised learning method
- Must provide a dataset that already contains the outcomes to train the model.

# Understanding the Logistic Regression

$$Y = f(x)$$

$$Y = b_0 + b_1 X$$

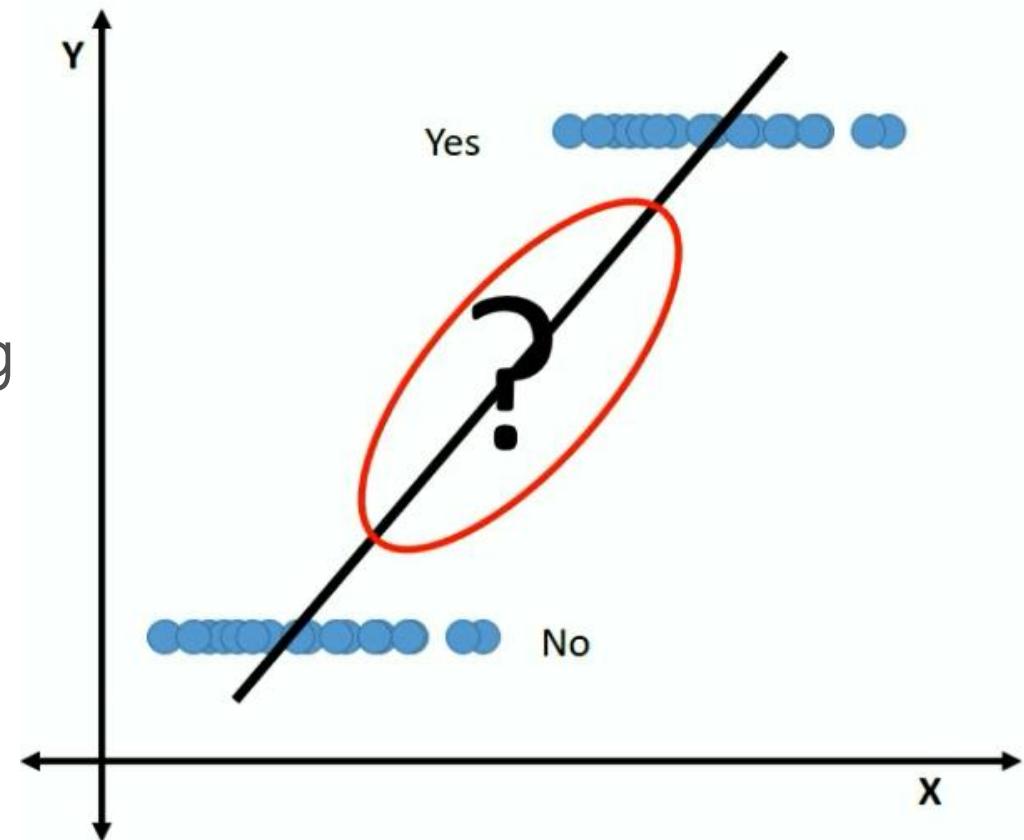
$$\text{No of claims} = 18 + b_1(\text{age})$$



Simple Linear Regression

# Logistic Regression?

- Outcome is categorical
- What is the probability of this customer buying this product?

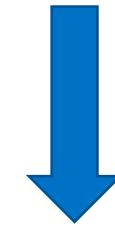


# Logistic Regression?

Probability needs to satisfy two basic conditions

- Always positive i.e.  $> 0$
- Always less than or equal to 1

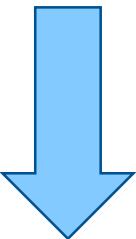
$$Y = b_0 + b_1 X \xrightarrow{\text{To make it positive}} \exp(Y) = \exp(b_0 + b_1 X)$$

 To make it less than 1 divide it with a number greater than P

# Logistic Regression?

$$\exp(Y) = \exp(b_0 + b_1 X)$$

$$\frac{\exp(Y)}{\exp(Y) + 1} = \frac{\exp(b_0 + b_1 X)}{\exp(b_0 + b_1 X) + 1}$$



$$e^y$$

$$P = \frac{e^y}{e^y + 1} \text{ probability of success}$$

# Logistic Regression?

$$P = \frac{e^y}{e^y + 1}$$

probability of success



$$Q = 1 - P = 1 - \frac{e^y}{e^y + 1} = \frac{e^y + 1 - e^y}{e^y + 1} = \frac{1}{e^y + 1}$$

Probability of Failure

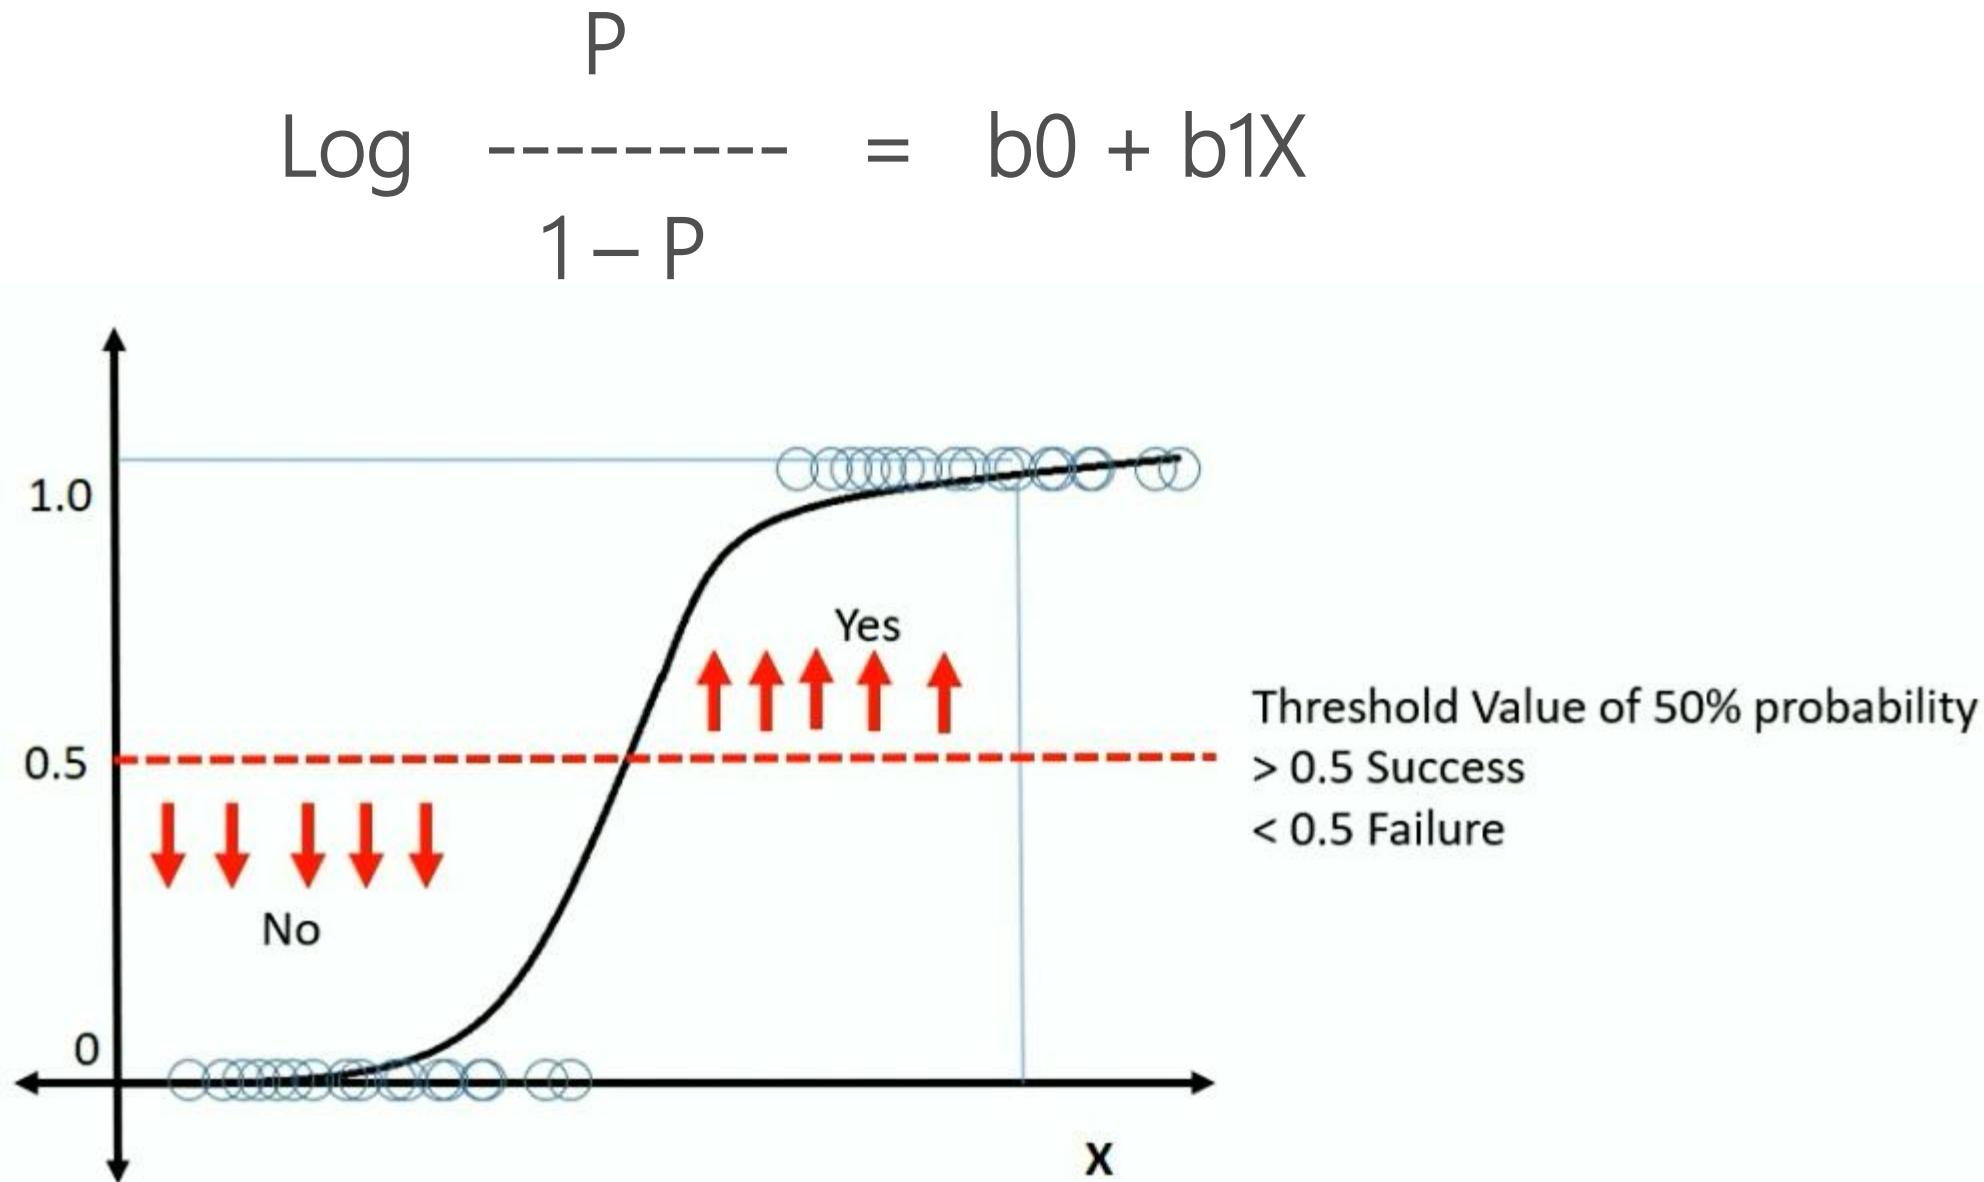
# Logistic Regression?

$$\frac{P}{1-P} = e^Y$$

$$\text{Log} \frac{P}{1-P} = Y$$

$$\text{Log} \frac{P}{1-P} = b_0 + b_1 X$$

# Plotting Logistic Regression?



Lets make the hands dirty

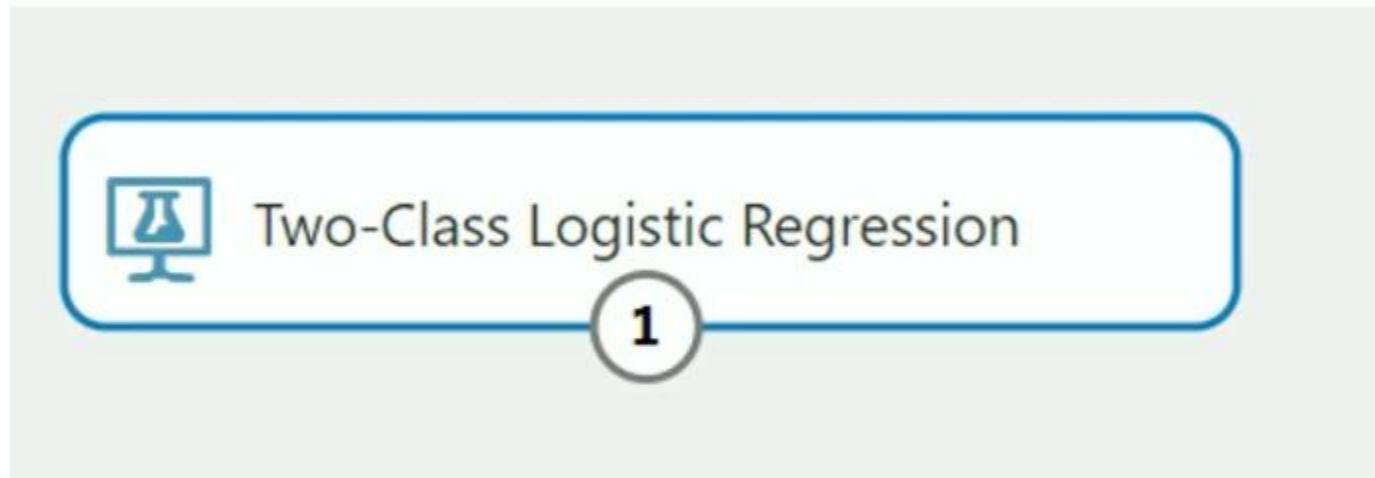
Demo -2. Loan Eligibility in Azure ML

# Problem Statement

- Automate loan eligibility process
- Identify customers whose loan will be approved

Loan_ID	Gender	Married	Dependents	Self_Employed	Income	LoanAmt	Term	CreditHistory	Property_Area	Status
LP001002	Male	No	0	No	\$5,849.00		60	1	Urban	Y
LP001003	Male	Yes	1	No	\$4,583.00	\$128.00	120	1	Rural	N
LP001005	Male	Yes	0	Yes	\$3,000.00	\$66.00	60	1	Urban	Y
LP001006	Male	Yes	2	No	\$2,583.00	\$120.00	60	1	Urban	Y

# Logistic Regression in Azure ML



Properties Project

## Two-Class Logistic Regression

Create trainer mode

Single Parameter

Optimization tolerance

1E-07

L1 regularization weight

1

L2 regularization weight

1

Memory size for L-BFGS

20

Random number seed

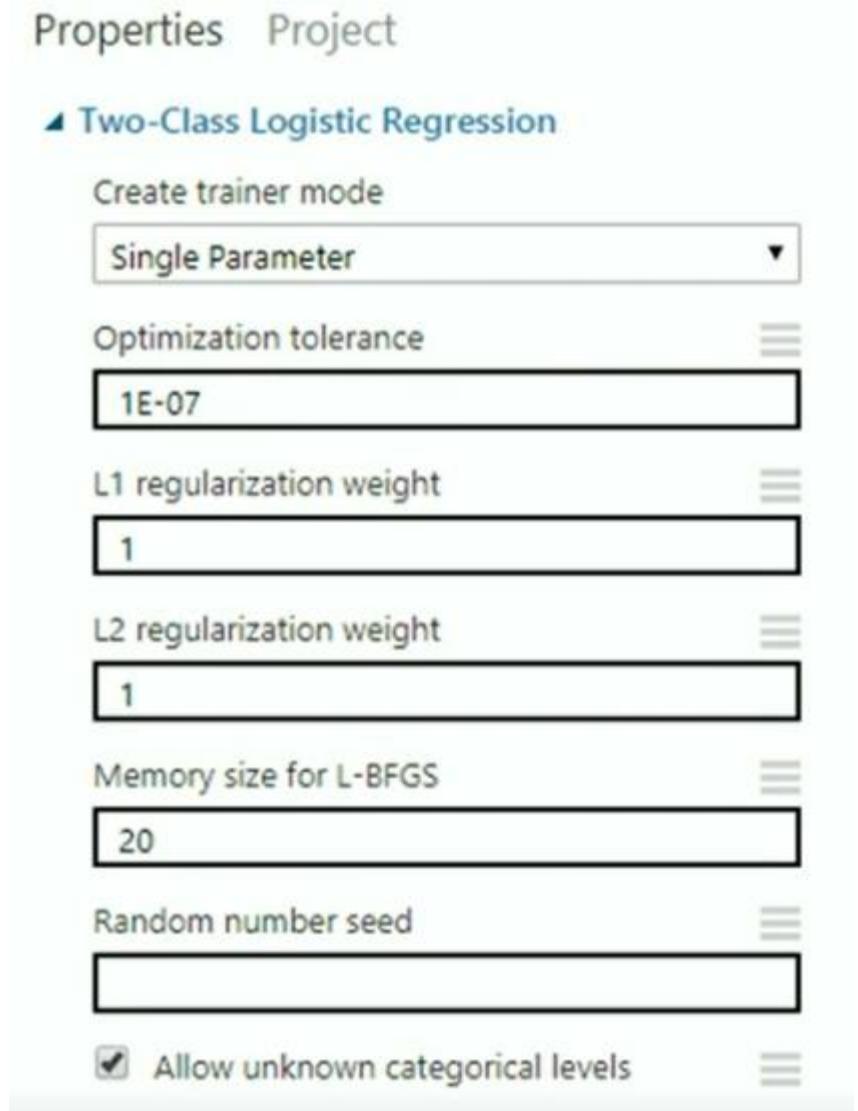
Allow unknown categorical levels

# What are Hyperparameters?

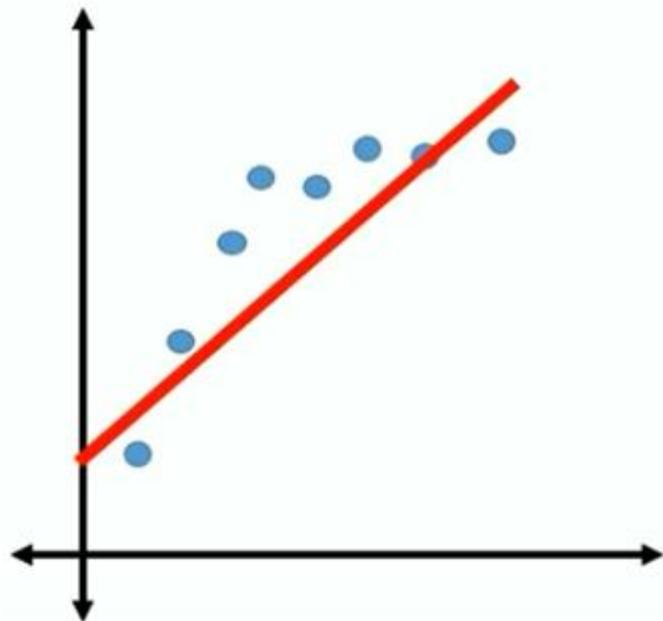


# Parameters to Logistic Regression

- Create Trainer Mode
  - Single Parameter - Provide specific set of values
  - Parameter Range - specify multiple values and get the optimum set for given configuration
- Optimization Tolerance - Threshold Value to stop the model iterations on trained dataset
- Memory Size for L-BFGS - Amount of memory to use for next steps and direction
- Random Number Seed - Random integer number that is used for reproducing the same results
- Allow Unknown Categorical Levels - Creates an additional "Unknown" level

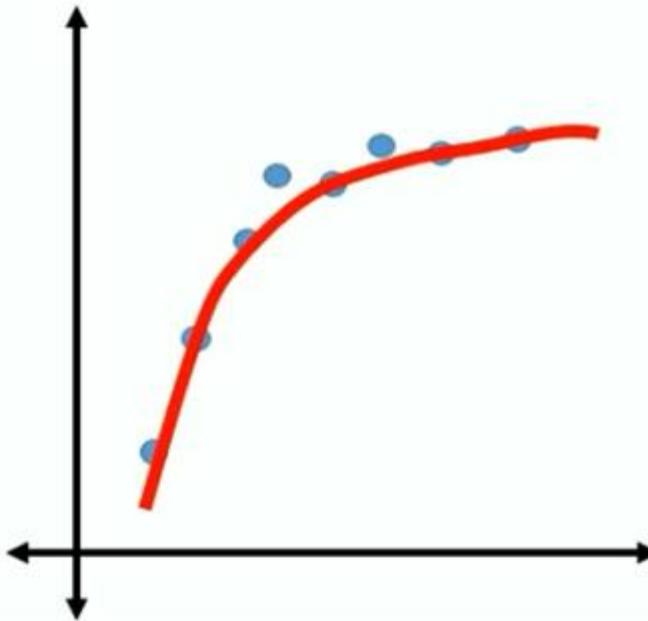


# Regularization Weight



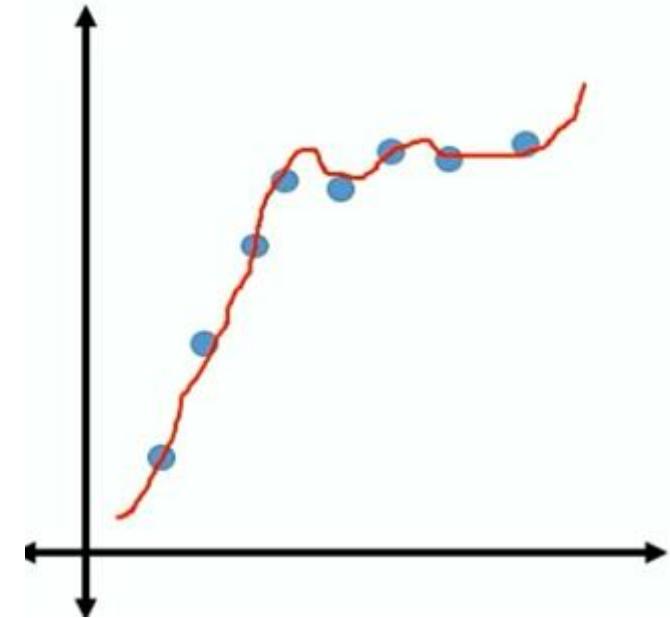
$b_0 + b_1 X$

Under-fit



$b_0 + b_1 X_1^2$

Right



$b_0 + b_1 X_1^2 + b_2 X_2^3$

Over Fit

What if the effect of such weights is reduced significantly  
Or reduced to zero

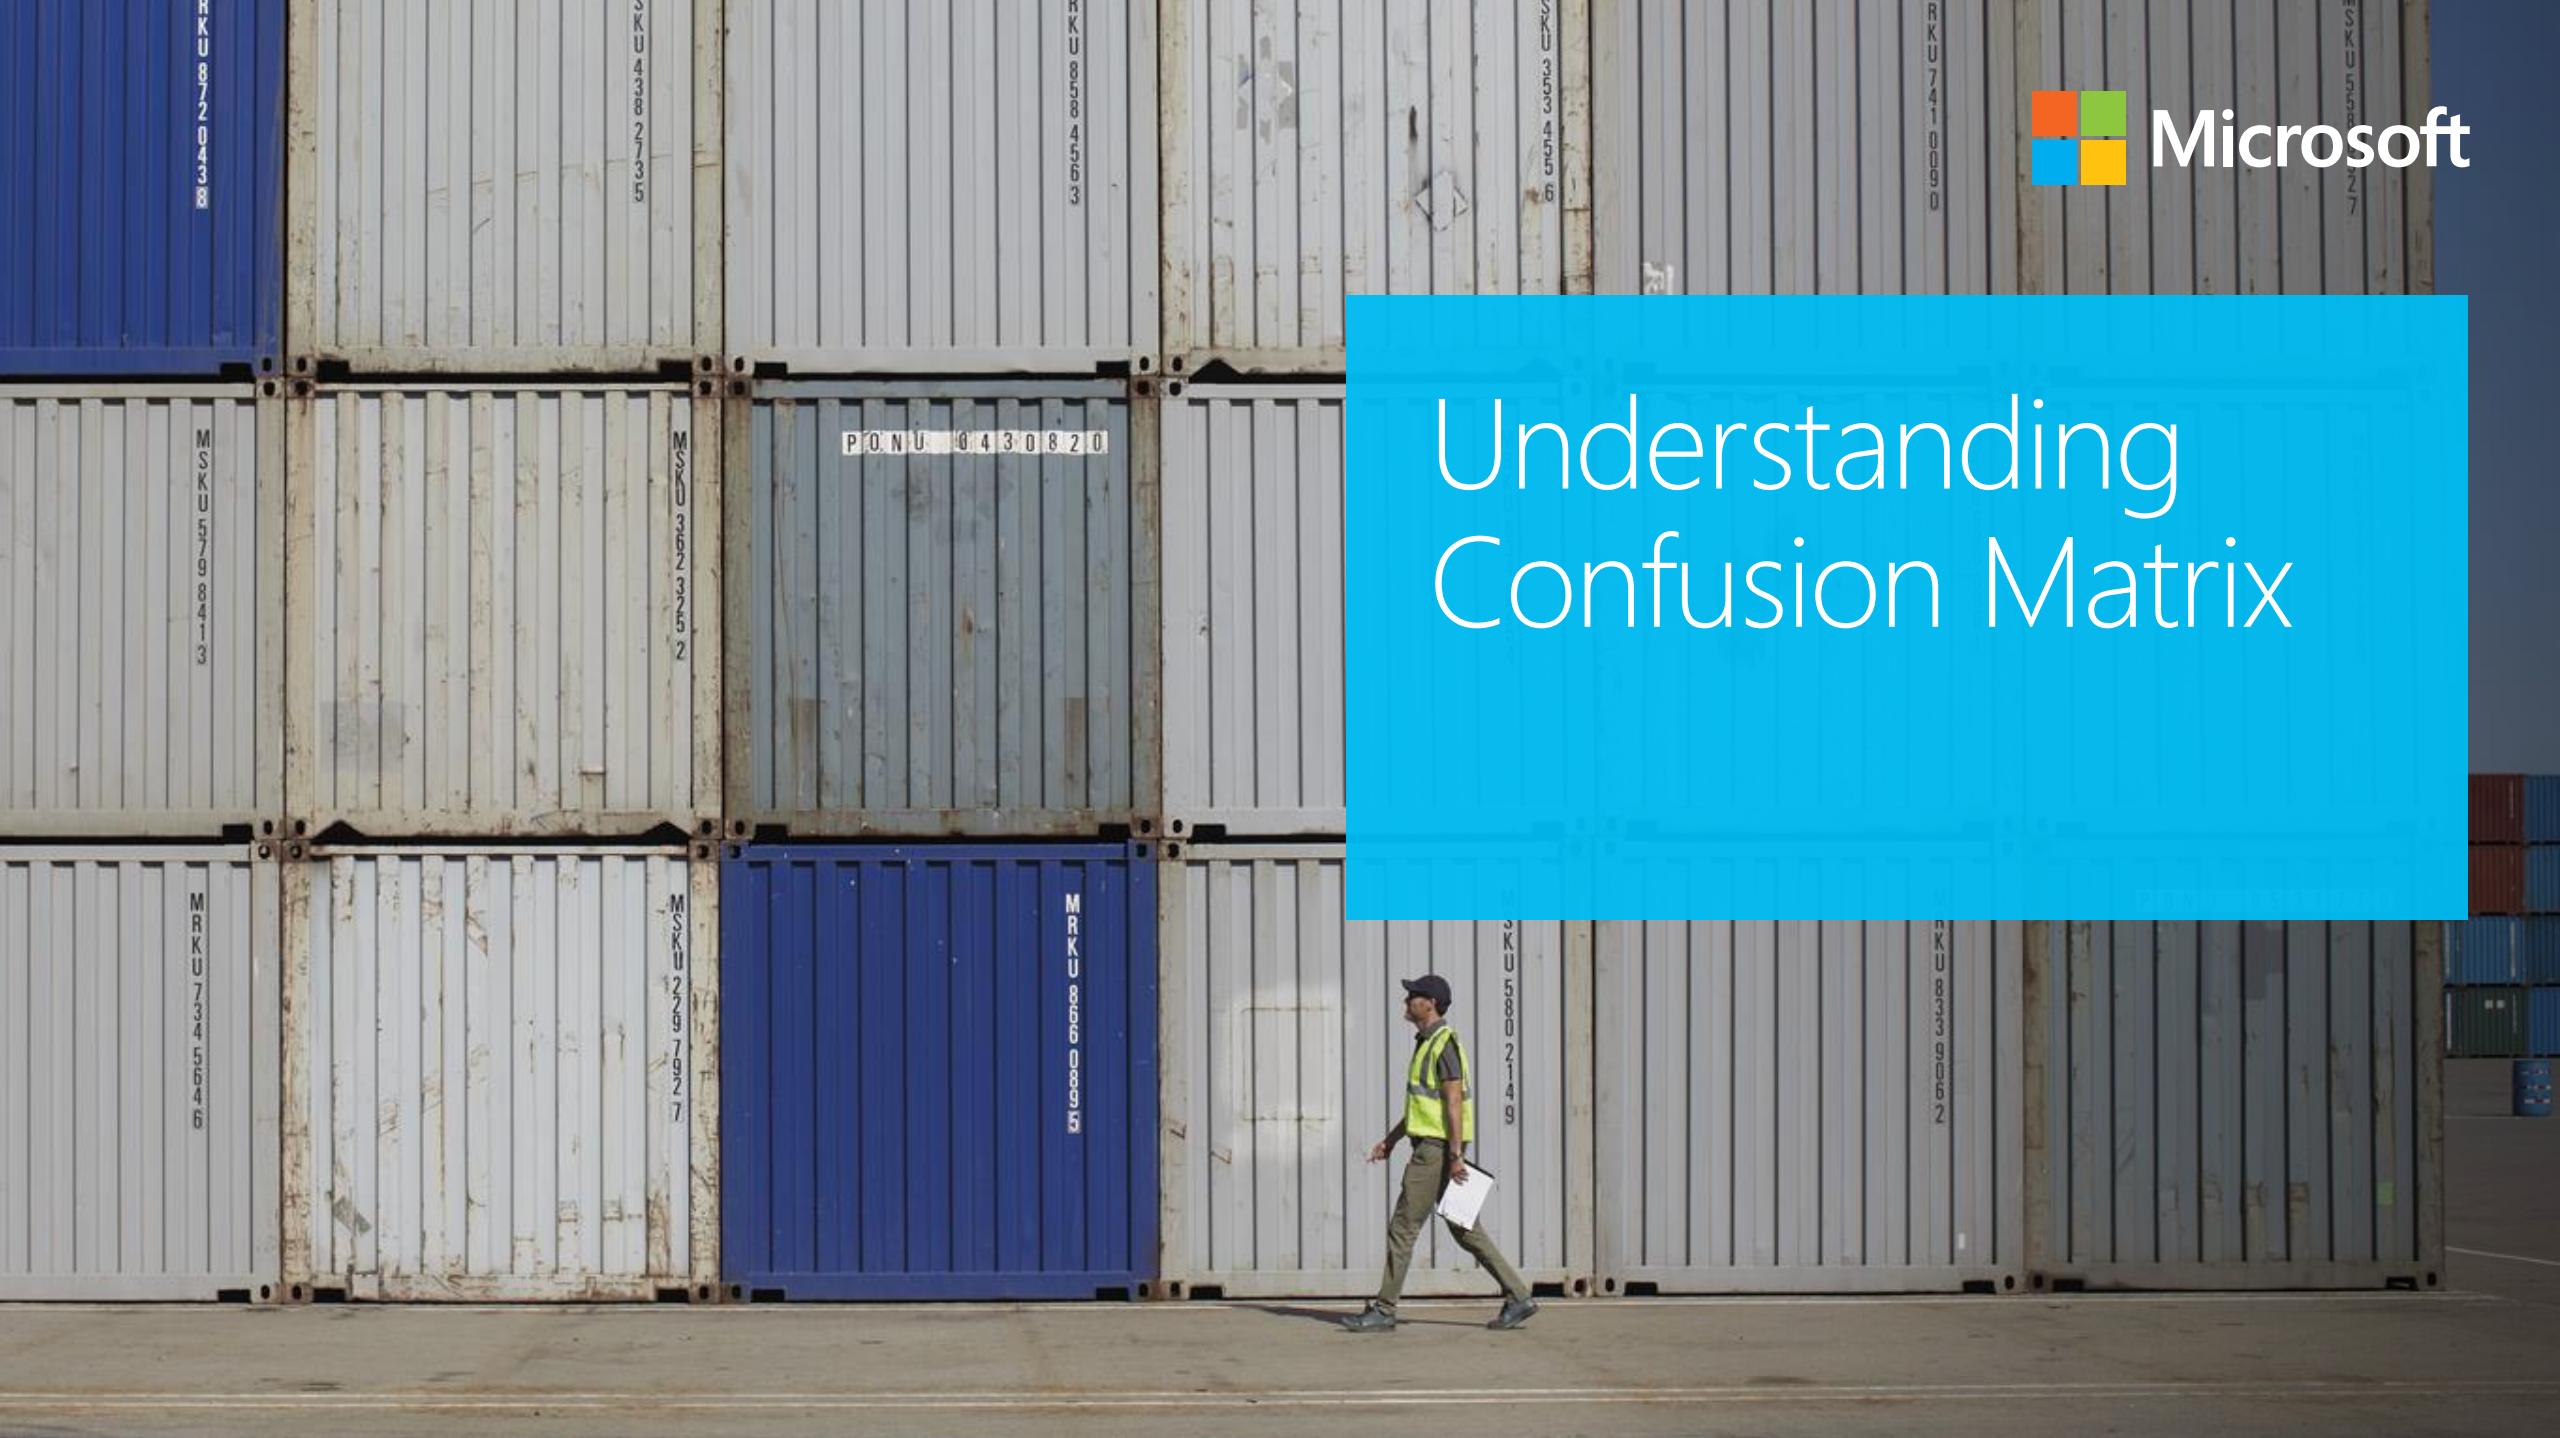
# Regularization Weight

- L2 (Ridge) shrinks all the coefficient by the same proportions but eliminates none
- L1(Lasso) can shrink some coefficients to zero, performing variable selection.
- Both L1 and L2 regularization prevents overfitting by shrinking (imposing a penalty) on the coefficients.
- L2 penalizes one big weight more than many small weights.
- With L2, you tend to end up with many small weights, while with L1, you tend to end up with larger weights, but more zeros.



Microsoft

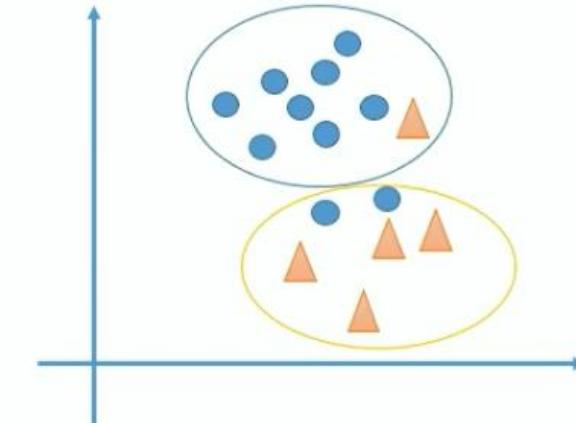
# Understanding Confusion Matrix



## Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	8	2	10
Actual Negative	1	4	5
	9	6	



Accuracy – Proportions of total number of correct results

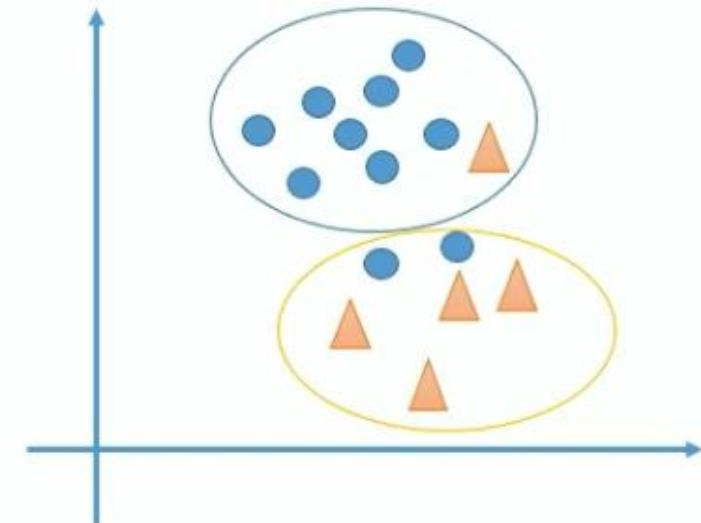
$$\text{Accuracy} = (8 + 4) / 15 = 0.8 \text{ or } 80\%$$

- **Accuracy:** Overall, how often is the classifier correct?

# Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	8	2	10
Actual Negative	1	4	5
	9	6	



Precision – Proportion of correct positive results out of all predicted positive results

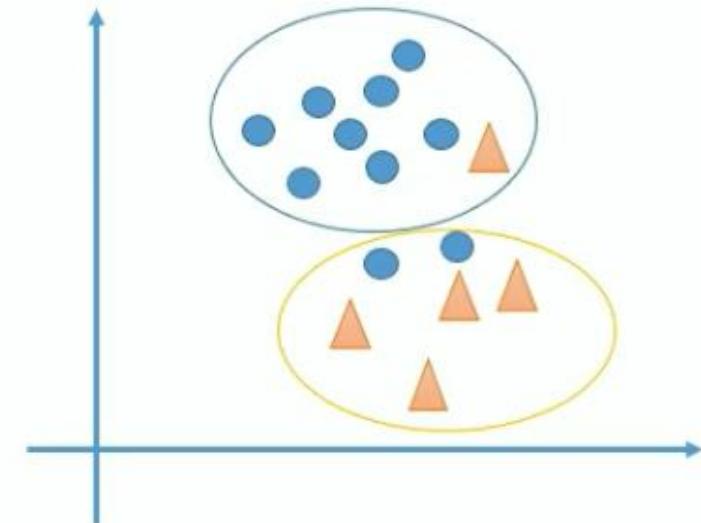
$$\text{Precision} = 8 / 9 = 0.889 \text{ or } 88.9\%$$

- **Precision:** When it predicts yes, how often is it correct?

# Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	8	2	10
Actual Negative	1	4	5
	9	6	



Recall – Proportion of actual positive cases

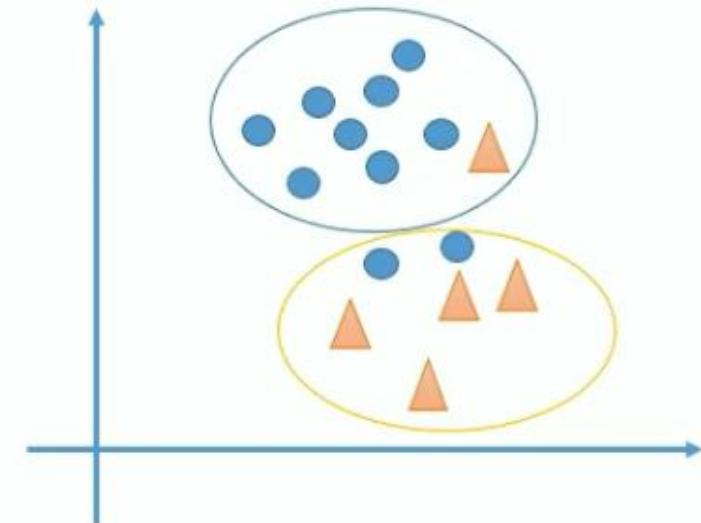
$$\text{Recall} = 8 / (8 + 2) = 0.8 \text{ or } 80\%$$

**True Positive Rate:** When it's actually yes, how often does it predict yes?  
also known as "Sensitivity" or "Recall"

# Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	8	2	10
Actual Negative	1	4	5
	9	6	



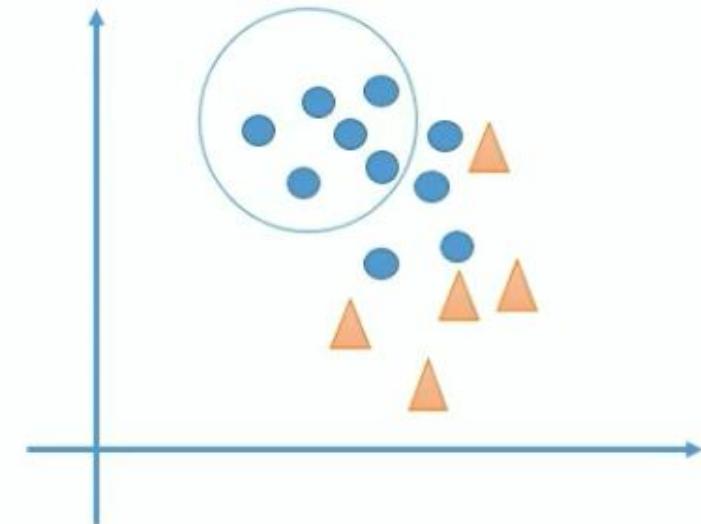
F1-Score – Weighted Average (Harmonic Mean) of Precision and Recall

$$\text{F1Score} = 2 * \text{Precision} * \text{Recall} / (\text{Precision} + \text{Recall}) = 0.84$$

# Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	6	4	10
Actual Negative	0	5	5
	6	6	



In the Previous example

$$\text{Precision} = 6 / 6 = 1 \text{ or } 100\%$$

$$\text{Recall} = 6 / (6 + 4) = 0.6 \text{ or } 60\%$$

$$\text{Average} = 0.8$$

May lead to false interpretation

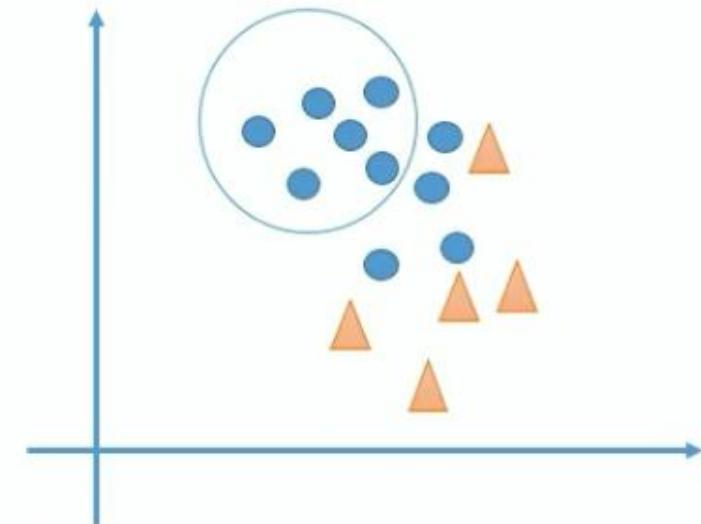
$$\text{Precision} = 0.889$$
$$\text{Recall} = 0.8$$

$$\text{Average} = 0.84$$

# Prediction Outcome

	Predicted Positives	Predicted Negatives
Actual Positives	True Positives	False Negatives
Actual Negative	False Positives	True Negatives

	Predicted Positives	Predicted Negatives	
Actual Positives	6	4	10
Actual Negative	0	5	5
	6	6	



In the first example

$$\text{Precision} = 6 / 6 = 1 \text{ or } 100\%$$

$$\text{Recall} = 6 / (8 + 2) = 0.6 \text{ or } 60\%$$

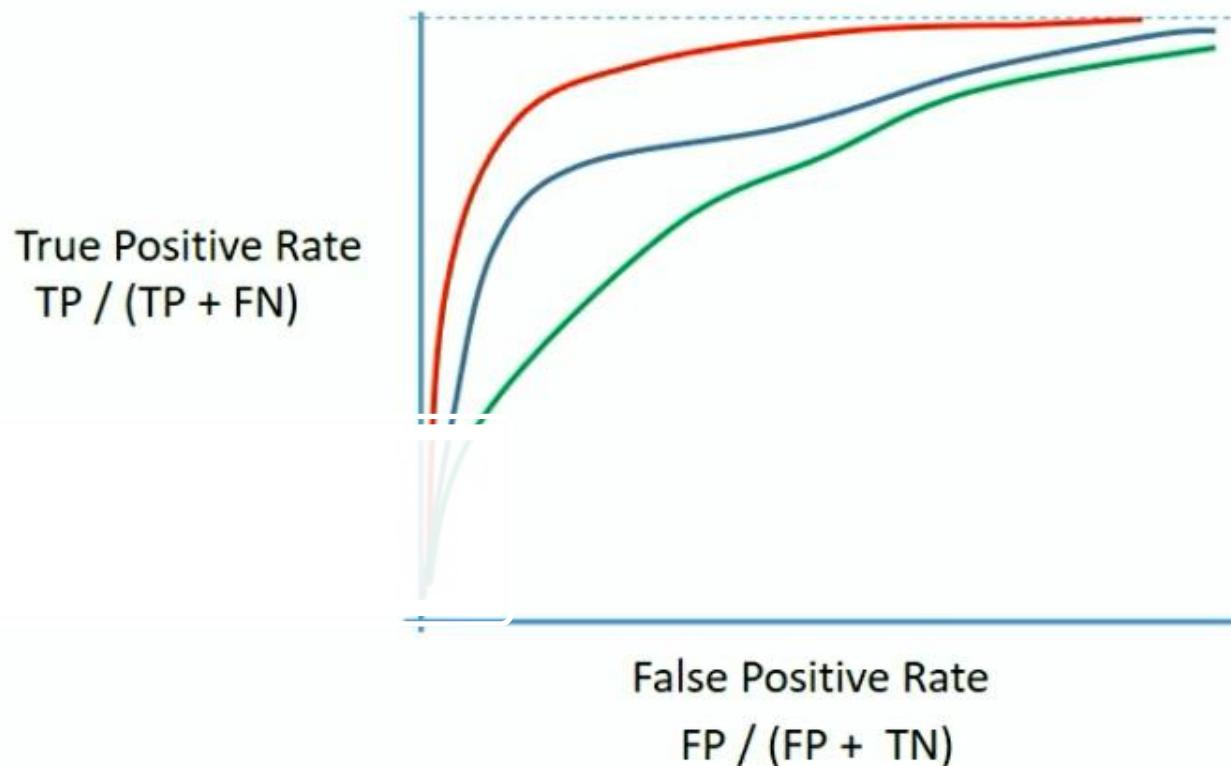
$$\text{F1Score} = 0.75$$

$$\text{Precision} = 0.889$$

$$\text{Recall} = 0.8$$

$$\text{F1Score} = 0.84$$

# AUC ROC



AUC – Area Under the Curve

ROC – Receiver Operating Characteristics

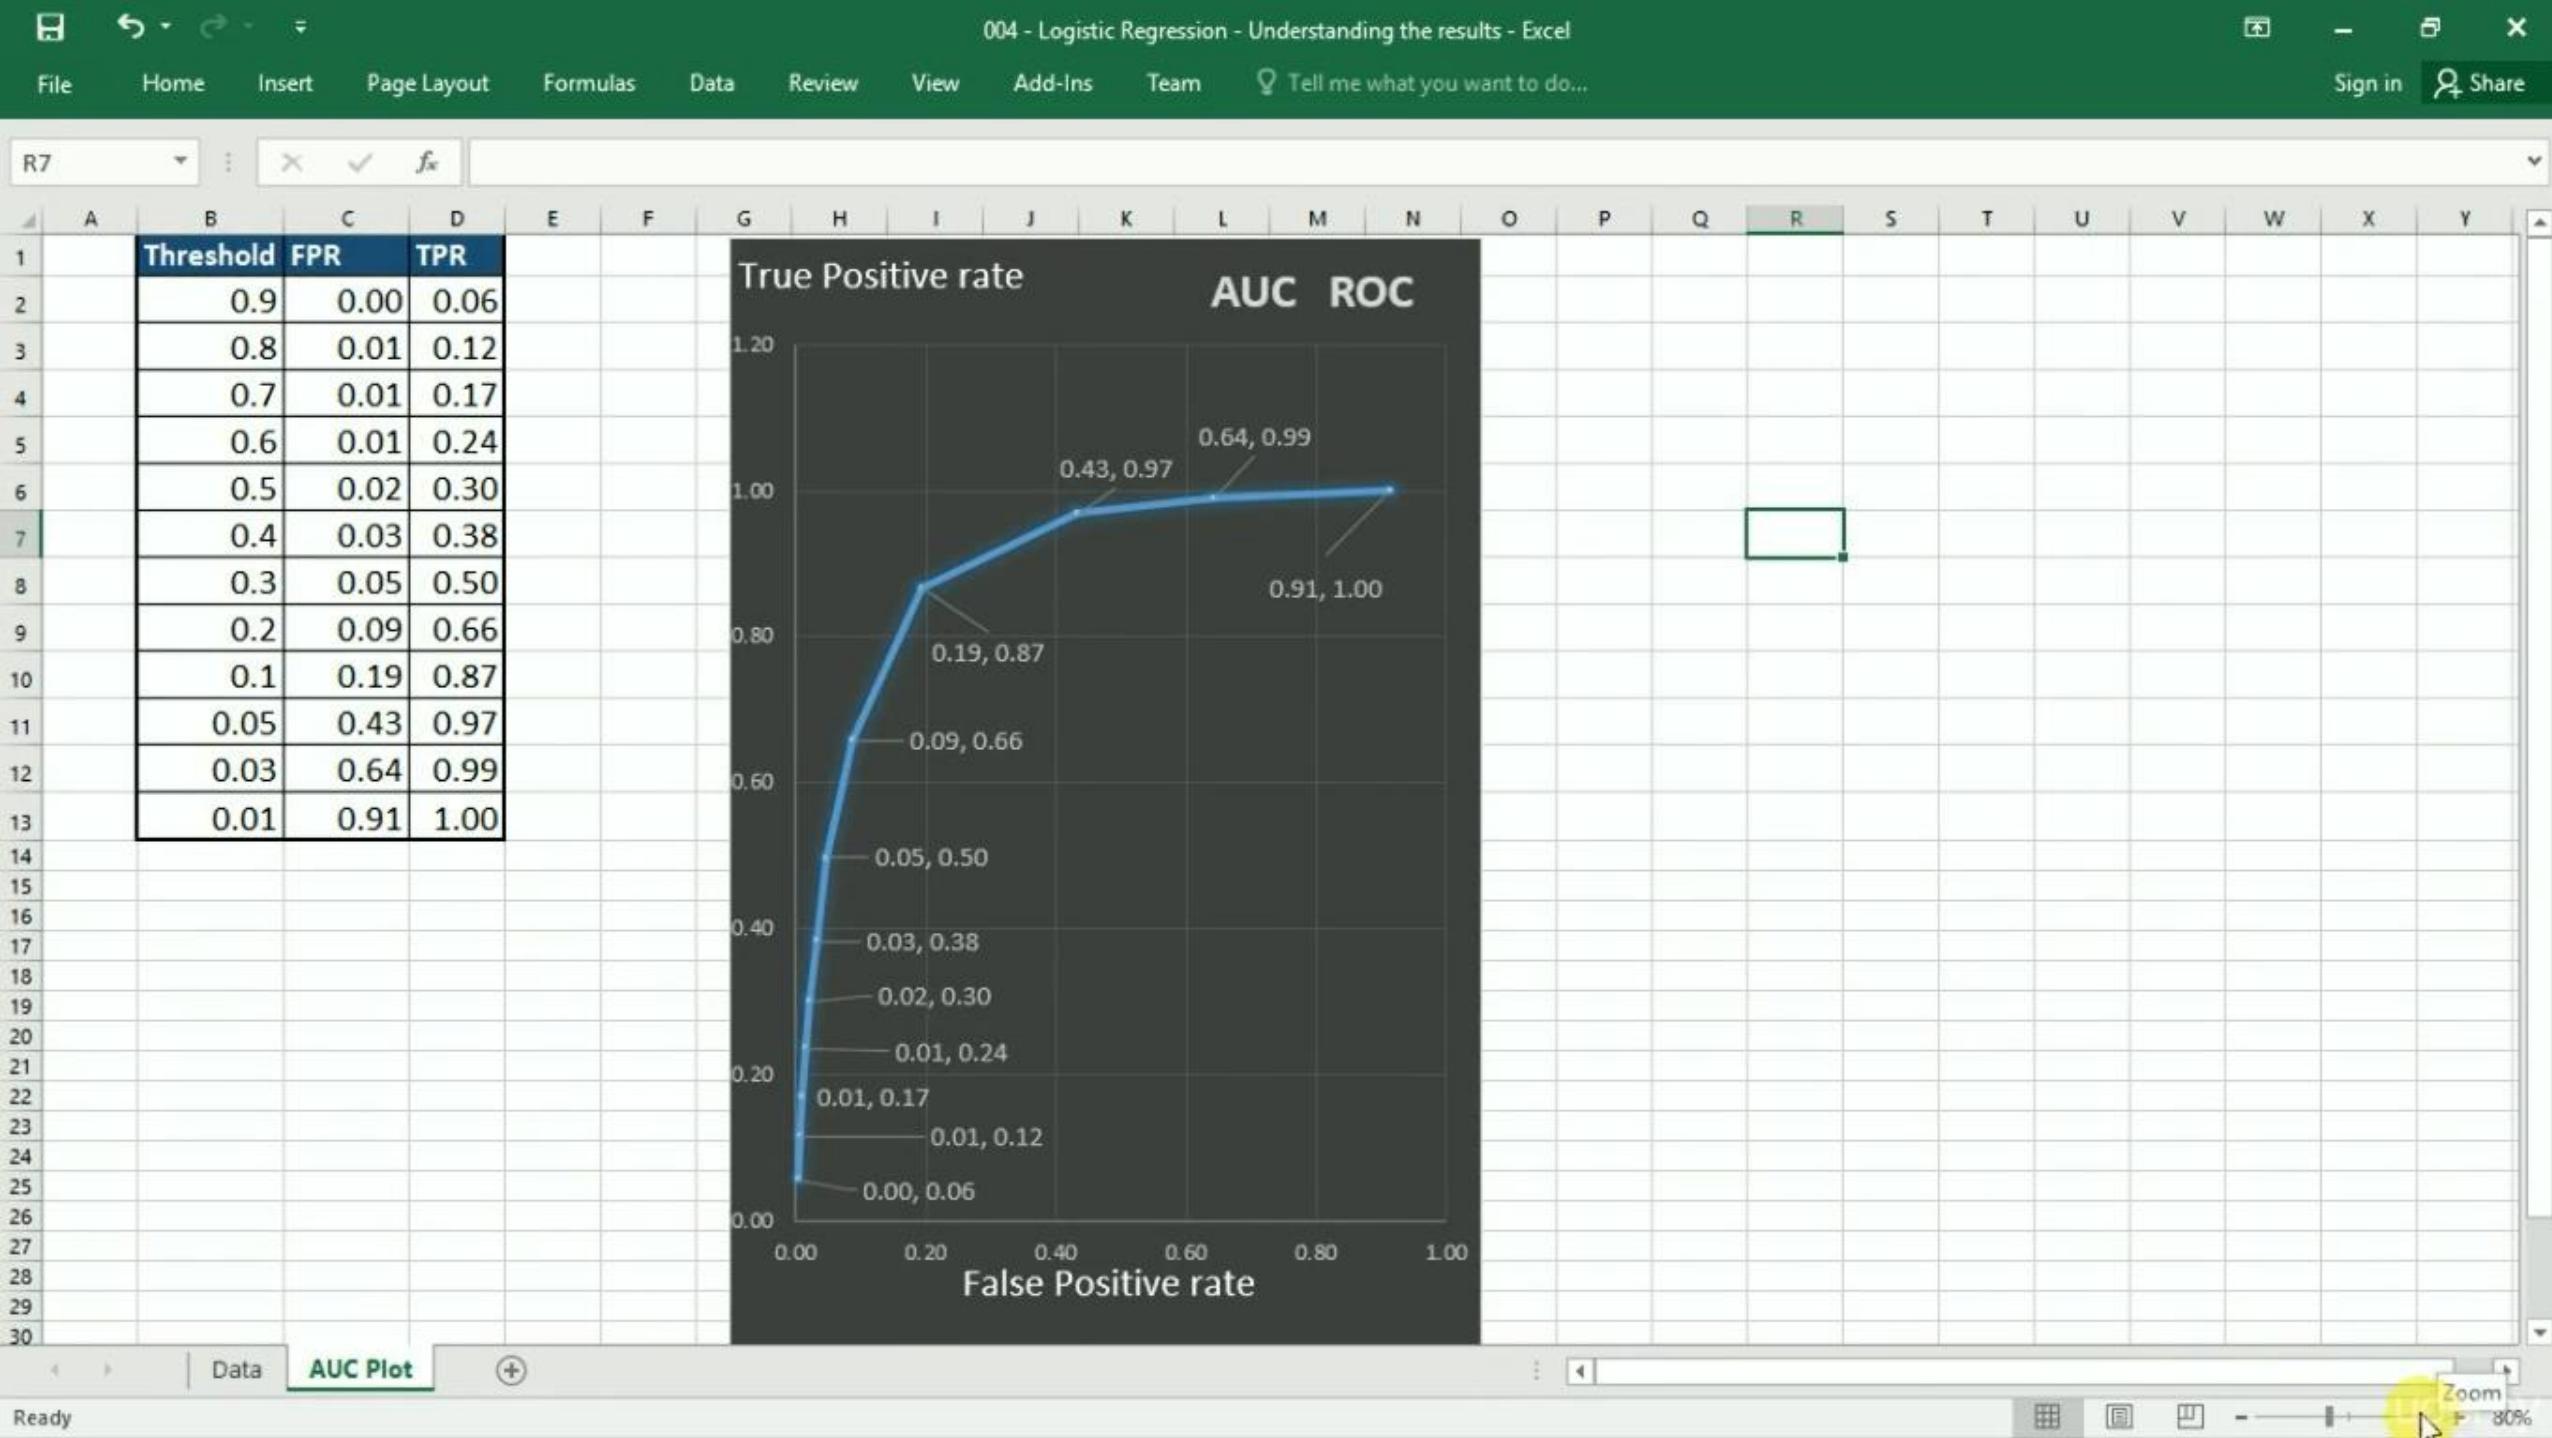
First used during World War II for the analysis of radar signals

Following the attack on Pearl Harbor in 1941, the United States army began new research to increase the prediction of correctly detected Japanese aircraft from their radar signals.

For this purposes they measured the ability of radar receiver operators to make these important distinctions, which was called the Receiver Operating Characteristics

Provides a single number that lets you compare models of different types.

004 - Logistic Regression - Understanding the results - Excel									
File	Home	Insert	Page Layout	Formulas	Data	Review	View	Add-Ins	Team
H17	X	Y	f(x)	0.5					Tell me what you want to do...
3	B	C	D	E	F	G	H	I	J
4	FPR = FP/(FP + TN)								
5	Threshold	0.05				Threshold	0.03		
6	TP	FN				TP	FN		
7	1019	33	TPR	0.96863		1042	10	TPR	0.99049
8	FP	TN	FPR	0.43342		FP	TN	FPR	0.64431
9	3463	4527				5148	2842		
11	Threshold	0.1				Threshold	0.2		
12	TP	FN				TP	FN		
13	913	139	TPR	0.86787		692	360	TPR	0.65779
14	FP	TN	FPR	0.19312		FP	TN	FPR	0.08811
15	1543	6447				704	7286		
17	Threshold	0.4				Threshold	0.5		
18	TP	FN				TP	FN		
19	404	648	TPR	0.38403		317	735	TPR	0.30133
20	FP	TN	FPR	0.03179		FP	TN	FPR	0.0219
21	254	7736				175	7815		
23	Threshold	0.7				Threshold	0.8		
24	TP	FN				TP	FN		
25	179	873	TPR	0.17015		122	930	TPR	0.11597
26	FP	TN	FPR	0.00914		FP	TN	FPR	0.00526
27	73	7917				42	7948		
28	Data	AUC Plot	+			Threshold	0.9		
						TP	FN		
						60	992	TPR	0.05703
						FP	TN	FPR	0.00275
						22	7968		



# Impact Analysis

# Multiple Model Analysis

- Scenario 1 - Impact of Split Percentage
- Scenario 2 - Impact of Stratification
- Scenario 3 - Low L1 (0.0001) and High L2 (1)
- Scenario 4 - High L1 (1) and Low L2 (0.0001)
- Scenario 5 - High L1 (1) and High L2 (1)
- Scenario 6 - Low L1 (0.0001) and Low L2 (0.0001)

# Multiclass Logistic Regression

# Wine Quality Prediction



- Fixed and Volatile Acidity
- Citric acid
- Residual sugar
- Chlorides
- Free and Total Sulphur dioxide
- Density
- pH
- Sulphates
- Alcohol Content

P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis.

Modelling wine preferences by data mining from physicochemical properties. In Decision Support Systems, Elsevier, 47(4):547-553,2009

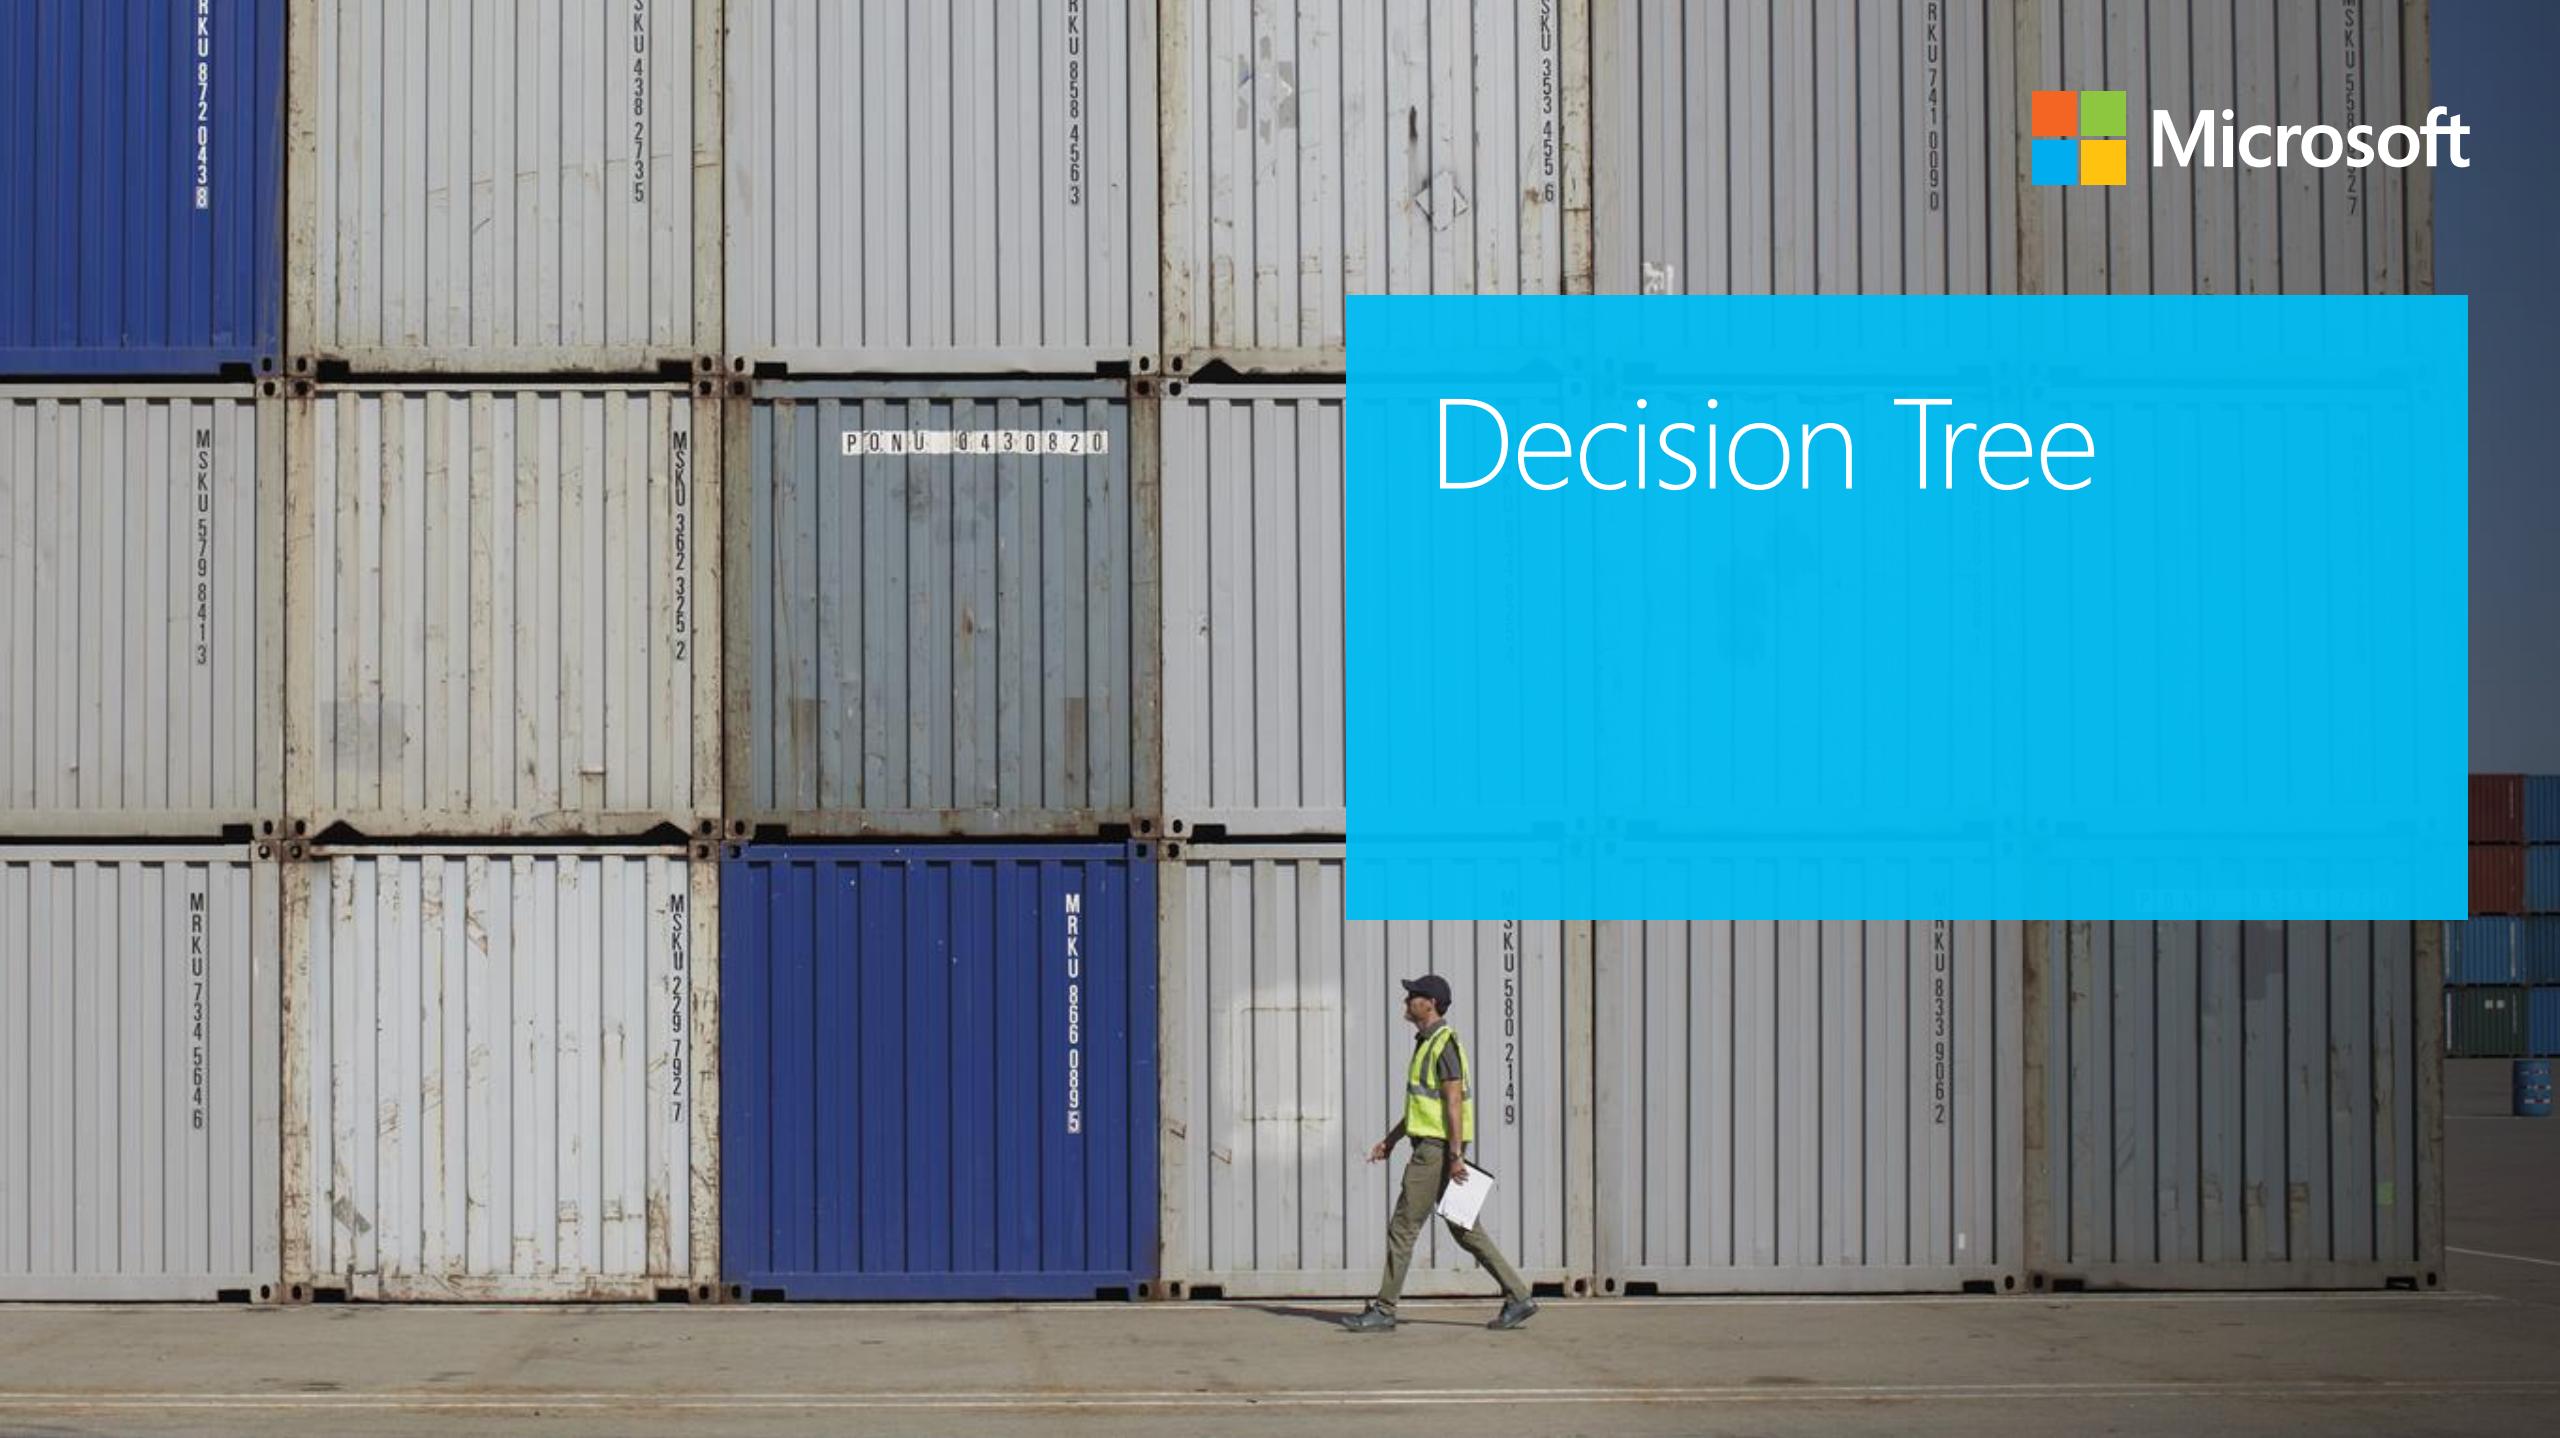
Lets make the hands dirty

Demo -3. Wine Quality in Azure ML



Microsoft

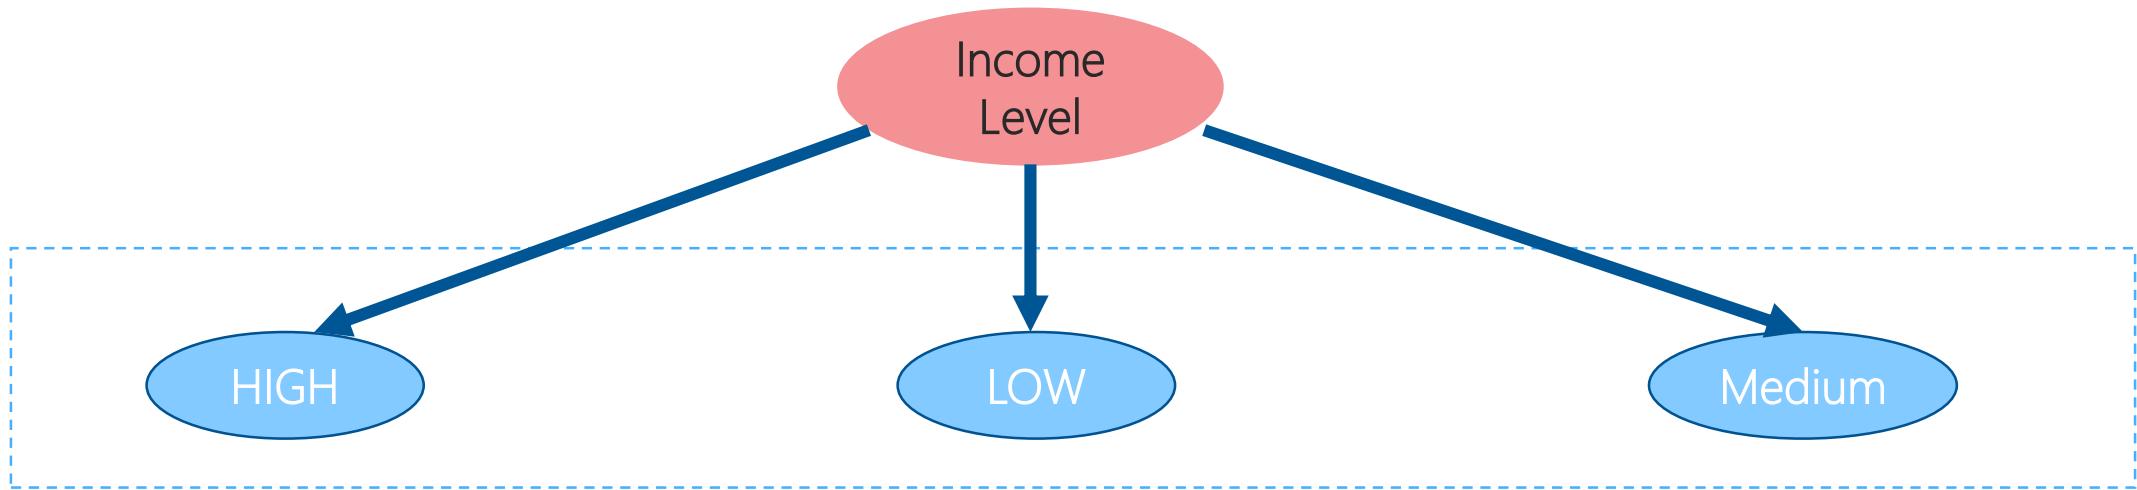
# Decision Tree



# What is Decision Tree?

- Supervised learning method
- Decision support tool that uses a tree-like graph or model of decisions and their possible consequences
- Various variations such as Boosted Decision Tree, Decision Forest, Decision Jungle
- Can be used for categorical as well as continuous variables

Loan ID	Income Level	Credit Score	Employment	Approved?
L1	Medium	Low	Self-Employed	No
L2	High	Low	Self-Employed	Yes
L3	High	High	Salaried	Yes
L4	Medium	Low	Salaried	Yes
L5	Low	High	Salaried	No
L6	Low	Low	Self-Employed	No
L7	High	Low	Salaried	Yes
L8	Medium	Low	Self-Employed	No
L9	High	High	Self-Employed	Yes
L10	Medium	High	Self-Employed	Yes
L11	High	Low	Salaried	Yes
L12	Medium	High	Salaried	Yes
L13	Medium	High	Self-Employed	Yes



LID	IL	CS	ET	Status
L2	High	Low	SE	Yes
L3	High	High	Salaried	Yes
L7	High	Low	Salaried	Yes
L9	High	High	SE	Yes
L11	High	Low	Salaried	Yes

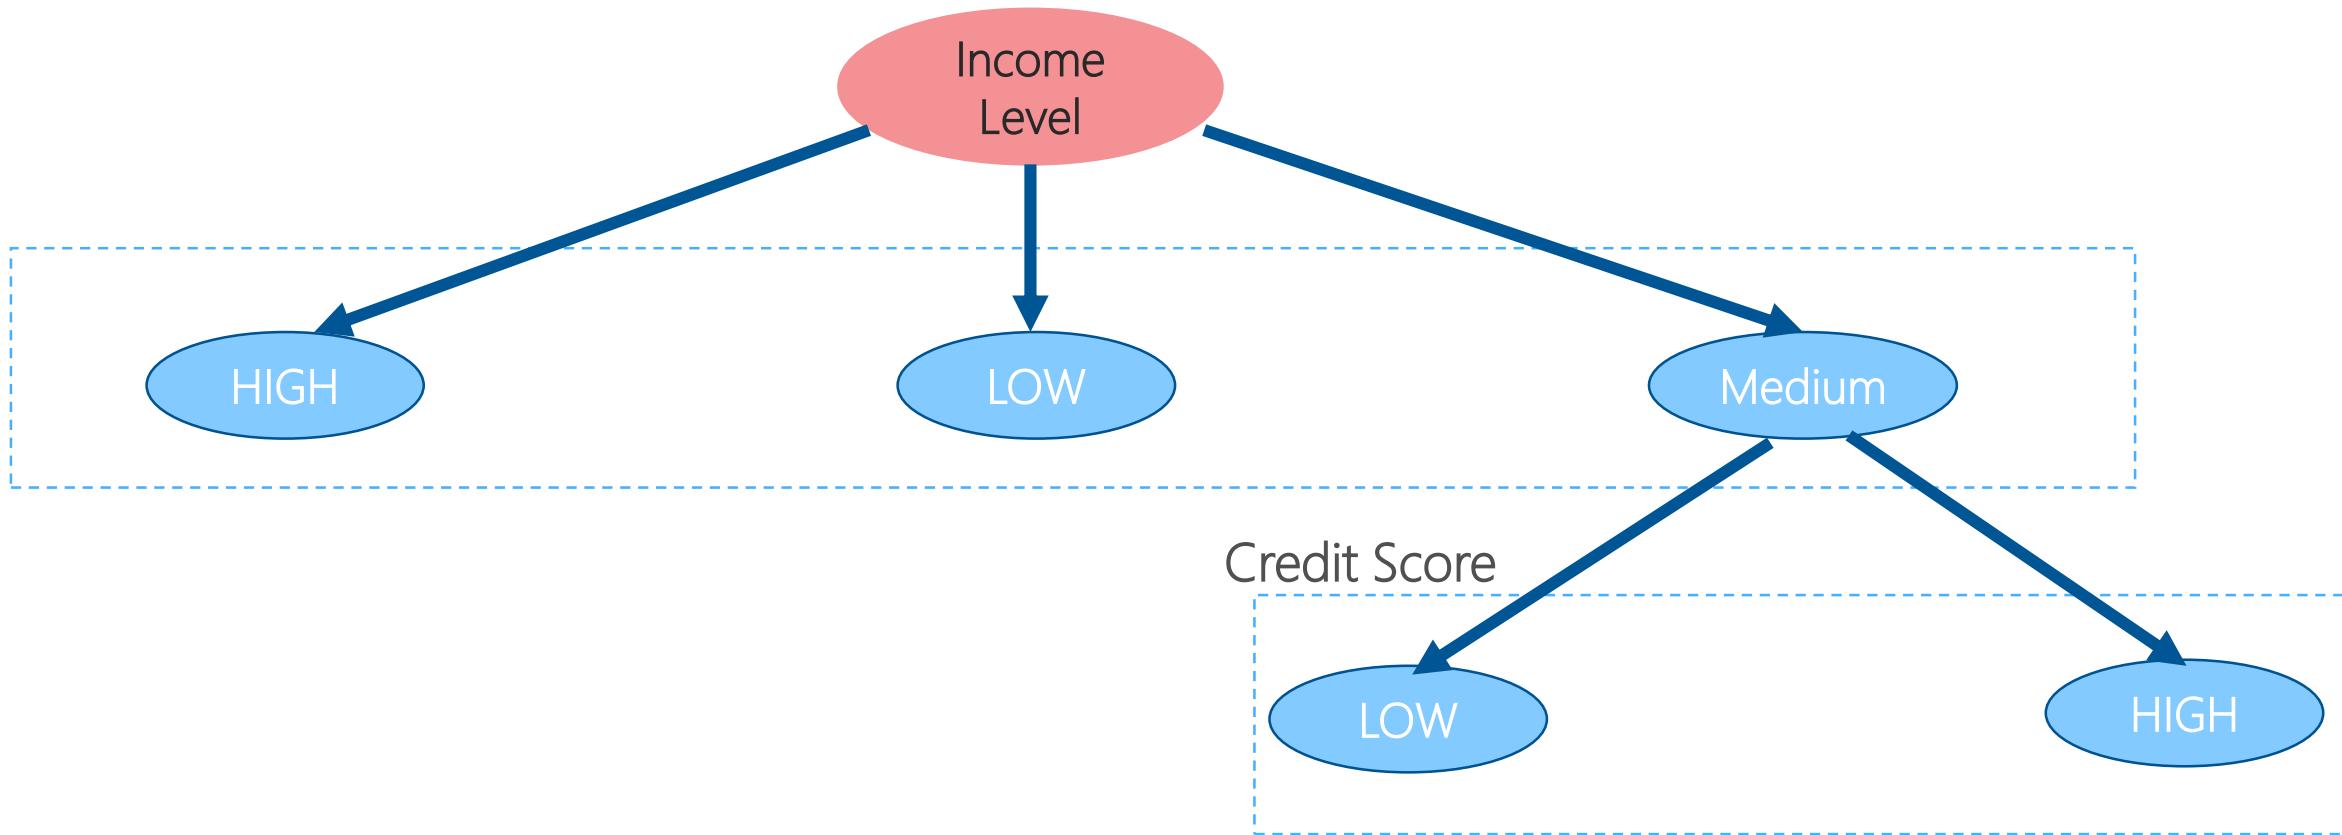
Pure Subset

LID	IL	CS	ET	Status
L5	Low	High	Salaried	No
L6	Low	Low	SE	No
L14	Low	Low	SE	No
L15	Low	High	SE	No

Pure Subset

LID	IL	CS	ET	Status
L1	Medium	Low	SE	No
L4	Medium	Low	Salaried	Yes
L8	Medium	Low	SE	No
L10	Medium	High	SE	Yes
L12	Medium	High	Salaried	Yes
L13	Medium	High	SE	Yes

Split Further

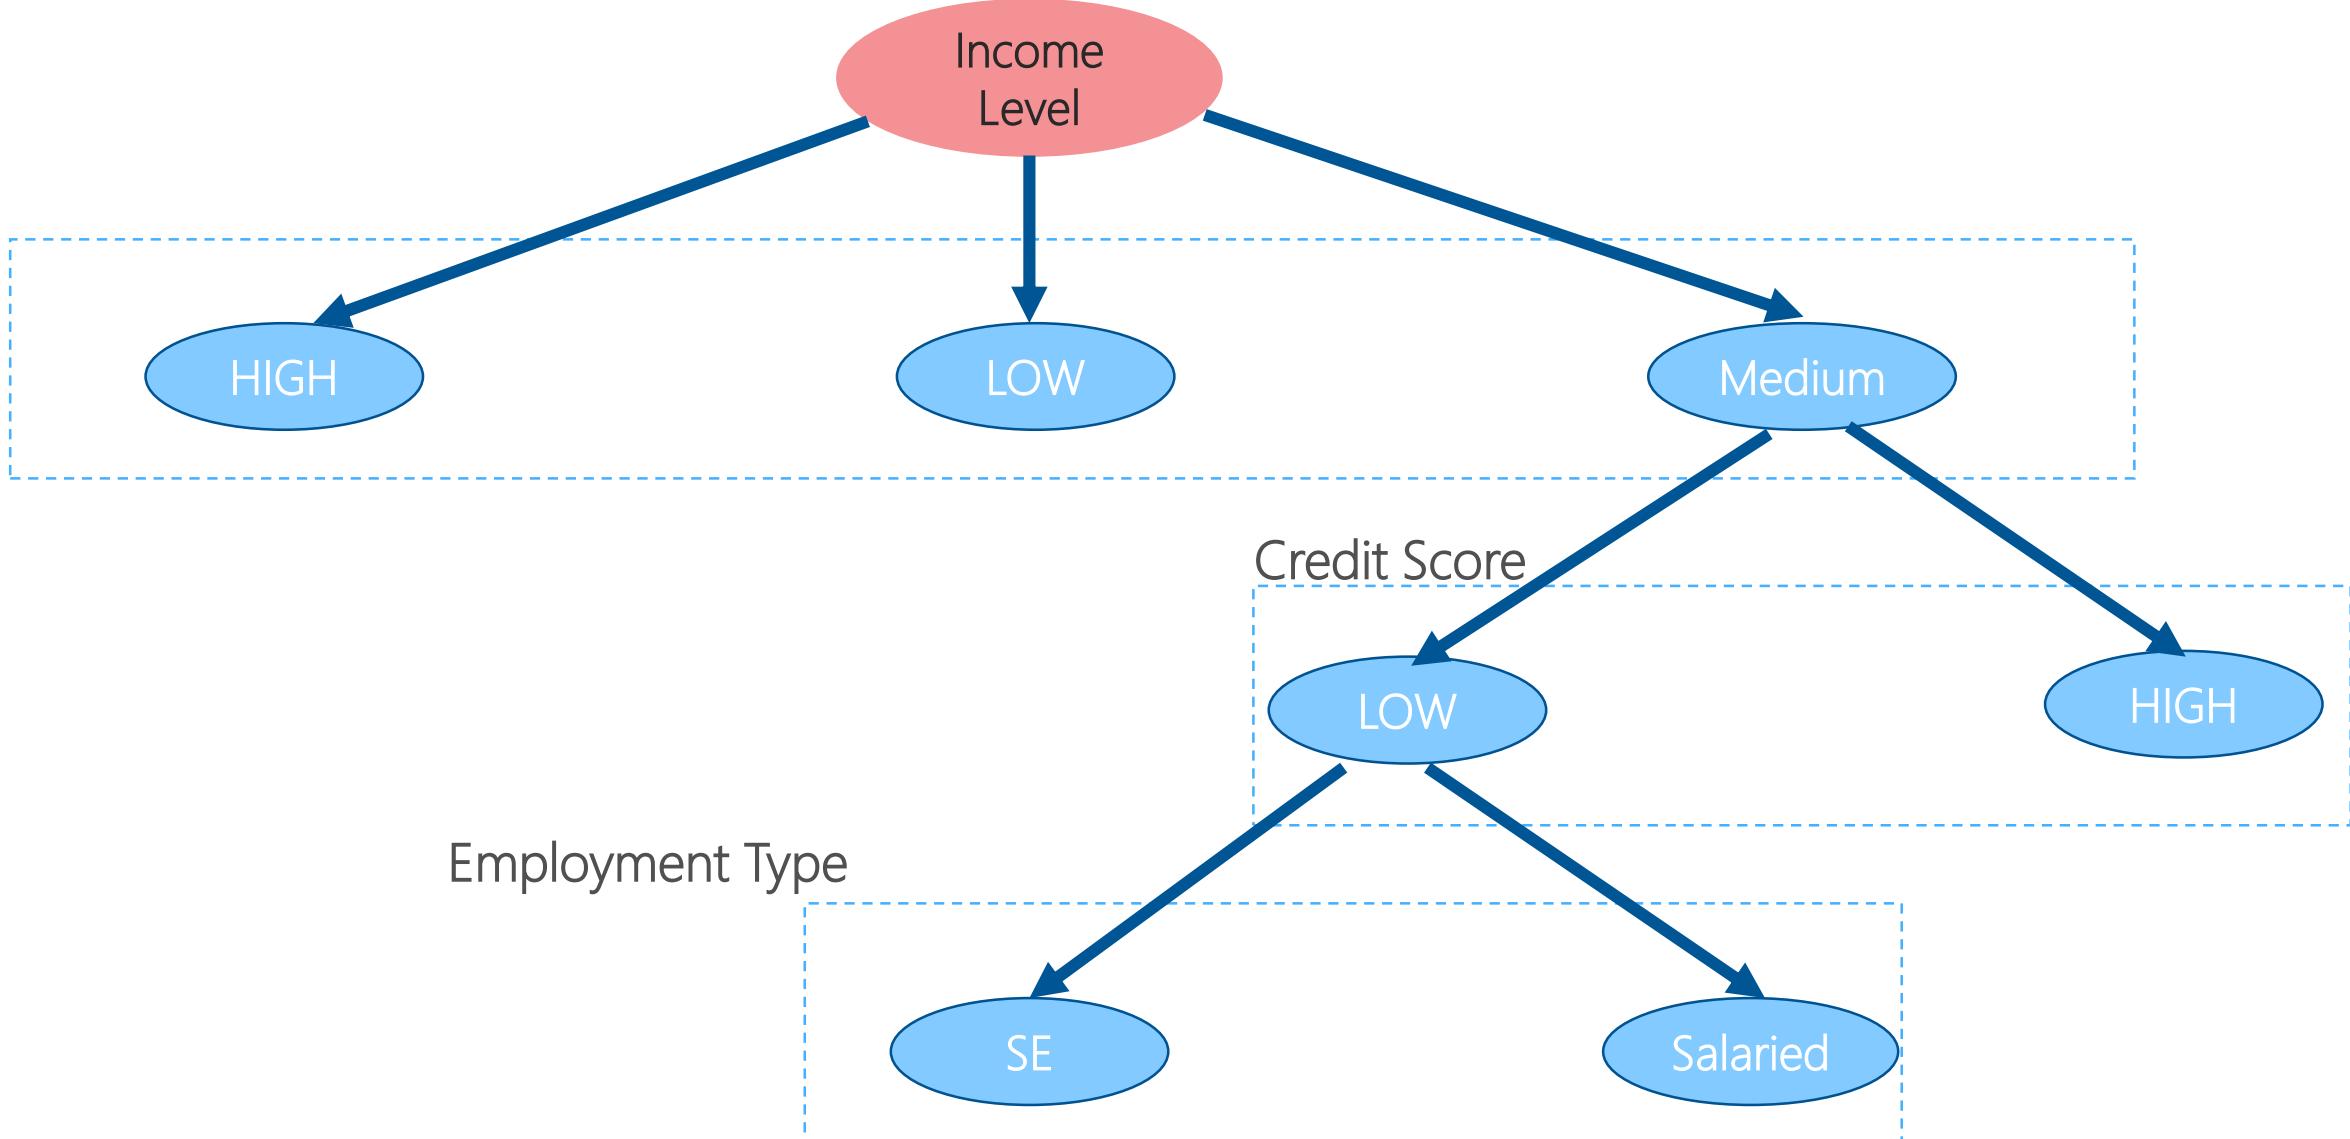


LID	IL	CS	ET	Status
L1	Medium	Low	SE	No
L4	Medium	Low	Salaried	Yes
L8	Medium	Low	SE	No

Split Further

LID	IL	CS	ET	Status
L10	Medium	High	SE	Yes
L12	Medium	High	Salaried	Yes
L13	Medium	High	SE	Yes

Pure Subset



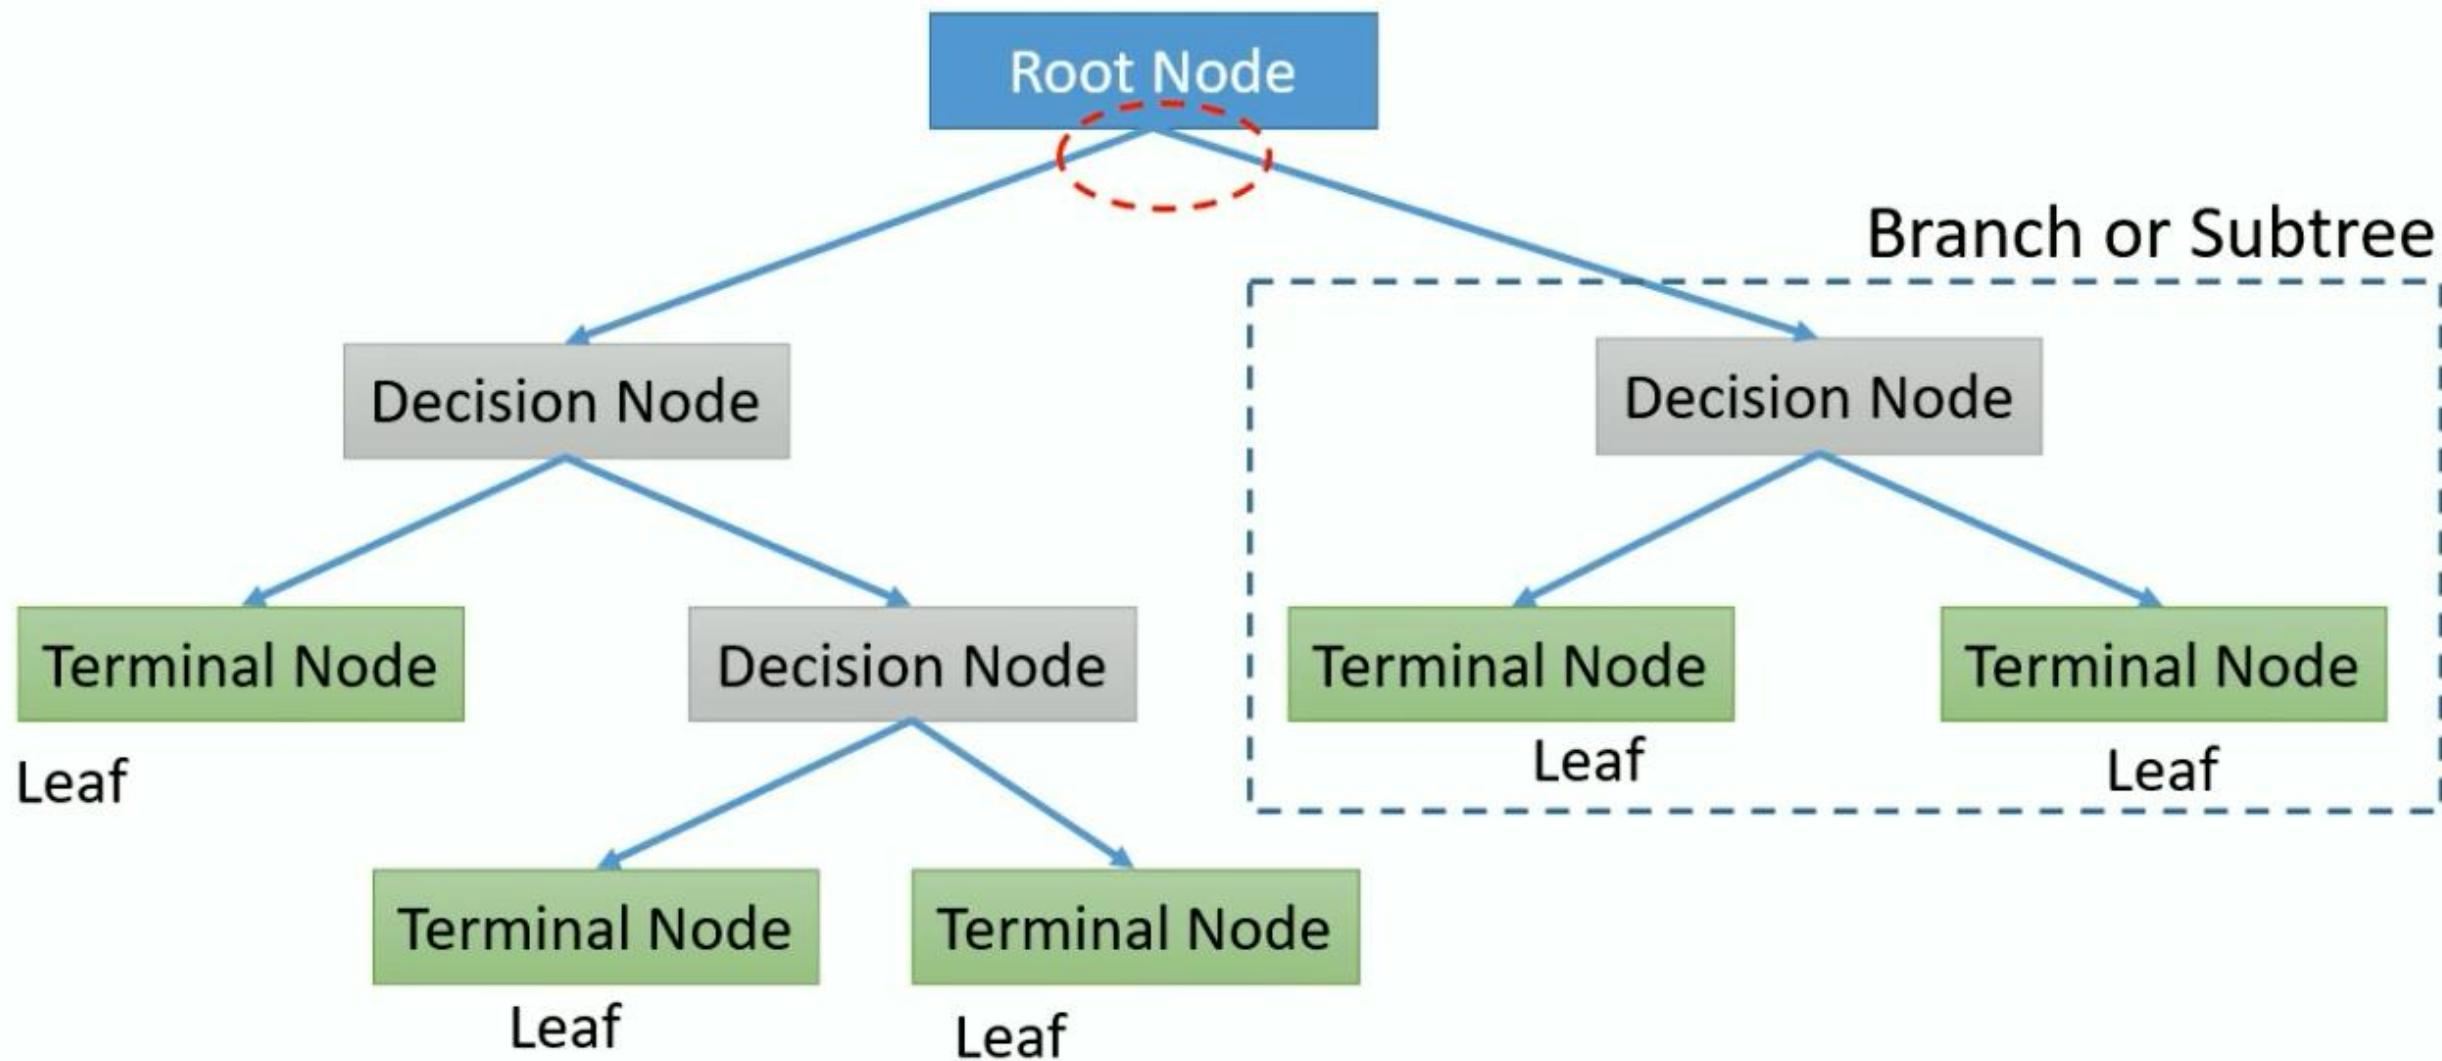
LID	IL	CS	ET	Status
L1	Medium	Low	SE	No
L8	Medium	Low	SE	No

Pure Subset

LID	IL	CS	ET	Status
L4	Medium	Low	Salaried	Yes

Pure Subset

# Decision Tree Terms



# Ensemble learning

# Everyday Ensemble Learning



# Decision?

Is this price fair?

Appreciation of price?

Construction Quality?

Location appropriate?

Neighbourhood?



# Decision?



Broker or real estate portal to check fair price, price appreciation

Friend or colleague who stays nearby or stayed in the neighbourhood

Inspection by an architect for quality checks and structural defects.

# Decision?

Is this price fair?



Appreciation of price?



Construction Quality?



Majority



Weighted Average

Location appropriate?



Neighbourhood?

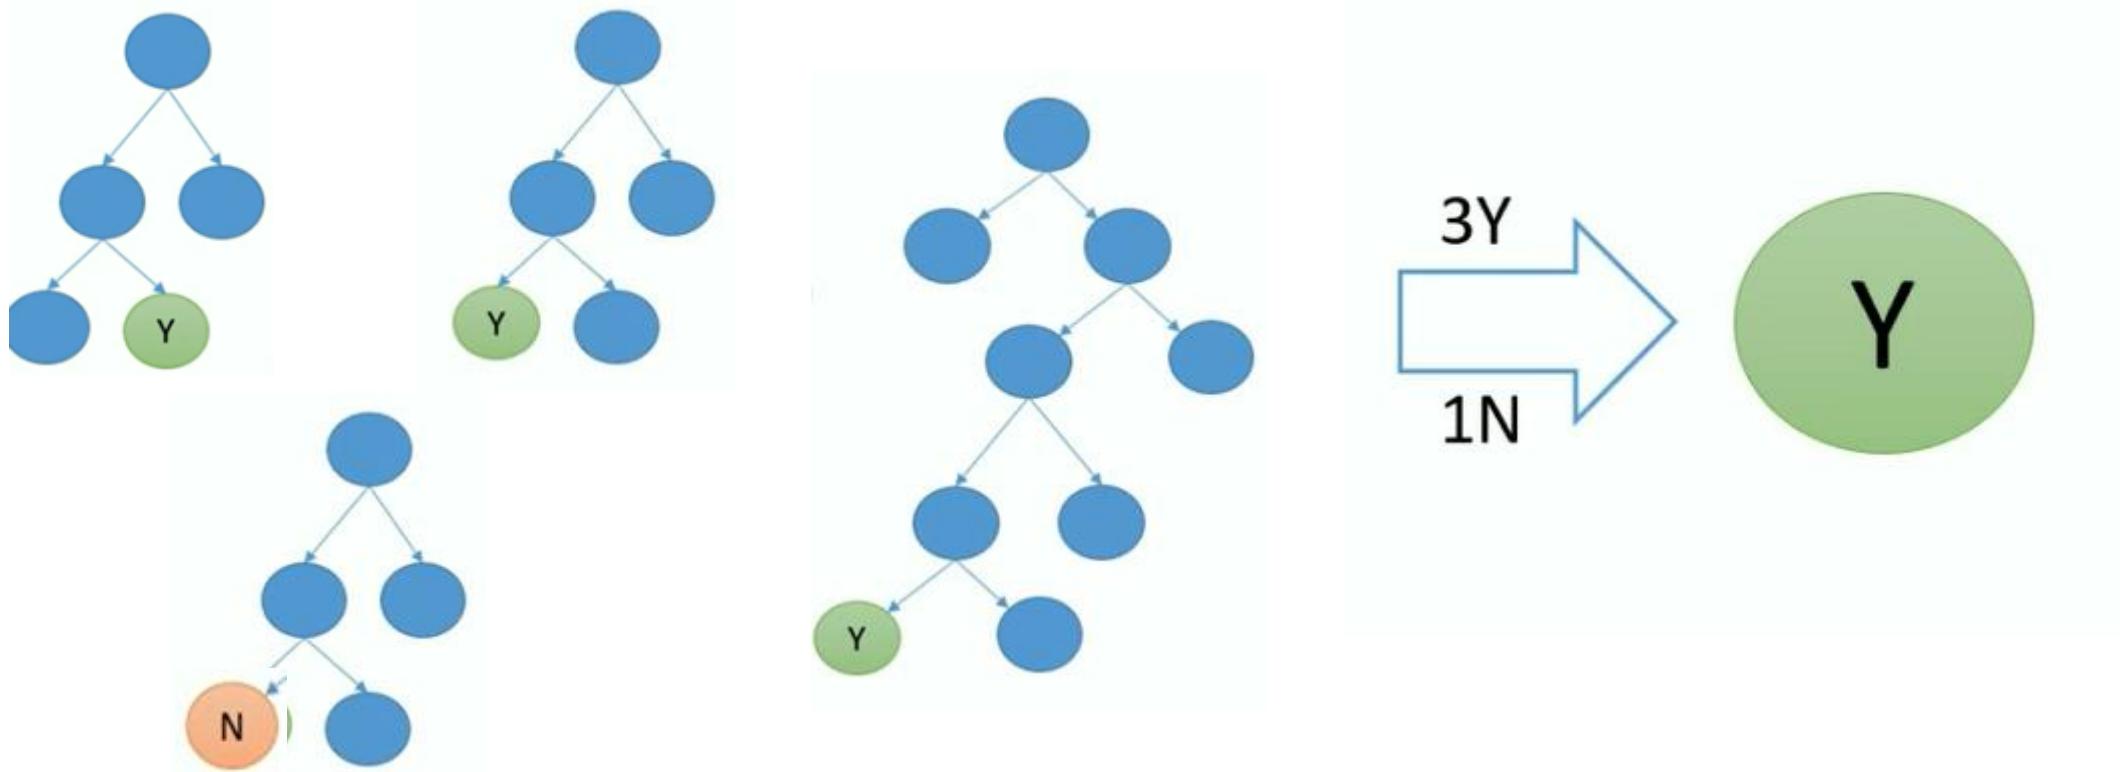


# Ensemble Learning

- All algorithms have errors
- Collective wisdom is higher than the individual intelligence
- Generate a group of base learners and combined result gives higher accuracy
- Different base learners can use different,
  - Parameters
  - Sequence
  - Training sets etc
- Two major Ensemble Learning Methods
  - Bagging
  - Boosting

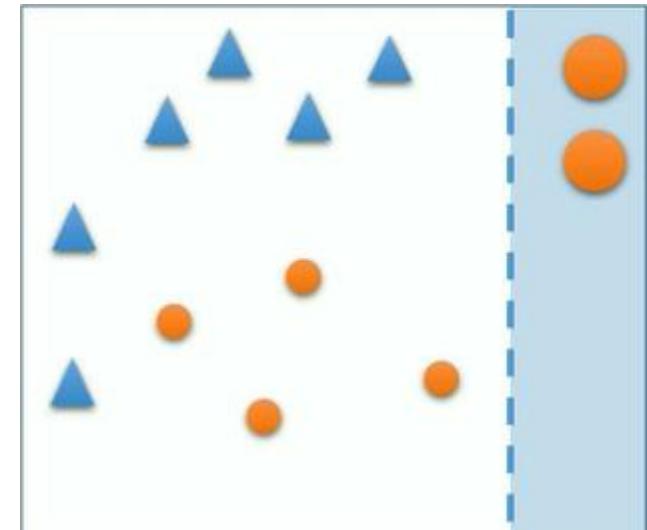
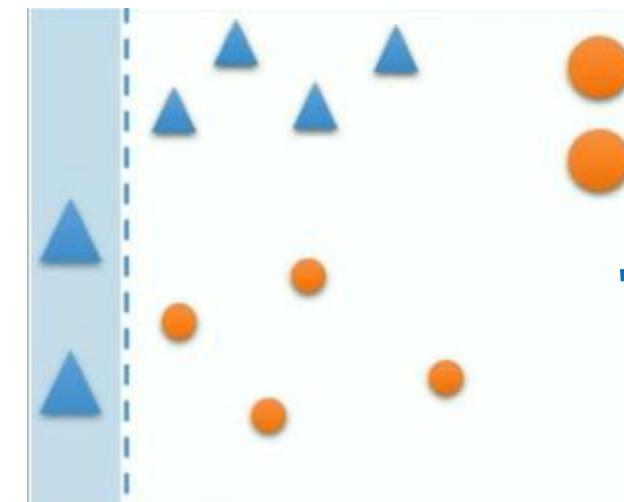
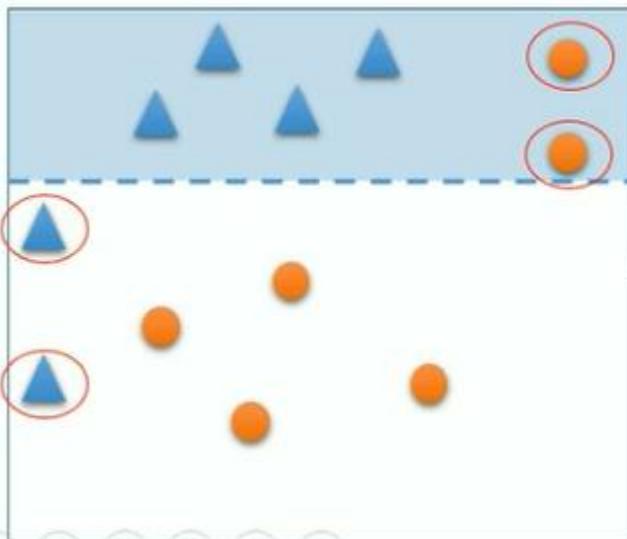
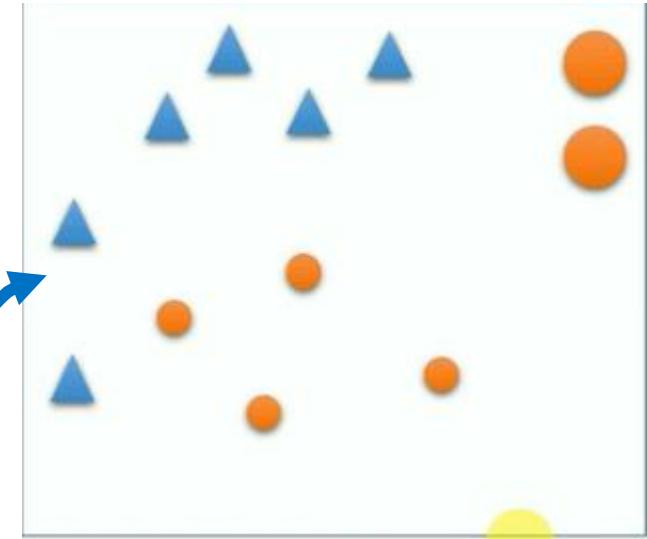
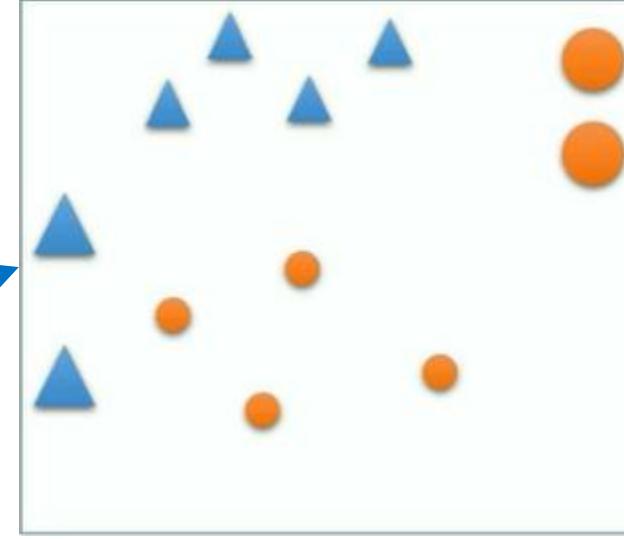
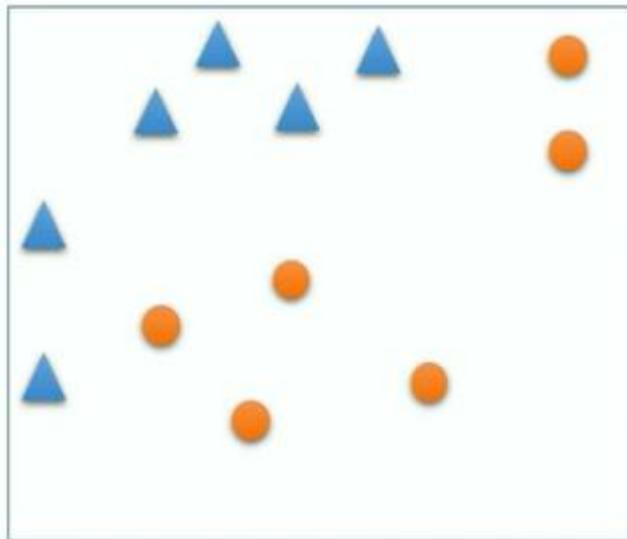
# Bagging

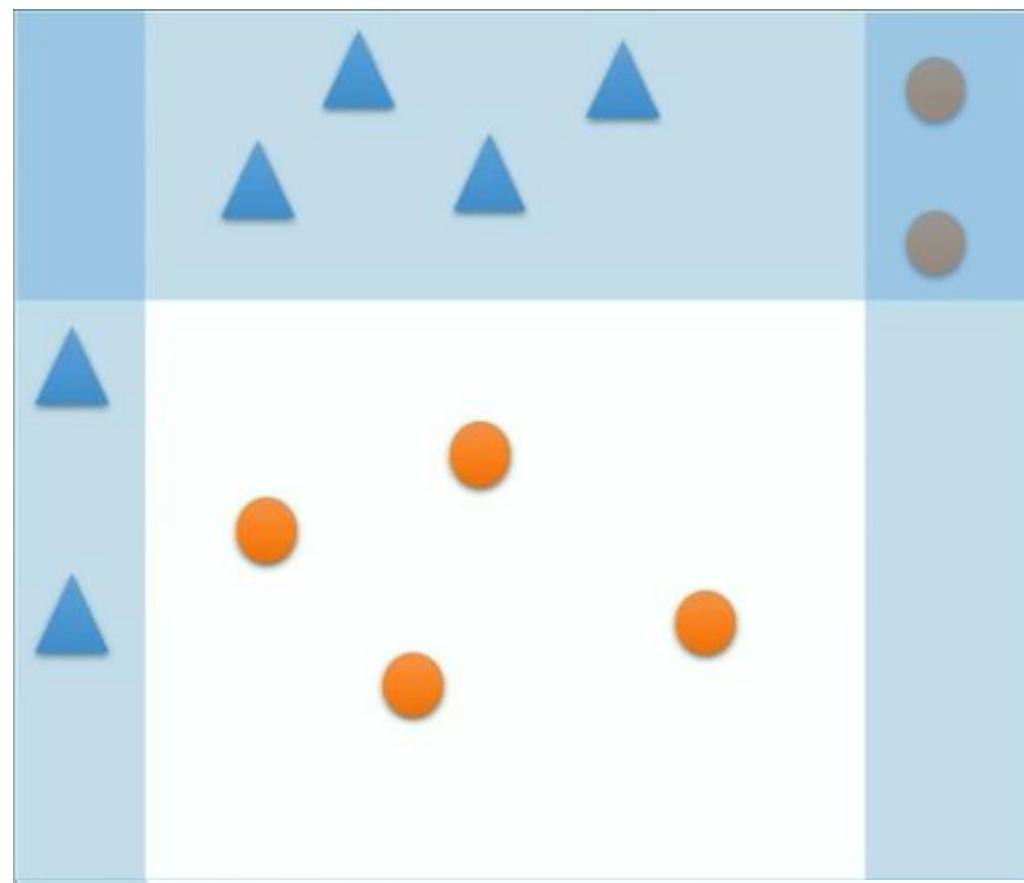
- Various models are built in parallel
- All models vote to give the final prediction



# Boosting

- Train the Decision Tree in a sequence
- Learn from the previous tree by focussing on incorrect observations
- Build new model with higher weight for incorrect observations from previous sequence





Lets make the hands dirty

Demo -4. Two Class Boosted Decision Tree

# Bank Telemarketing

- Goal is to predict if the client will subscribe to a product or not

Number of instances — 45, 211

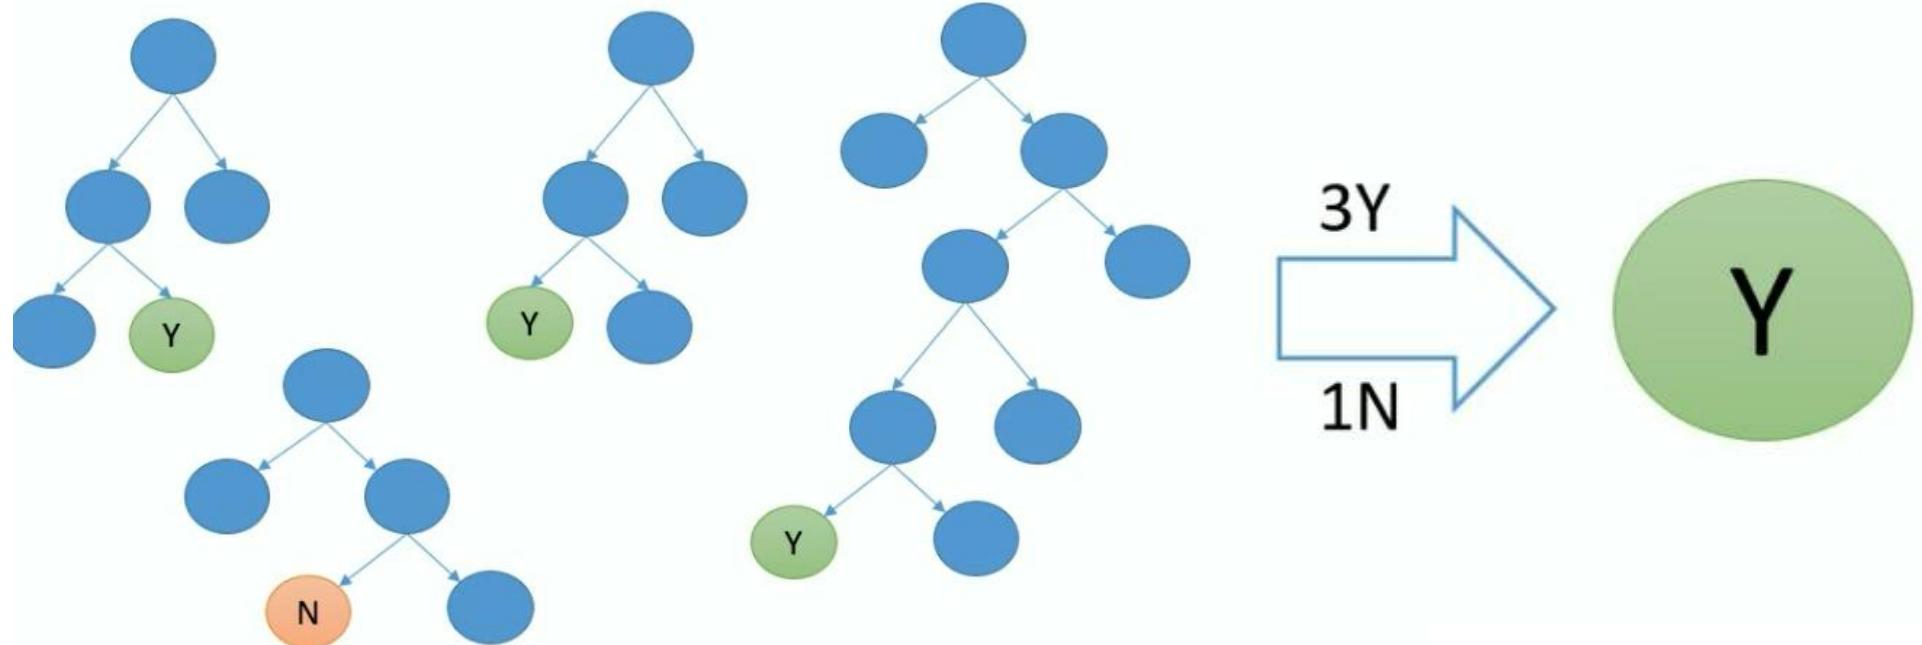
1. Age
2. Job Type
3. Marital Status
4. Education Level
5. Credit Default?
6. Housing Loan?
7. Personal Loan
8. Contacted Type
9. Contacted Month
10. Last Contacted day
11. Contact Duration
12. Campaign Type
13. P-Days
14. Previous
15. P-Outcome
16. Emp-Var-Rate
17. Consumer Price Index
18. Consumer Confidence Index
19. Euribor 3 Month Rate
20. Number of employees
21. Subscribed?

# Two Class Boosted Decision Tree?

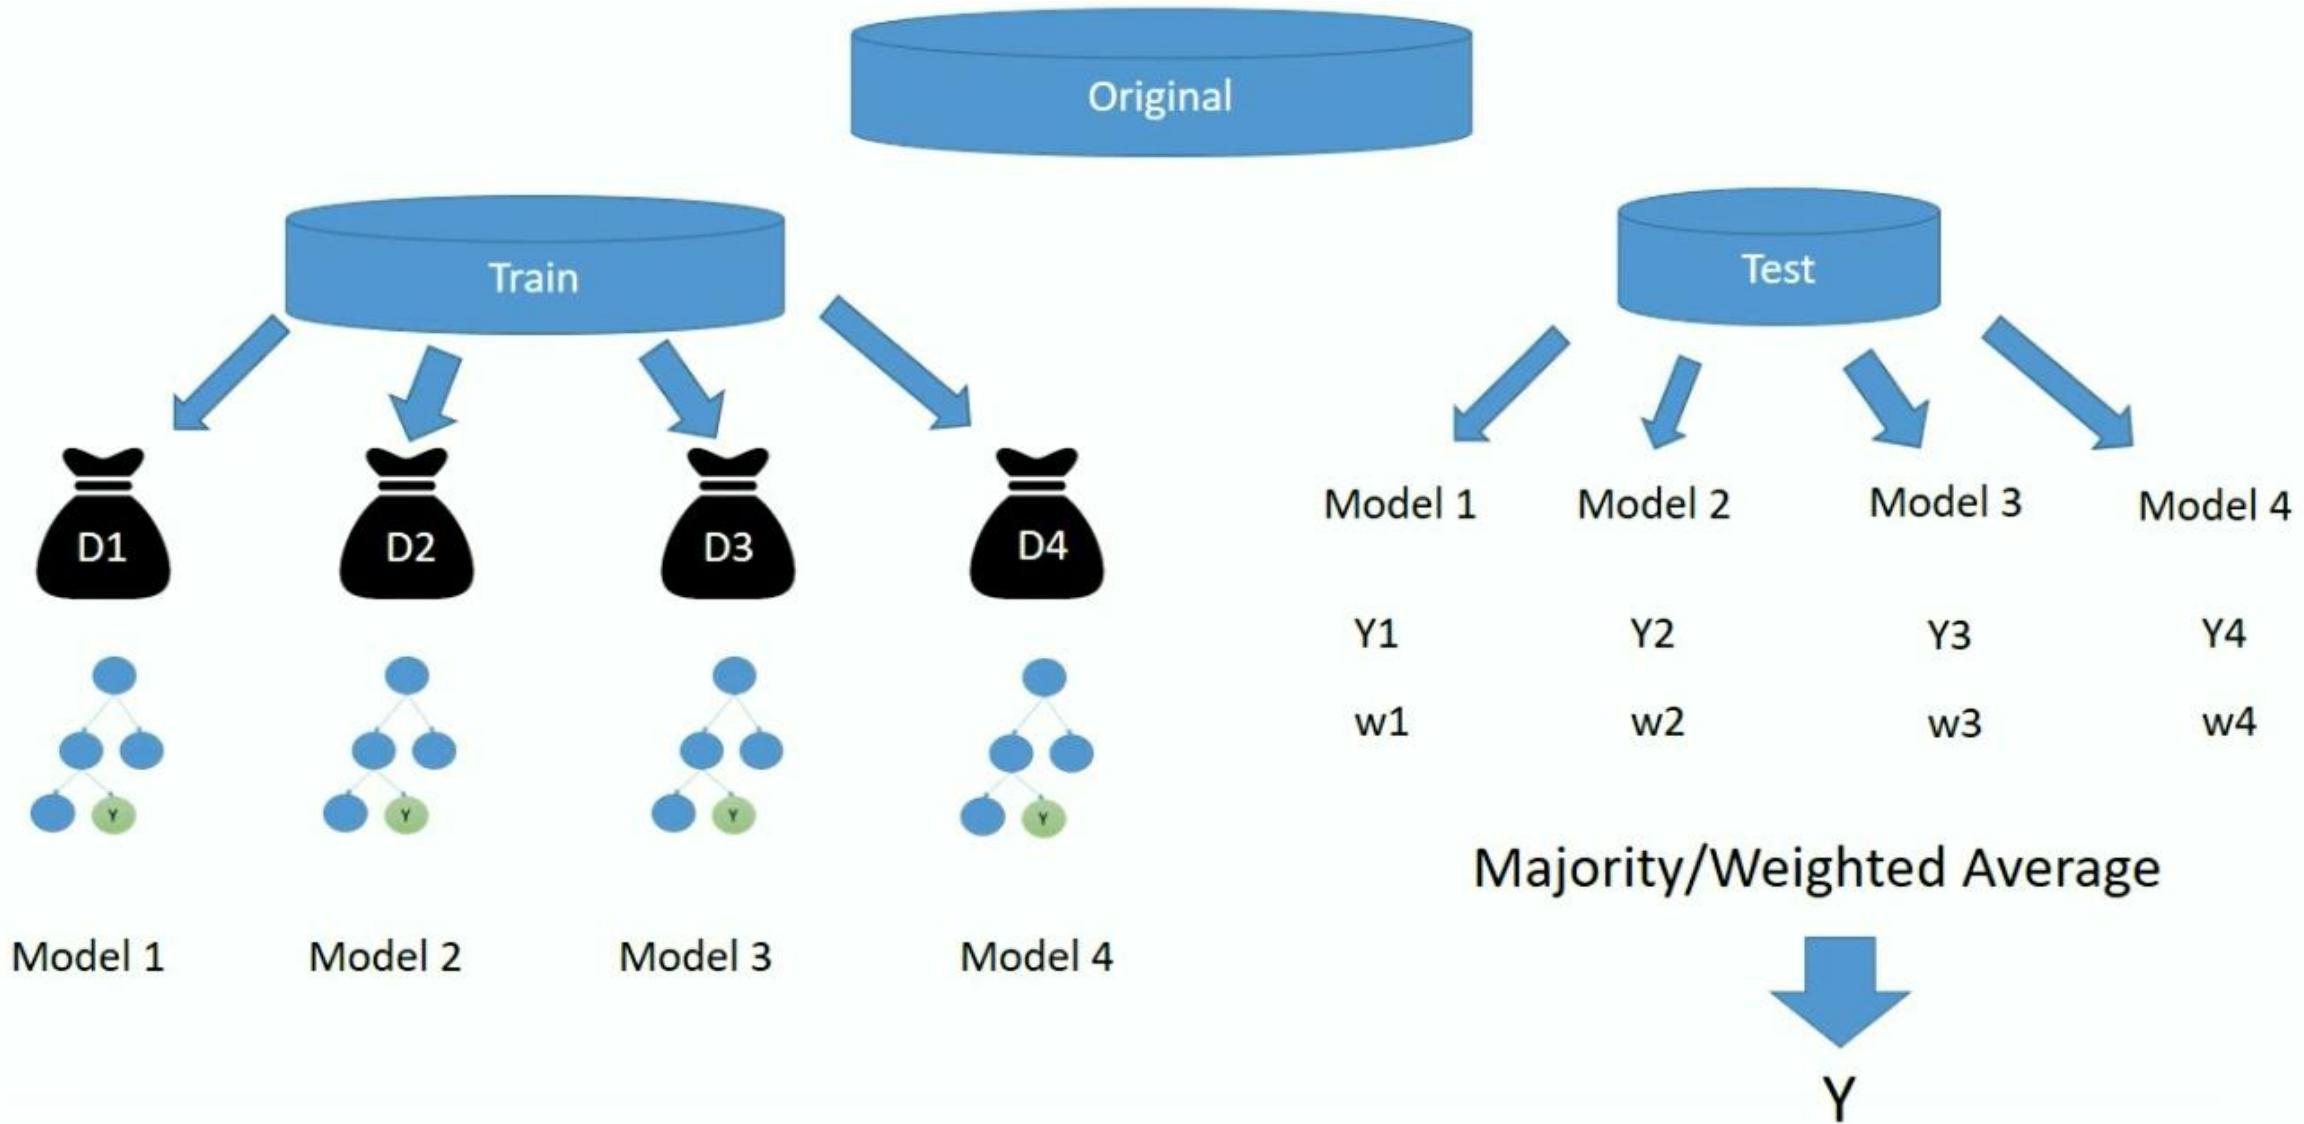
- Machine learning model based on the boosted decision trees algorithm
- Based on ensemble learning method
- Among the easiest methods to get top performance
- One of the more memory-intensive learners

# Bagging

- Various models are built in parallel
- All models vote to give the final prediction

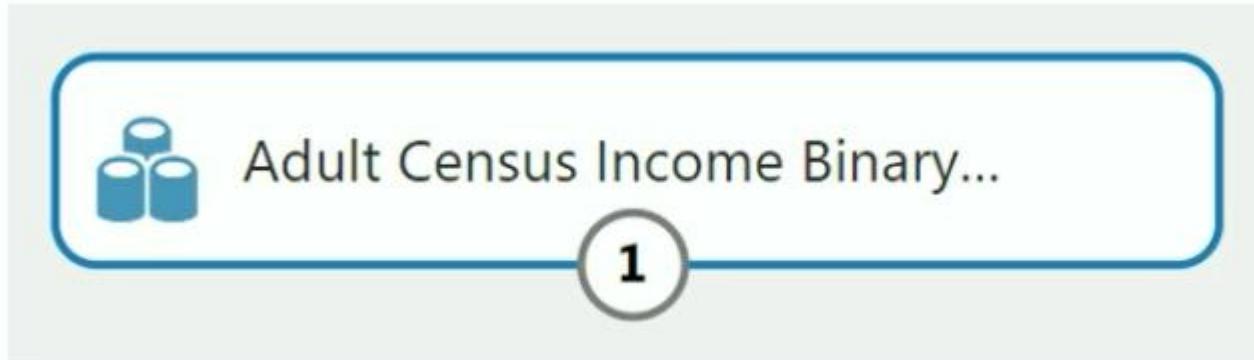


# Bagging



# Adult Census Data

Problem statement: Predict whether income exceeds \$50K/yr based on census data



1. Age
2. Workclass
3. Fnlwgt
4. Education
5. Education-Num
6. Marital Status
7. Occupation
8. Relationship
9. Race
10. Sex
11. Capital Gains
12. Capital Losses
13. Hours per week
14. Native Country
15. Income

# Demo -5. Decision Forest

# Demo -6. Support Vector Machine & Evaluate 2 Models

# Tune Model Hyperparameters

# Models and Parameters

## Two-Class Logistic Regression

1

Create trainer mode

Single Parameter

Optimization tolerance

1E-07

L1 regularization weight

1

L2 regularization weight

1

Memory size for L-BFGS

20

Random number seed

Allow unknown categorical levels

## Two-Class Boosted Decision...

1

Create trainer mode

Single Parameter

Maximum number of leaves per tree

20

Minimum number of samples per leaf node

10

Learning rate

0.2

Number of trees constructed

100

Random number seed

Allow unknown categorical levels

## Two-Class Support Vector M...

1

Create trainer mode

Single Parameter

Number of iterations

1

Lambda

0.001

Normalize features

Project to the unit-sphere

Random number seed

Allow unknown categorical levels

# What are Hyperparameters?



# Model parameters?

- Decision Trees
  - Maximum number of leaves per tree
  - Minimum number of samples per leaf node
  - Learning rate
  - Number of trees to construct
- Logistic regression
  - Optimization tolerance
  - L1 regularization weight
  - L2 regularization weight
  - Memory size for L-BFGS

# Tune Model Hyperparameters?

- Helps in determining the best possible combination of hyperparameters
- Also known as hyperparameter optimization
- Performance metric to measure
  - Accuracy (Overall, how often is the classifier correct?)
  - Precision (When it predicts yes, how often is it correct?)
  - Recall (When it's actually yes, how often does it predict yes?)
  - AUC
  - F1Score

# Parameter Sweeping Modes

- Random Grid
- Entire Grid
- Random Sweep

# What is a Grid?

- Cartesian Product of Parameters
- Parameter 1 → 1, 2, 3
- Parameter 2 → A, B, C, D

		Parameter 1 →		
		1	2	3
← Parameter 2	A	A, 1	A, 2	A, 3
	B	B, 1	B, 2	B, 3
	C	C, 1	C, 2	C, 3
	D	D, 1	D, 2	D, 3

# Random Grid

Parameter 1 →

← Parameter 2

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

# Entire Grid

← Parameter 2

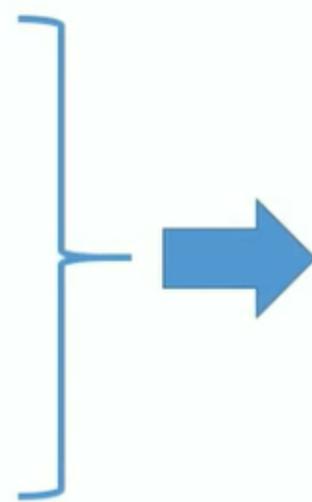
Parameter 1 →

	1	2	3
A	A, 1	A, 2	A, 3
B	B, 1	B, 2	B, 3
C	C, 1	C, 2	C, 3
D	D, 1	D, 2	D, 3

# Random Sweep

Parameter 1 → Range 1.....4

Parameter 2 → Range A.....D



Iterations

P1, P2  
P1, P2  
P1, P2  
. .  
. .  
. .  
P1, P2

# DEMOS

# Tuning Hyper Parameters



Microsoft

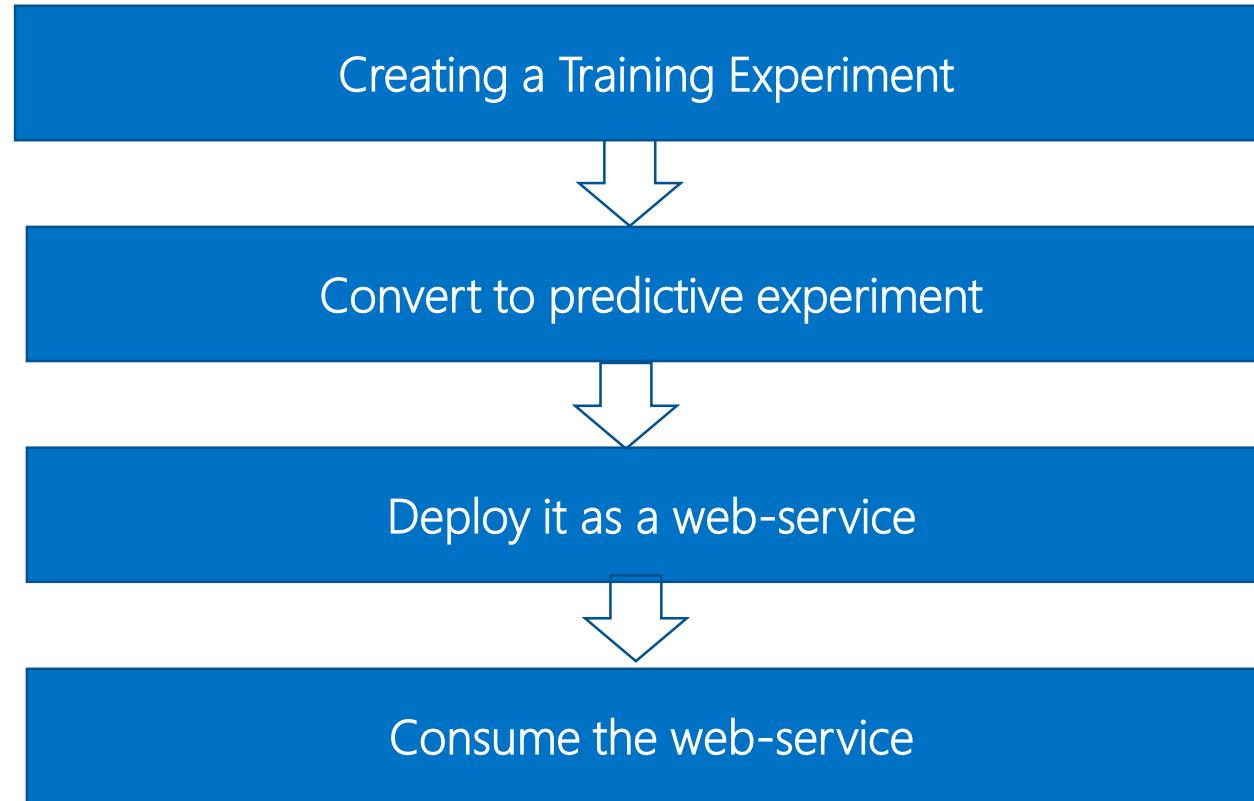
# Deploying Models



# I have a Model. Now What?

- Develop but how to deploy?
- Model language is not supported
- Difficult to deploy in the current architecture
- Tedious Environment set-up
- Many more...

# Steps for Azure ML webservices Deployment



# DEMOS

## Deploying & Creating Webservice



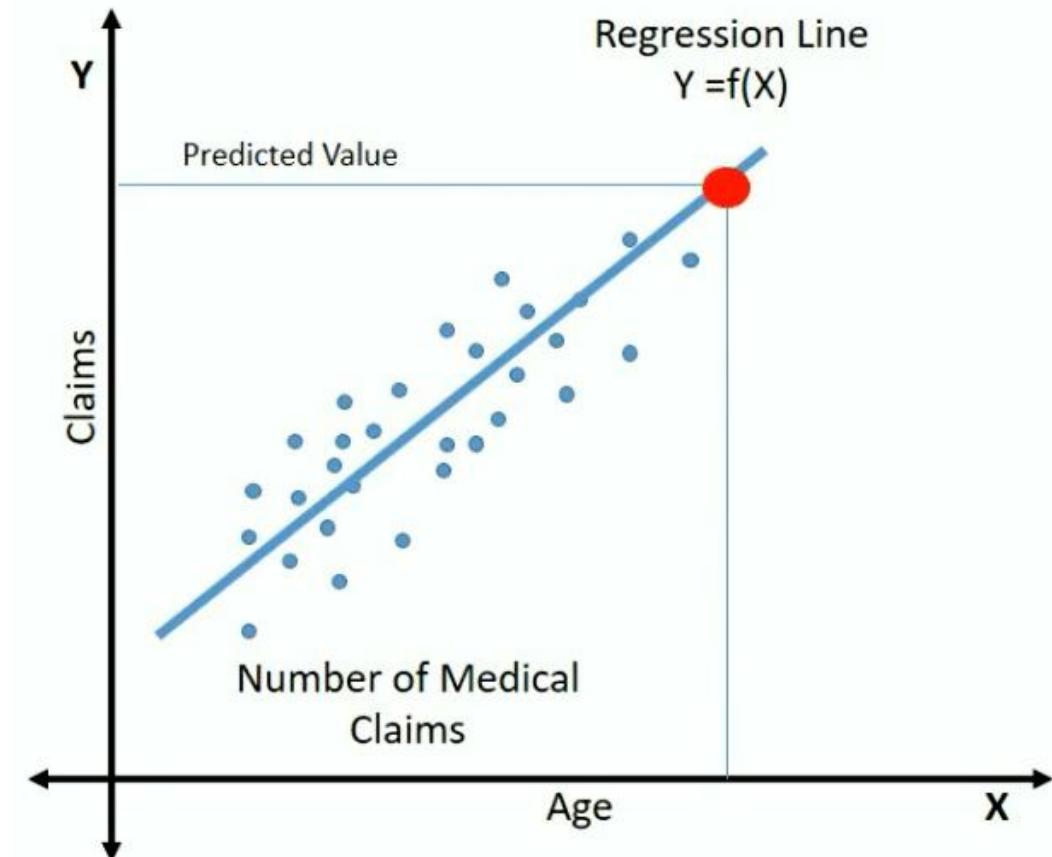
Microsoft

# Regression in AML



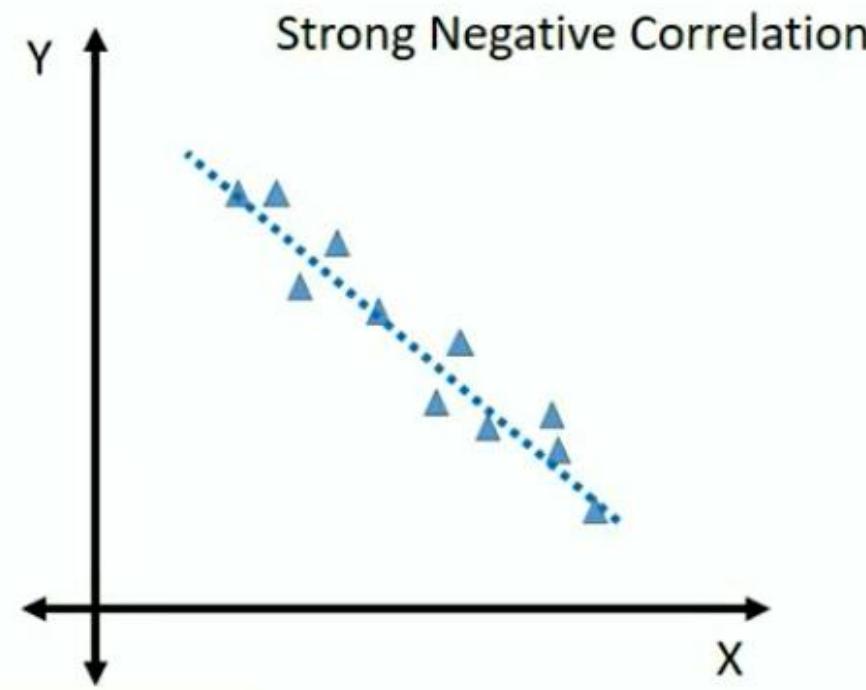
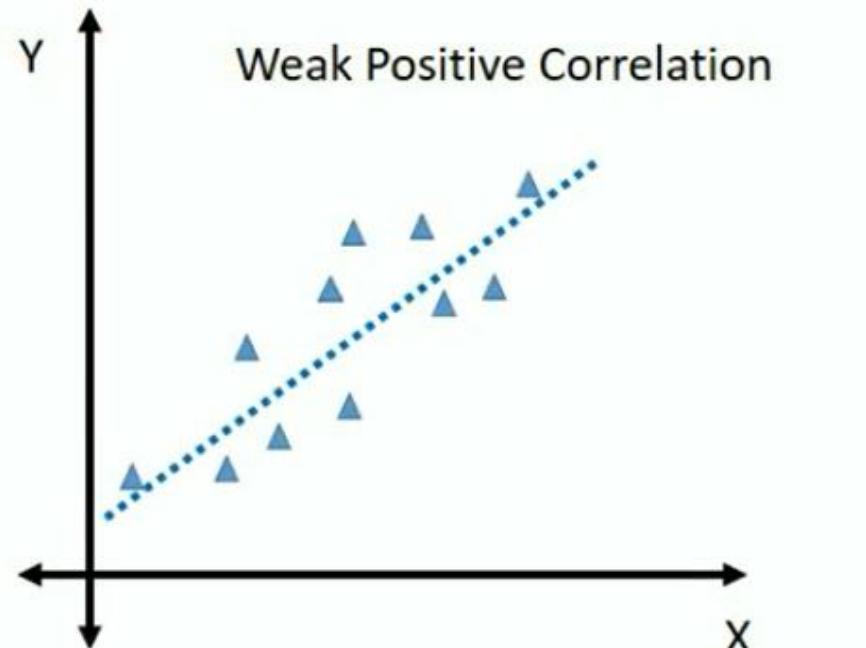
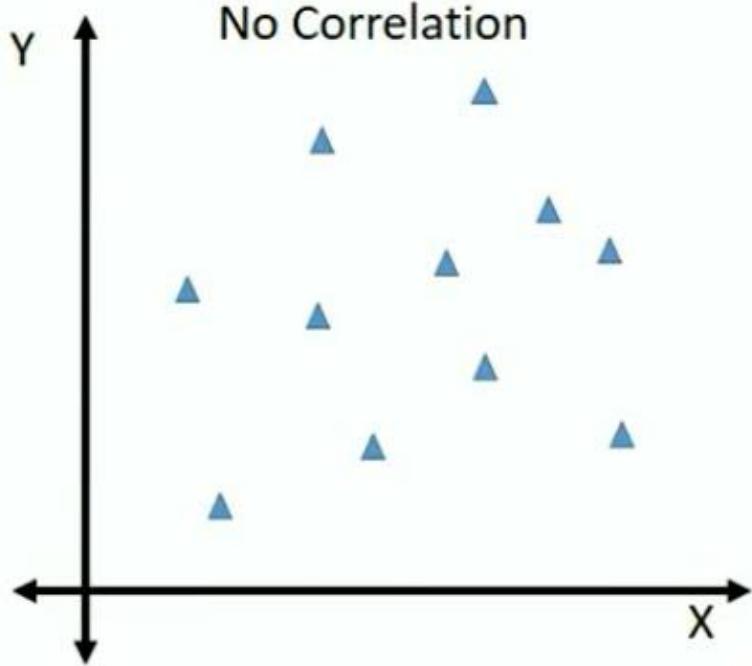
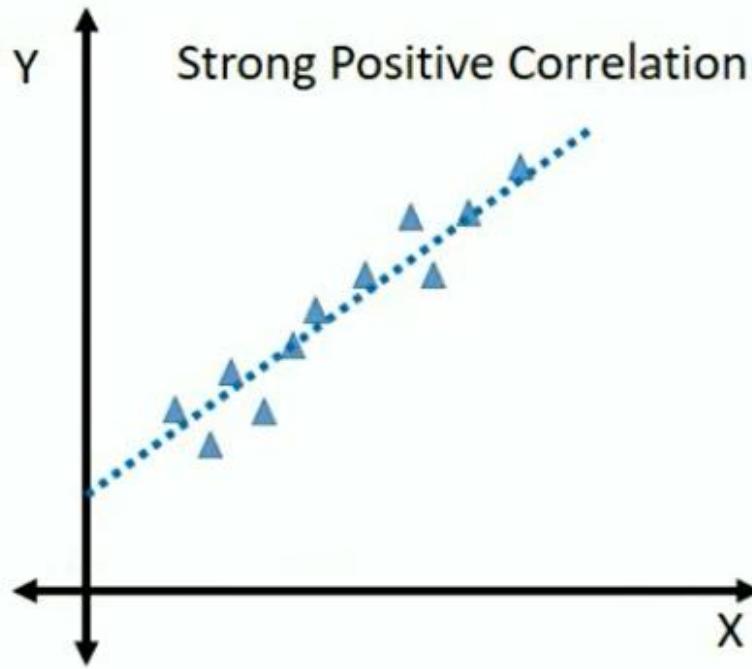
# Regression Analysis

- Statistical process for estimating the relationships among variables
- Relationship between a dependent variable and one or more independent variables (or 'predictors')
- The predictor is a continuous variable
- Can also be used to infer causal relationships between dependent and independent variables.



# Casual Relationship?





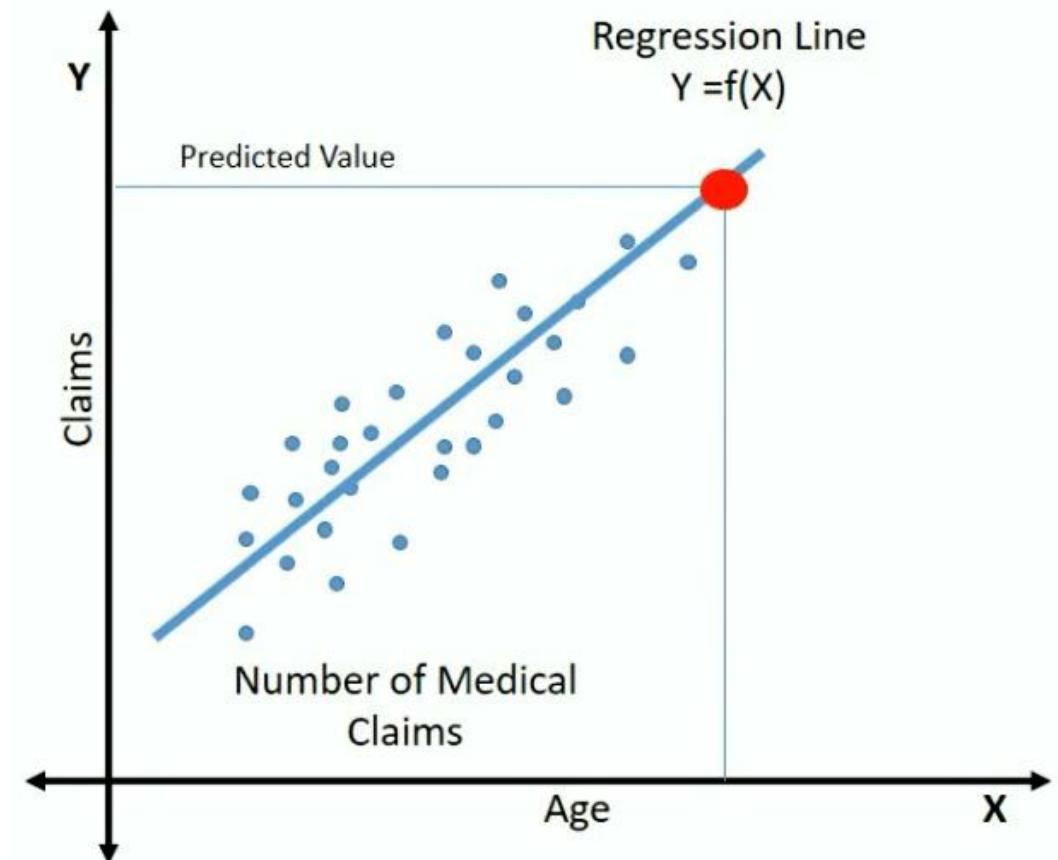
# Linear Regression

Simple Regression

$$Y = B_0 + B_1 X$$

Multivariate linear regression.

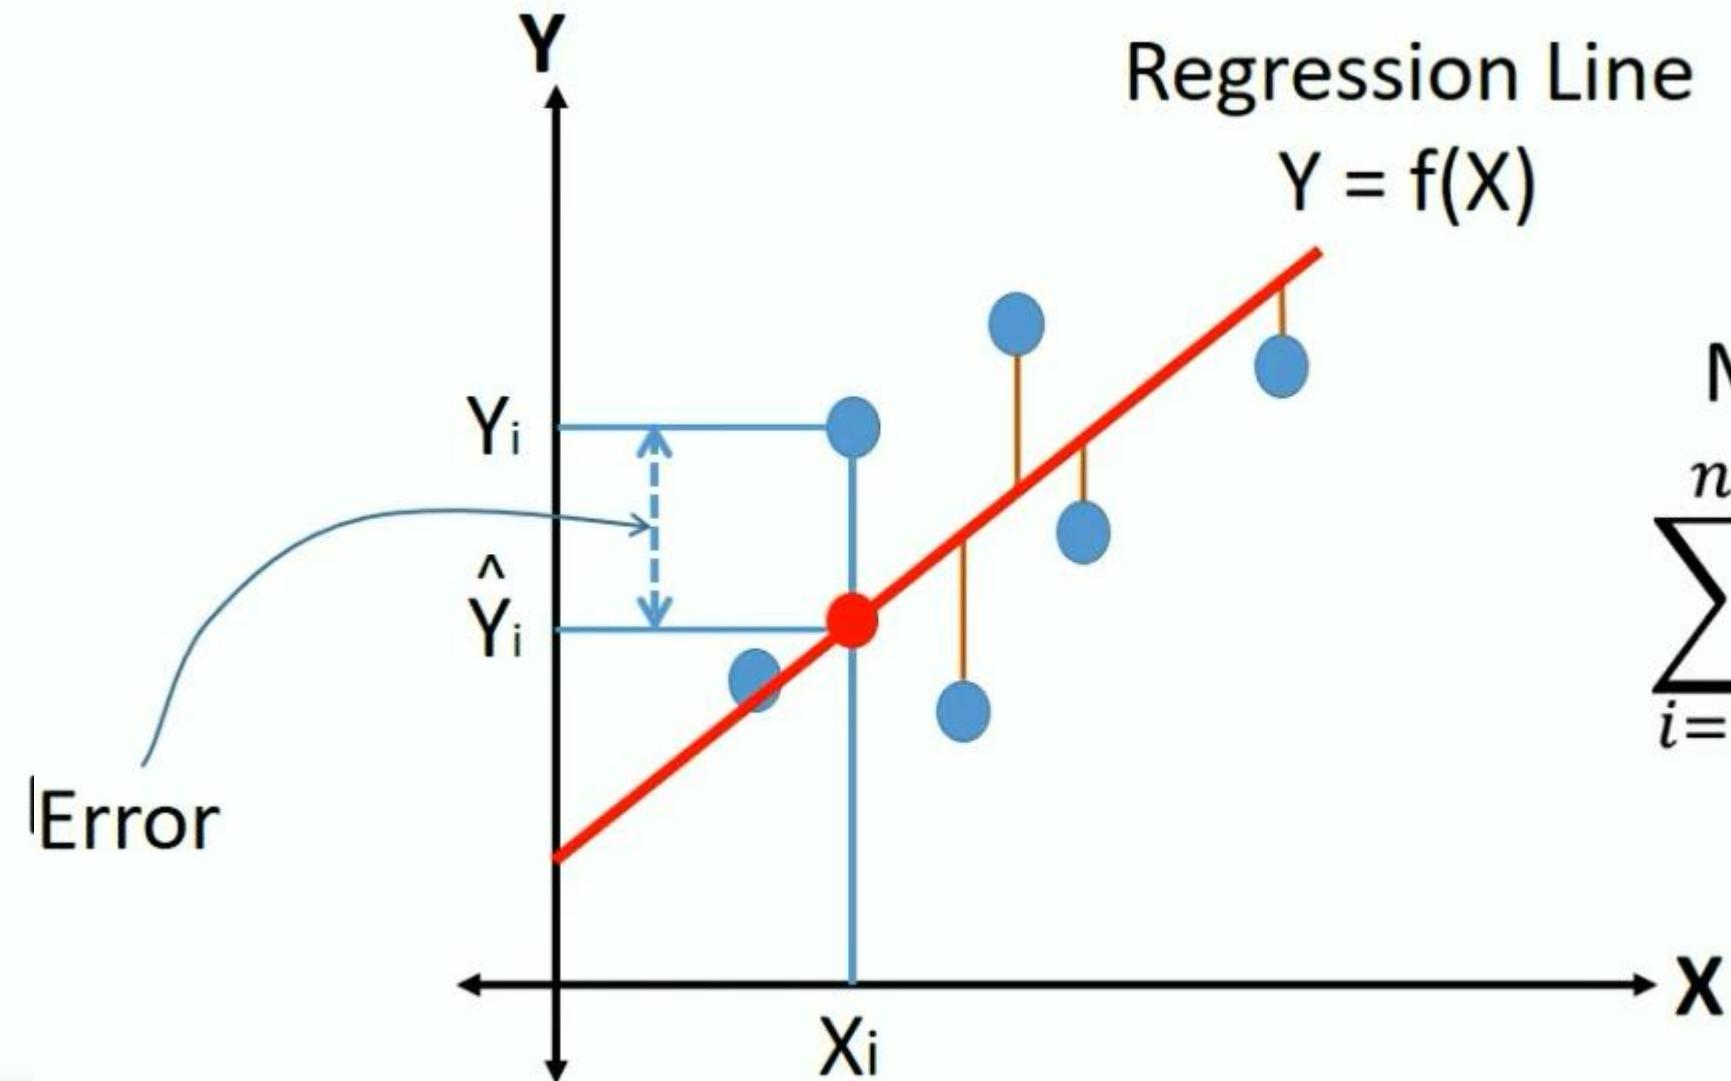
$$Y = B_0 + B_1 X_1 + B_2 X_2 + B_3 X_3 + \dots + B_n X_n$$



# Simple Linear Regression

# Common Regression Terms

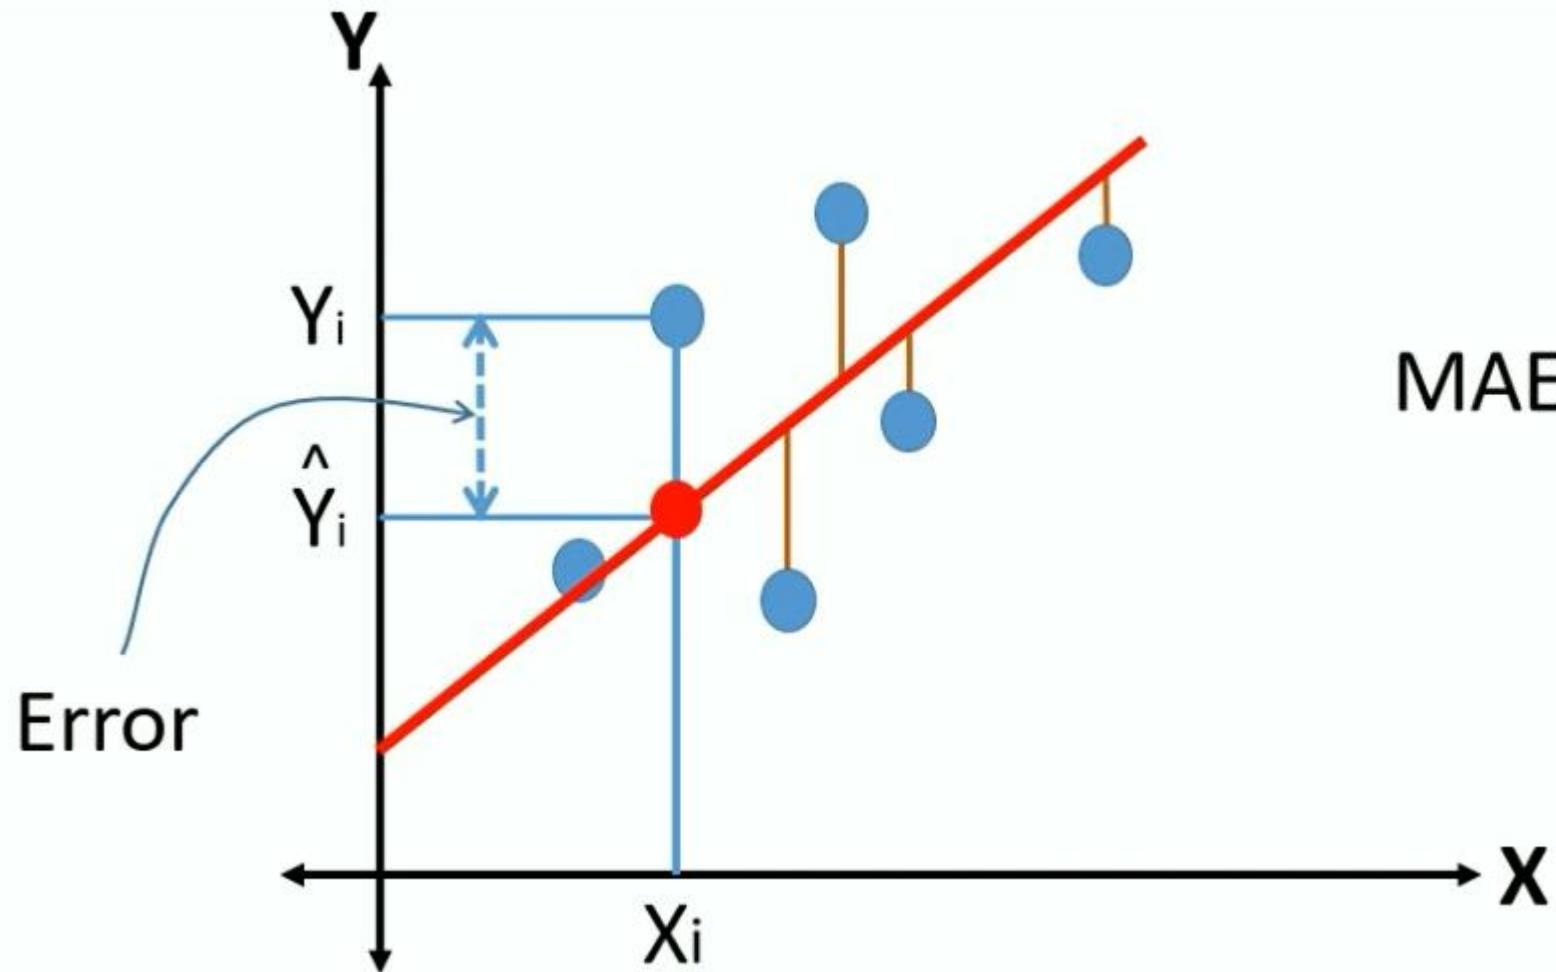
# Ordinary Least Square



Minimum

$$\sum_{i=1}^n (y_i - \hat{y}_i)^2$$

# Mean Absolute Error



$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

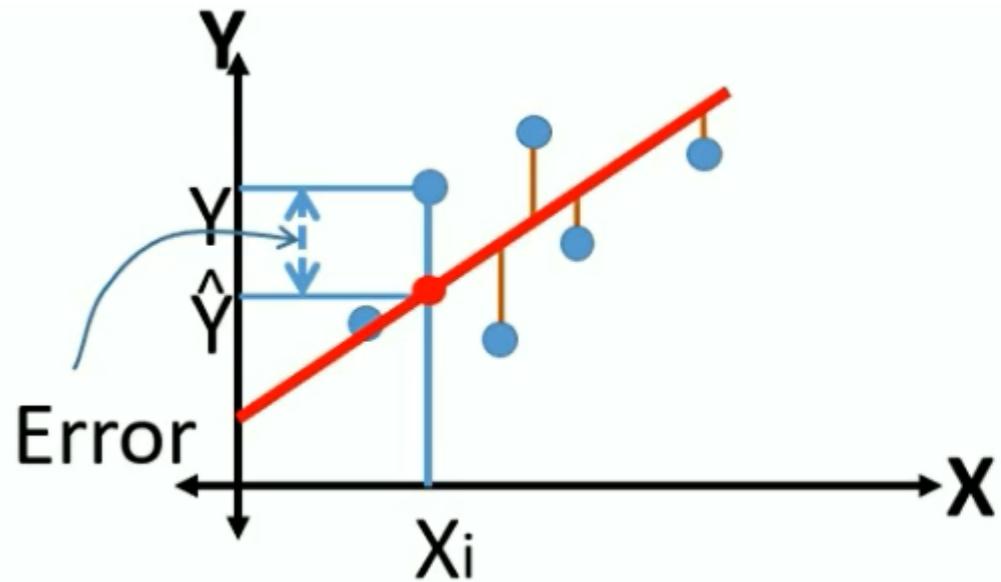
Mean absolute error (MAE) is a quantity used to measure how close forecasts or predictions are to the eventual outcomes.

# Root Mean Square Error

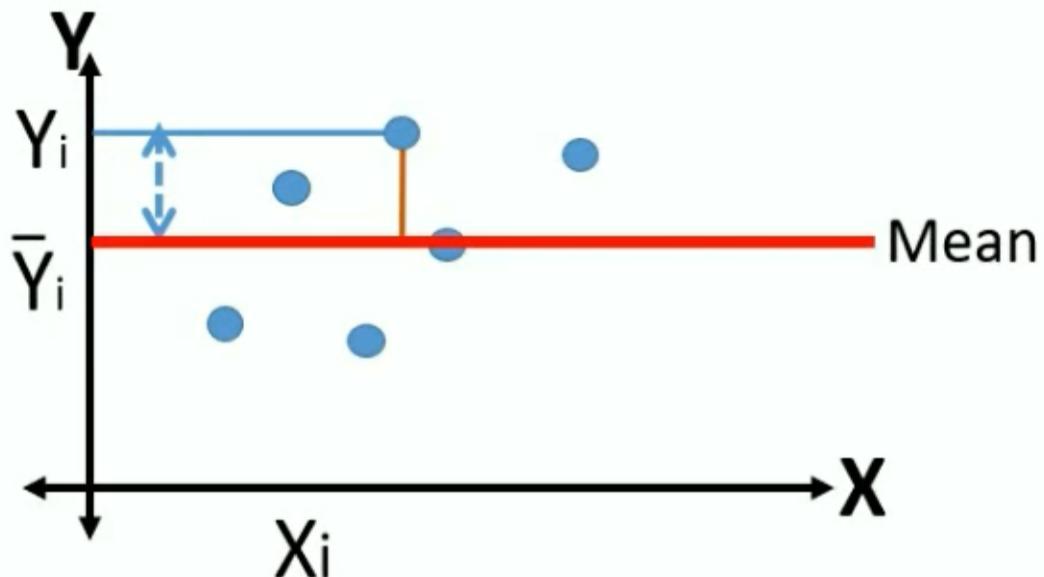
$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

- Very commonly used and makes for an excellent general purpose error metric for numerical predictions.
- Compared to the similar Mean Absolute Error, RMSE amplifies and severely punishes large errors.

# Relative Absolute Error

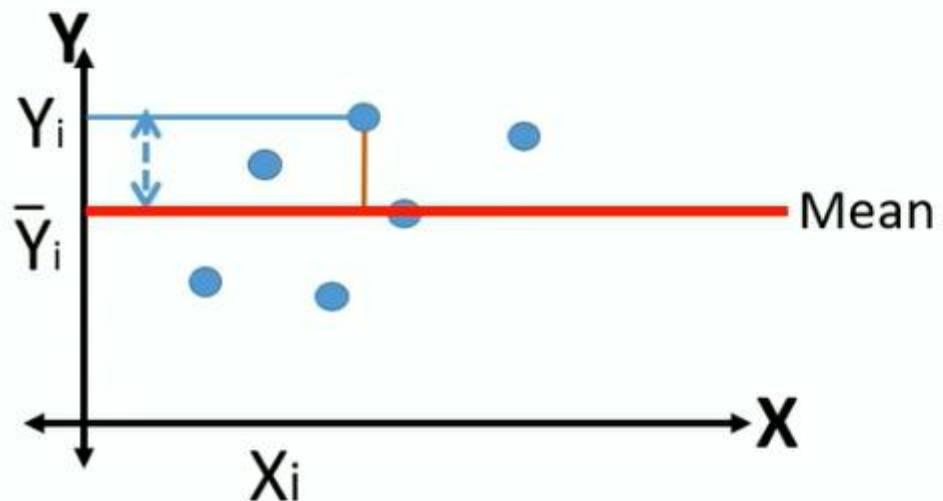
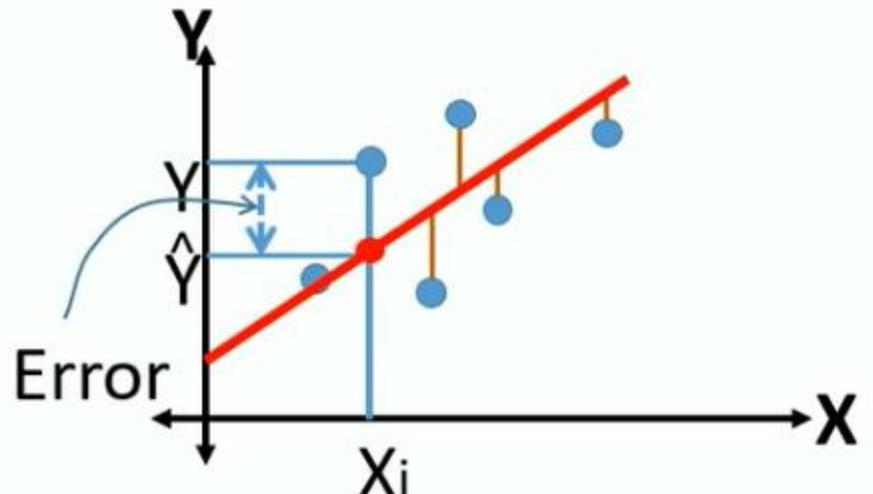


$$\sum_{i=1}^n |y_i - \hat{y}_i|$$



$$\sum_{i=1}^n |y_i - \bar{y}_i|$$

# Relative Absolute Error



$$\text{RAE} = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{\sum_{i=1}^n |y_i - \bar{y}_i|}$$

# Business Problem

Build a model to predict the price of the vehicle based on the available historic data

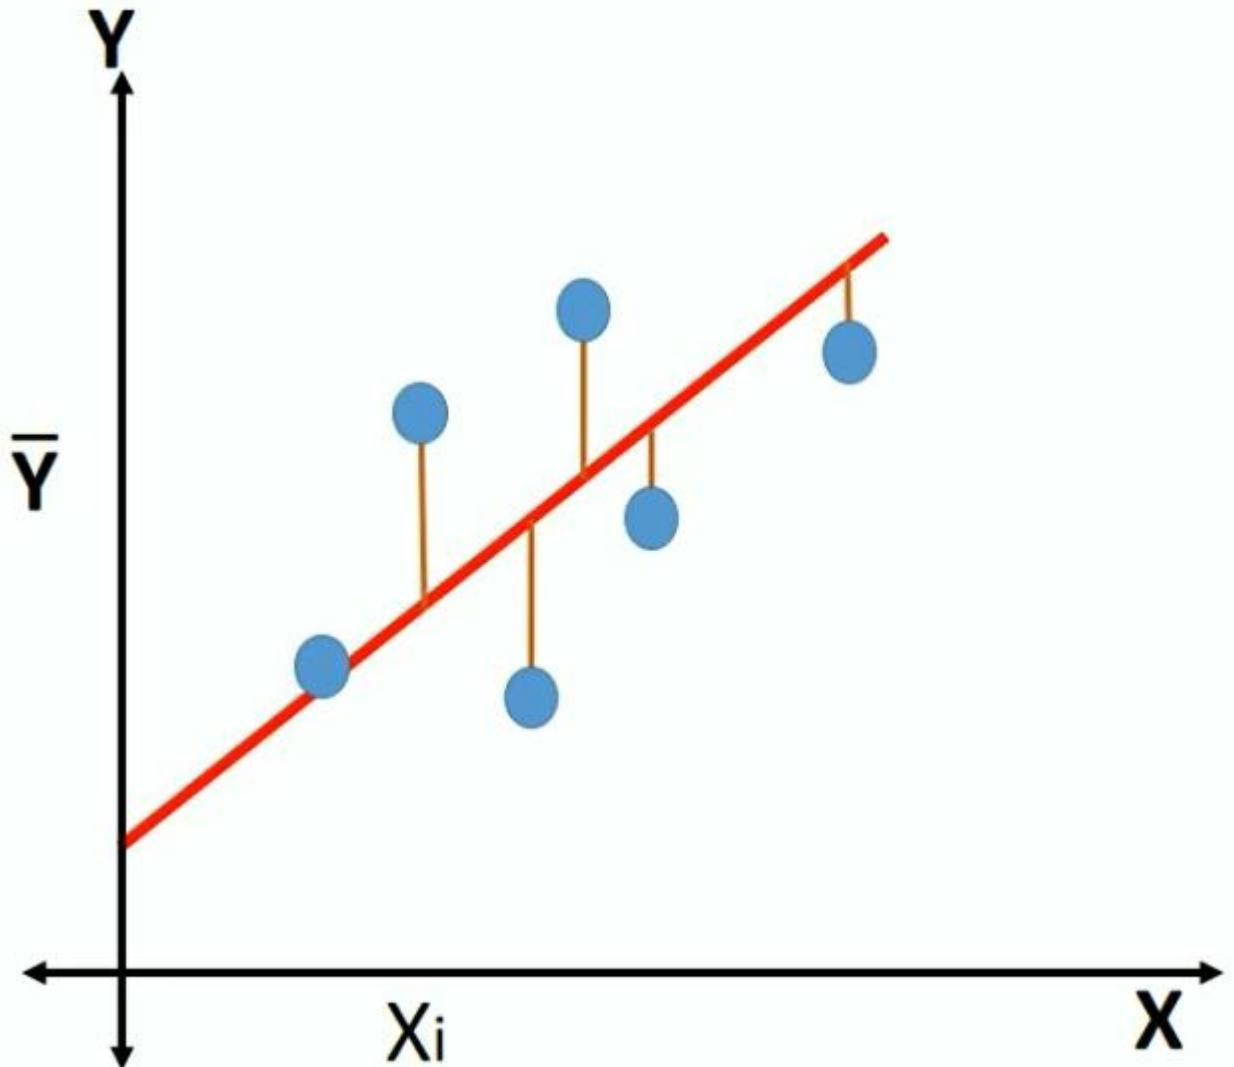
# DEMOS

## Linear Regression for Automobile Data

Coefficient of Determination  
or R Squared

# Coefficient of Determination

How much (what % ) of variation in Y is described by the variation in X?



# R-Square with an Example

Hrs Studied (X)	Marks (Y)
0	40
2	52
3	53
4	55
4	56
5	72
6	71
6	88
7	56
7	74
8	89
9	67
9	89
5.38	66.31
Mean	

$$Y = 41.8 + 4.55X$$

Predicted Marks $\hat{Y}$
41.80
50.90
55.45
60.00
60.00
64.55
69.10
69.10
73.65
73.65
78.20
82.75
82.75

$(Y - \bar{Y})^2$	$(\hat{Y} - \bar{Y})^2$
692.22	600.74
204.78	237.47
177.16	117.94
127.92	39.82
106.30	39.82
32.38	3.10
22.00	7.78
470.46	7.78
106.30	53.88
59.14	53.88
514.84	141.37
0.48	270.27
514.84	270.27
3028.77	1844.12
SST	SSR

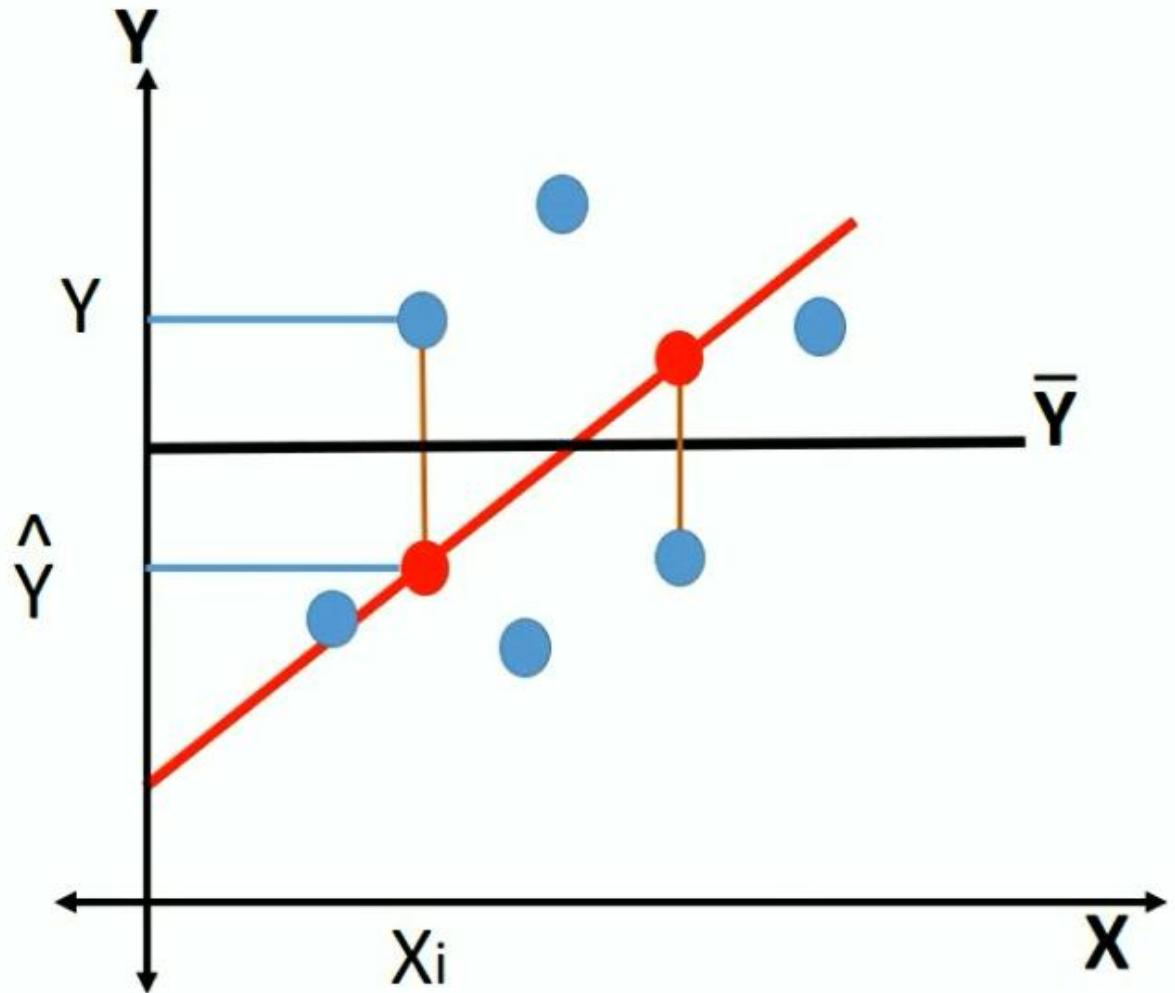
# Coefficient of Determination

Sum of Squares Due to Regression

$$SSR = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$$

Total Sum of Squares

$$SST = \sum_{i=1}^n (y_i - \bar{y})^2$$



# Coefficient of Determination

$$\begin{aligned} R^2 &= \text{SSR/SST} = 1844.12/3028.77 \\ &= 0.60886 \end{aligned}$$

Higher the value → Variation in Y is explained by variation in X.

# Gradient Descent

# Hypothesis

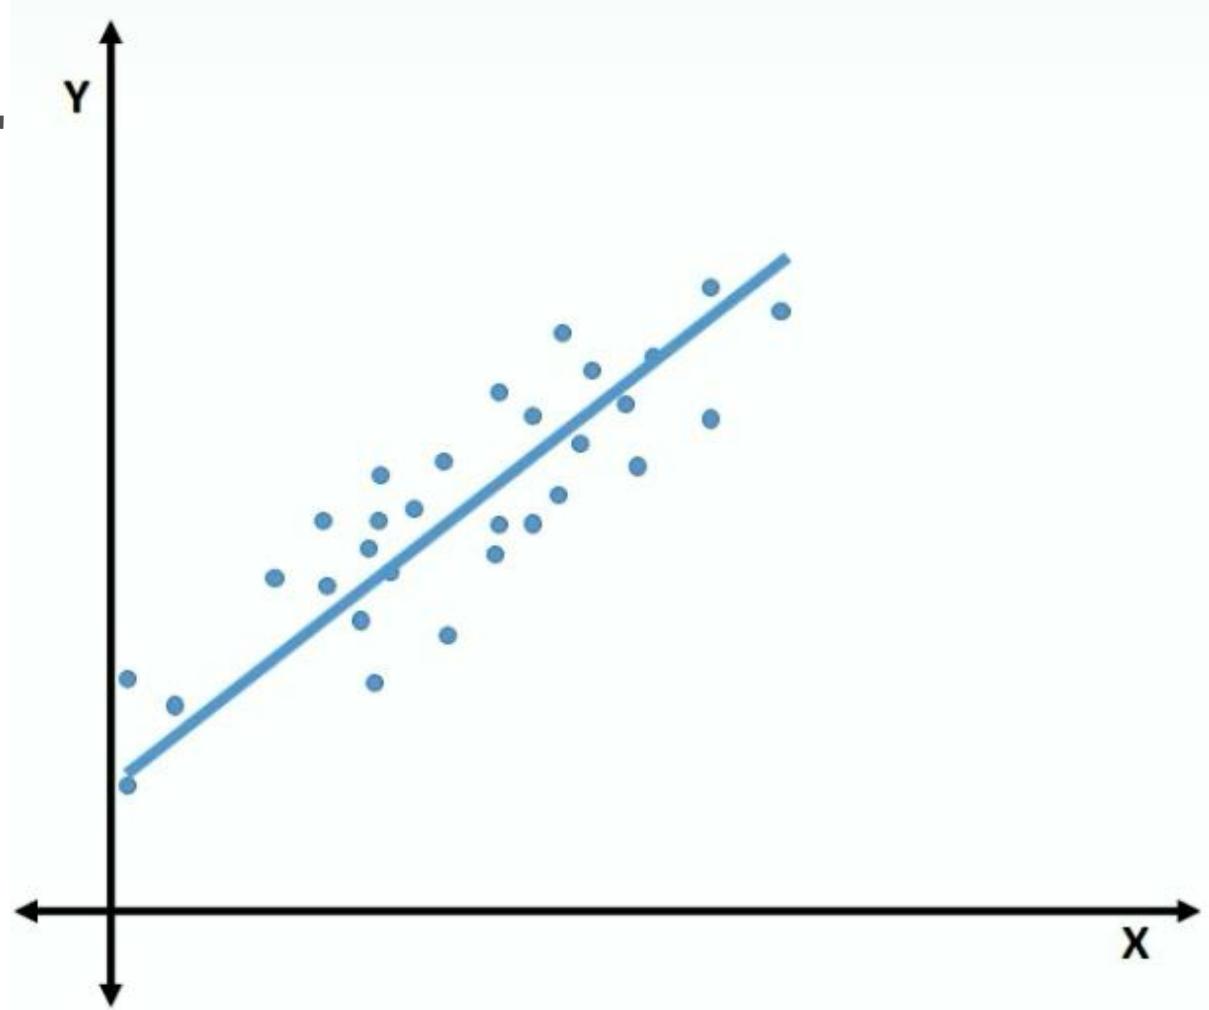
*"Proposed explanation made on the basis of limited evidence as a starting point for further investigation"*

$$h(x) = b_0 + b_1 x$$

Find out value of  $b_0$  and  $b_1$  such that

$$Y \sim N(h(x))$$

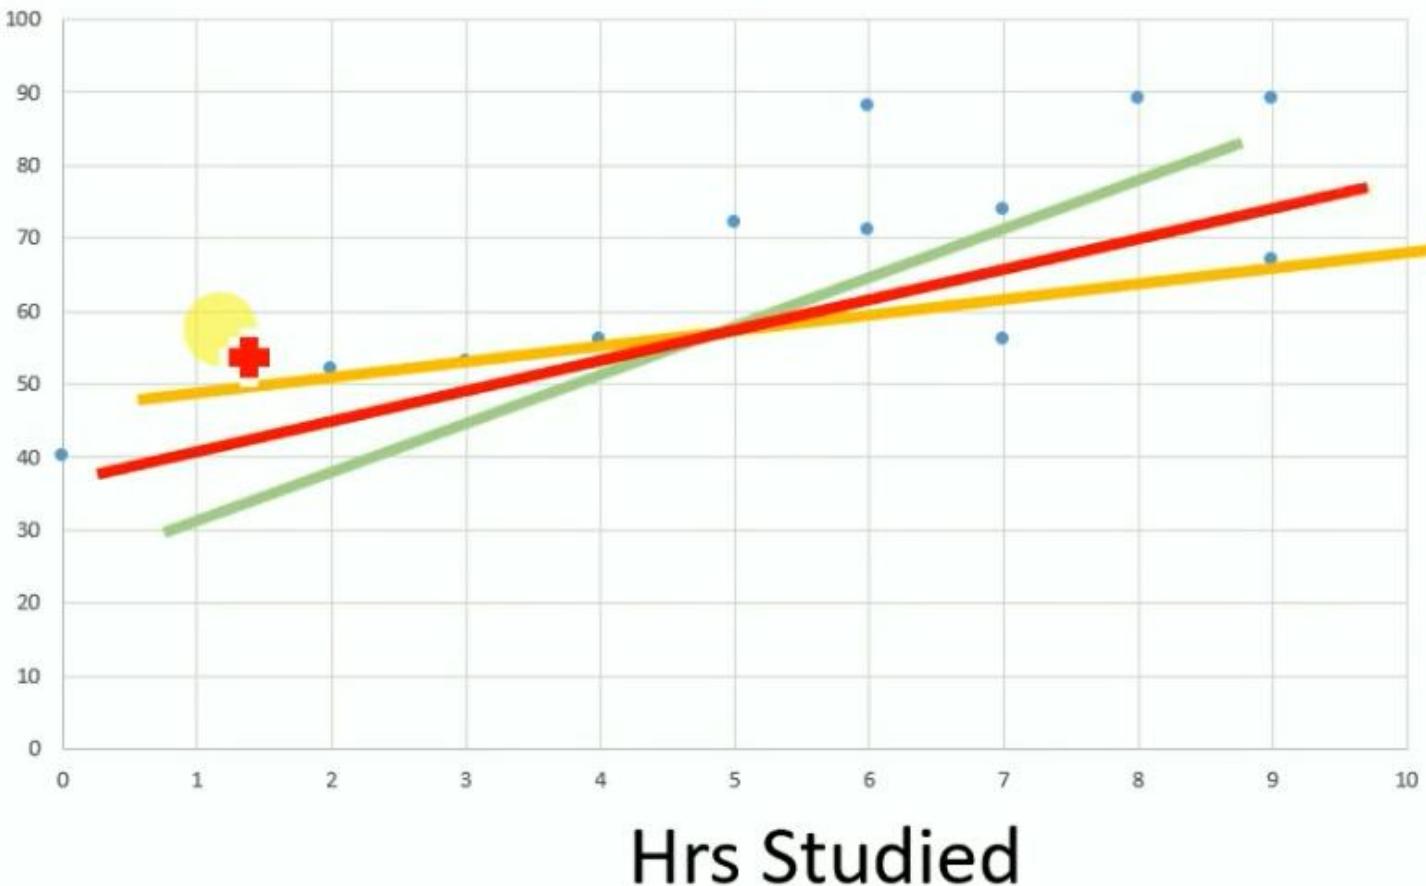
for the given observations



# Example of Linear Regression

Hrs Studied	Marks
0	40
2	52
3	53
4	55
4	56
5	72
6	71
6	88
7	56
7	74
8	89
9	67
9	89

Marks



# Cost Function

Hypothesis:  $h(x) = b_0 + b_1x$

Hrs Studied	Marks	$b_0 = 0; b_1 = 1$ Marks Predicted
0	40	0
2	52	2
3	53	3
4	55	4
4	56	4
5	72	5
6	71	6
6	88	6
7	56	7
7	74	7
8	89	8
9	67	9
9	89	9

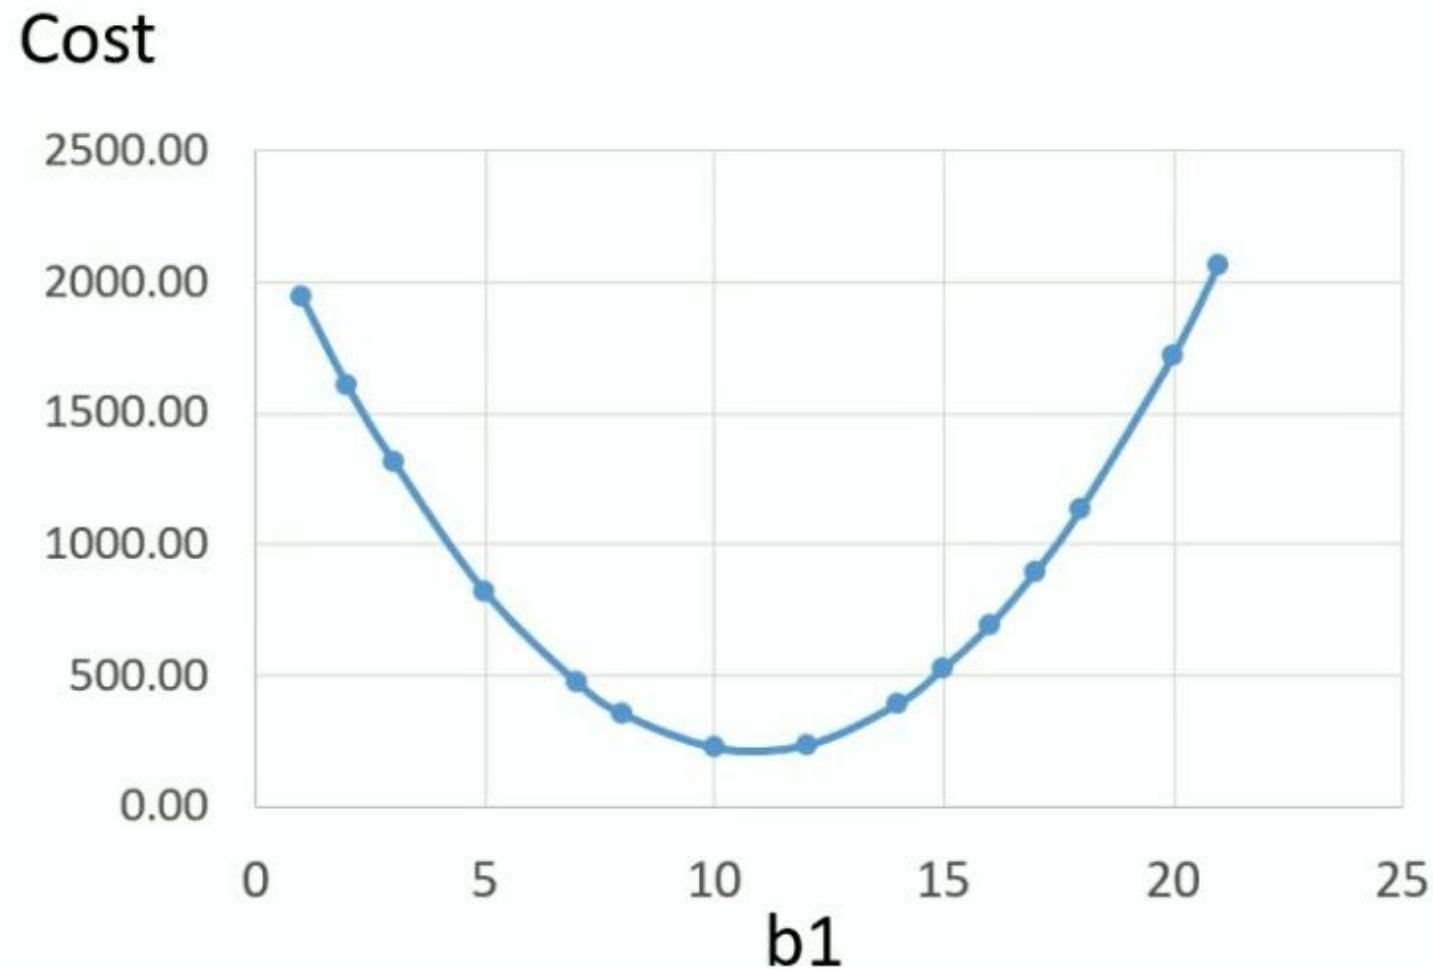
$(Y_i - \hat{Y}_i)^2$
1600
2500
2500
2601
2704
4489
4225
6724
2401
4489
6561
3364
6400

$$\frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

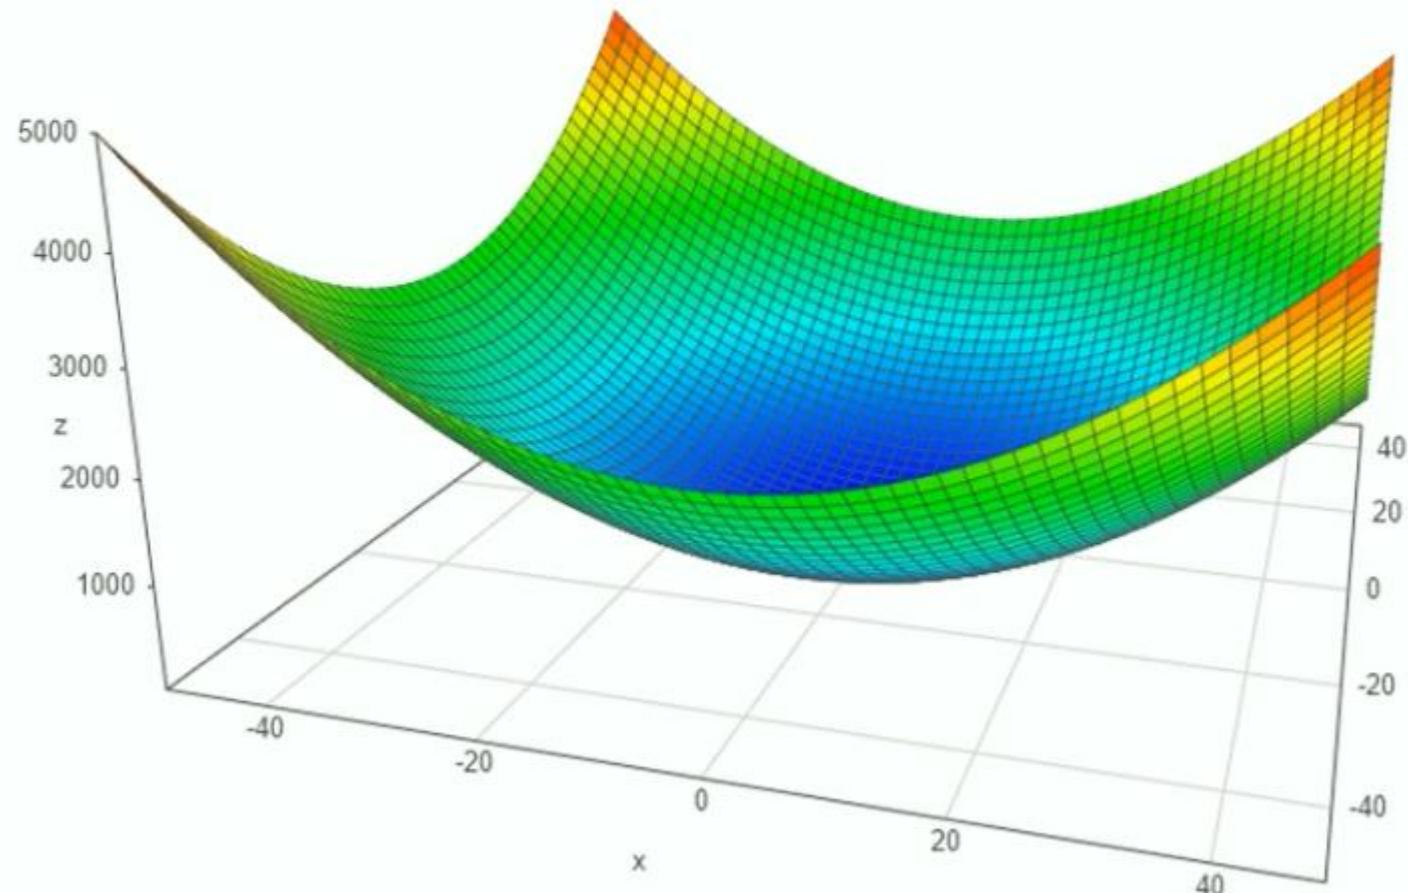
$b_0$	$b_1$	Cost
0	1	1944.538

# Cost Function Plot

b0	b1	cost
0	1	1944.54
0	2	1610.08
0	3	1311.46
0	5	821.77
0	7	475.46
0	8	356.08
0	10	224.85
0	12	237.00
0	14	392.54
0	15	524.08
0	16	691.46
0	17	894.69
0	18	1133.77
0	20	1719.46
0	21	2066.08



# Cost Function with b0 and b1

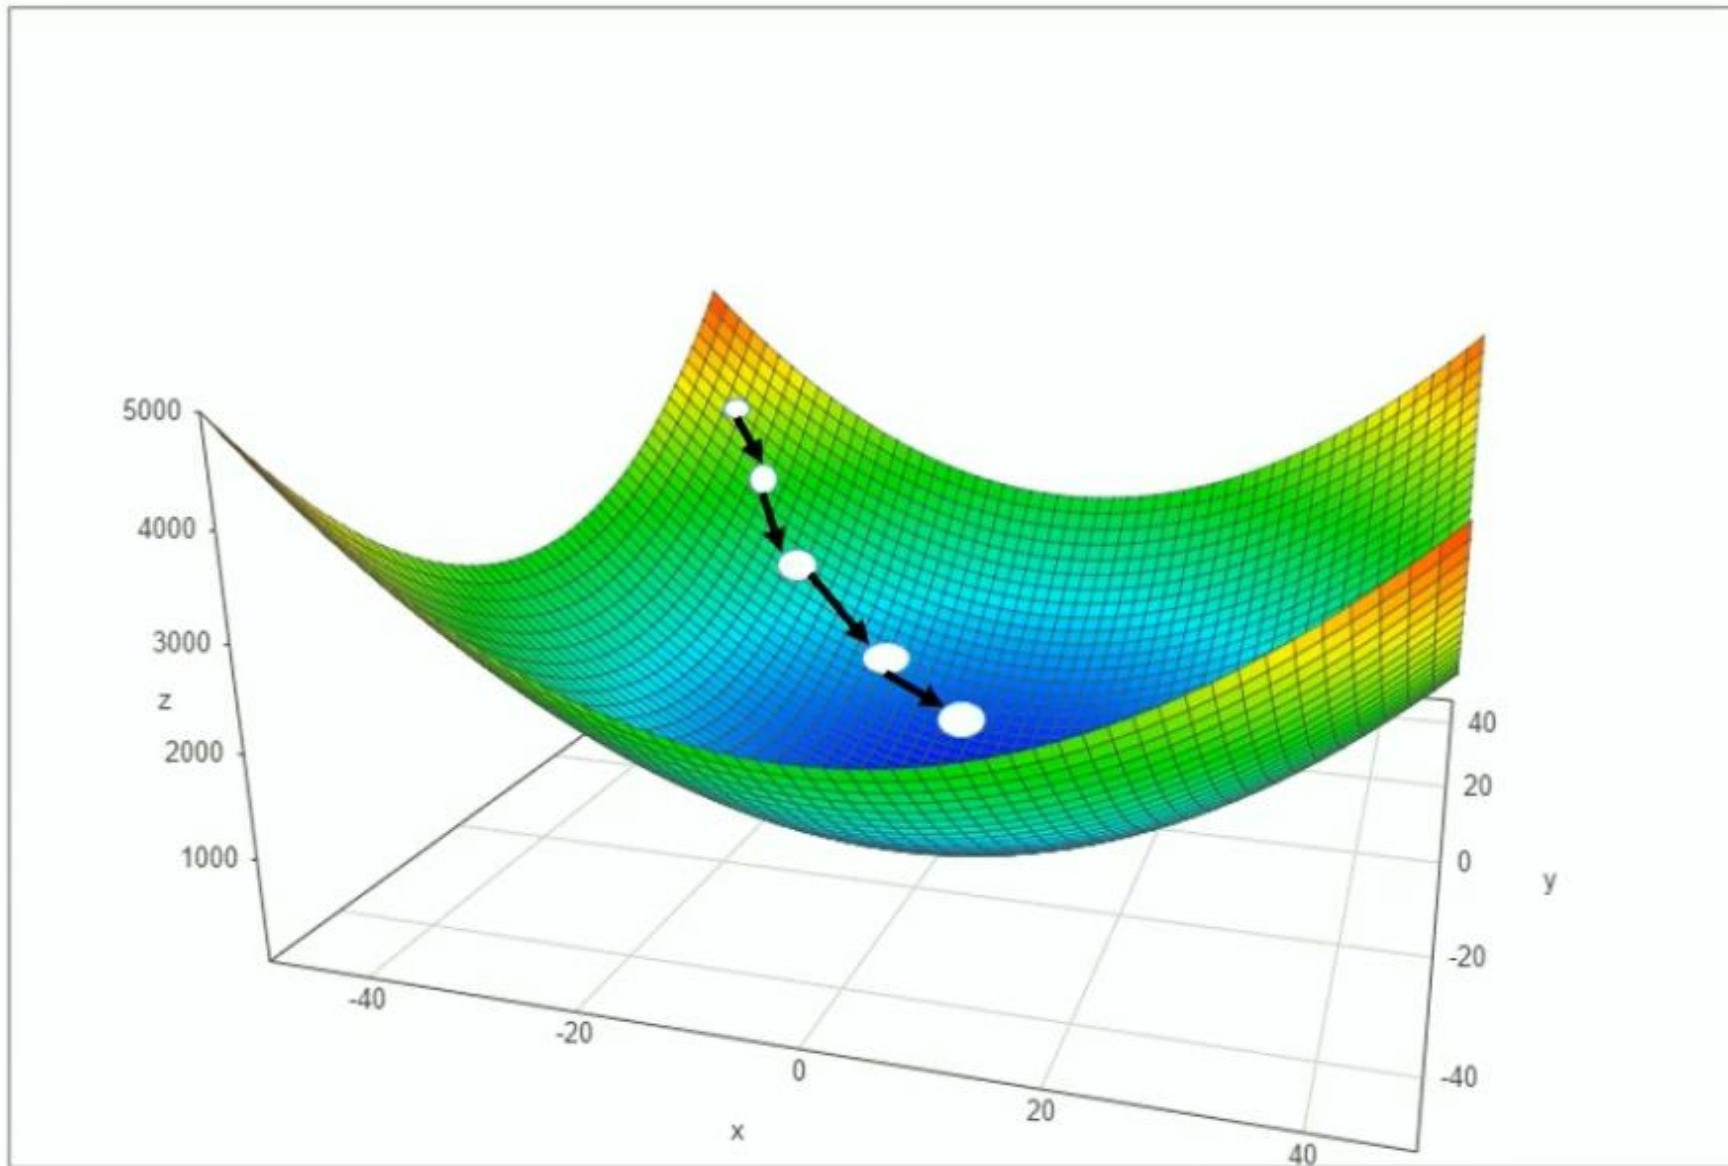


X axis –  $b_1$

Y axis –  $b_0$

$Z = C(b_0, b_1)$

# Gradient Descent

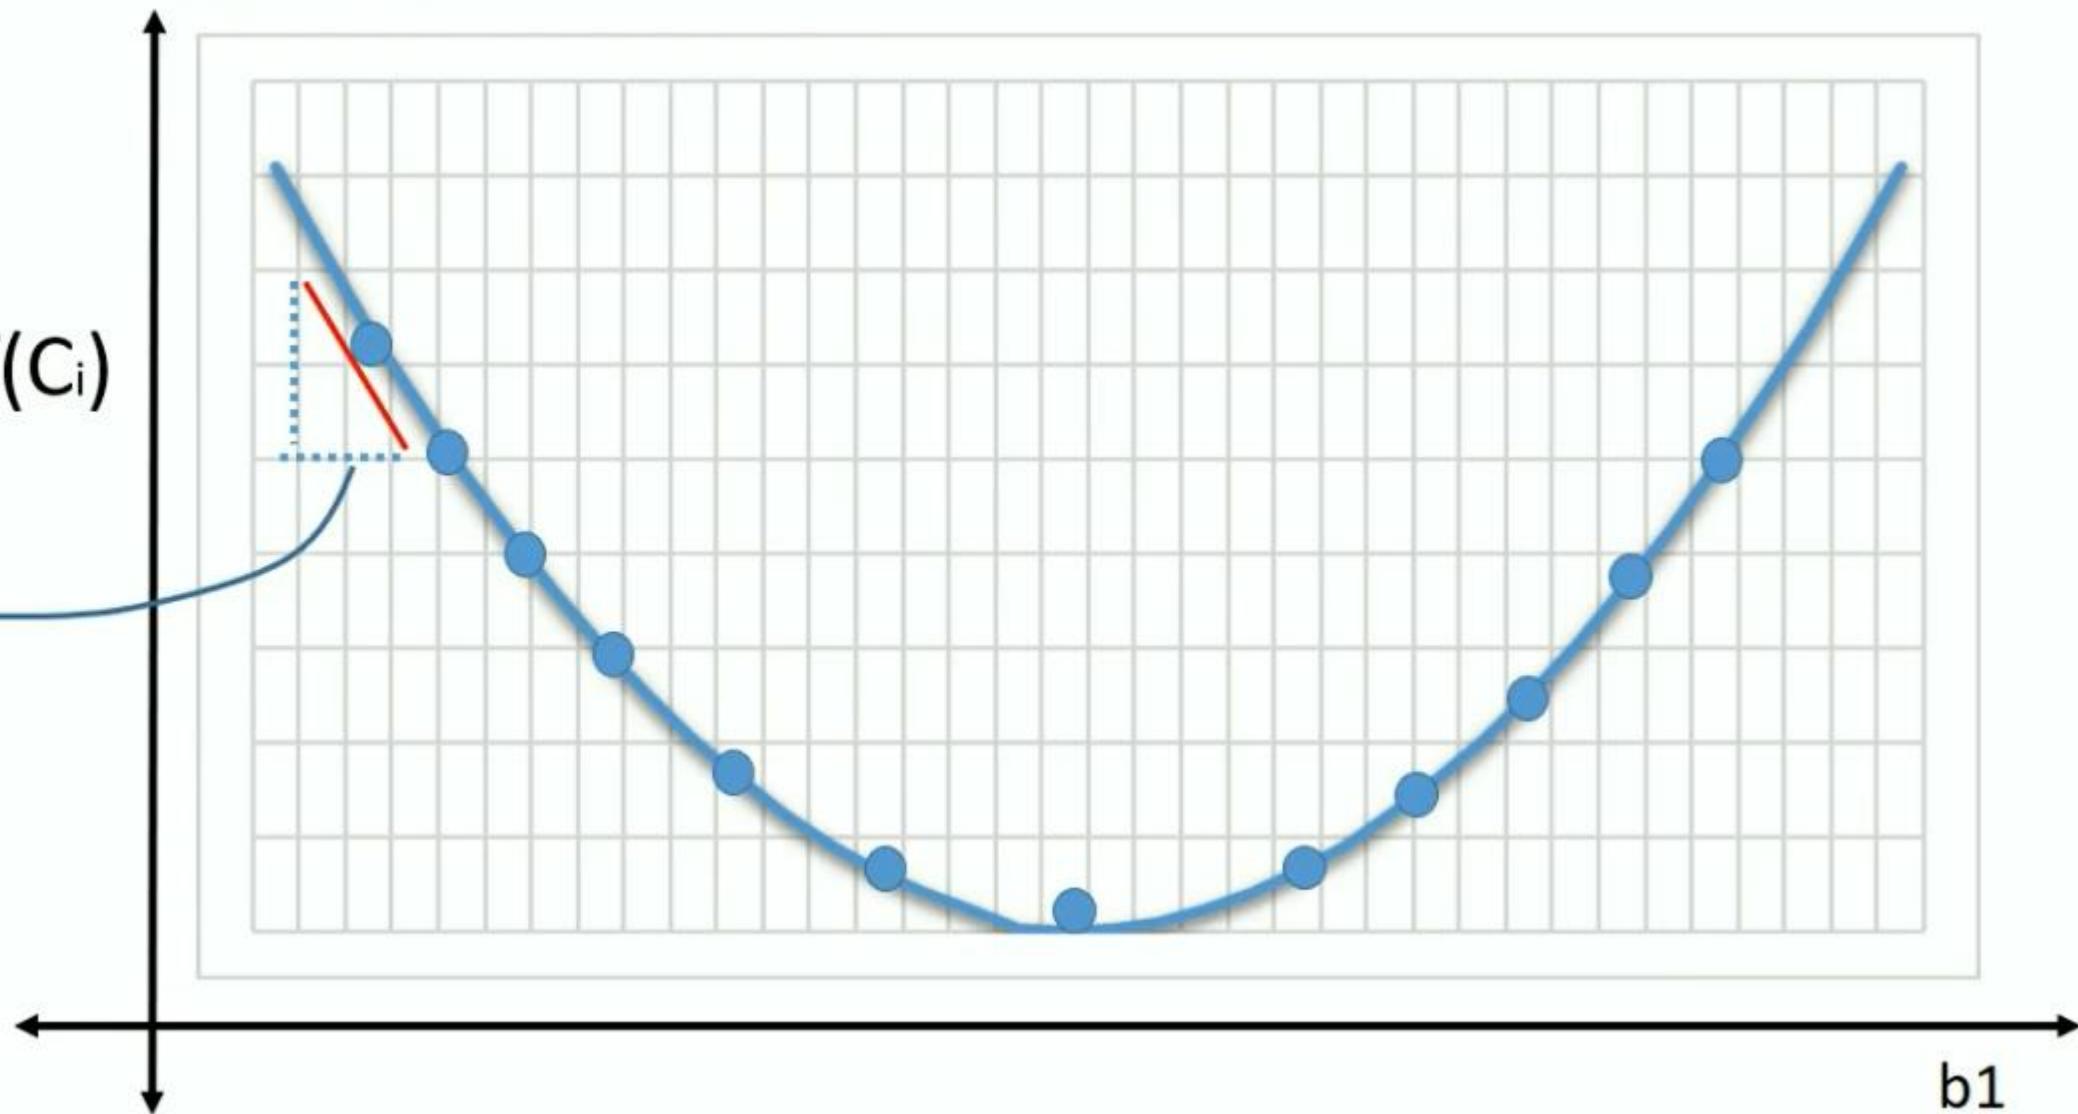


# Gradient Descent

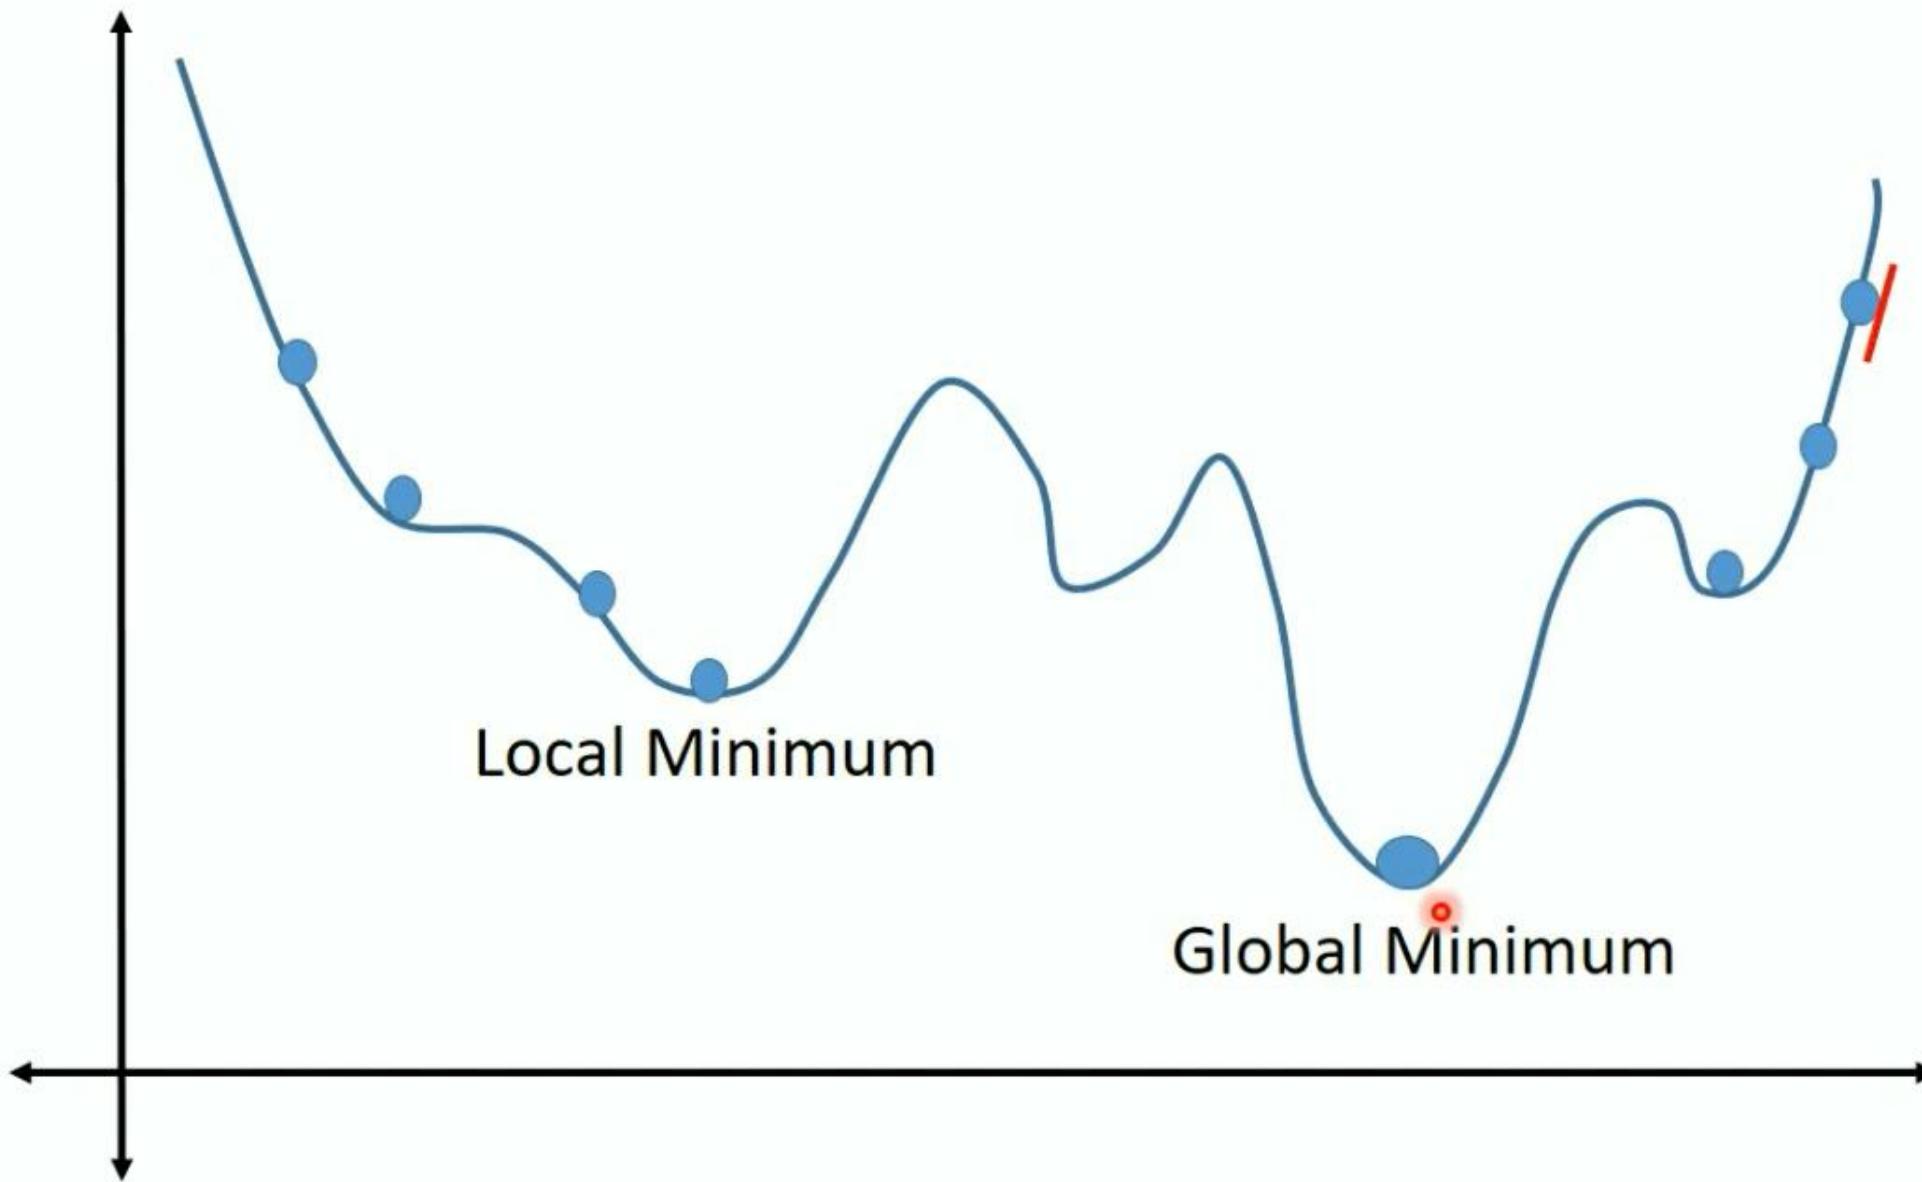
Cost Function:  $C(b_1)$

$$b_j := b_i - \alpha f(C_i)$$

$\alpha$



# Gradient Descent



# Batch Gradient Descent

$$b_j := b_j - \alpha f(C_i)$$

Does it for  
number of examples  
number of features  
learning rate

Sum of All before taking one step (epoch)  
Long time to reach the bottom

## Batch Gradient Descent

# Batch Vs Stochastic Gradient Descent

$$b_j := b_j - \alpha f(C_i)$$

Does it for  
number of examples  
number of features  
learning rate

## Batch Gradient Descent

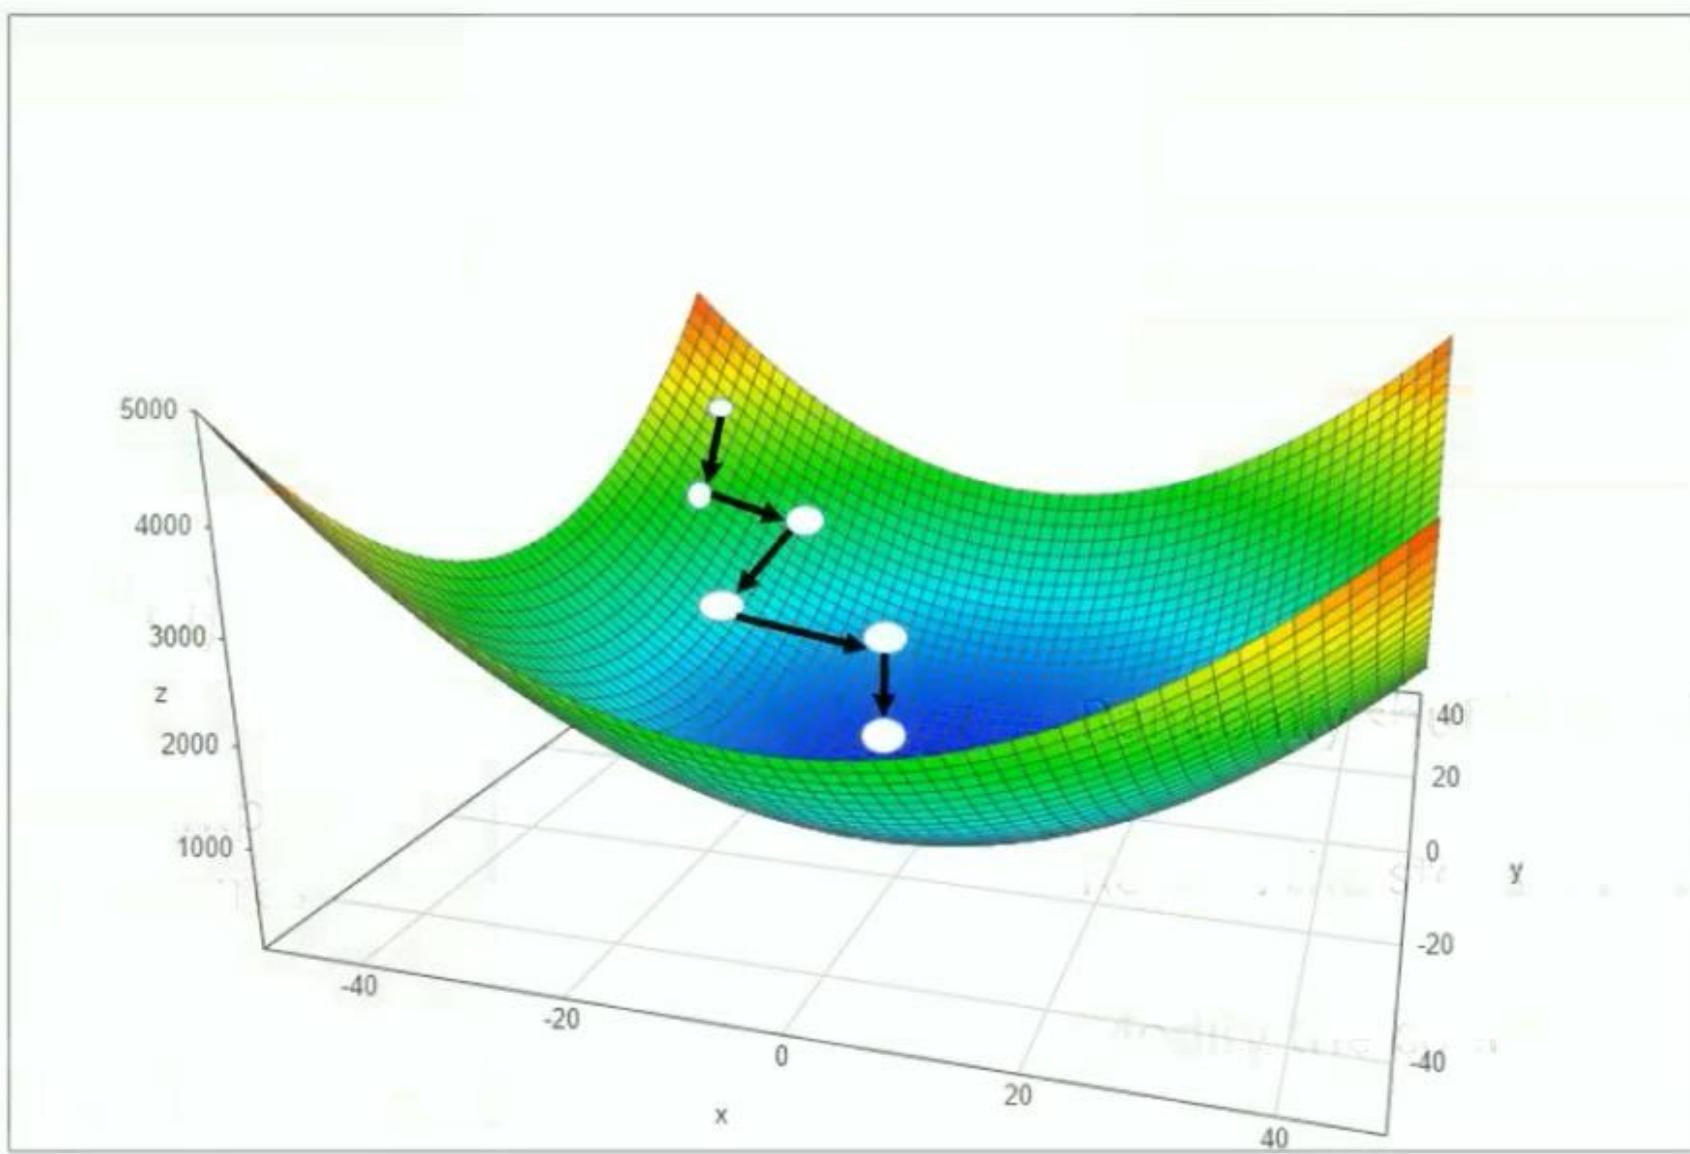
X1	X2	...	Xn

## Randomly shuffle the dataset

Repeat the steps for every example

Modify the coefficient at every step

# Stochastic Gradient Descent

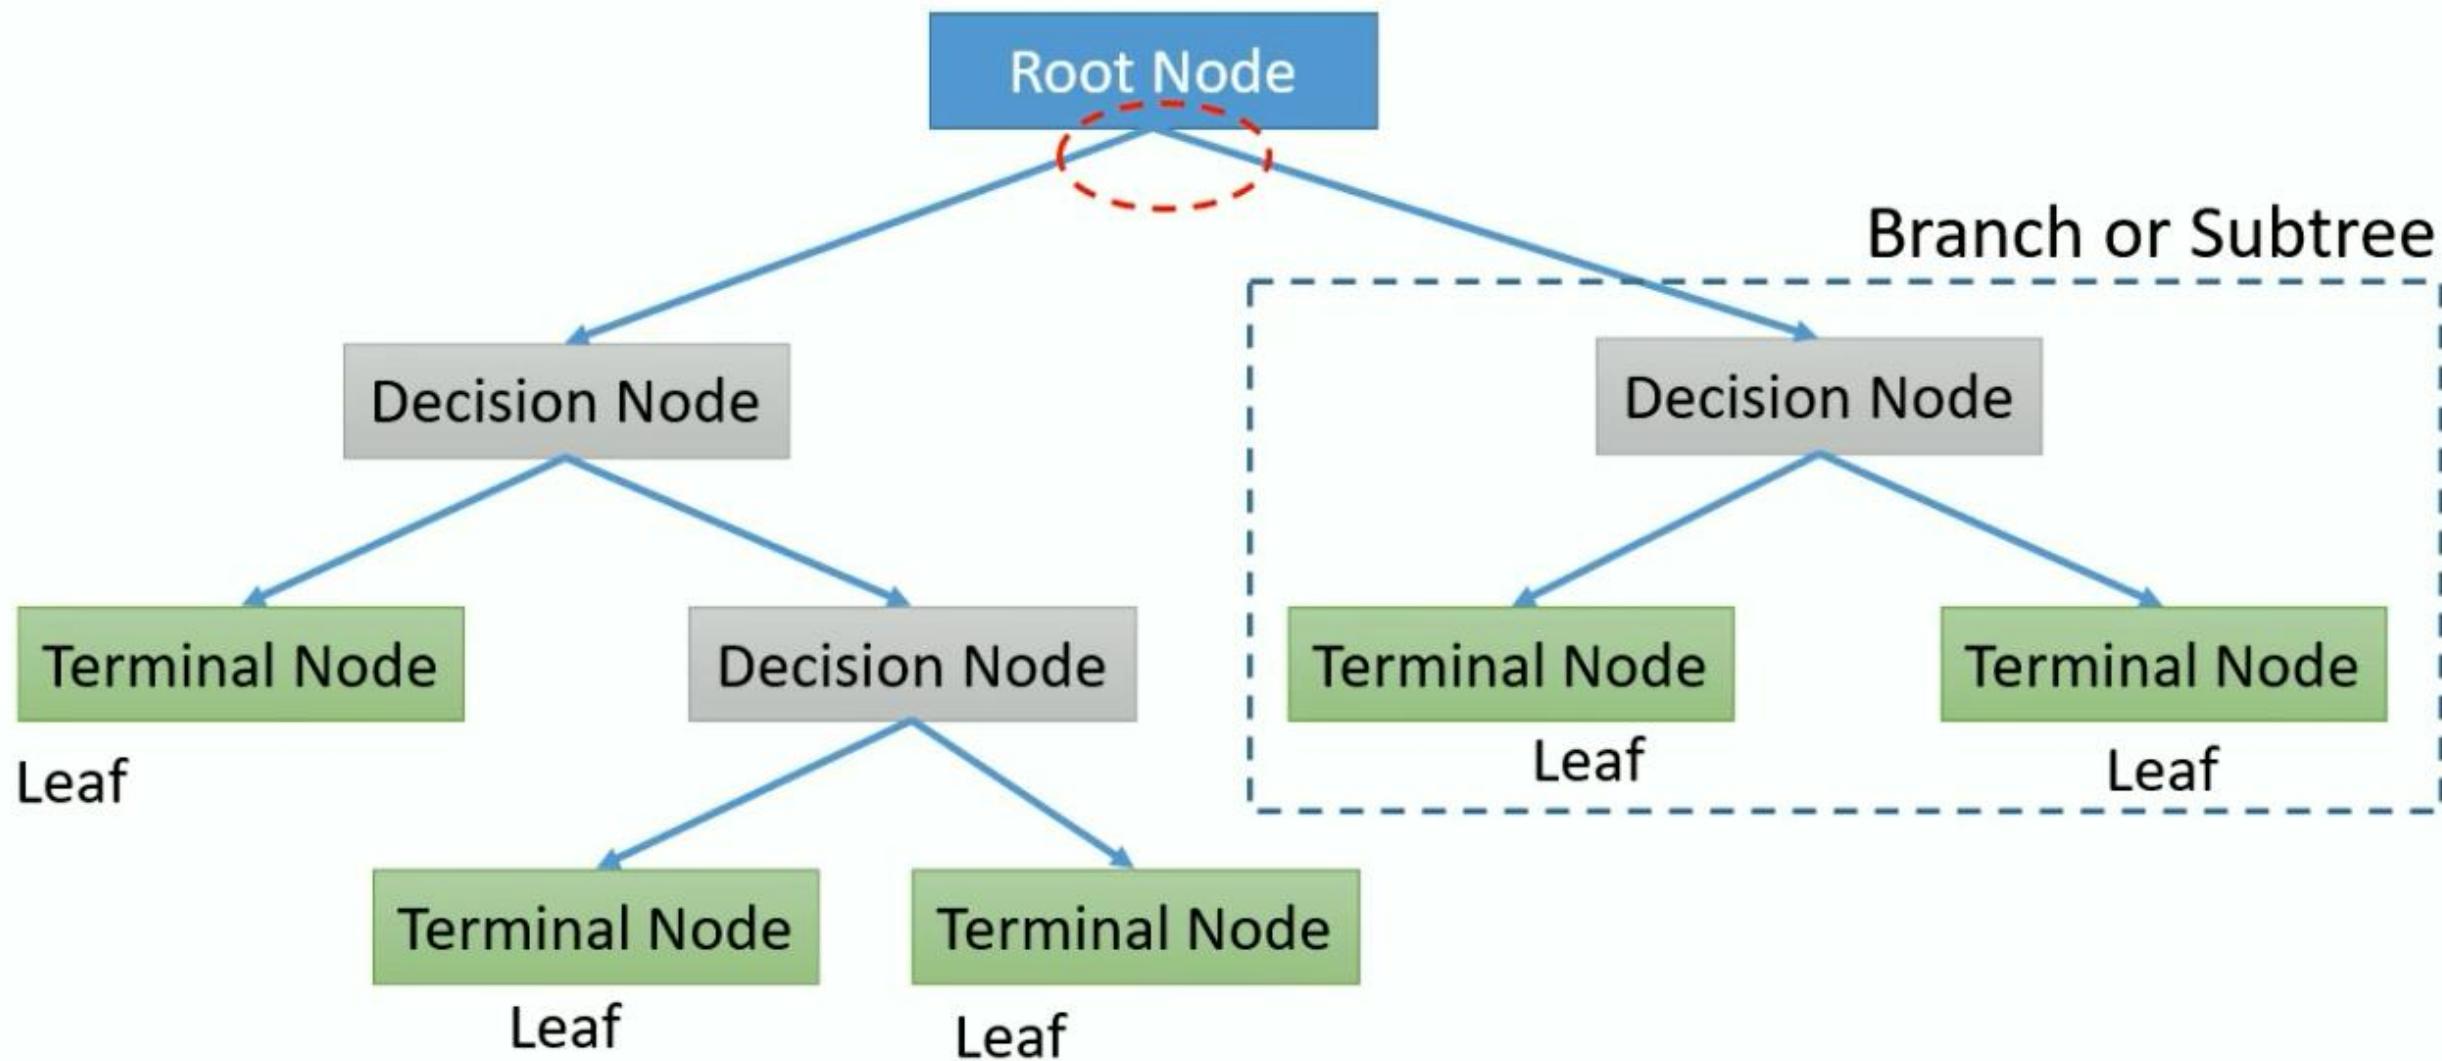


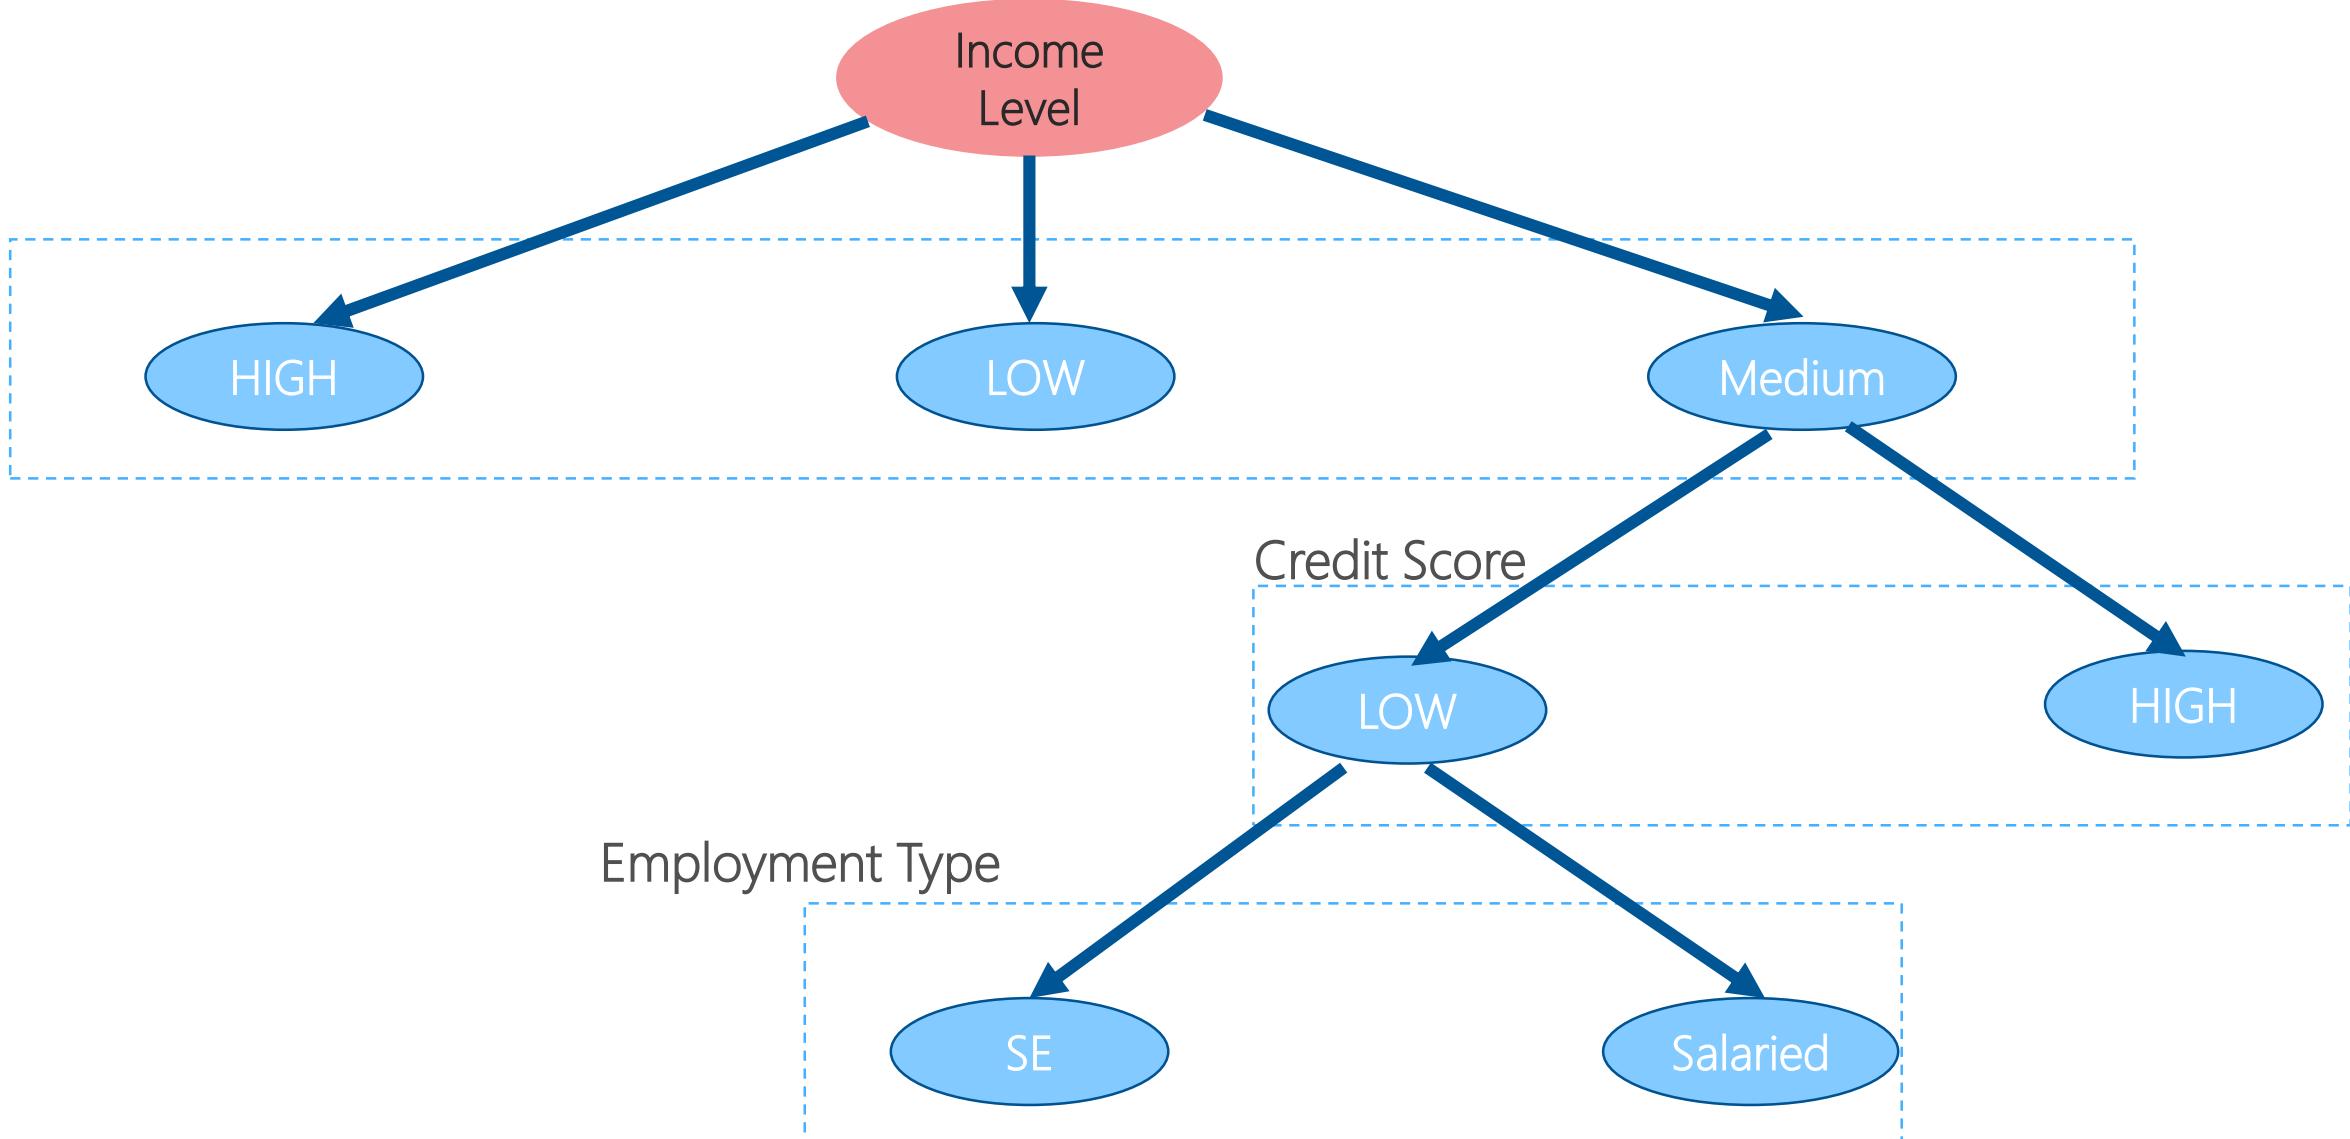
# DEMO

## Online Gradient Descent in Linear Regression

# Regression Using Decision Tree?

# Decision Tree Terms





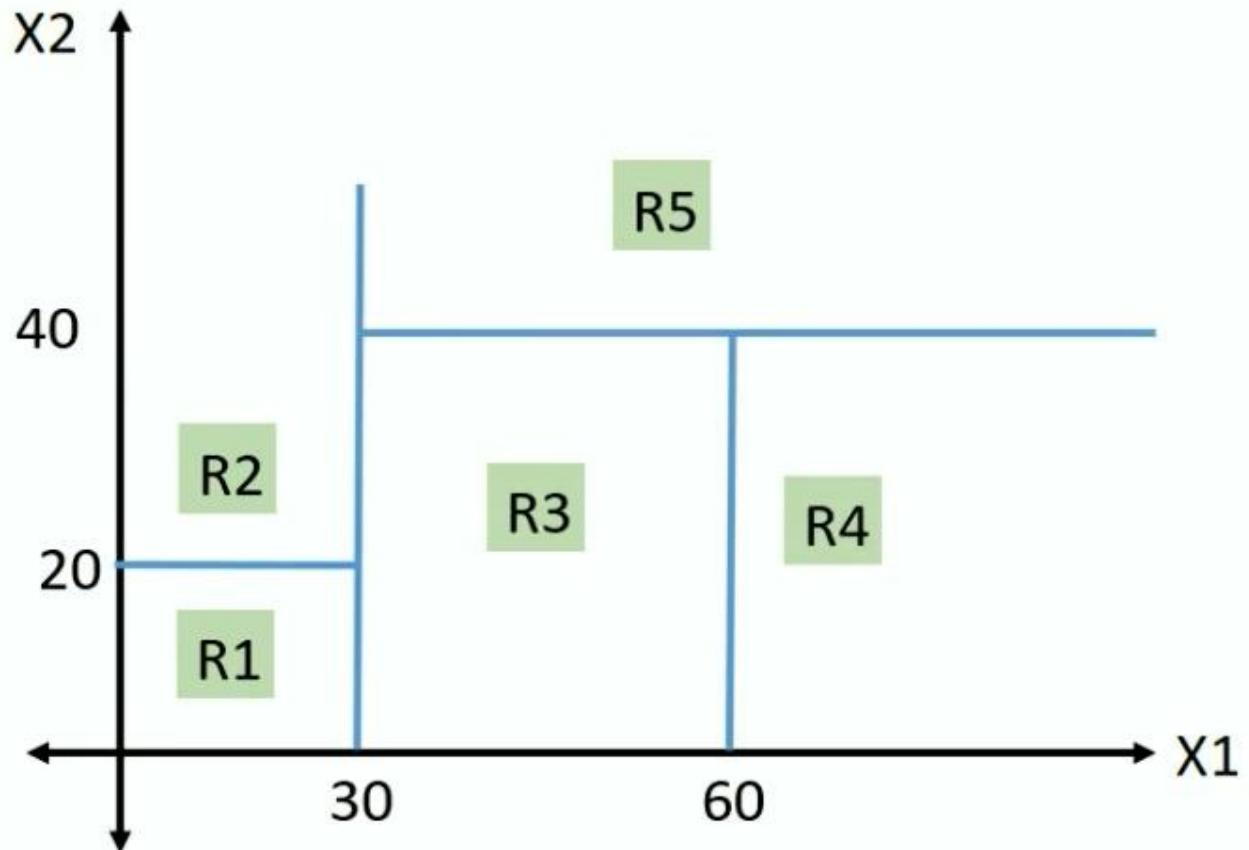
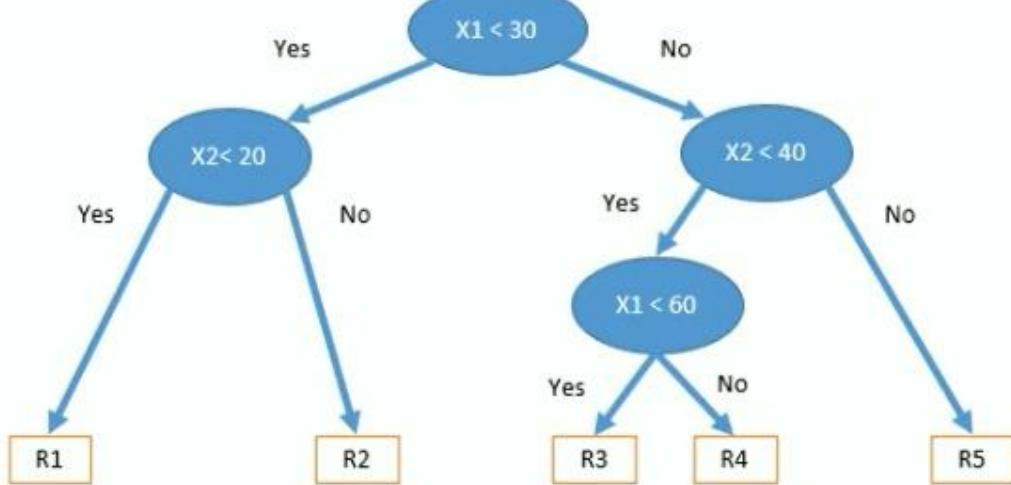
LID	IL	CS	ET	Status
L1	Medium	Low	SE	No
L8	Medium	Low	SE	No

Pure Subset

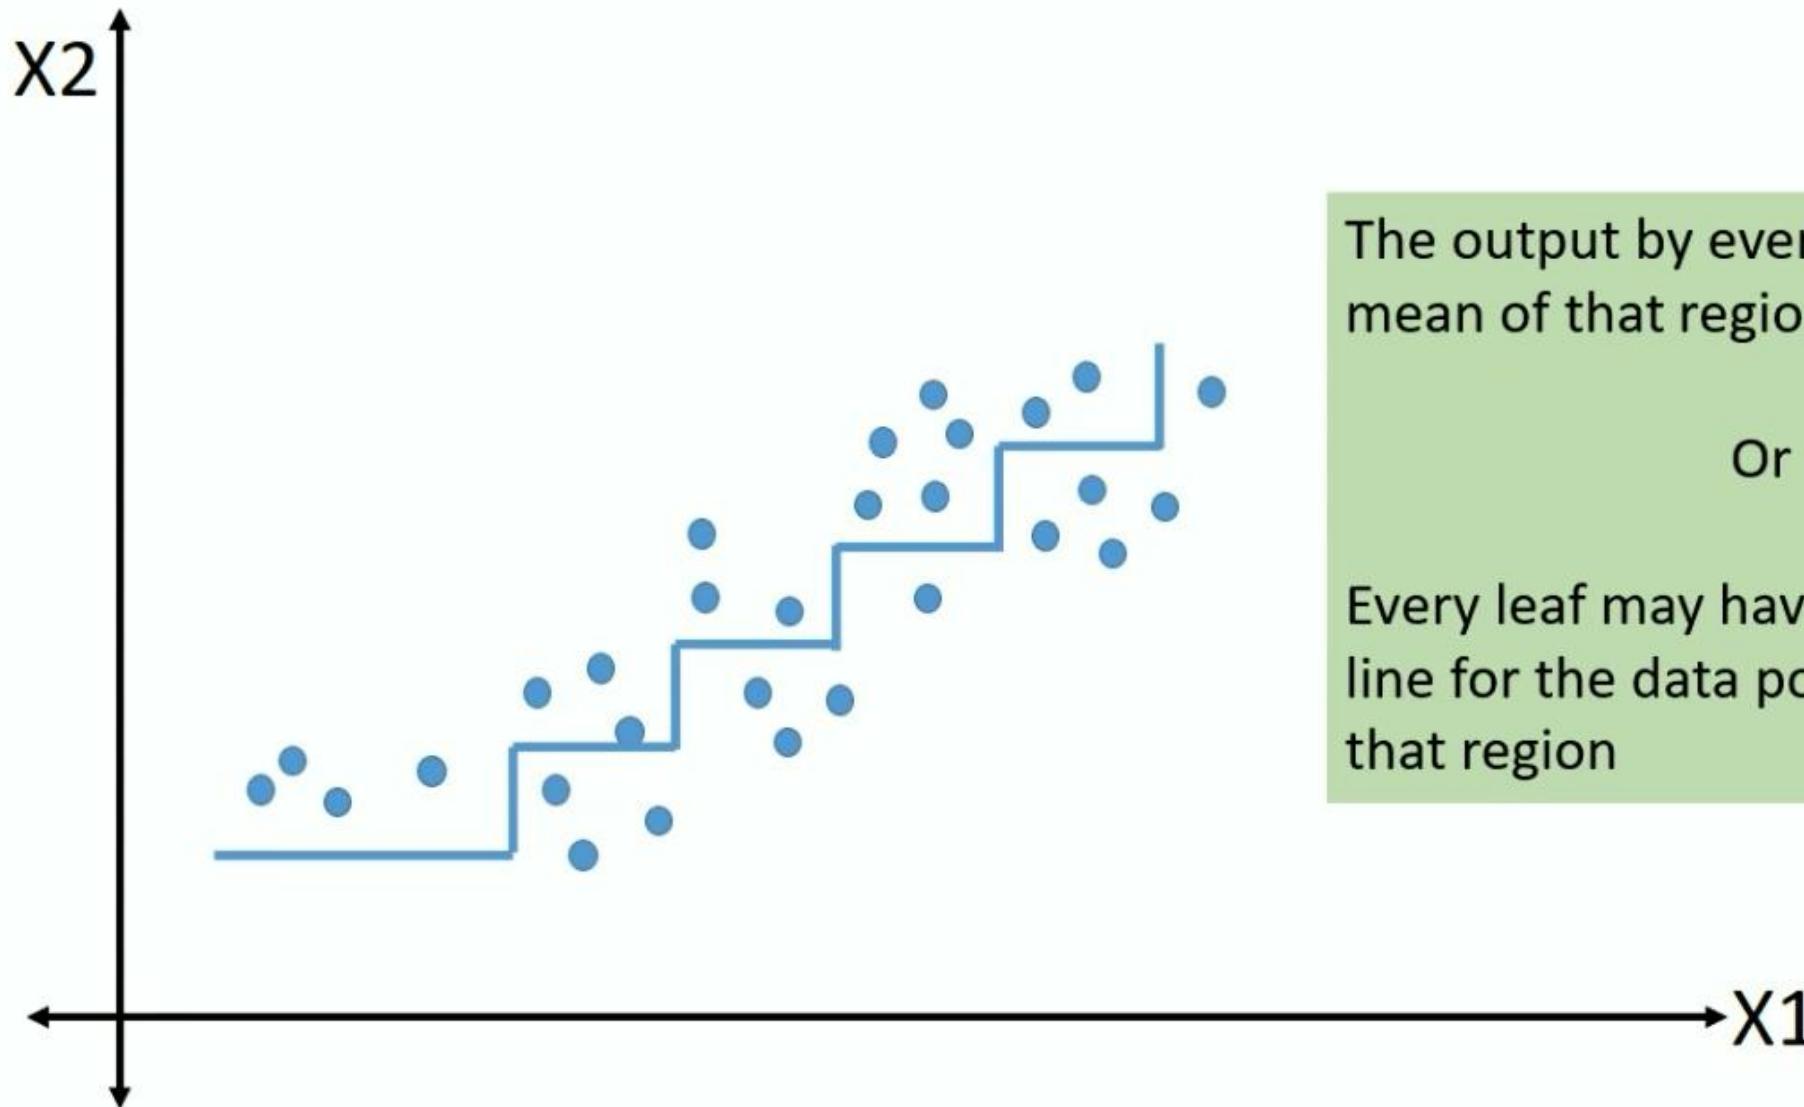
LID	IL	CS	ET	Status
L4	Medium	Low	Salaried	Yes

Pure Subset

# Decision Tree Regression



# Decision Tree Regression



The output by every region is the mean of that region

Or

Every leaf may have a regression line for the data points within that region

# Boosted Decision Tree Regression

- MART gradient boosting algorithm.
- Builds each regression tree in a step-wise fashion
- Predefined loss function to measure the error in each step

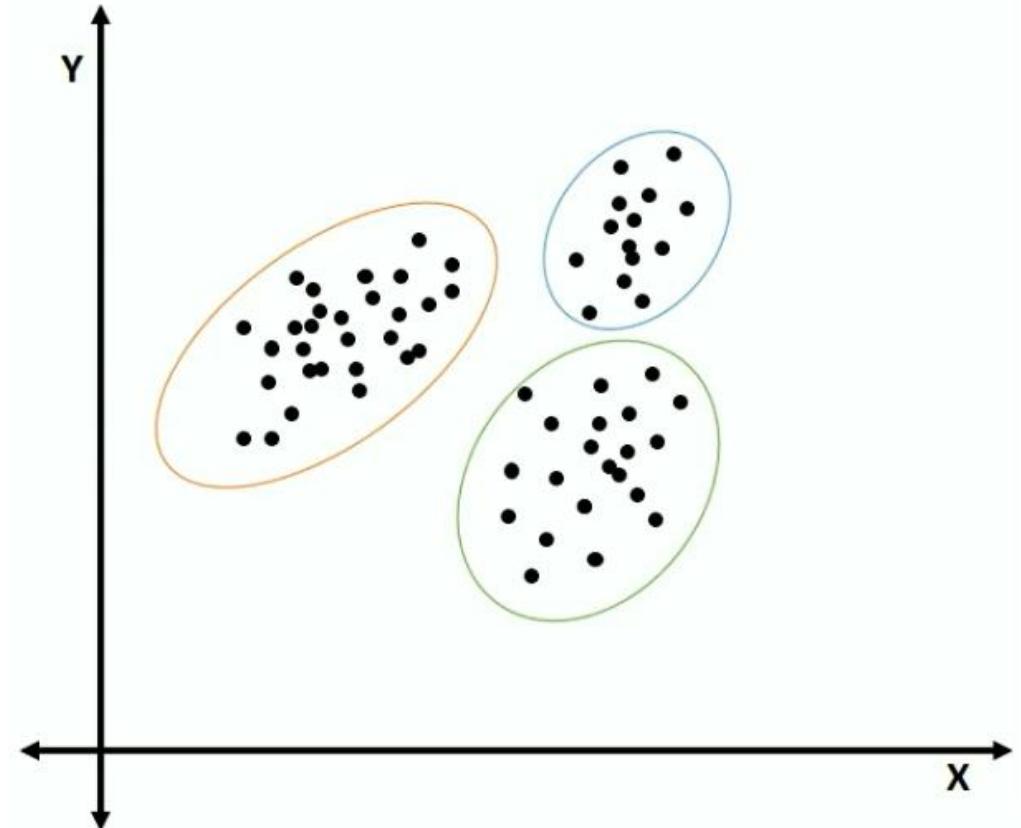
# DEMO

## Boosted Decision Tree Regression

# Cluster Analysis

# Clustering or Cluster Analysis

- Clustering is the task of grouping a set of objects
- Unsupervised Learning model
- Discovering distinct groups in customer databases
- Used for creating strategies to adopt for certain segments

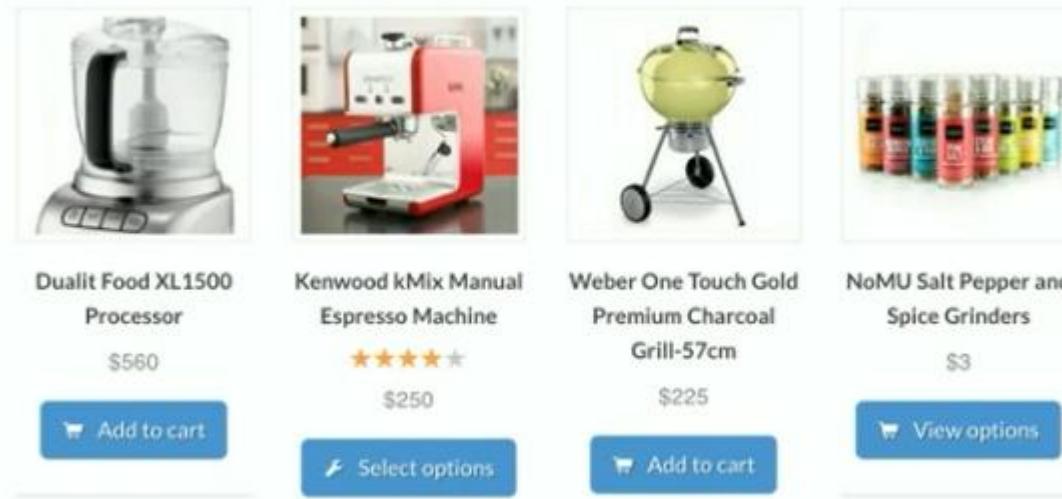


# Examples of Clustering

- Recommendation engines
- Market Segmentation
- Social Network Analysis



Customers who viewed this item also viewed these products

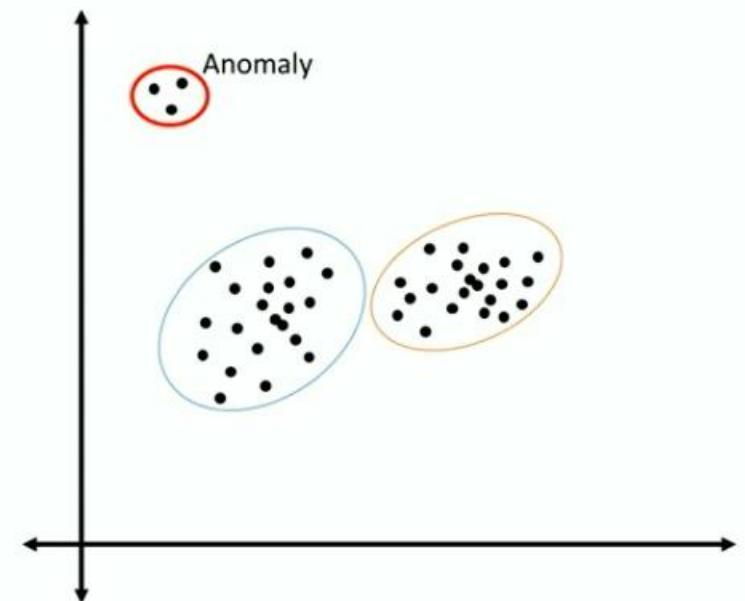


# Examples of Clustering

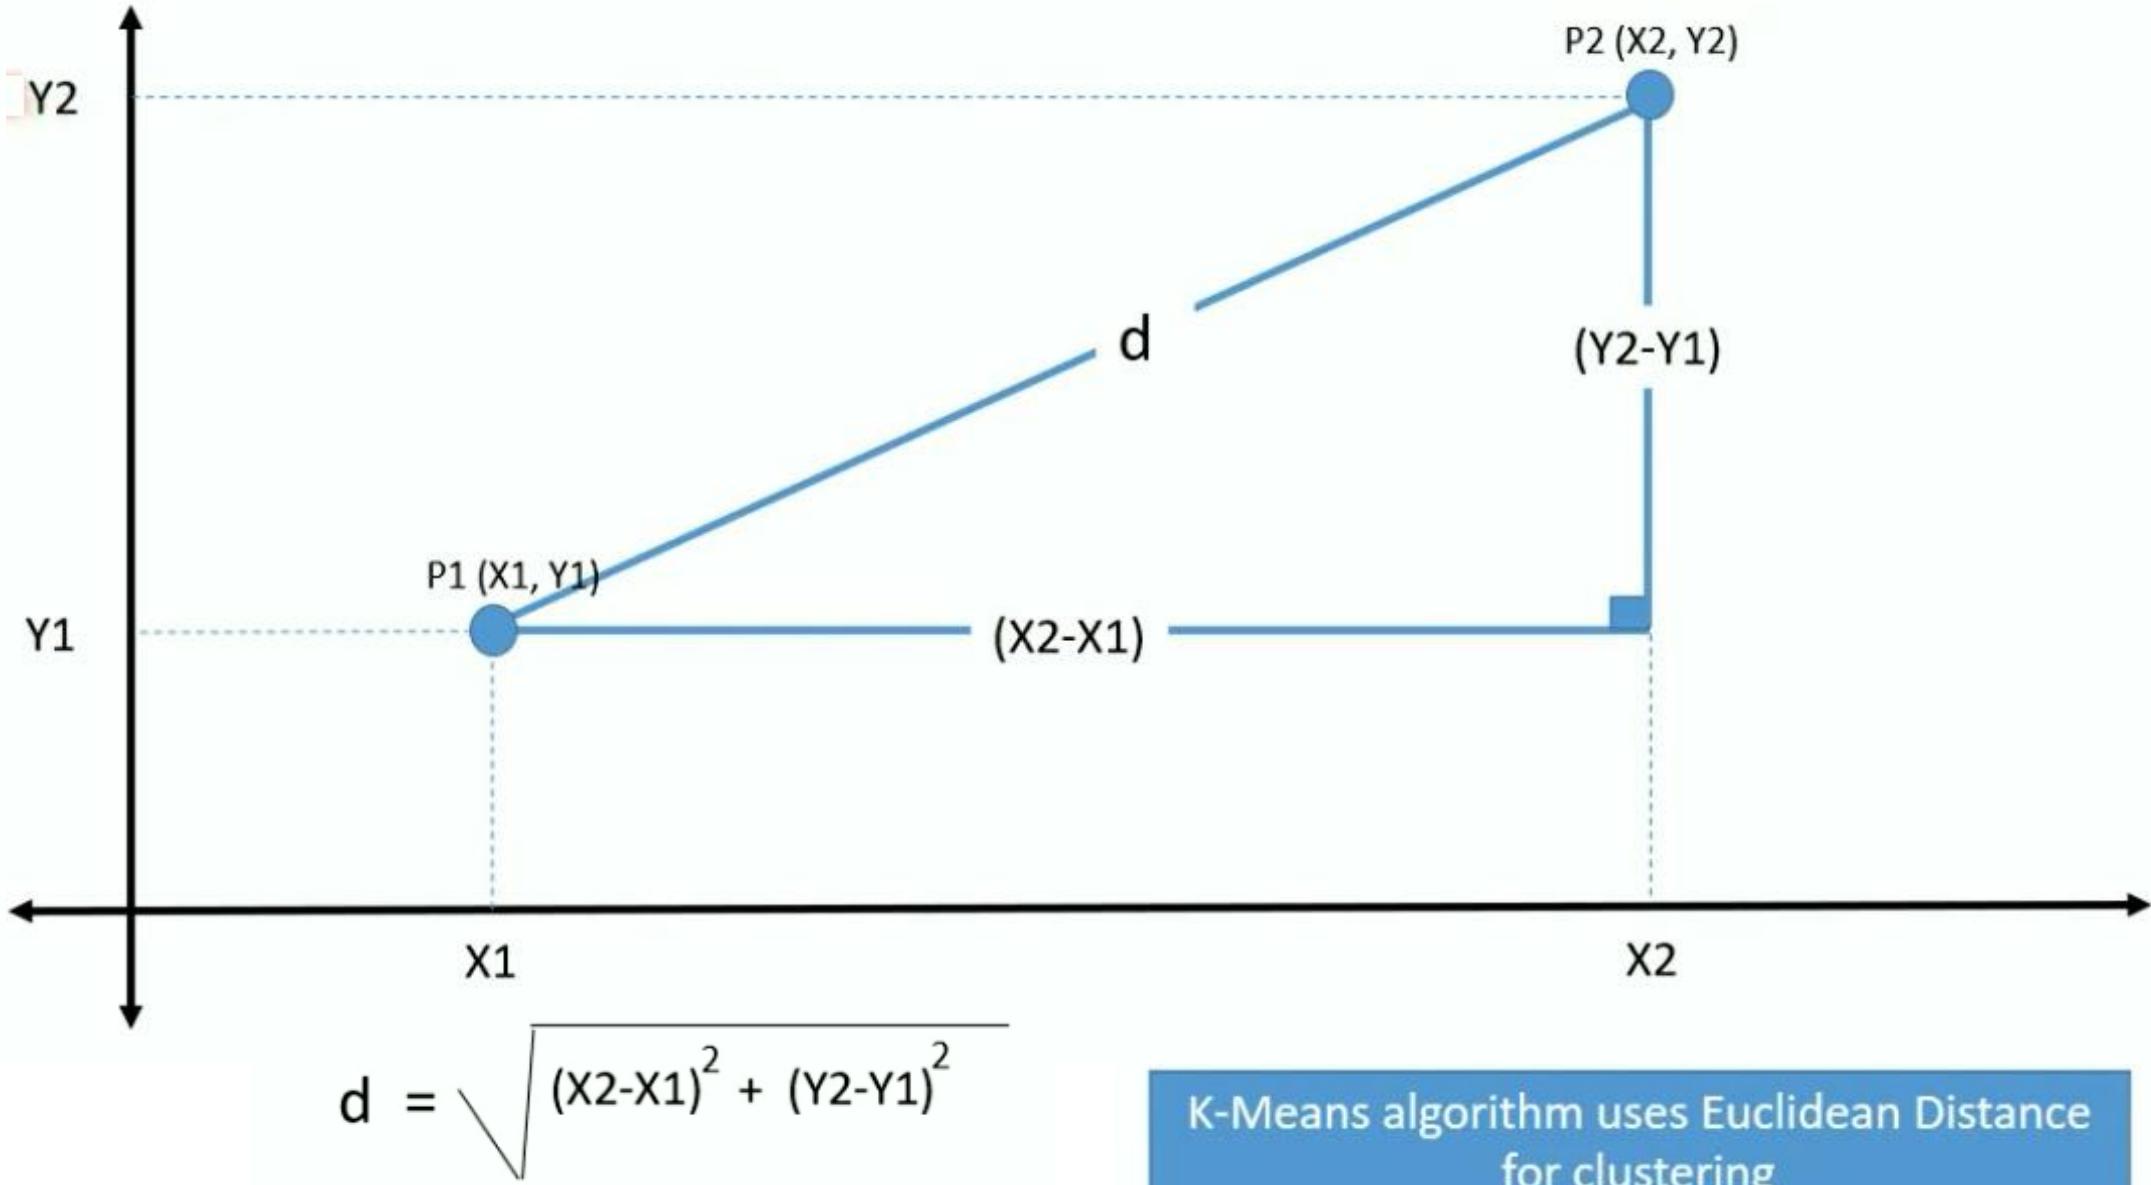
- Medical /Health Science
- Image Segmentation



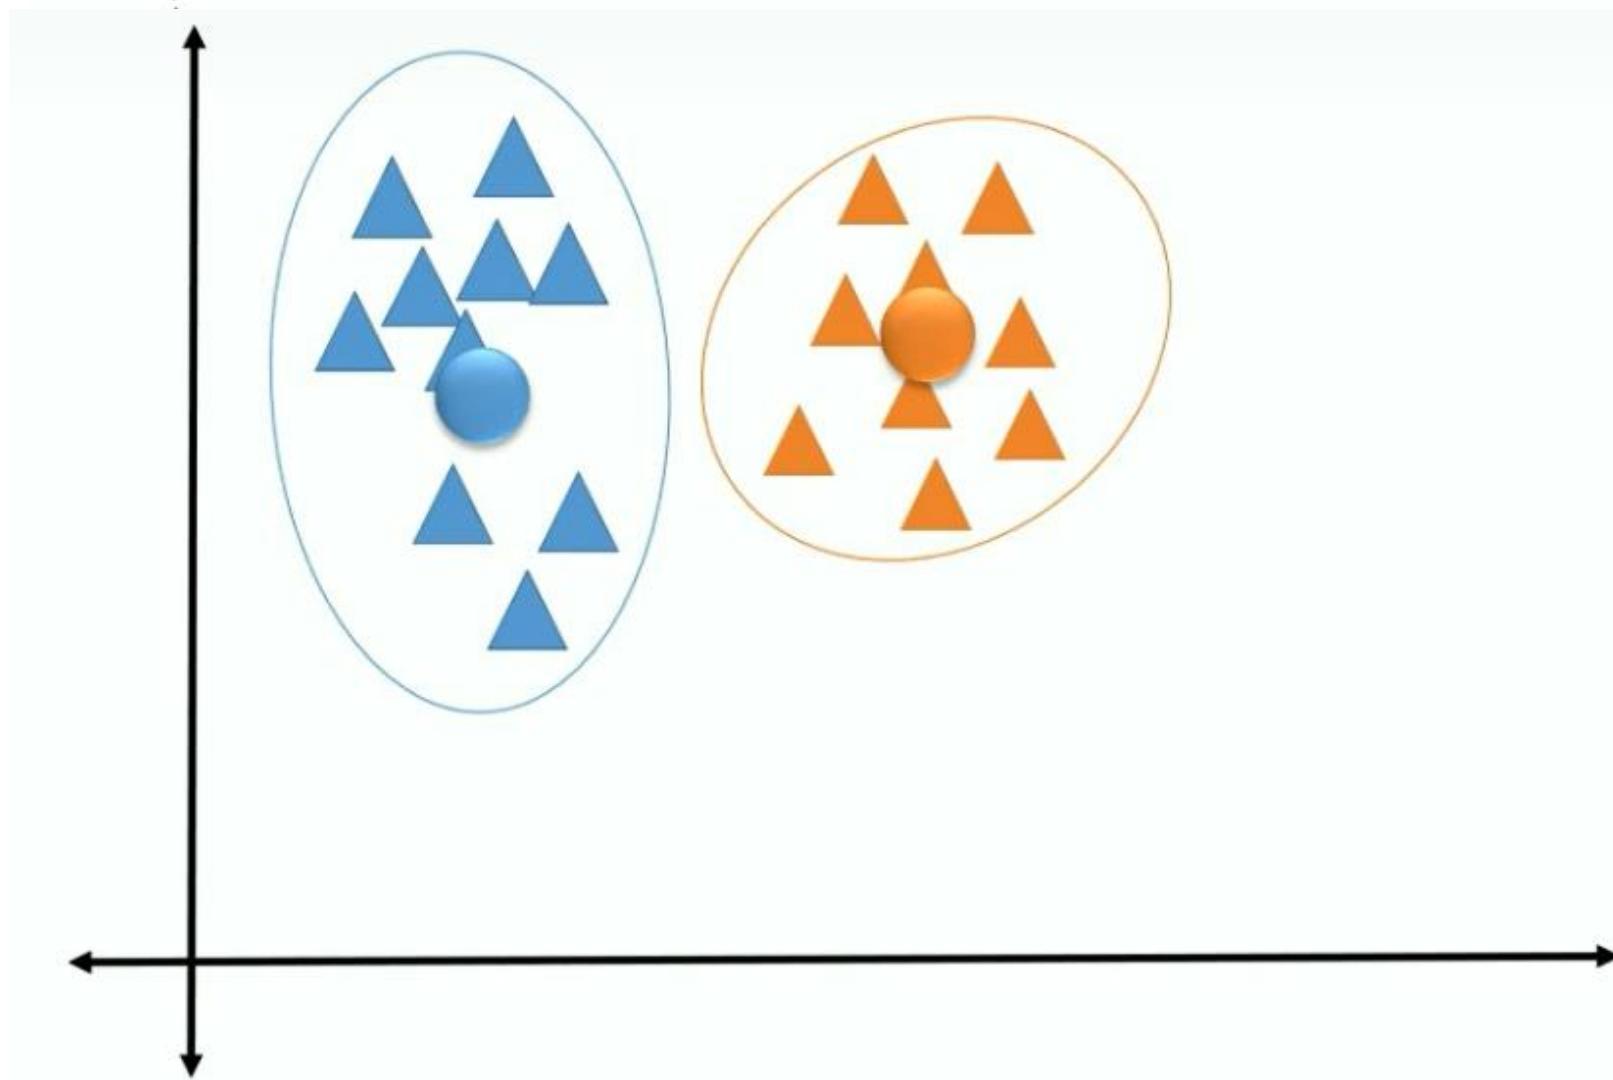
- Anomaly detection



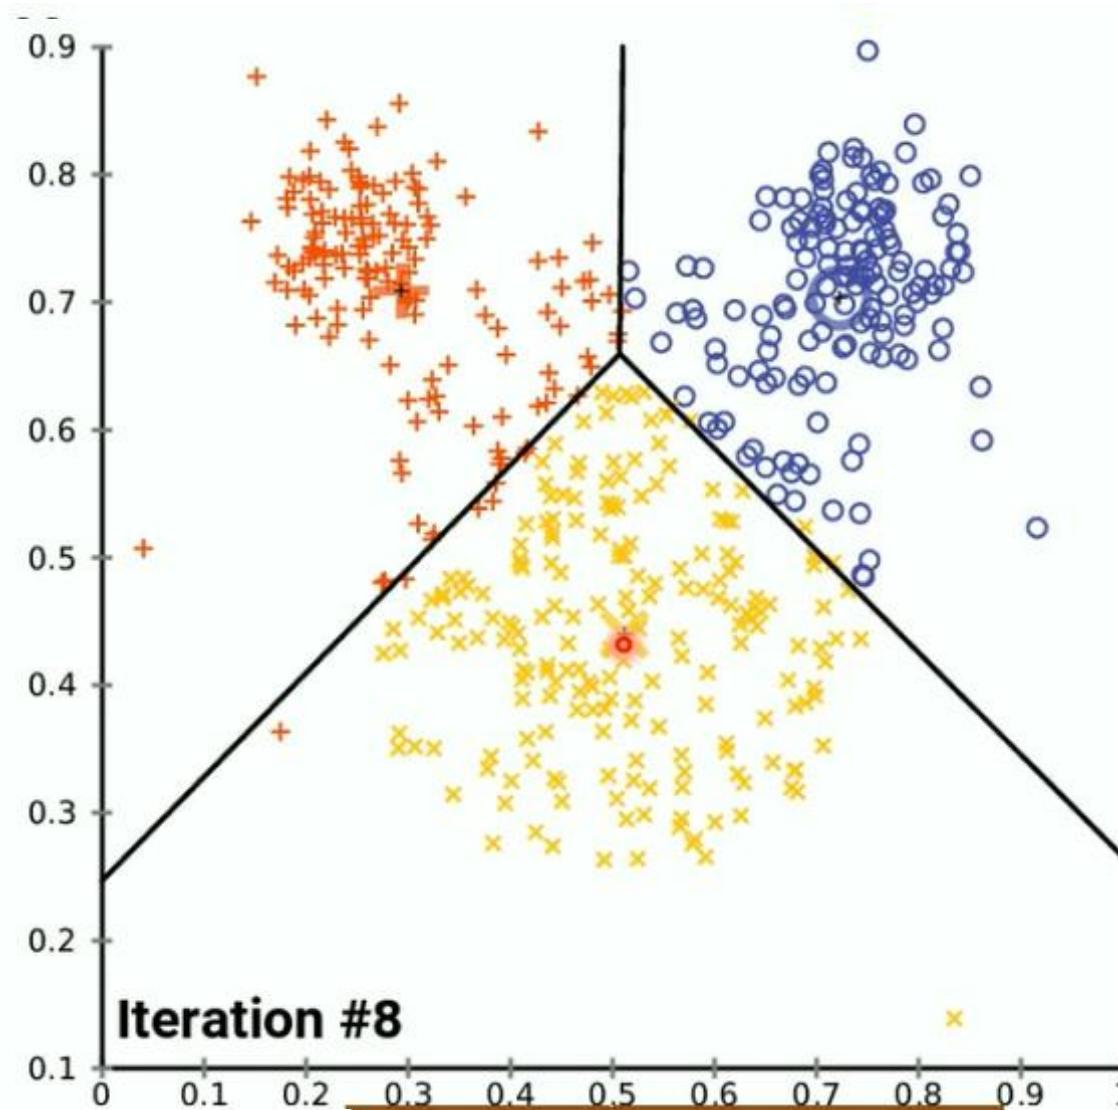
# Euclidean Distance



# How Clusters are formed?



[https://en.wikipedia.org/wiki/K-means\\_clustering](https://en.wikipedia.org/wiki/K-means_clustering) User: Chire



# Good Cluster Analysis

- Observations in the same group share similar characteristics
- Clusters have proportionate number observations

# Cluster Initialization

- Random - Random placement of data points into clusters
- First N or Forgy Method - First Data points at Random
- K-Means++ - Default method and an improvement over finding the initial means
- K-Means++ Fast — Optimised for faster clustering
- Evenly
- Use Label Column

# DEMO

# Clustering using K-Means



Microsoft

# Recommendation



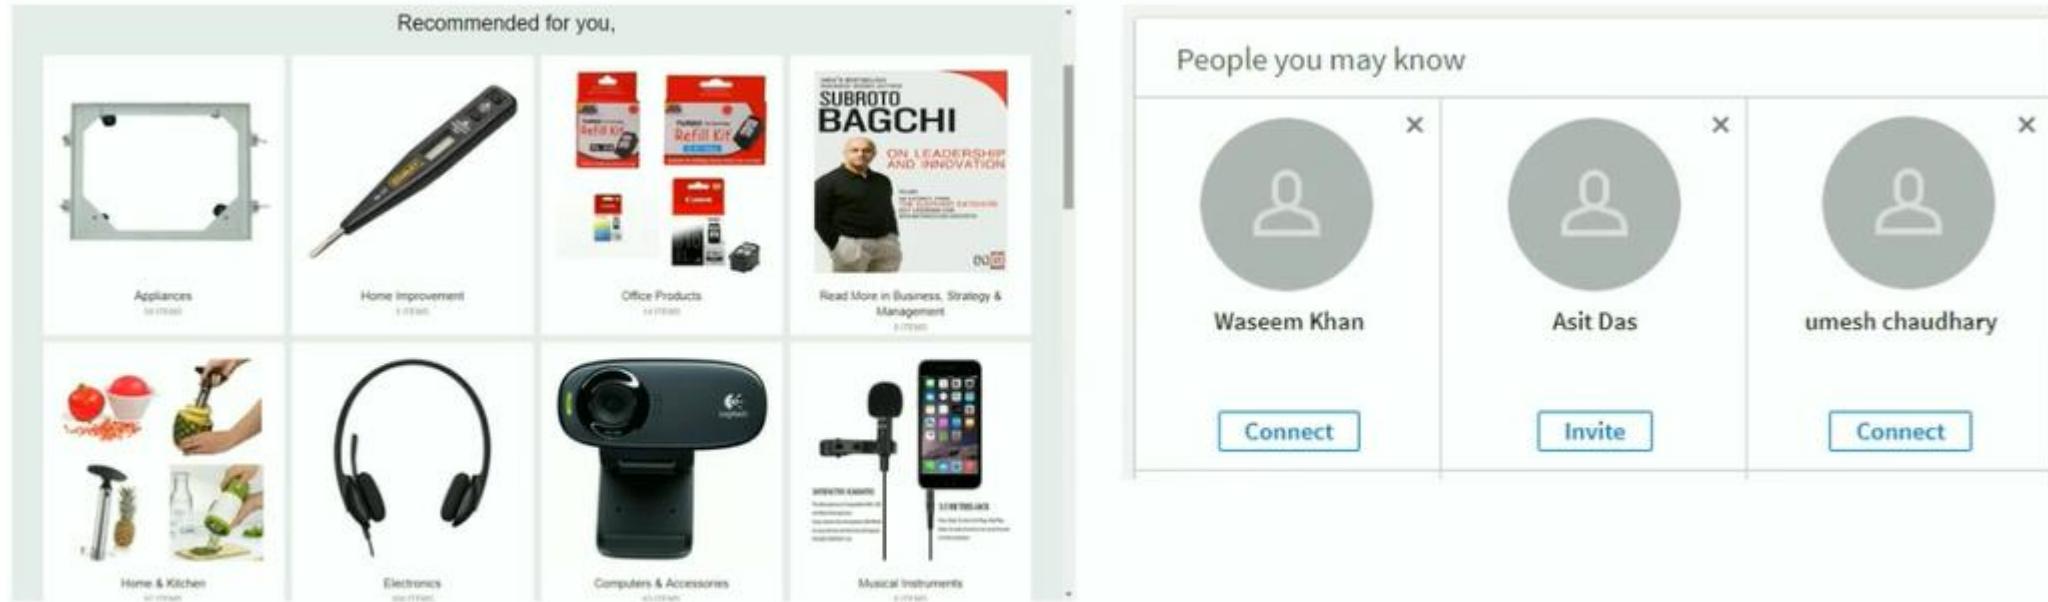
# AzureML Matchbox Recommender

# Recommendation System

- What is a Recommendation System?
- Types of Recommendation Systems
  - Collaborative Filtering
  - Content Based Filtering
- How a recommendation system works?

# What is Recommendation System?

"A recommender system or a recommendation system (platform or engine) seeks to predict the "rating" or "preference" that a user would give to an item."



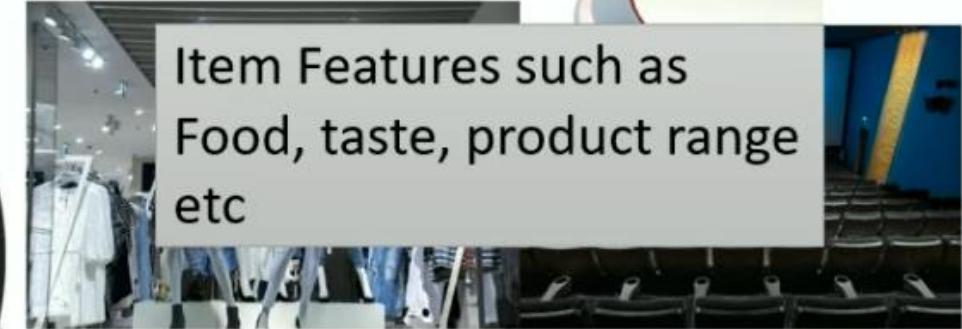
# Everyday Recommendations



History of preferences or behaviour



Friends, Family, Colleagues, Professors



Item Features such as Food, taste, product range etc



Recommend

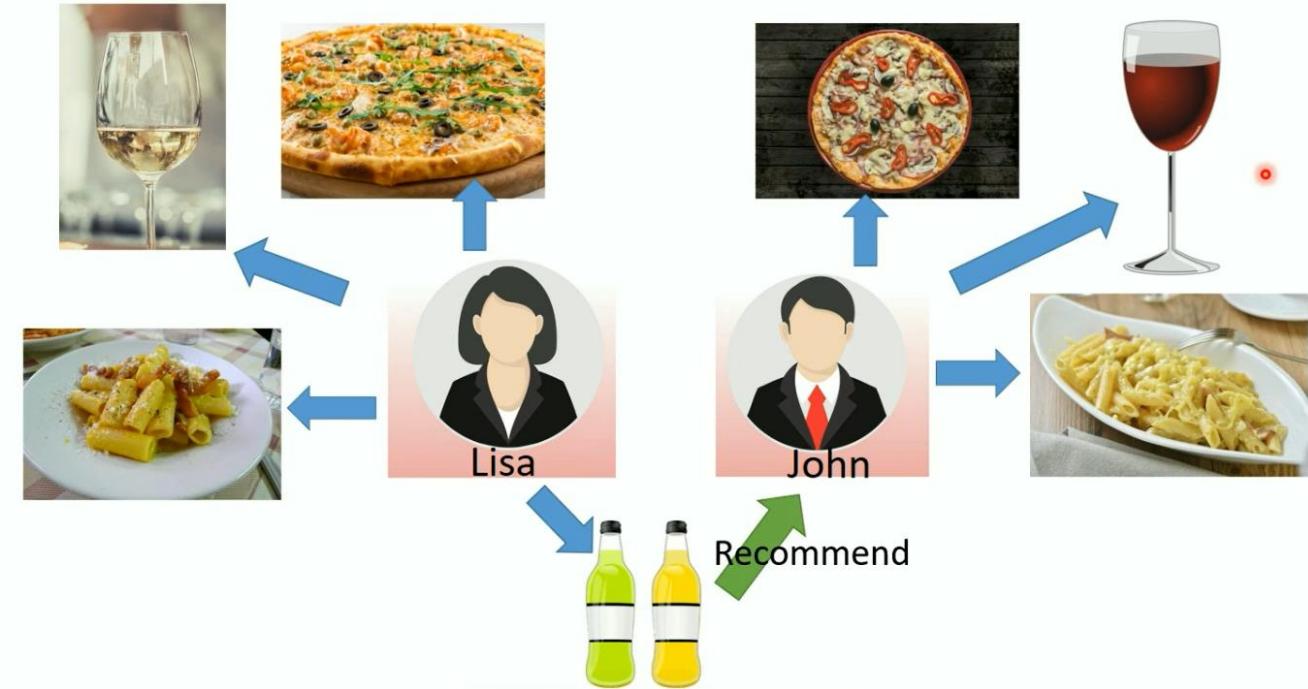


# Types of Recommendations Systems

- Collaborative Filtering
- Content Based Filtering
- Hybrid - Combination of the Collaborative and Content Based
- Popularity based - Most bought, most watched, most downloaded, most heard

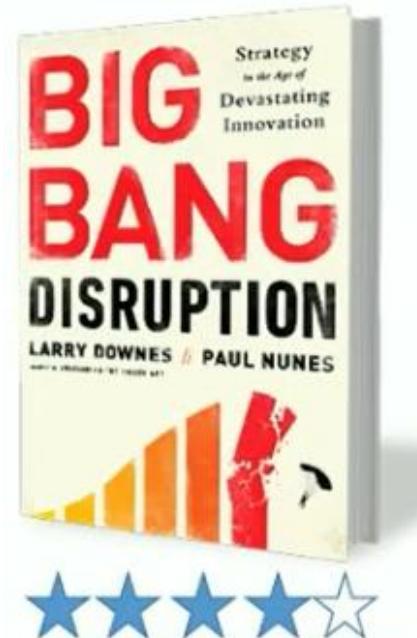
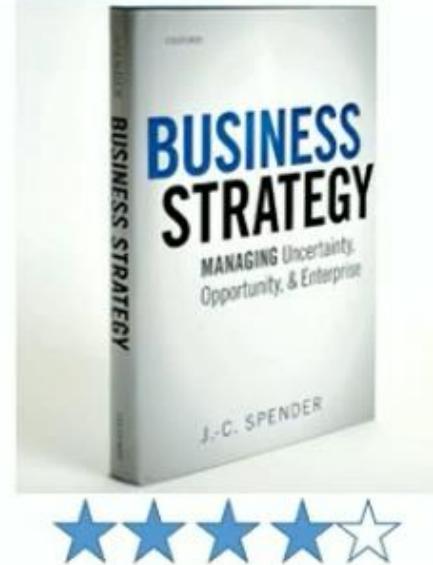
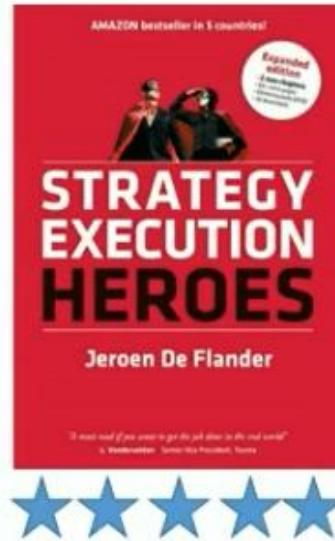
# Collaborative Filtering

- Analyse User behaviour, Activities and preferences
- Recommend based on similarity to other user
- People who agreed in the past will agree in the future, and that they will like similar kinds of items as they liked in the past.

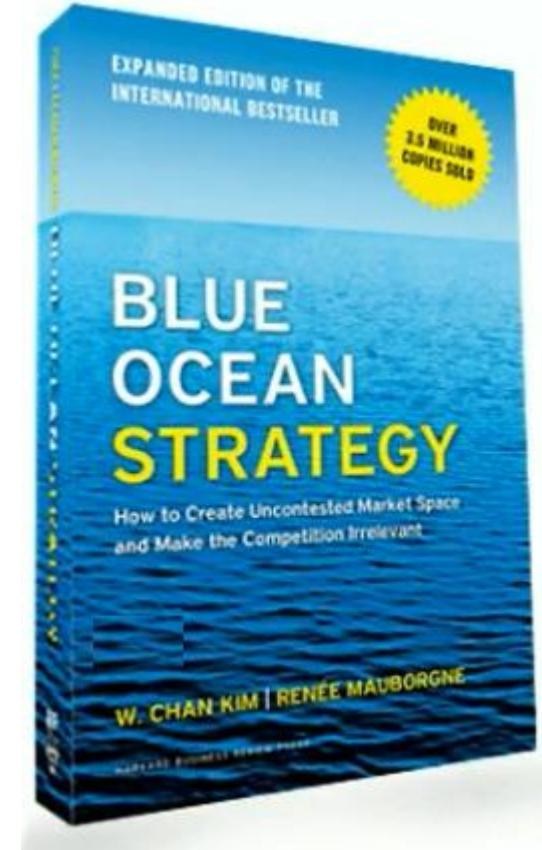


People who buy “x” also buy “y”.

# Content based Filtering



Recommend



# Business Problem

A travel booking website books thousands of hotel rooms every day.

With such a huge growth, the task at hand is to improve the booking per customer by showing the best hotels as per the taste and preferences of the user.

# Sample Set

- Hotel 1 -Good Gym, Small Pool
- Hotel 2 - Amazing Gym, No Pool
- Hotel 3 - Great Gym, Small Pool
- Hotel 4- Basic Gym, Large Pool
- Hotel 5-No Gym, Large and Beautifully designed pool

- John - Prefers Gym over swimming pool
- Kavin - Likes Gym more than pool
- Bill - Needs a pool with basic Gym
- Frans - A pool is a must compared to Gym

**Item – Feature Matrix**

Hotel	Gym	Pool
Hotel 1	0.8	0.2
Hotel 2	1	0
Hotel 3	0.9	0.1
Hotel 4	0.1	0.9
Hotel 5	0	1

**User – Feature Matrix**

User	Gym	Pool
John	0.9	0.1
Kavin	0.8	0.2
Bill	0.3	0.7
Frans	0	1

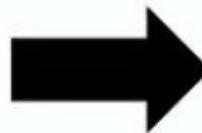
# Creating Feature Vector

Hotel	Gym	Pool
Hotel 1	0.8	0.2
Hotel 2	1	0
Hotel 3	0.9	0.1
Hotel 4	0.1	0.9
Hotel 5	0	1



Hotel	Feature Vector
H1	[ 0.8, 0.2 ]
H2	[ 1, 0 ]
H3	[ 0.9, 0.1 ]
H4	[ 0.1, 0.9 ]
H5	[ 0, 1 ]

User	Gym	Pool
John	0.9	0.1
Kavin	0.8	0.2
Bill	0.3	0.7
Frans	0	1



User	Feature Vector
U1	[ 0.9, 0.1 ]
U2	[ 0.8, 0.2 ]
U3	[ 0.3, 0.7 ]
U4	[ 0, 1 ]

# Creating Feature Vector

Hotel	Feature Vector	User	Feature Vector
H1	[ 0.8, 0.2 ]	U1	[ 0.9, 0.1 ]
H2	[ 1, 0 ]	U2	[ 0.8, 0.2 ]
H3	[ 0.9, 0.1 ]	U3	[ 0.3, 0.7 ]
H4	[ 0.1, 0.9 ]	U4	[ 0, 1 ]
H5	[ 0, 1 ]		

- Which Hotels should be recommended to John (U1)?
- MAX (Uj \* Hi)

MAX (U1\*H1, U1\*H2, U1\*H3, U1\*H4, U1\*H5)

MAX (0.74, 0.9, 0.82, 0.18, 0.1)

**Recommended Hotels in the order of preference**

**H2, H3, H1**

# Collaborative Filtering

Workflow of CF, (Wikipedia)

- A user expresses his or her preferences by rating items.
- Ratings can be viewed as an approximate representation of the user's interest in the corresponding domain.
- The system matches this user's ratings against other users' and finds the people with most "similar" tastes.
- With similar users, the system recommends items that the similar users have rated highly but not yet being rated by this user (presumably the absence of rating is often considered as the unfamiliarity of an item)

User	Hotel 1	Hotel 2	Hotel 3	Hotel 4	Hotel 5
John	4	5	?	?	1
Kavin	?	5	4.5	?	1
Bill	2	3	?	5	4
Frans	1	?	?	?	5

# Collaborative Filtering

User	Hotel 1	Hotel 2	Hotel 3	Hotel 4	Hotel 5
John	4	5	?	?	1
Kavin	?	5	4.5	?	1
Bill	2	3	?	5	4
Frans	1	?	?	?	5

# Matchbox Recommender

# Recommendation Data Issues

- Cold Start
- Same Scale Judgement

# Cold Start

- Content based filtering constructs user profile before system can recommend
- Collaborative Filtering needs like-minded users for recommendations
- Completely new profile or item
- Can be tackled using Hybrid recommendation

# Same Scale judgement

- A user rates all the items on a same scale
- An item is rated on the same scale by all users

# Recommender Split

- Fraction of training only users - Fraction of users assigned only to the training dataset. The rows would never be used to test the model.
- Fraction of test user ratings for training - Portion of the user ratings that can be used for training.
- Fraction of cold users - Cold users are users that the system has not previously encountered
- Fraction of cold items - Cold items are items that the system has not previously encountered.
- Fraction of ignored users - Specify the percentage of users that should be ignored.
- Fraction of ignored items - Specify the percentage of items to ignore.
- Remove occasionally produced cold items

# DEMOS

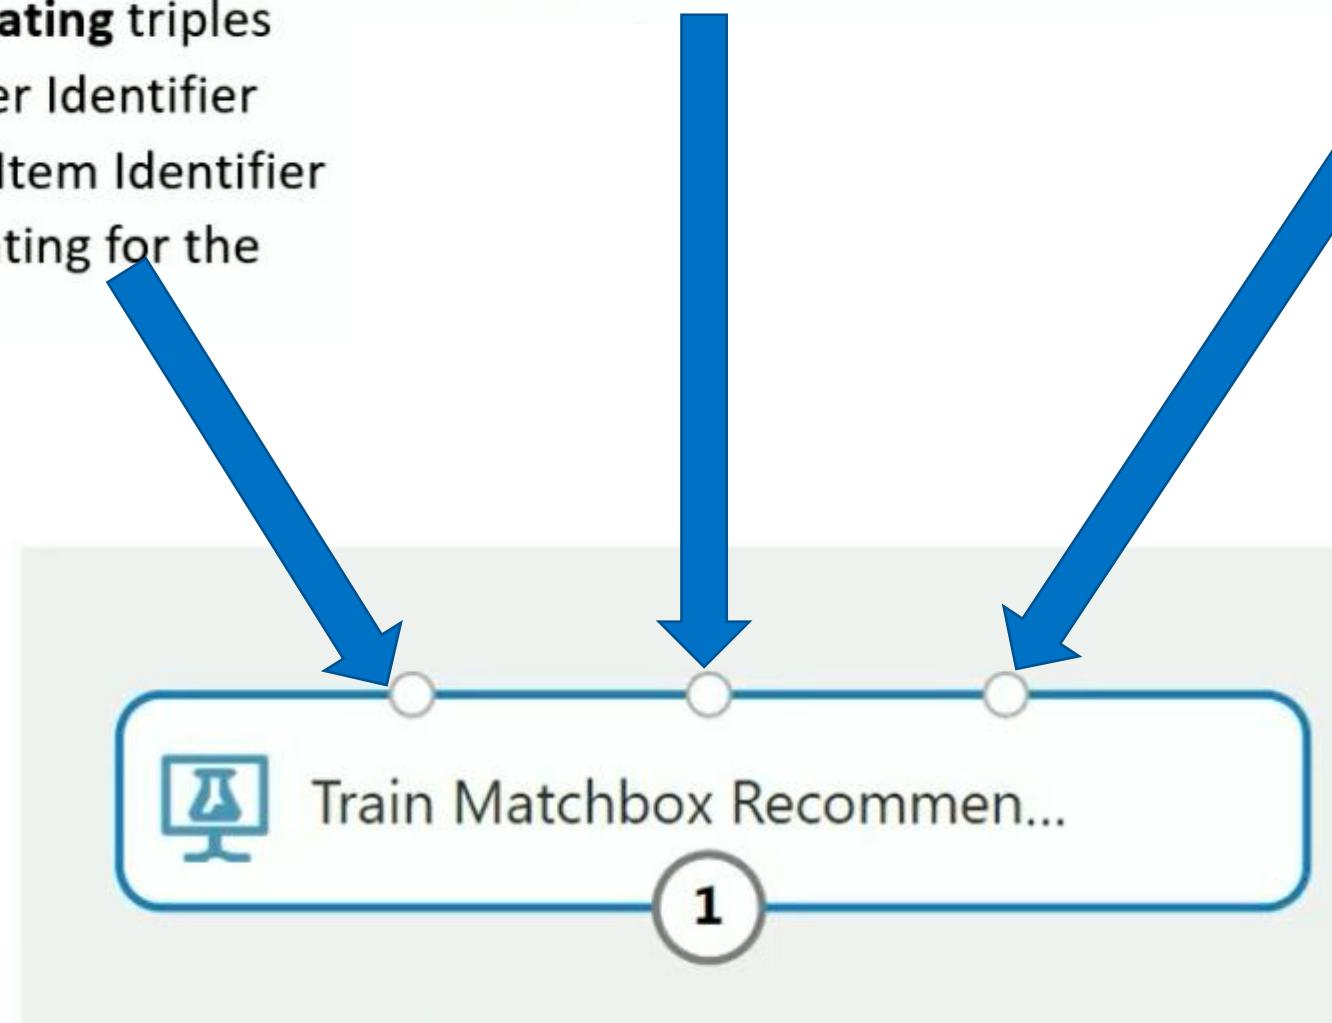
# Matchbox Recommendation System

# What is Matchbox Recommendation System

- Recommendation system developed by Microsoft Research
- Based on a the Hybrid Approach of Content and Collaborative filtering
- Uses Train Matchbox Recommender and Score Matchbox Recommender Modules
- Train Matchbox Recommender reads a dataset of user-item-rating triples and, optionally, some user and item features.

# Configuring Train Matchbox Recommender

- User-Item-Rating Data - Prepare the data used for training, containing **user-item-rating** triples
  - First Column – User Identifier
  - Second Column – Item Identifier
  - Third Column – Rating for the user-item pair
- User-Feature – userID and the userfeatures
- Item-Feature – itemID and the Itemfeatures



# Parameters to Train Matchbox Recommender

- **Number of Traits**
  - How many traits to learn for each user and item
  - Each feature is associated with a latent “trait” vector
- **Number of recommendation algorithm iterations**
  - how many times the algorithm should process the input data
- **Number of training batches**
  - Number of batches for dividing the data during training

Properties Project

## Train Matchbox Recommender

Number of traits

10

Number of recommendation algorithm iterations

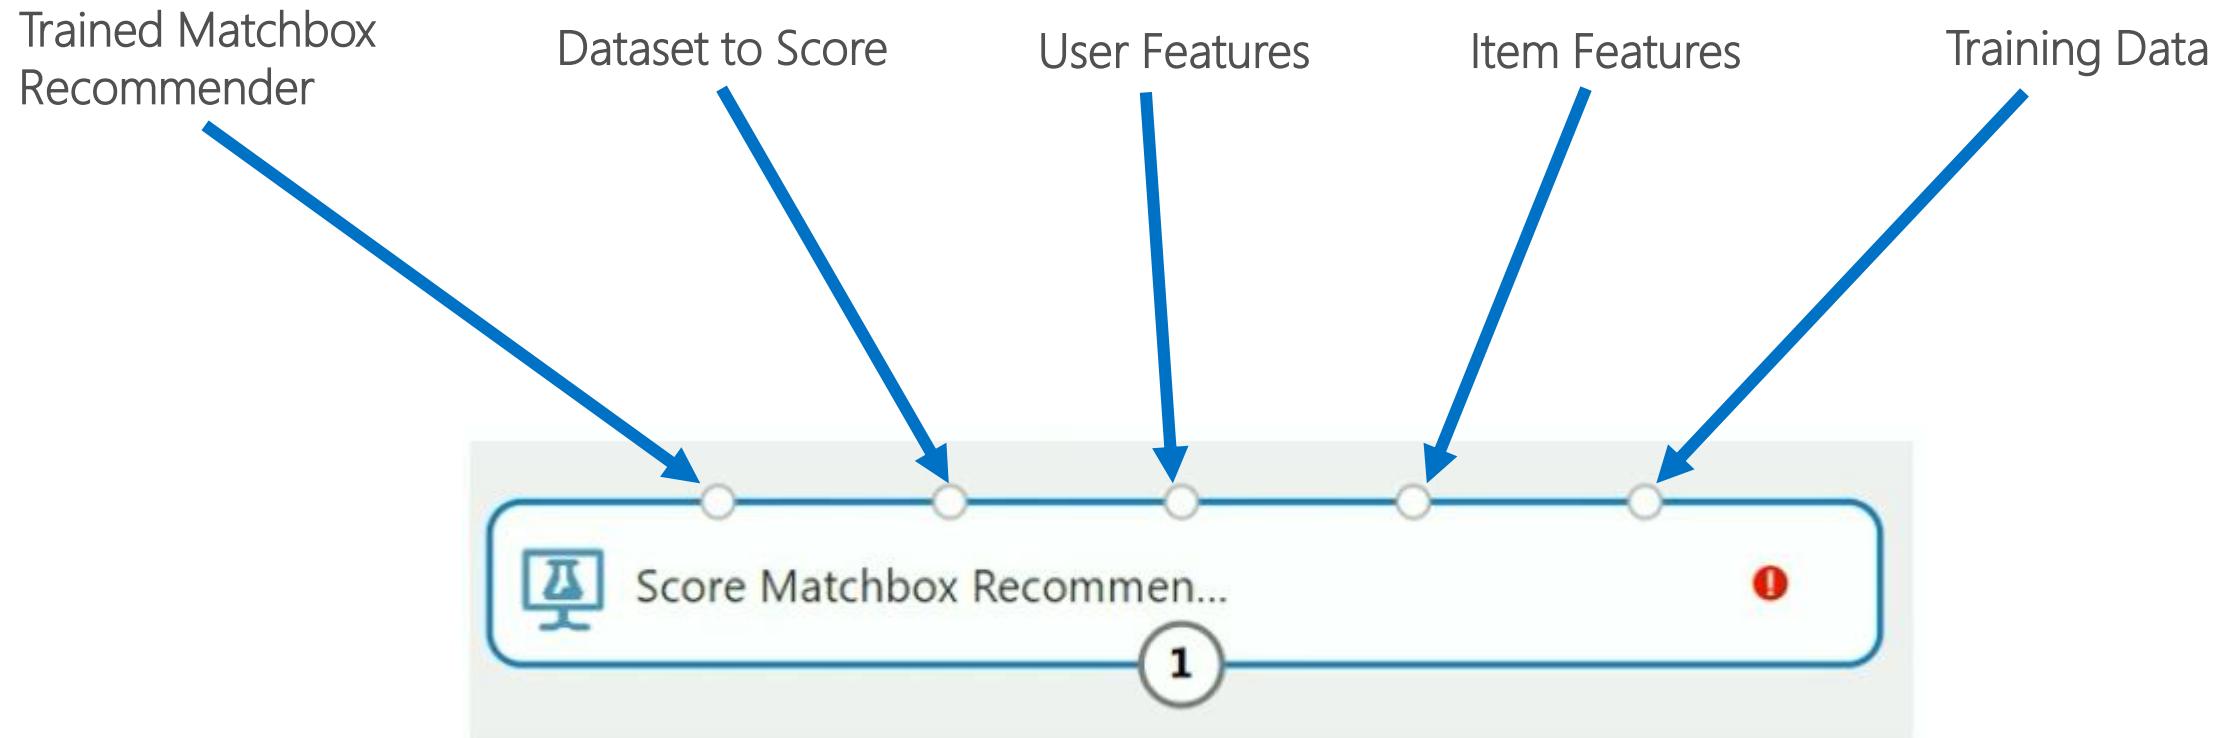
5

Number of training batches

4

# Score Matchbox Recommender

# Input to Score Matchbox Recommender



# Types of Prediction

- Predict Ratings
- Item Recommendation
- Related Users
- Related Items

# Ratings of Prediction

- Rating predictions by a user
- Input data must contain User and Item
- Does not require any parameter

# Item Recommendation

- Users and Items as input
- Uses its knowledge about existing items and users
- Generates a list of items that will appeal to each user

## Score Matchbox Recommender

Recommender prediction kind

Item Recommendation

Recommended item selection

From Rated Items (for model evaluation)

Maximum number of items to recommend to a user

5

Minimum size of the recommendation pool for a single user

2

Whether to return the predicted ratings of the items along with the labels

# Related Users and Items

- Can be used for "People like you" predictions
- Generate recommendations for users based on items that have already been rated

# Understanding the Recommender Result

- For Item Recommendations
- Normalized Discounted Cumulative Gain (NDCG)



# Ranking Quality

## Search Result

What is machine learning? - Definition from WhatIs.com

[whatis.techtarget.com › Topics › AppDev › Programming](https://whatis.techtarget.com/Topics/AppDev/Programming) ▾

Jun 24, 2017 - Machine learning is the field of study that gives computers the ability to learn without being explicitly programmed.

The 10 Algorithms Machine Learning Engineers Need to Know

[www.kdnuggets.com/2016/08/10-algorithms-machine-learning-engineers.html](https://www.kdnuggets.com/2016/08/10-algorithms-machine-learning-engineers.html) ▾

Aug 8, 2016 - It is no doubt that the sub-field of machine learning / artificial intelligence has increasingly gained more popularity in the past couple of years.

Machine Learning | SAP - SAP.com

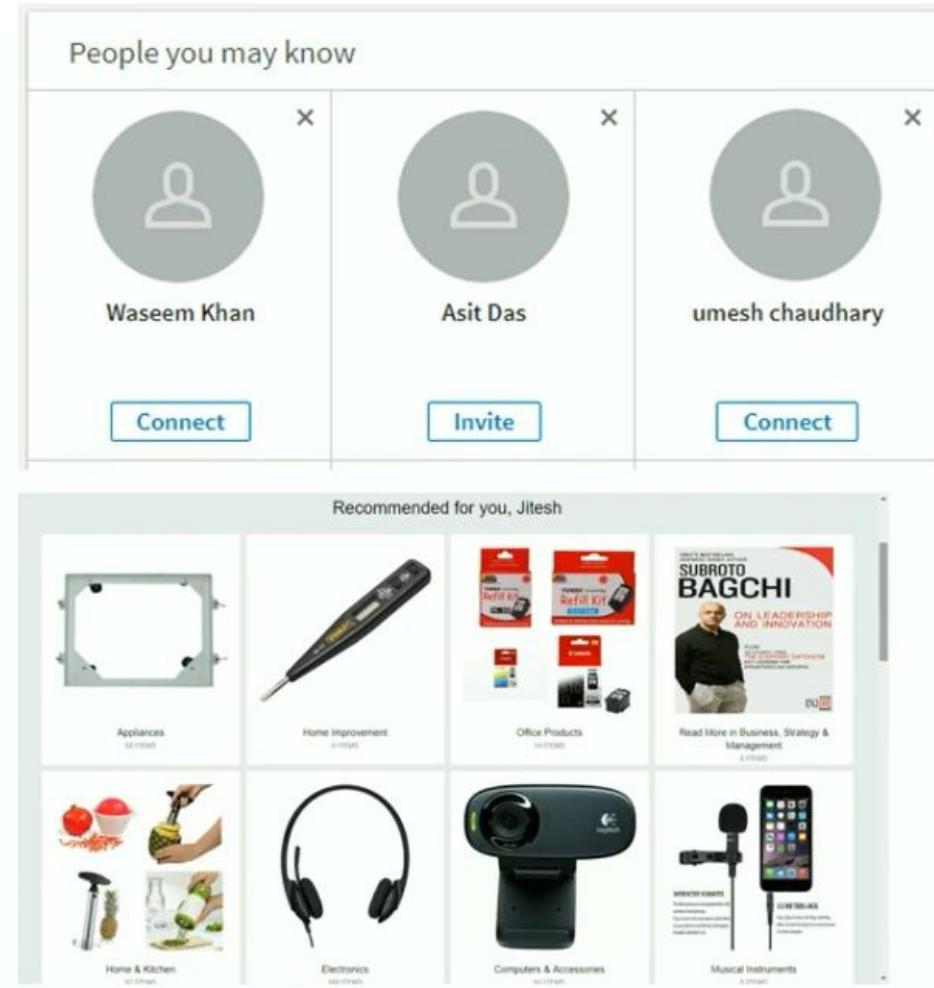
<https://www.sap.com/india/trends/machine-learning.html> ▾

Discover how AI, machine learning, and deep learning are powering a new breed of software that uses Big Data to drive radical changes to business.

Machine Learning | edX

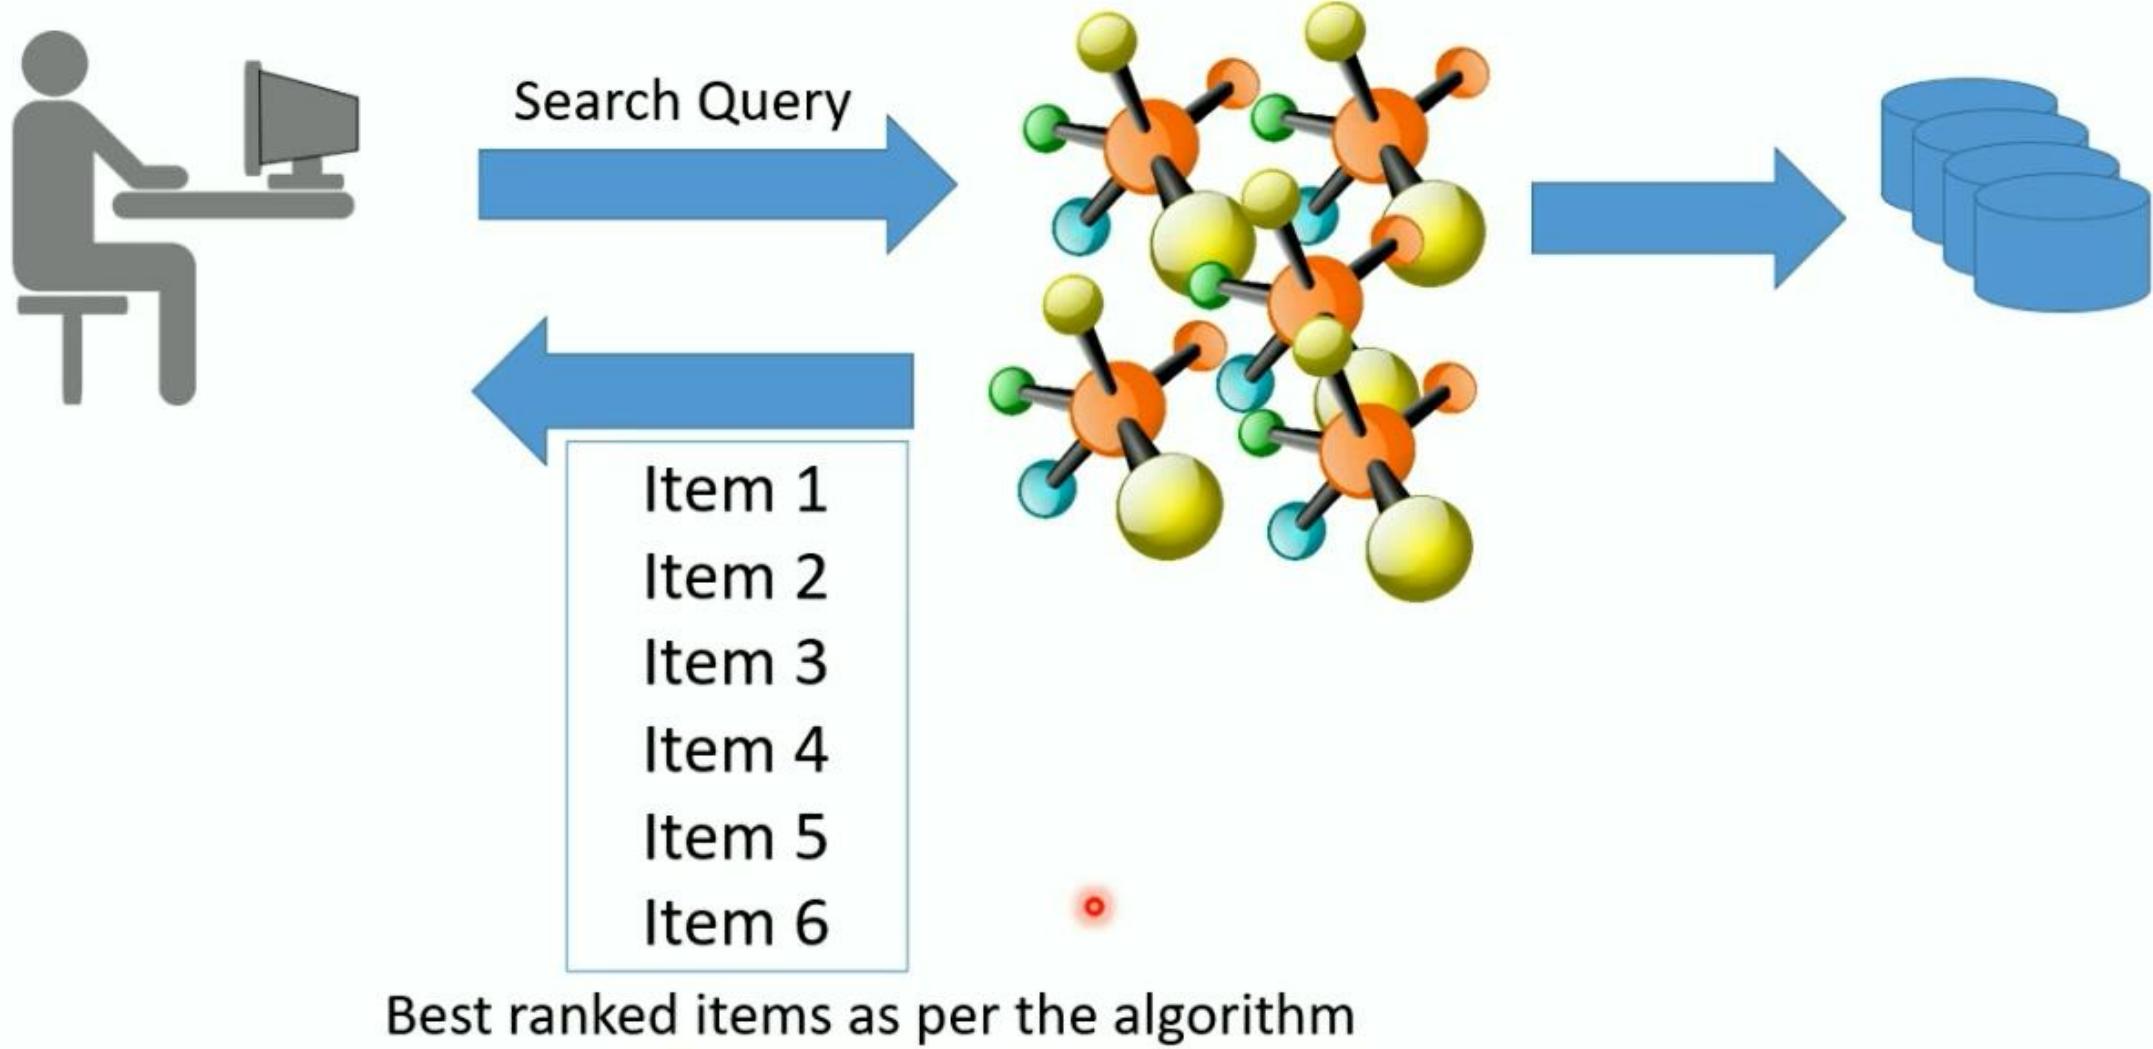
<https://www.edx.org/course/machine-learning-columbiasx-csmm-102x-0> ▾

Master the essentials of machine learning and algorithms to help improve learning from data without human intervention.



The highest ranked item should result in highest gain.

# Ranking Quality



# Discounted Cumulative Gain

Ranked items by the algorithms  I1, I2, I3, I4, I5, I6

$$\begin{aligned} \text{CG} &= 4 + 3 + 4 + 0 + 1 + 2 \\ &= 14 \end{aligned}$$

Gain perceived by the user  4, 3, 4, 0, 1, 2

$i$	$\text{rel}_i$	$\log_2(i+1)$	$\frac{\text{rel}_i}{\log_2(i+1)}$
1	4	1	4.00
2	3	1.585	1.89
3	4	2	2
4	0	2.322	0
5	1	2.585	0.39
6	2	2.81	0.71
<b>DCG</b>			<b>8.99</b>

# Ideal Discounted Cumulative Gain

Ideal Ranking by algorithm



I1, I3, I2, I6, I5, I4

Ideal Gain perceived by the user



4, 4, 3, 2, 1, 0

$$\begin{aligned} \text{CG} &= 4 + 4 + 3 + 2 + 1 + 0 \\ &= 14 \end{aligned}$$

$i$	$\text{rel}_i$	$\log_2(i+1)$	$\frac{\text{rel}_i}{\log_2(i+1)}$
1	4	1	4.00
2	4	1.585	2.52
3	3	2	1.5
4	2	2.322	0.86
5	1	2.585	0.38
6	0	2.81	0
IDCG			<b>9.27</b>

# Normalised DCG

**DCG = 8.99**

**IDCG = 9.27**

$$\text{NDCG} = \frac{\text{DCG}}{\text{IDCG}} = \frac{8.99}{9.27} = 0.9697$$

Highly relevant documents are more useful than marginally relevant documents, which are in turn more useful than non-relevant documents

# Understanding the Result

- Normalized Discounted Cumulative Gain (NDCG)



# Computer Vision Tasks

## Image Classification

- Is there a deer in the image?



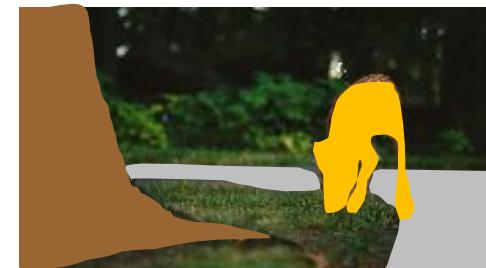
## Object detection

- Where in the image is the deer?



## Image segmentation

- Where exactly is the deer? What pixels?



## Image Similarity

- Which images are similar to the query image?

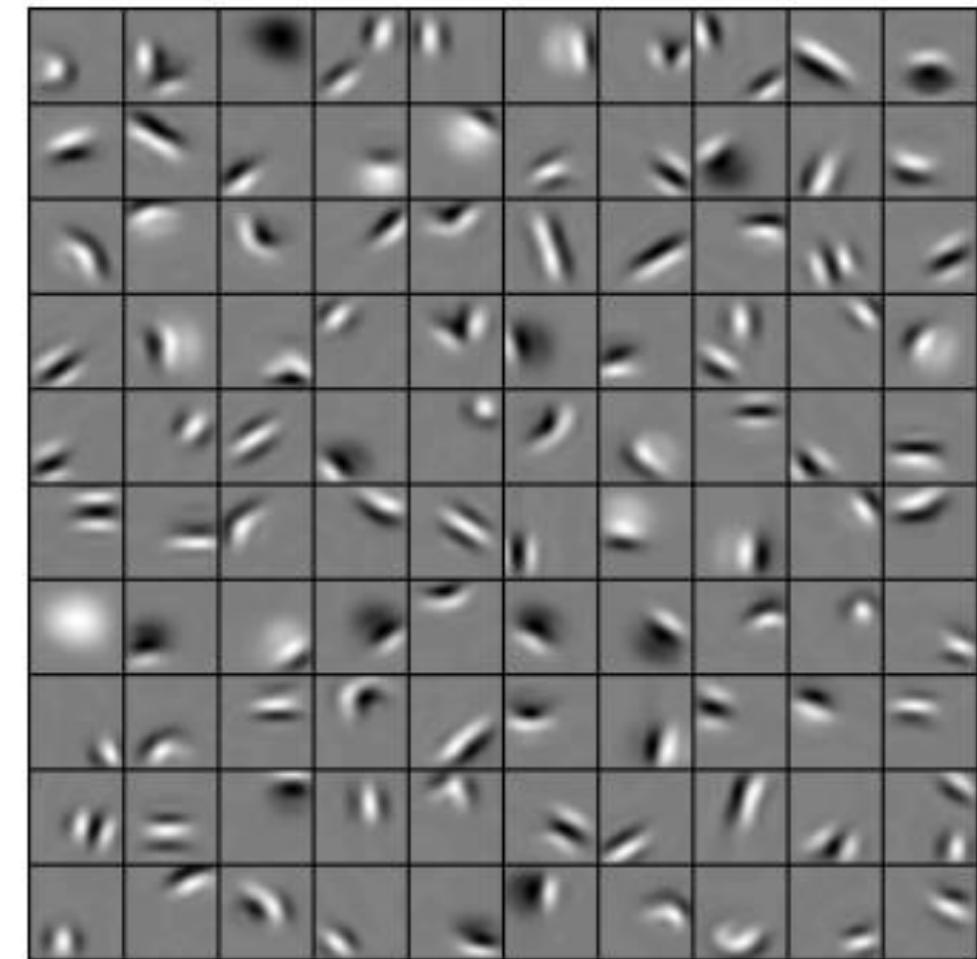


# Before Deep Learning

- Researchers took a traditional machine learning approach
  - Manual creation of a variety of different visual feature extractors
  - Followed by traditional ML classifiers
- Example: HoG Detectors
  - Histogram of oriented gradients (HoG) features
  - Sliding window detector
  - SVM Classifier
  - Very fast OpenCV implementation (<100ms)



# How it Works: Convolutional Neural Networks



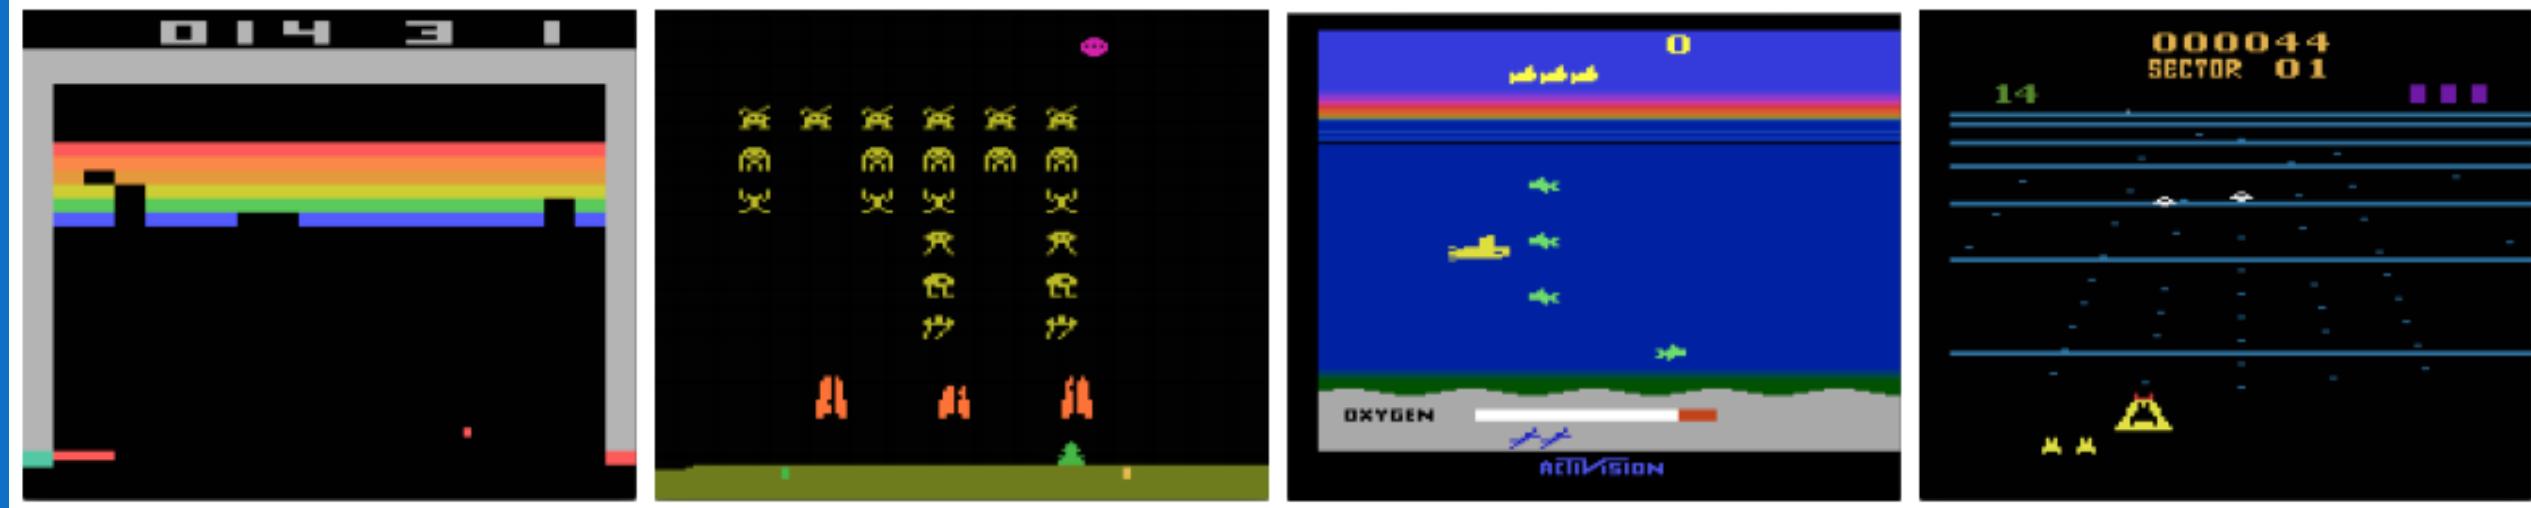
faces



cars

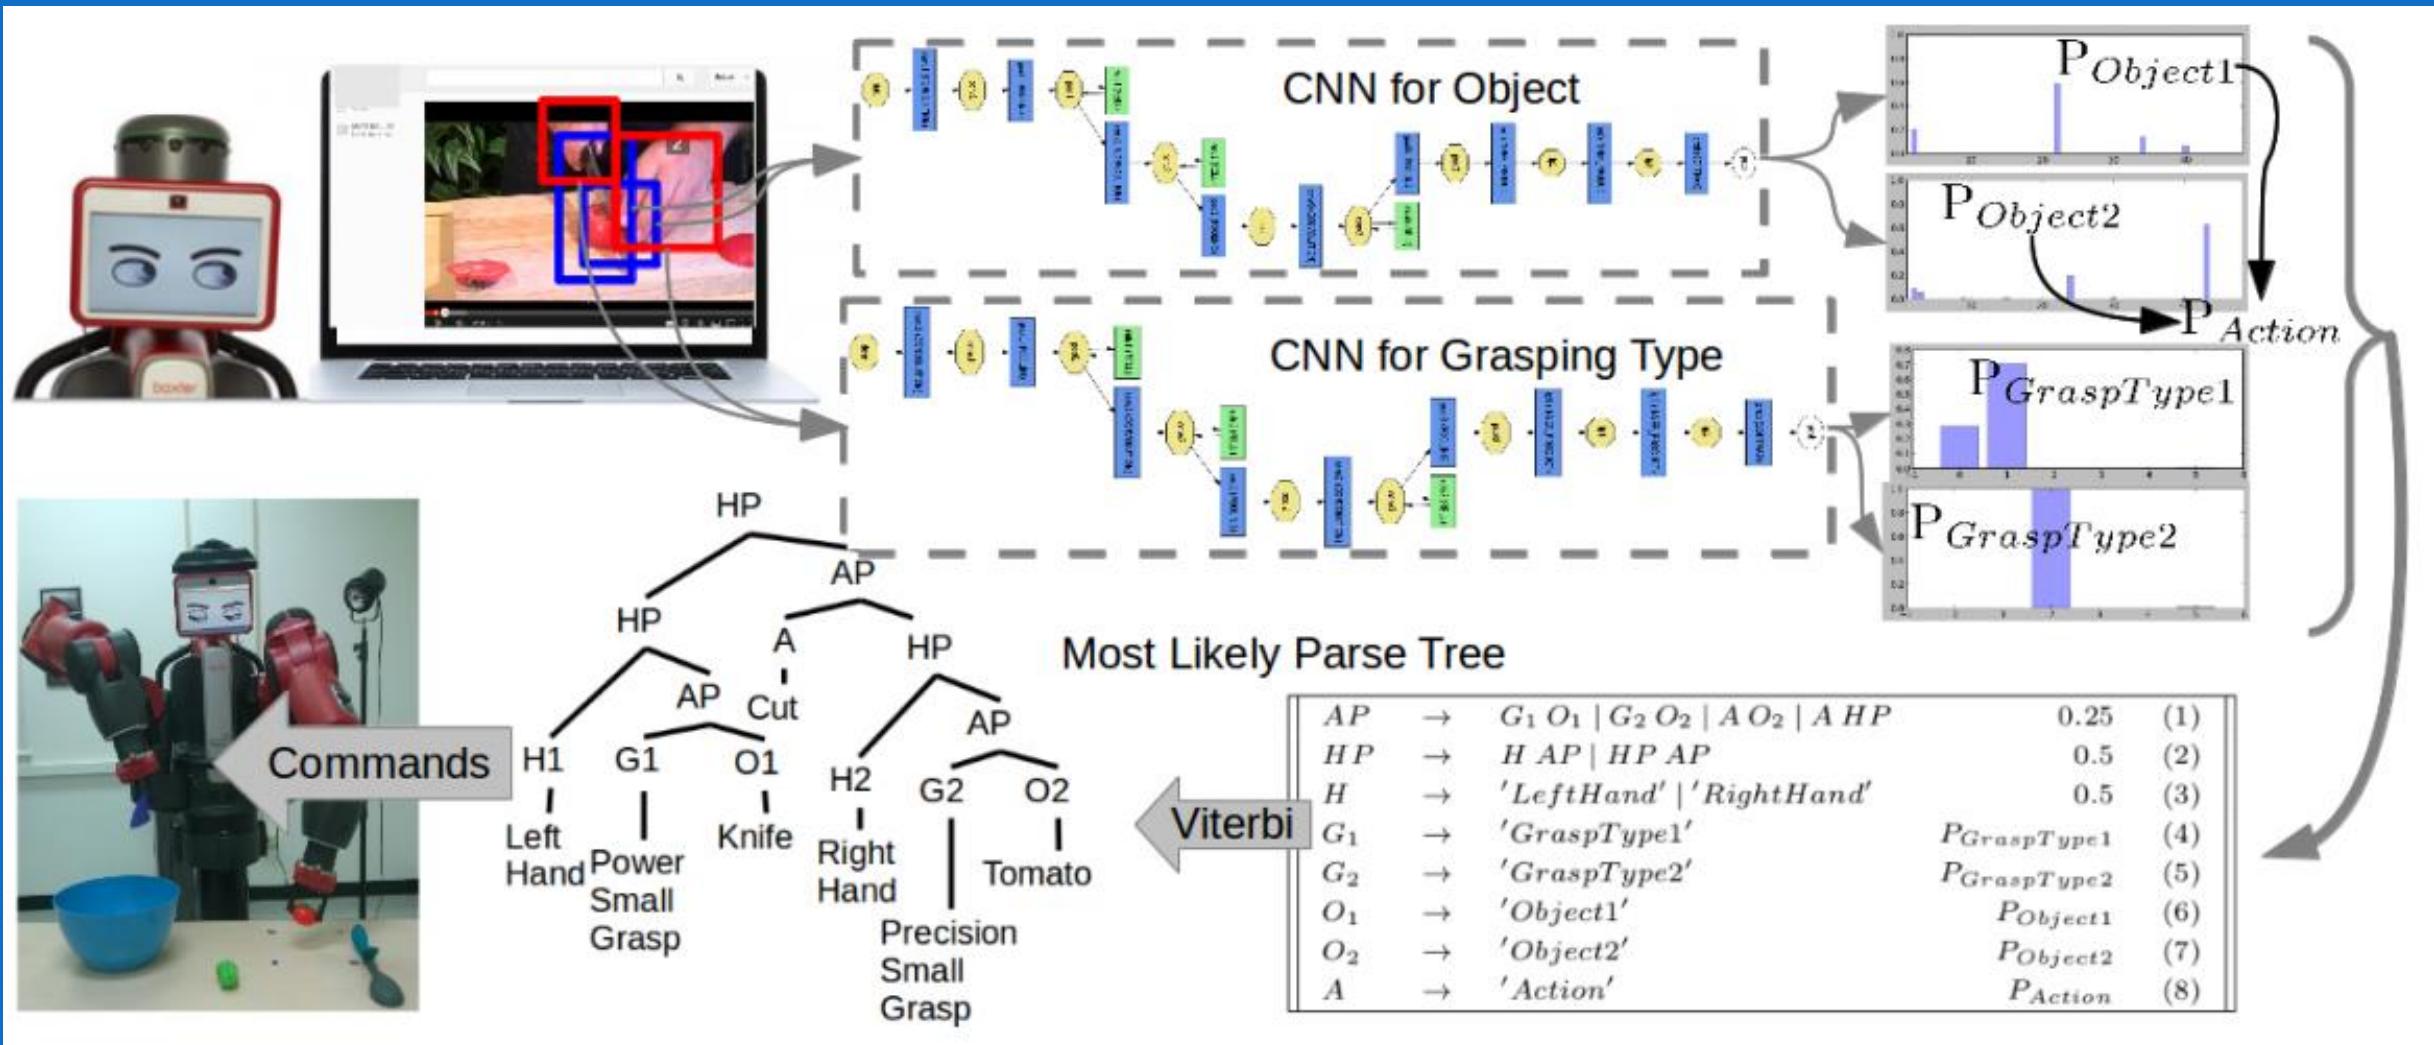


Convolutional Deep Belief Networks for Scalable  
Unsupervised Learning of Hierarchical Representations  
Honglak Lee, Roger Grosse, Rajesh Ranganath, Andrew Y.  
Ng



Playing Atari with Deep Reinforcement Learning.

Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, Martin Riedmiller



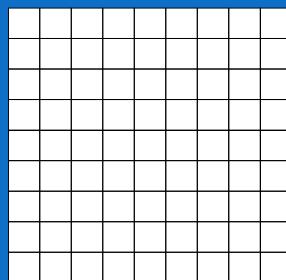
Robot Learning Manipulation Action Plans by “Watching” Unconstrained Videos from the World Wide Web.

Yezhou Yang, Cornelia Fermuller, Yiannis Aloimonos

# A toy ConvNet: X's and O's

Says whether a picture is of an X or an O

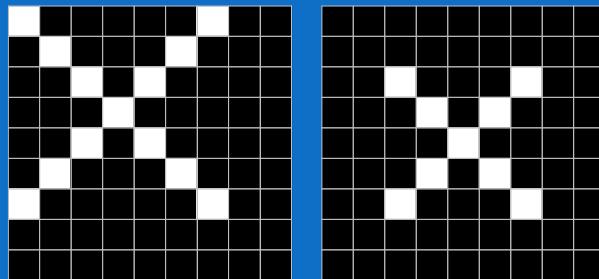
A two-dimensional  
array of pixels



For example



# Trickier cases

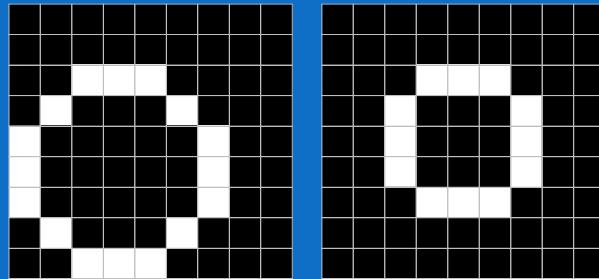


translation

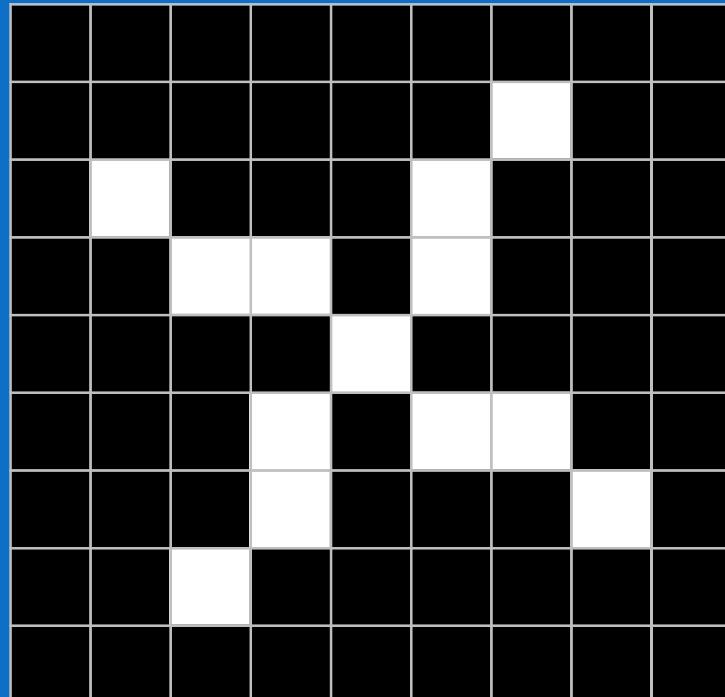
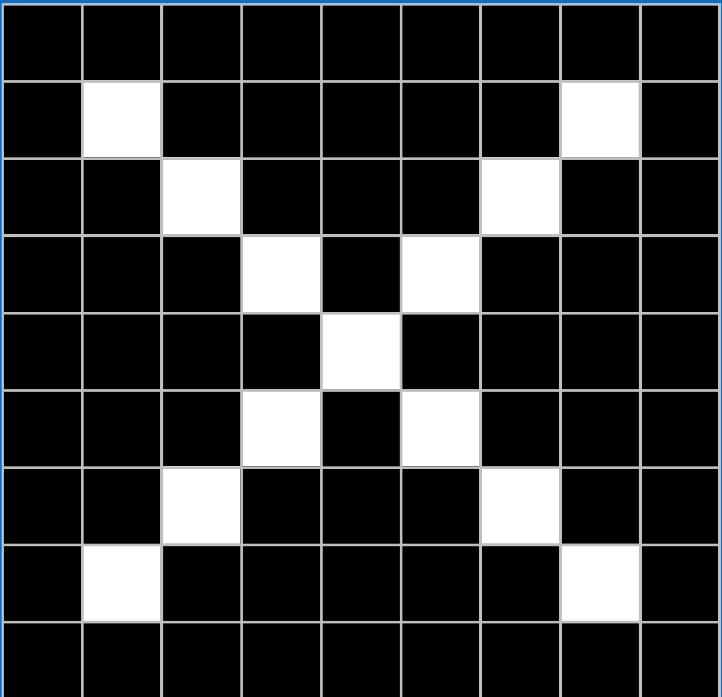
scaling

rotation

weight



# Deciding is hard



# What computers see



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	-1	1	1	-1	-1
-1	-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

# What computers see

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	X	-1	-1	-1	-1	X	X	-1
-1	X	X	-1	-1	X	X	-1	-1
-1	-1	X	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	X	-1	-1
-1	-1	X	X	-1	-1	X	X	-1
-1	X	X	-1	-1	-1	-1	X	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

# Computers are literal

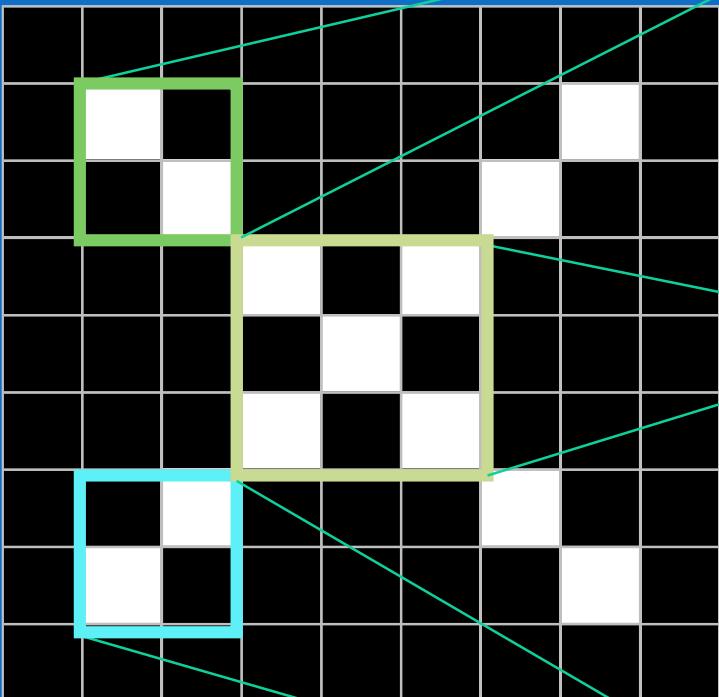
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



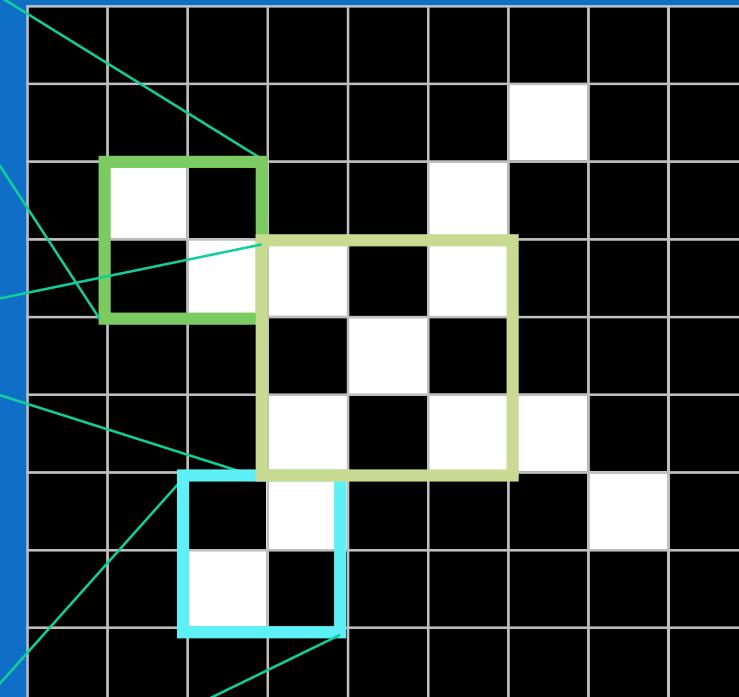
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	1	1	-1	-1
-1	-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

# ConvNets match pieces of the image

=



=



=

# Features match pieces of the image

1	-1	-1
-1	1	-1
-1	-1	1

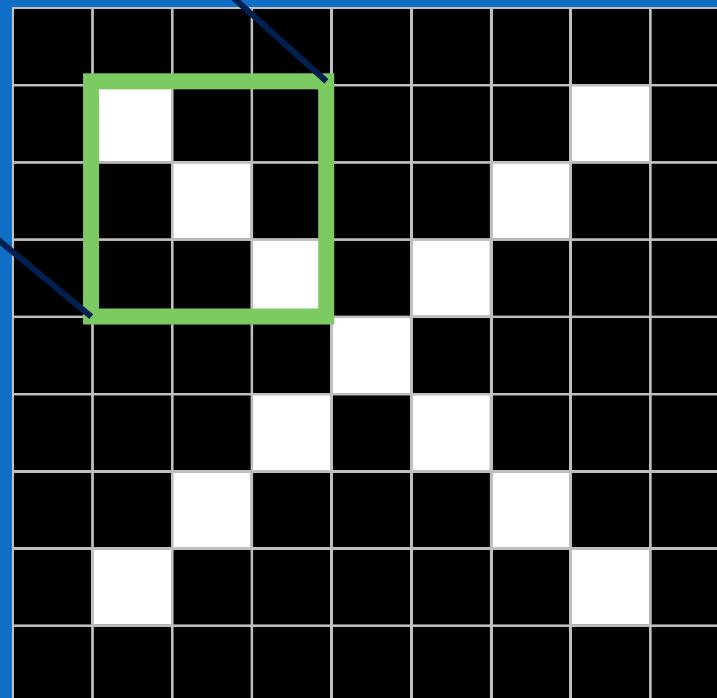
1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1

1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

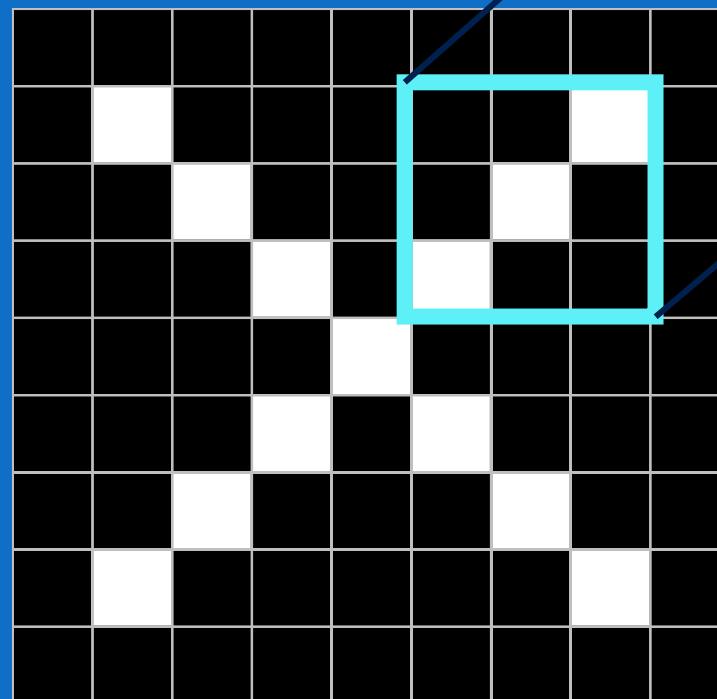
-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

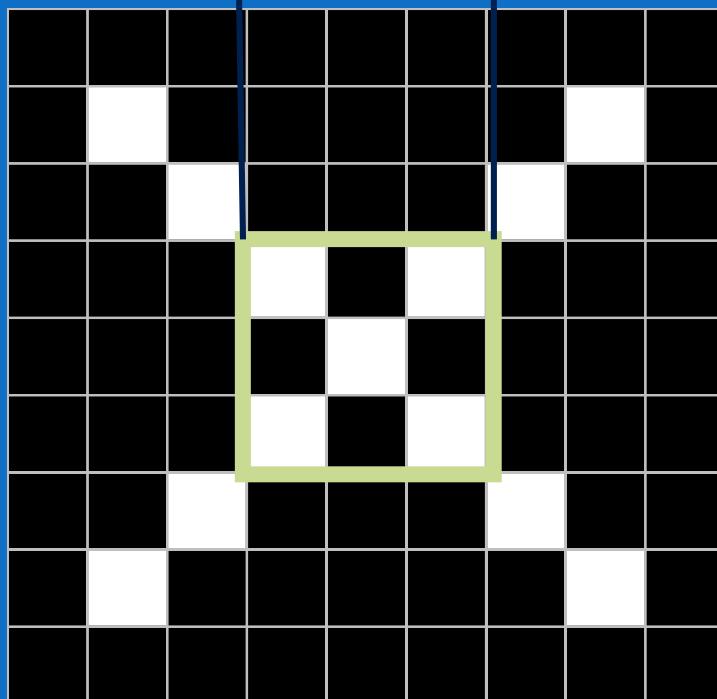
-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

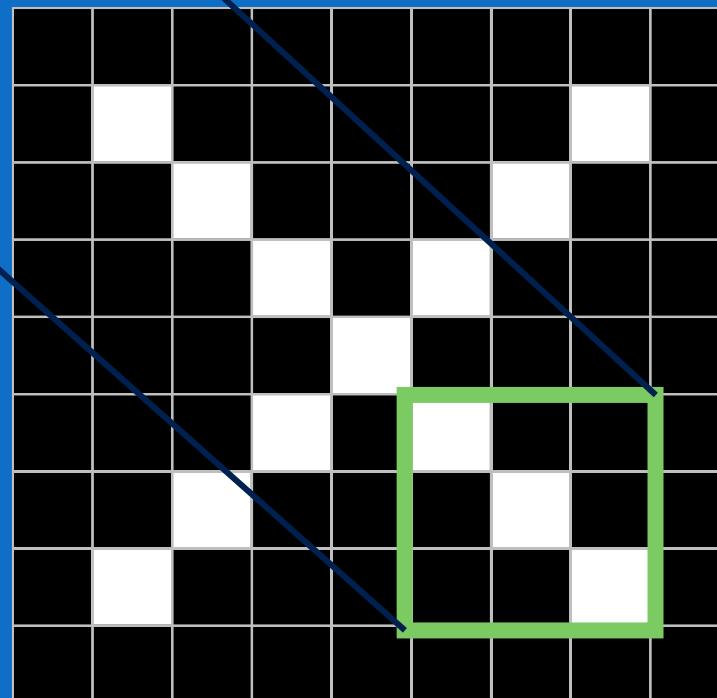
-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

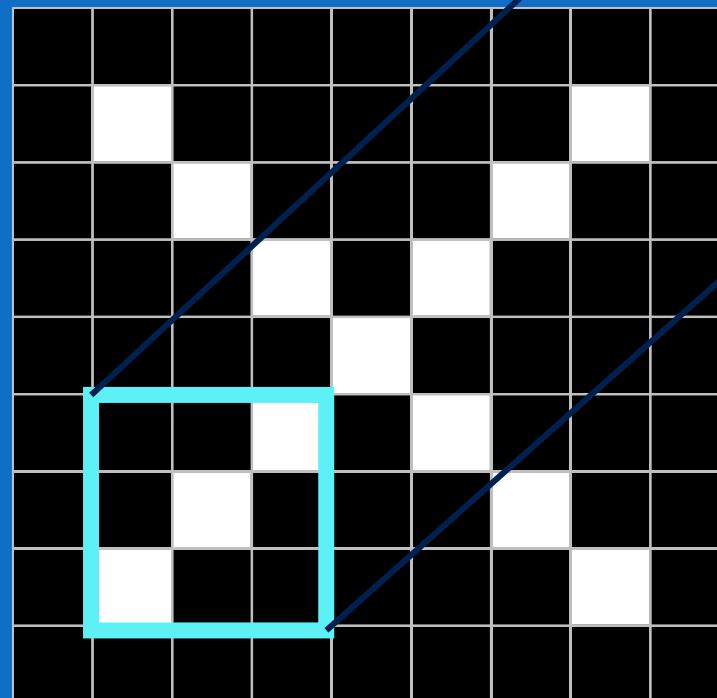
-1	-1	1
-1	1	-1
1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

1	-1	1
-1	1	-1
1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1



# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

**DON'T  
PANIC**

# Filtering: The math behind the match

1. Line up the feature and the image patch.
2. Multiply each image pixel by the corresponding feature pixel.
3. Add them up.
4. Divide by the total number of pixels in the feature.

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1		

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$-1 \times -1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$-1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$-1 \times -1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1		

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$-1 \times -1 = 1$$

1	1	1
1	1	1
1	1	1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$-1 \times -1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1		

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$[-1] \times [1] = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	1	-1	1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	1
1	1	1
1	1	1

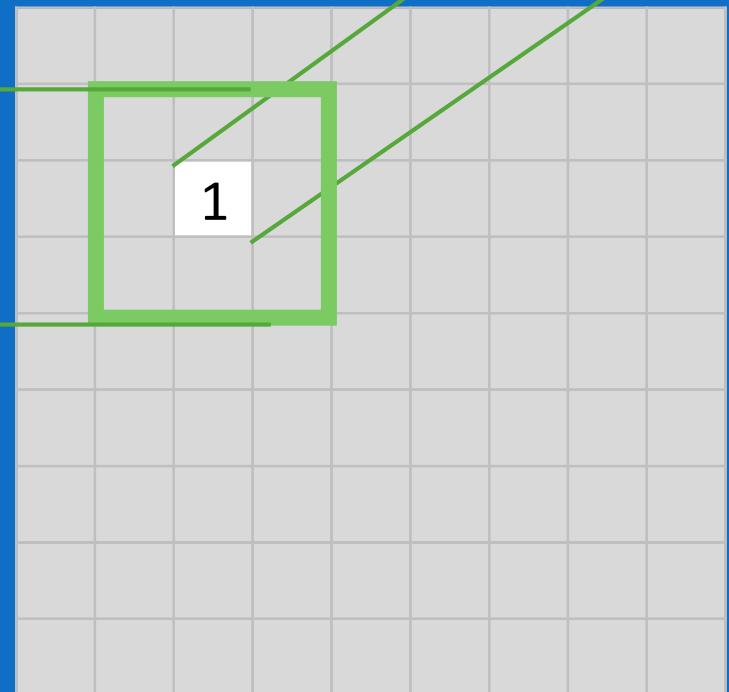
# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

1	1	1
1	1	1
1	1	1

$$\frac{1 + 1 + 1 + 1 + 1 + 1 + 1 + 1}{9} = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

$$1 \times 1 = 1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1		

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

-1
----

1
---

$$-1 \times 1 = -1$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	-1

# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

1	1	-1
1	1	1
-1	1	1

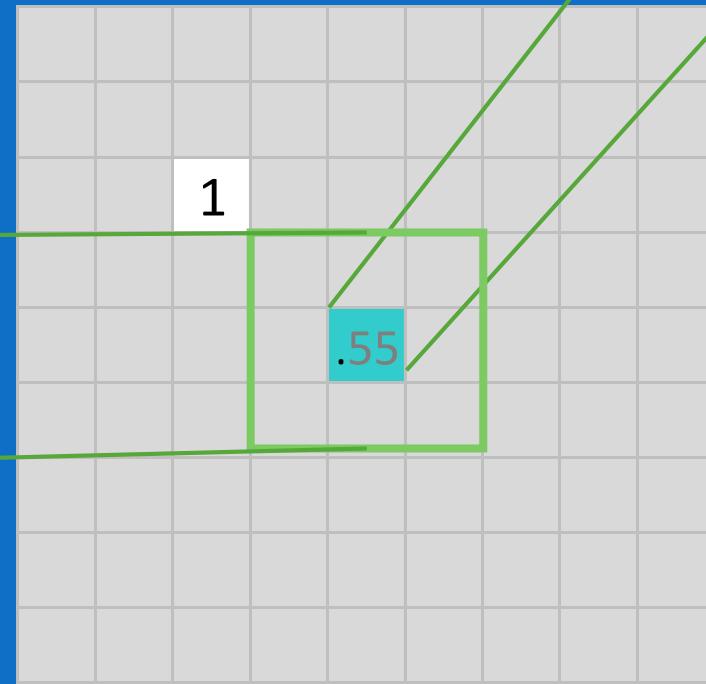
# Filtering: The math behind the match

1	-1	-1
-1	1	-1
-1	-1	1

1	1	-1
1	1	1
-1	1	1

$$\frac{1 + 1 - 1 + 1 + 1 + 1 - 1 + 1 + 1}{9} = .55$$

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



# Convolution: Trying every possible match

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

# Convolution: Trying every possible match

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	-1	-1	1	-1	-1	-1	-1	
-1	-1	-1	-1	1	-1	-1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	-1	-1	-1	-1	-1	-1	-1	



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	1
-1	1	-1
1	-1	1

=

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



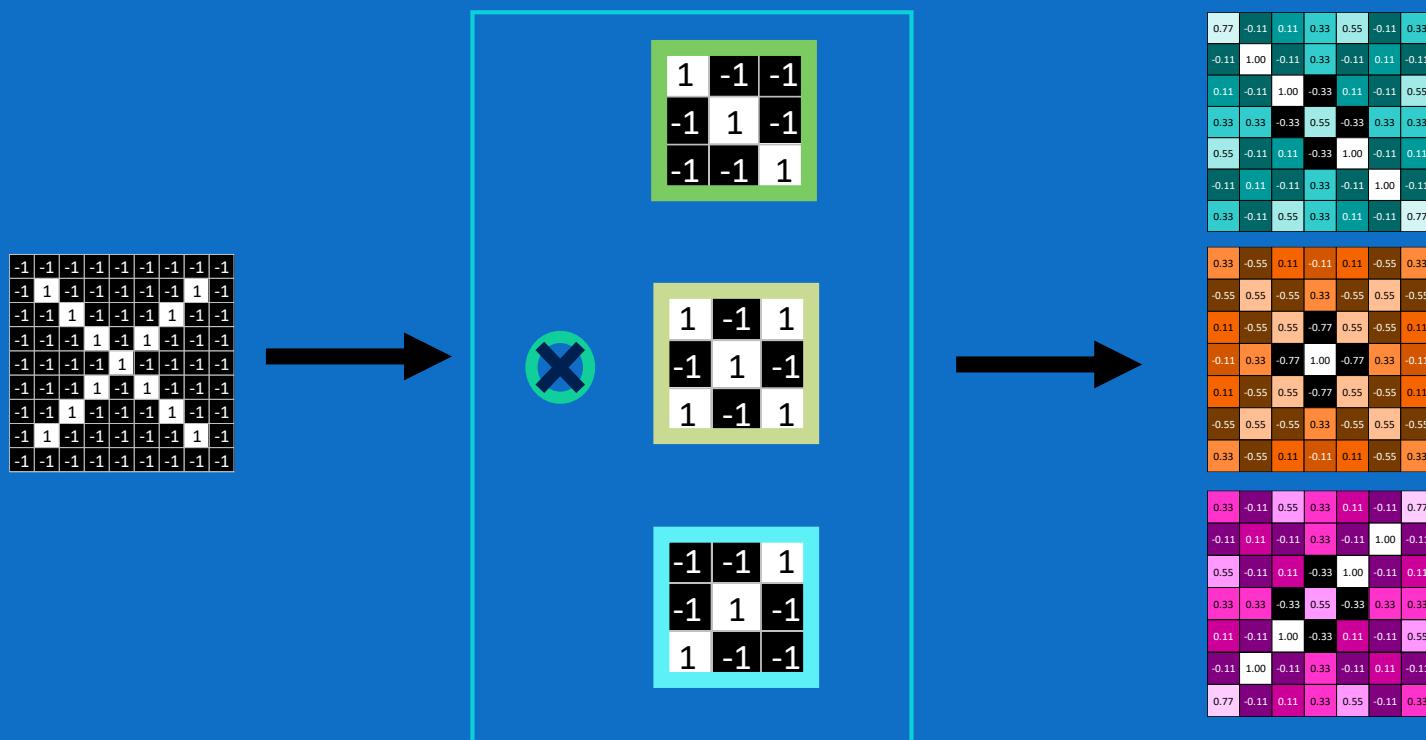
-1	-1	1
-1	1	-1
1	-1	-1

=

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

# Convolution layer

One image becomes a stack of filtered images



# Convolution layer

One image becomes a stack of filtered images

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	-1	-1	1	-1	-1	-1	-1	
-1	-1	-1	-1	-1	1	-1	-1	-1	
-1	-1	-1	-1	-1	-1	1	-1	-1	
-1	1	-1	-1	-1	-1	-1	1	-1	
-1	-1	-1	-1	-1	-1	-1	-1	1	
-1	-1	-1	-1	-1	-1	-1	-1	-1	1



0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

# Pooling: Shrinking the image stack

1. Pick a window size (usually 2 or 3).
2. Pick a stride (usually 2).
3. Walk your window across your filtered images.
4. From each window, take the maximum value.

# Pooling

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

maximum

1.00			

# Pooling

maximum

0.77	-0.11	0.11	0.33	0.55	0.11	0.33	
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11	
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55	
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33	
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11	
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11	
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77	

1.00	0.33		

# Pooling

maximum

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	0.33	0.55	

# Pooling

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33	
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11	
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55	
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33	
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11	
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11	
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77	

maximum

1.00	0.33	0.55	0.33

# Pooling

maximum

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

1.00	0.33	0.55	0.33
0.33			

# Pooling

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

max pooling

1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33



0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33



0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33

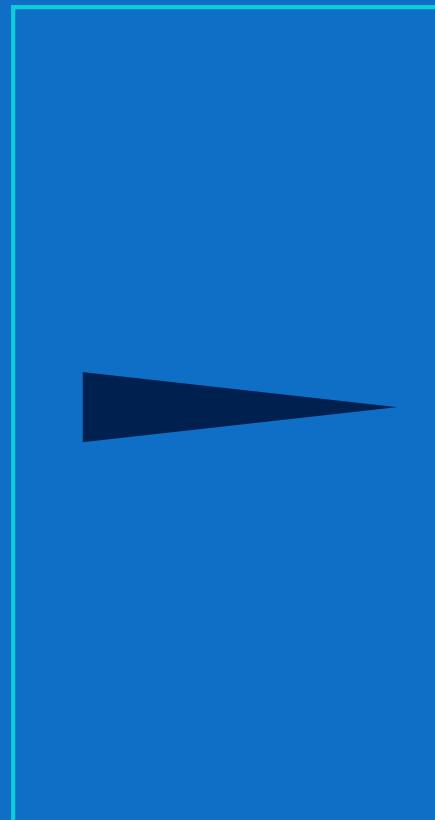
# Pooling layer

A stack of images becomes a stack of smaller images.

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

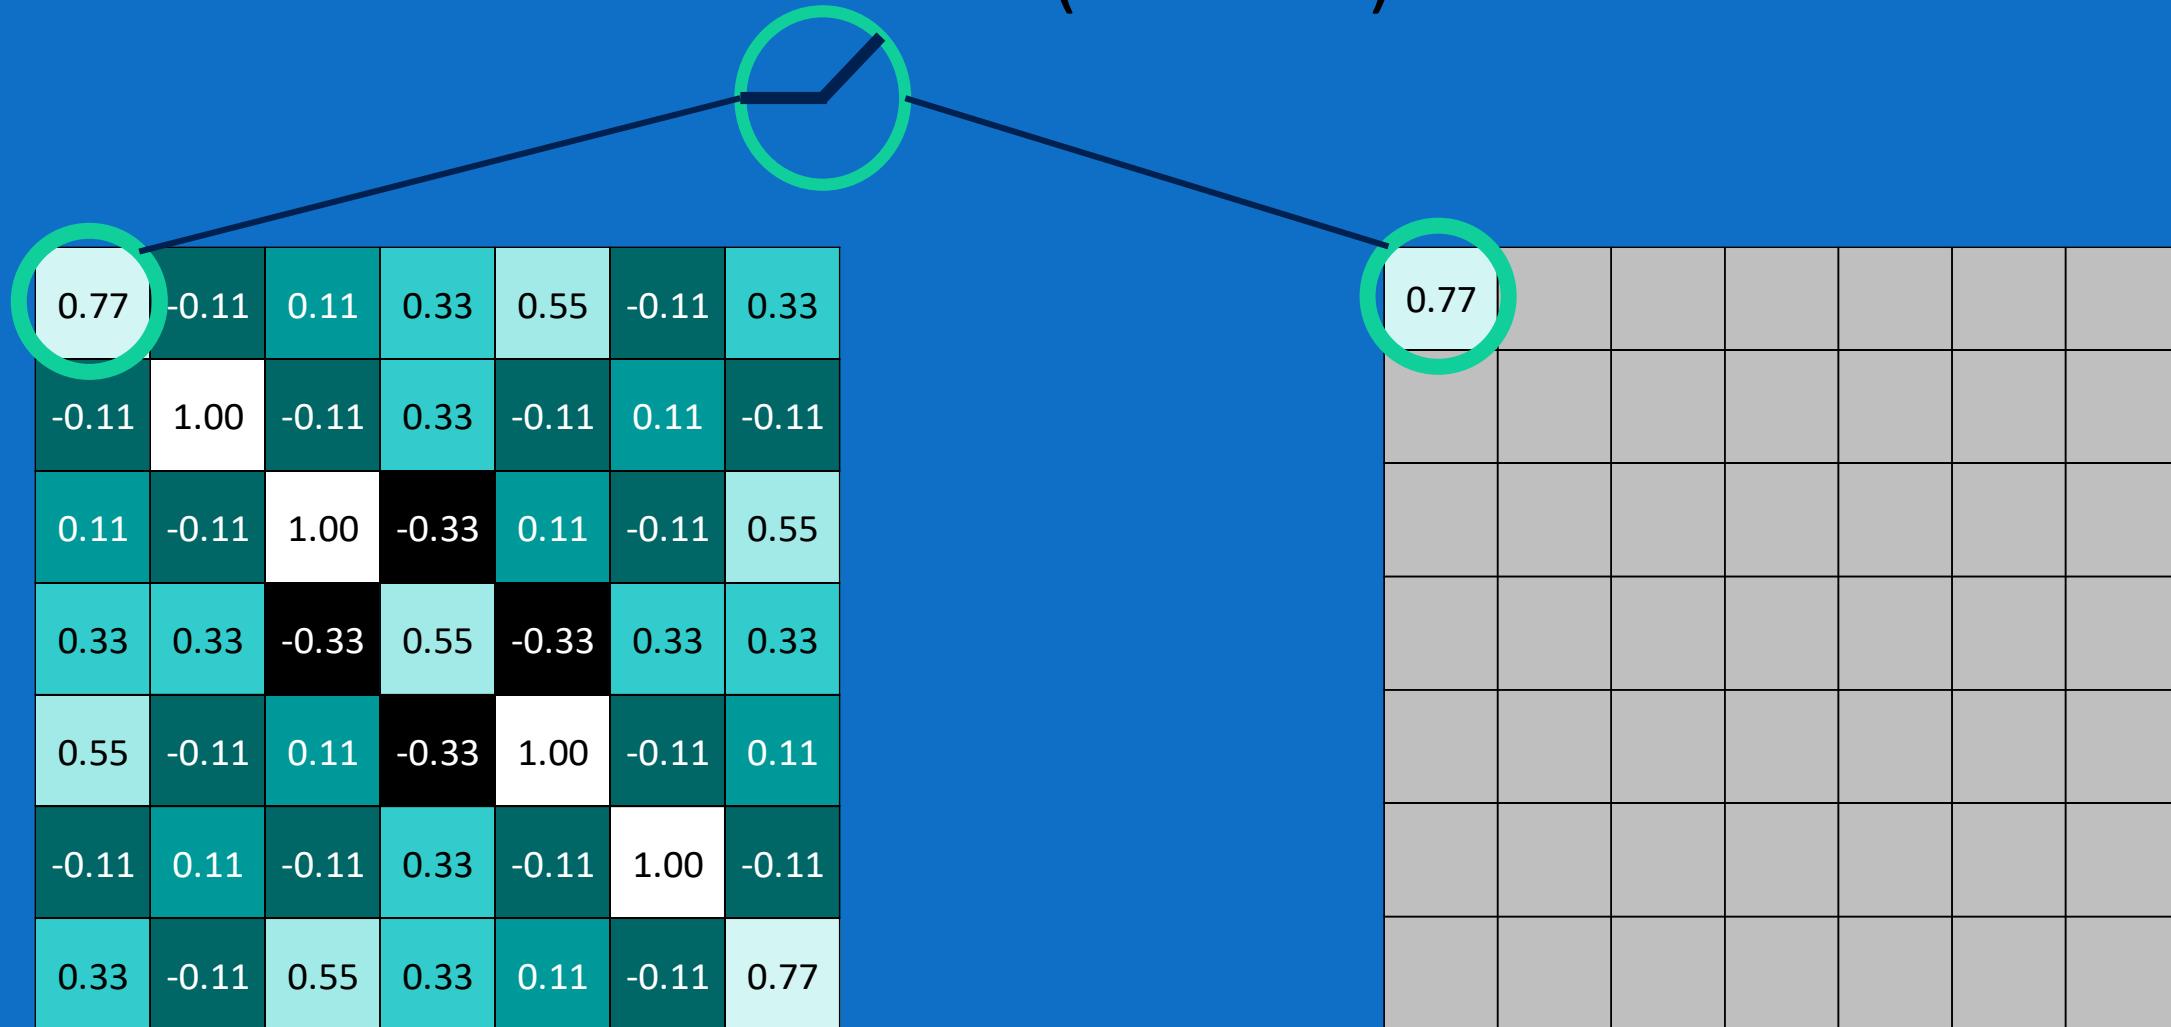
0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33

# Normalization

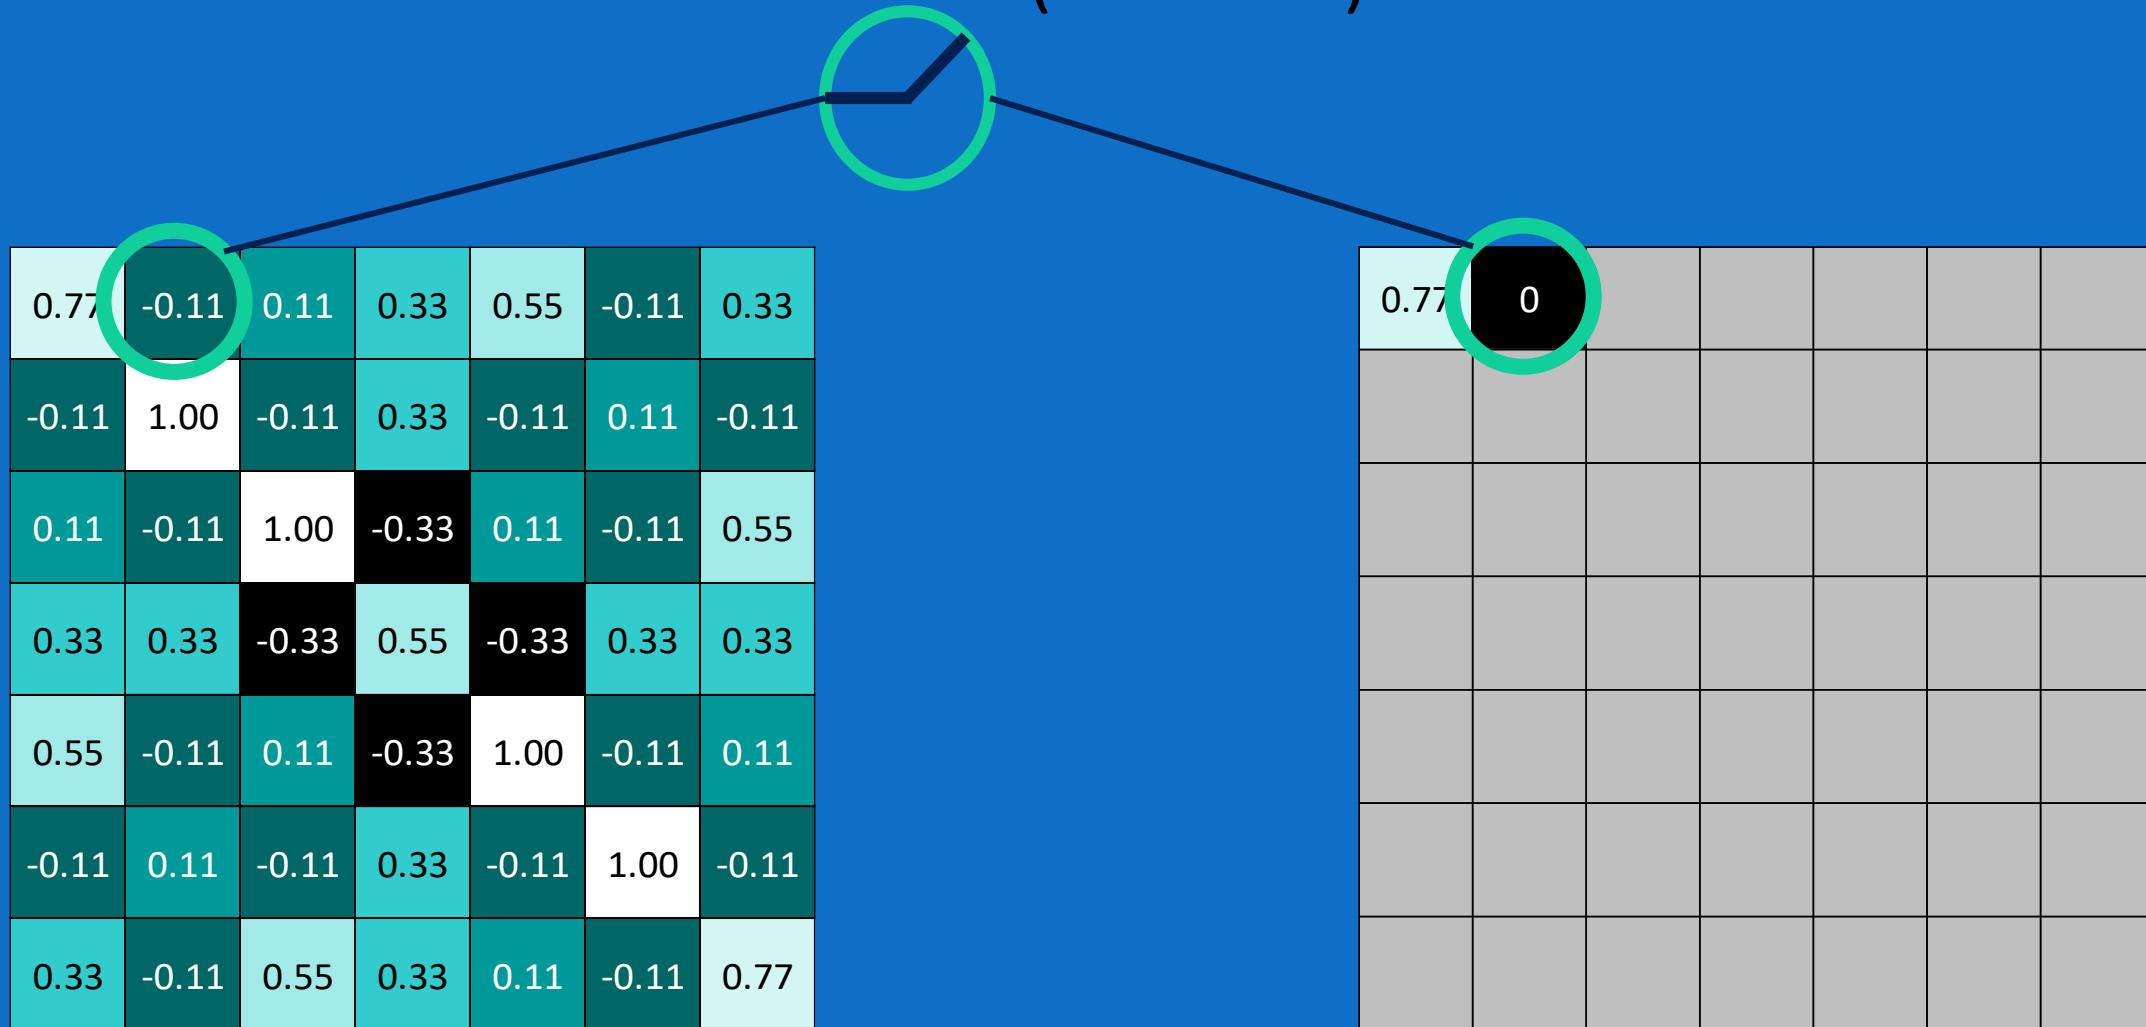
Keep the math from breaking by tweaking each of the values just a bit.

Change everything negative to zero.

# Rectified Linear Units (ReLUs)



# Rectified Linear Units (ReLUs)



# Rectified Linear Units (ReLUs)

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.77	0	0.11	0.33	0.55	0	0.33

# Rectified Linear Units (ReLUs)

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77



0.77	0	0.11	0.33	0.55	0	0.33
0	1.00	0	0.33	0	0.11	0
0.11	0	1.00	0	0.11	0	0.55
0.33	0.33	0	0.55	0	0.33	0.33
0.55	0	0.11	0	1.00	0	0.11
0	0.11	0	0.33	0	1.00	0
0.33	0	0.55	0.33	0.11	0	0.77

# ReLU layer

A stack of images becomes a stack of images with no negative values.

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33



0.77	0	0.11	0.33	0.55	0	0.33
0	1.00	0	0.33	0	0.11	0
0.11	0	1.00	0	0.11	0	0.55
0.33	0.33	0	0.55	0	0.33	0.33
0.55	0	0.11	0	1.00	0	0.11
0	0.11	0	0.33	0	1.00	0
0.33	0	0.55	0.33	0.11	0	0.77

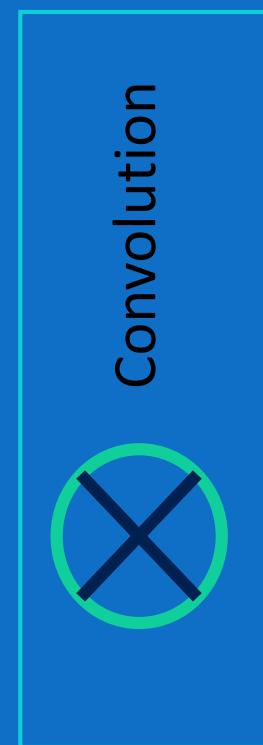
0.33	0	0.11	0	0.11	0	0.33
0	0.55	0	0.33	0	0.55	0
0.11	0	0.55	0	0.55	0	0.11
0	0.33	0	1.00	0	0.33	0
0.11	0	0.55	0	0.55	0	0.11
0	0.55	0	0.33	0	0.55	0
0.33	0	0.11	0	0.11	0	0.33

0.33	0	0.55	0.33	0.11	0	0.77
0	0.11	0	0.33	0	1.00	0
0.55	0	0.11	0	1.00	0	0.11
0.33	0.33	0	0.55	0	0.33	0.33
0.11	0	1.00	0	0.11	0	0.55
0	1.00	0	0.33	0	0.11	0
0.77	0	0.11	0.33	0.55	0	0.33

# Layers get stacked

The output of one becomes the input of the next.

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

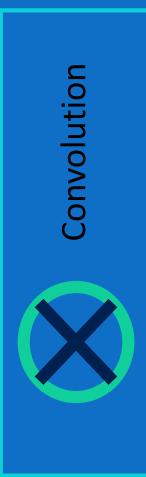
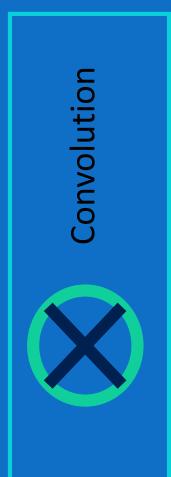
0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33

# Deep stacking

Layers can be repeated several (or many) times.

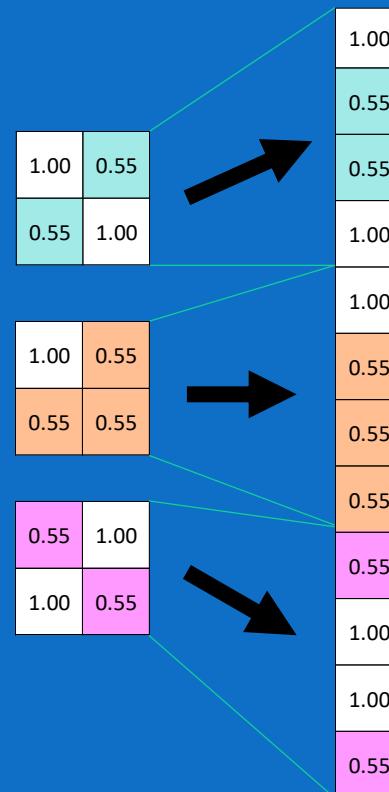
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1



1.00	0.55
0.55	1.00
1.00	0.55
0.55	0.55
0.55	1.00
1.00	0.55

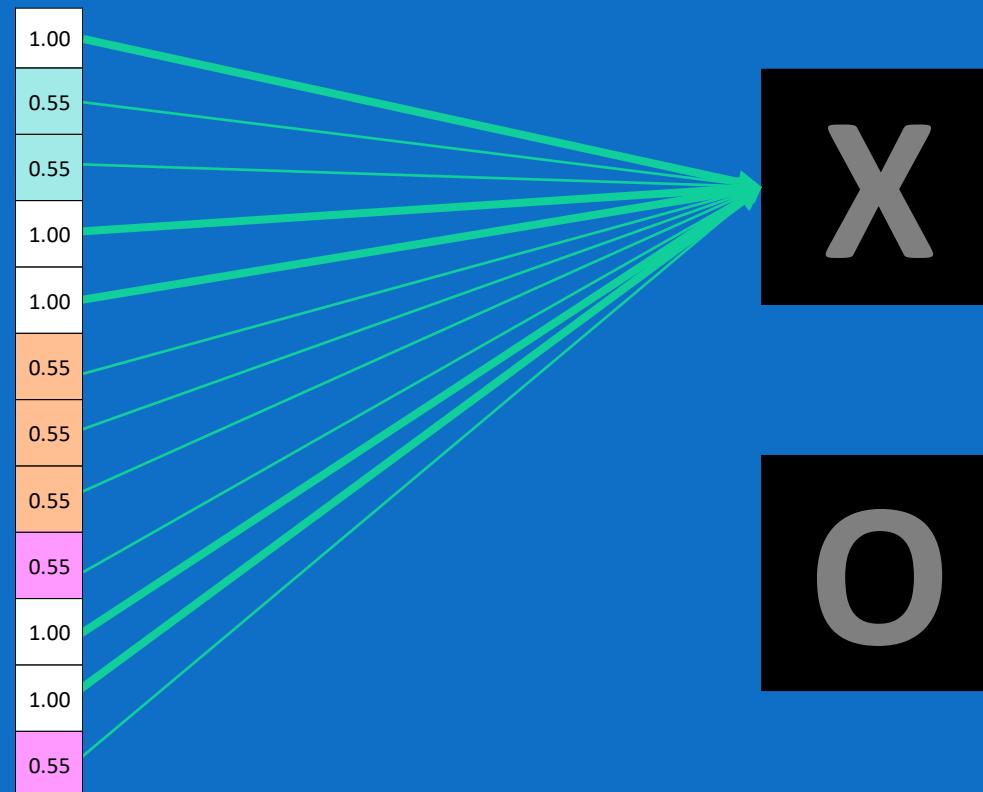
# Fully connected layer

Every value gets a vote



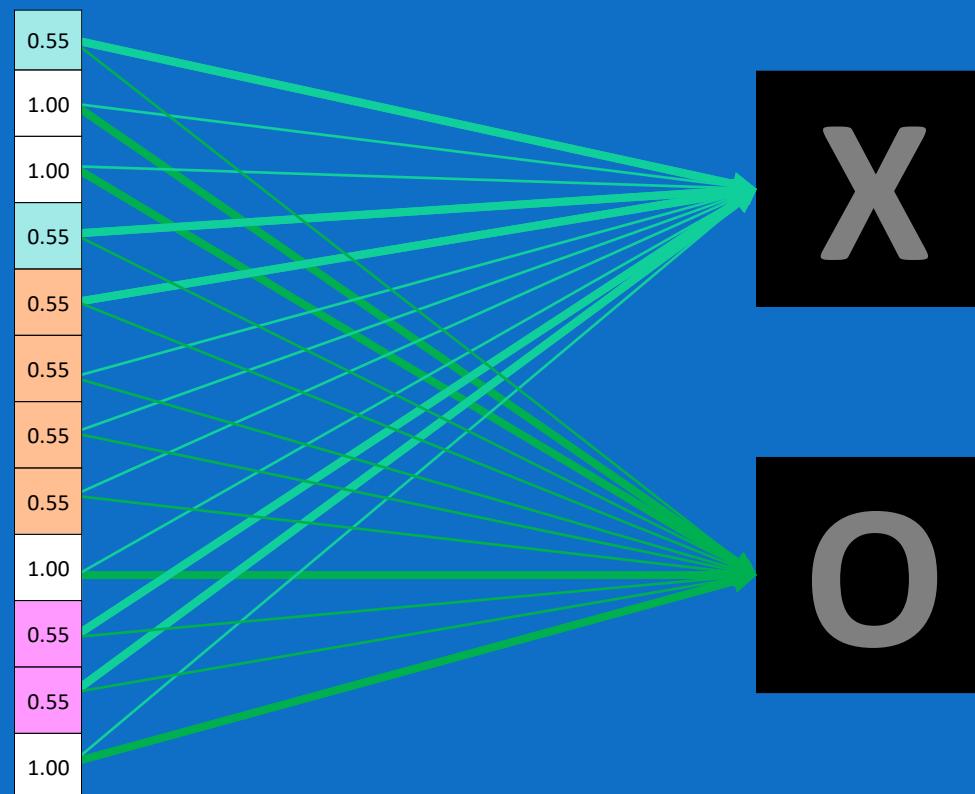
# Fully connected layer

Vote depends on how strongly a value predicts X or O



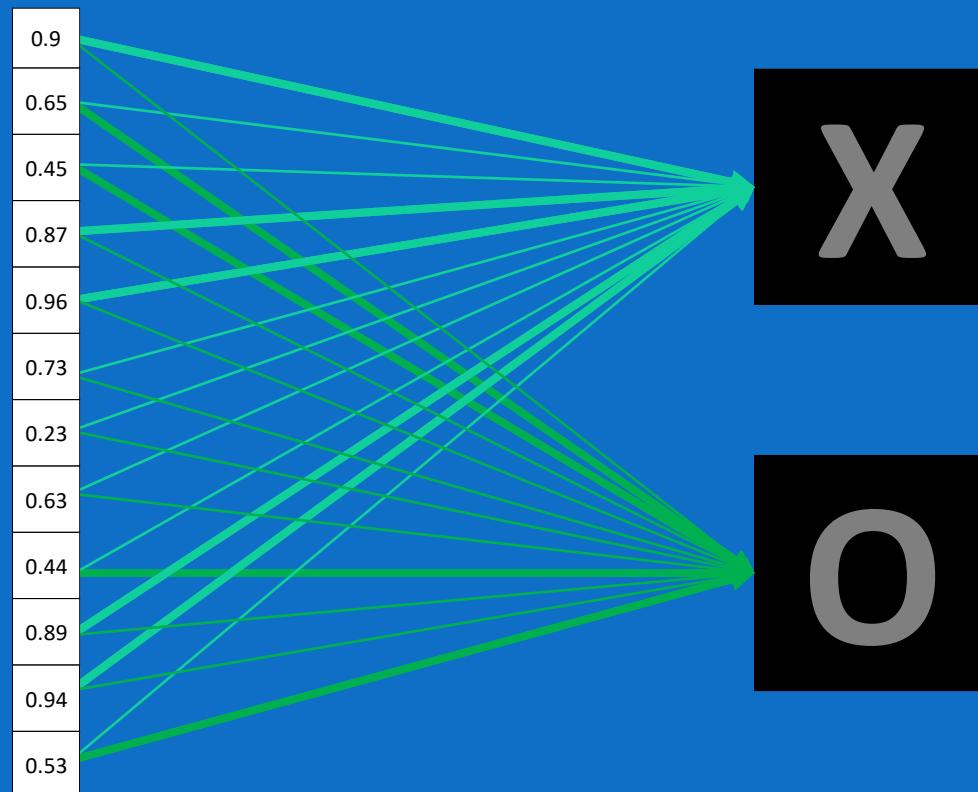
# Fully connected layer

Vote depends on how strongly a value predicts X or O



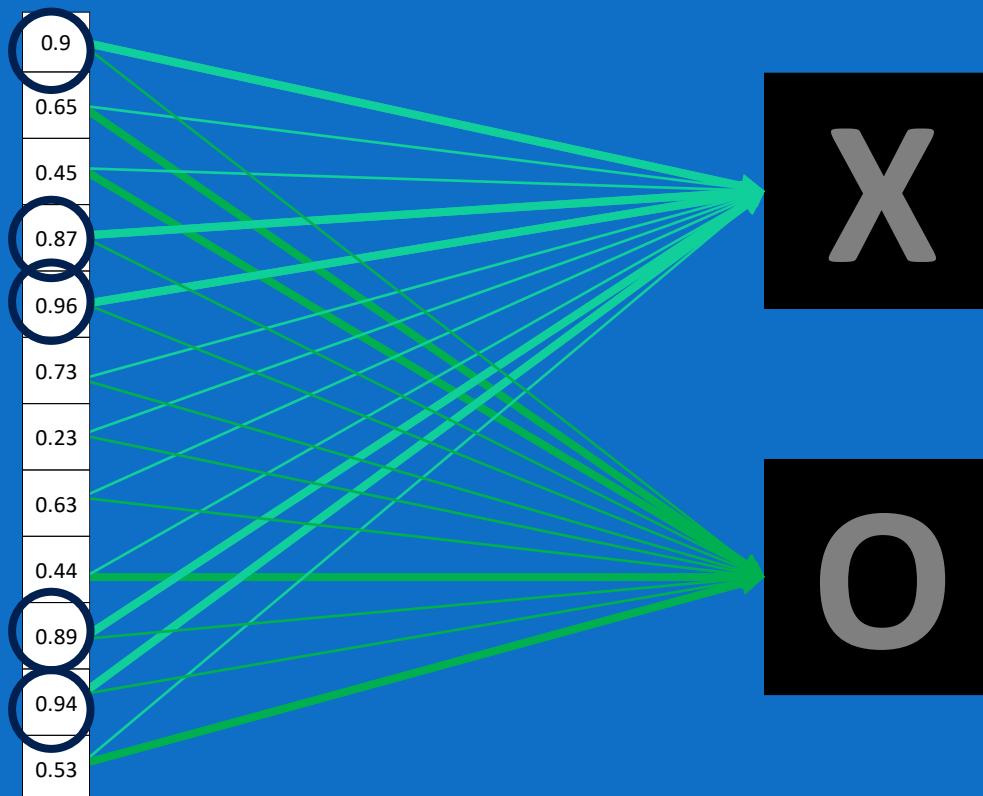
# Fully connected layer

Future values vote on X or O



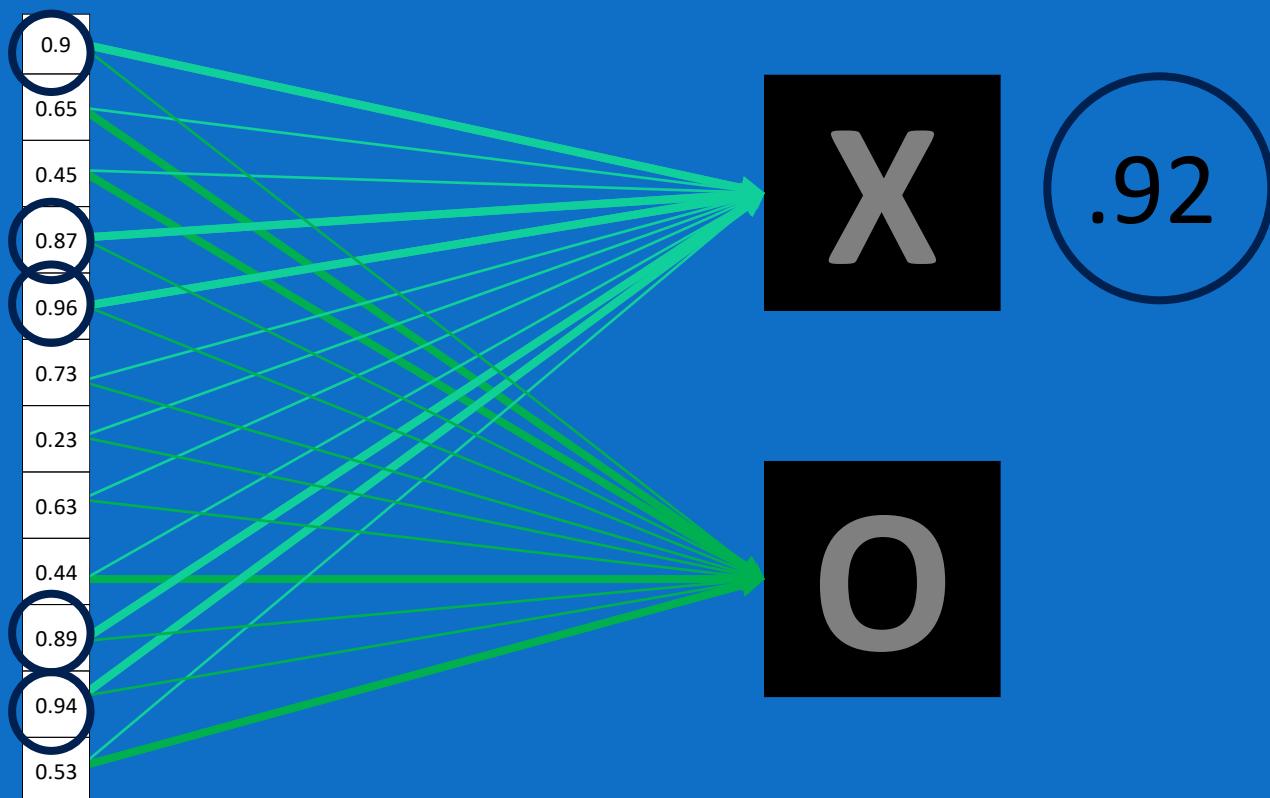
# Fully connected layer

Future values vote on X or O



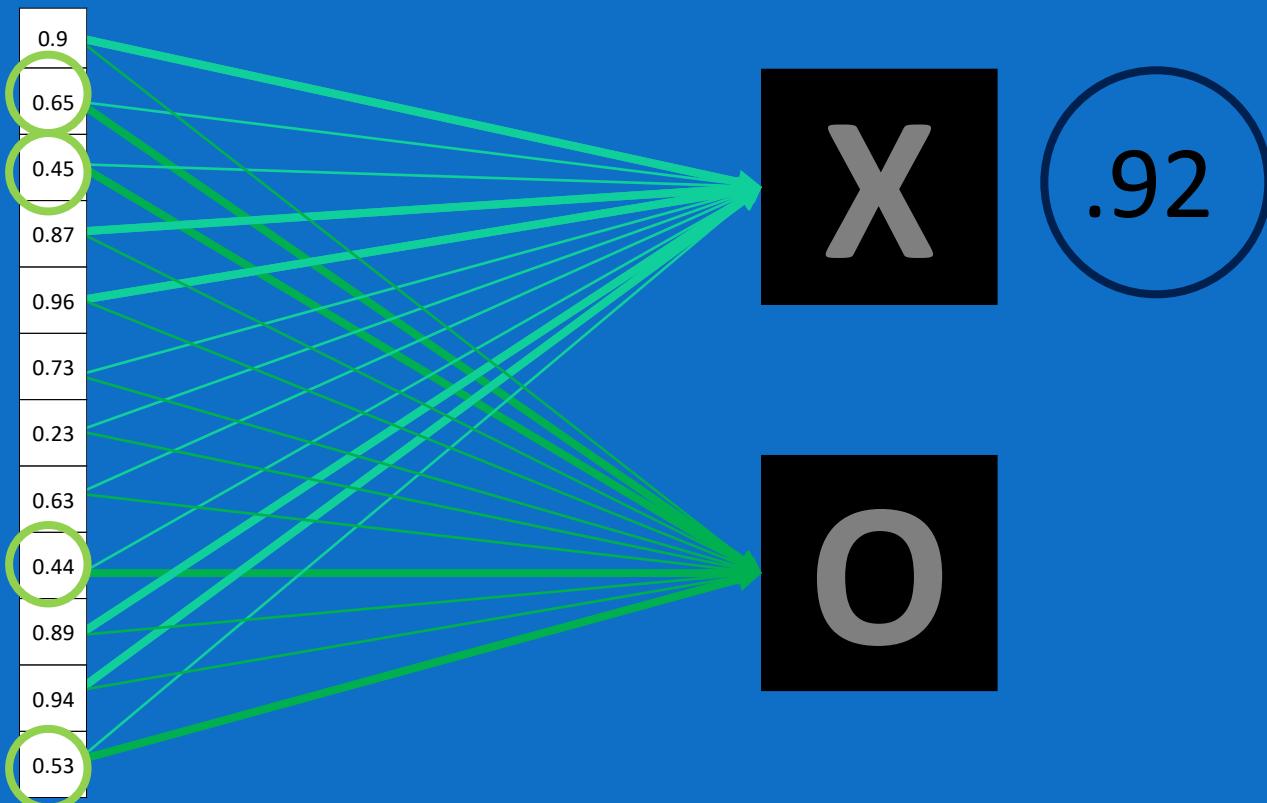
# Fully connected layer

Future values vote on X or O



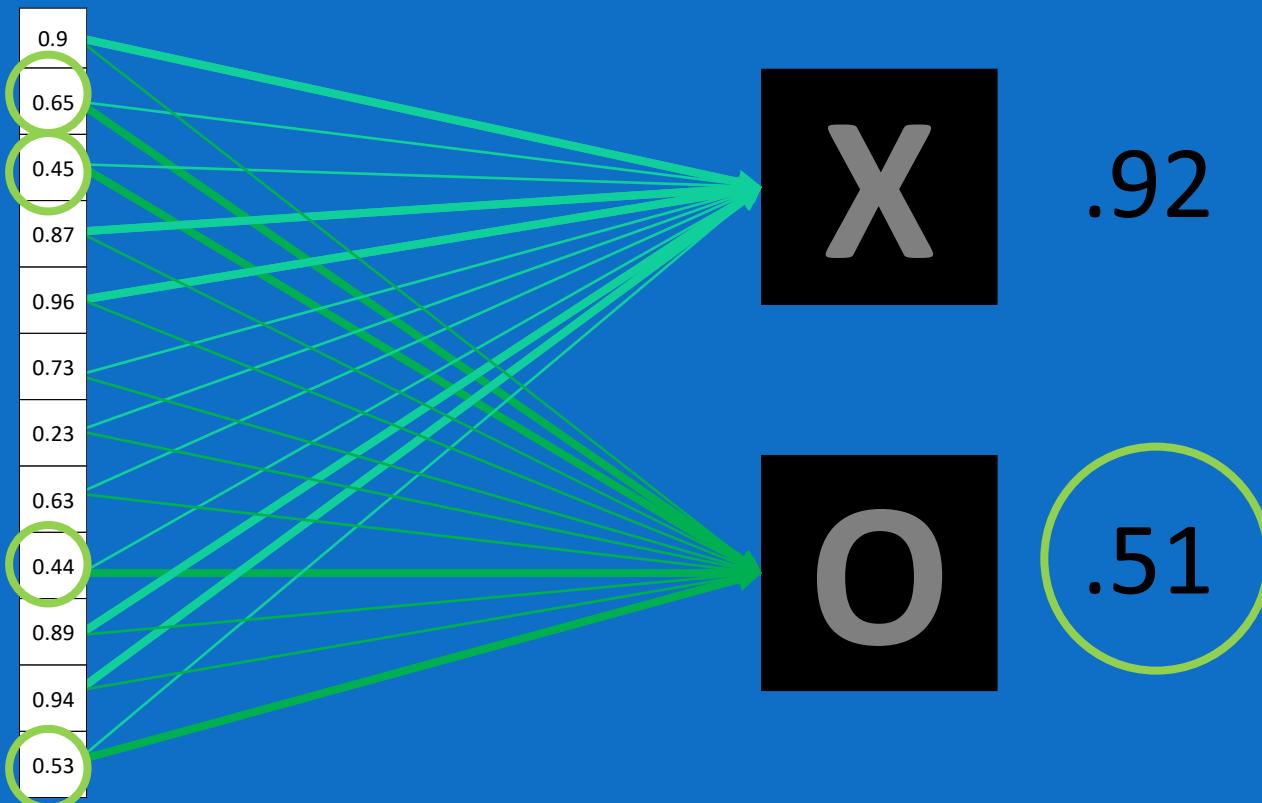
# Fully connected layer

Future values vote on X or O



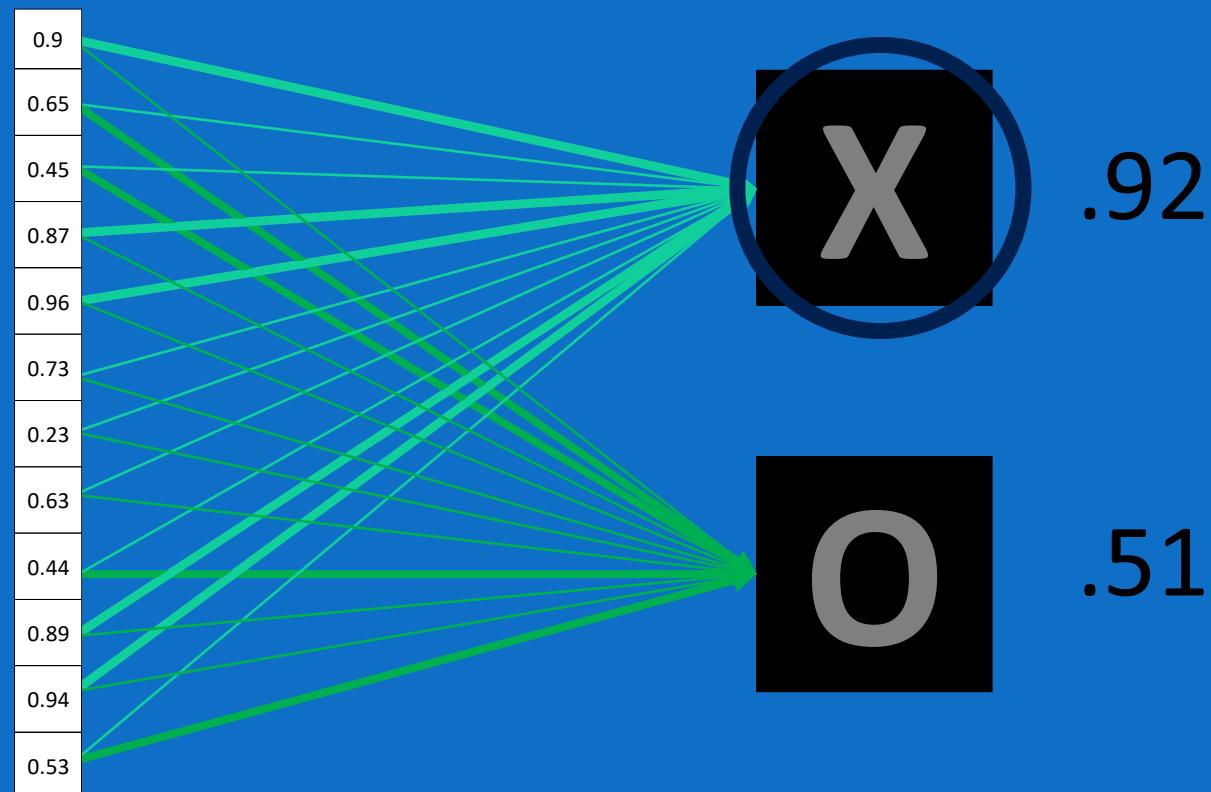
# Fully connected layer

Future values vote on X or O



# Fully connected layer

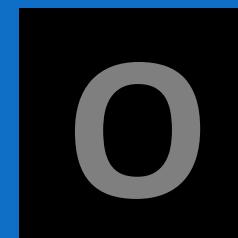
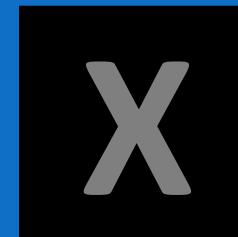
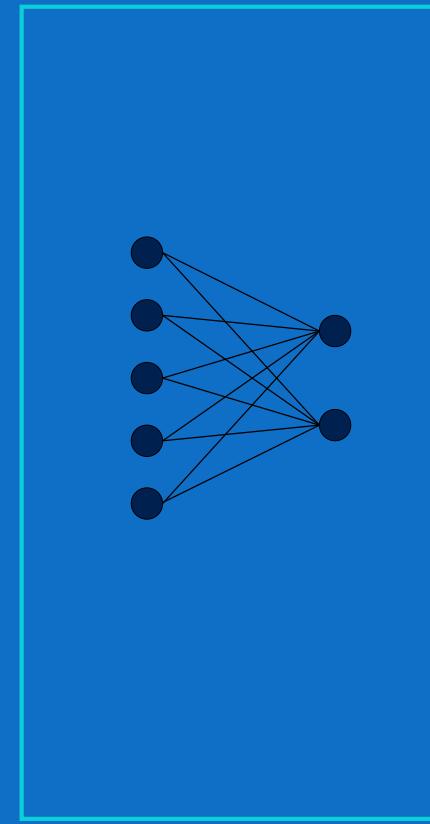
Future values vote on X or O



# Fully connected layer

A list of feature values becomes a list of votes.

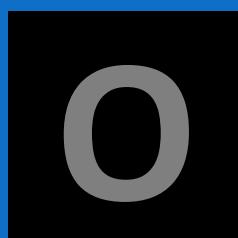
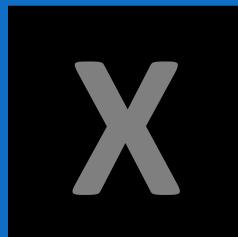
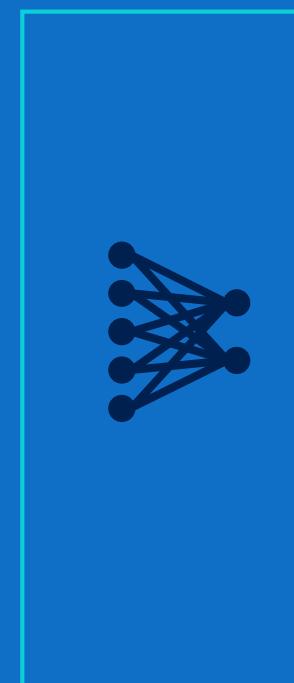
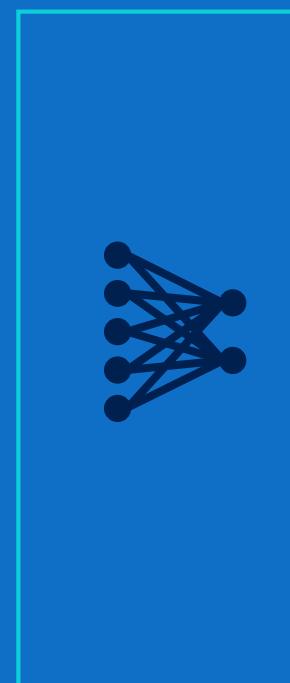
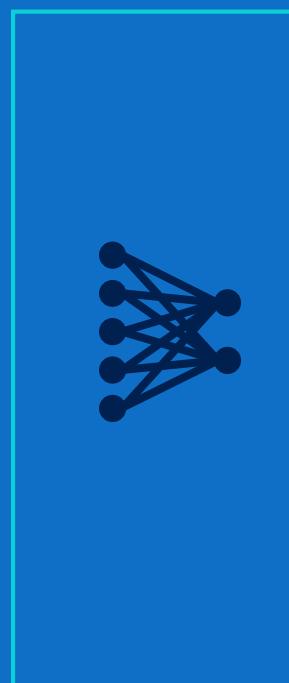
0.9
0.65
0.45
0.87
0.96
0.73
0.23
0.63
0.44
0.89
0.94
0.53



# Fully connected layer

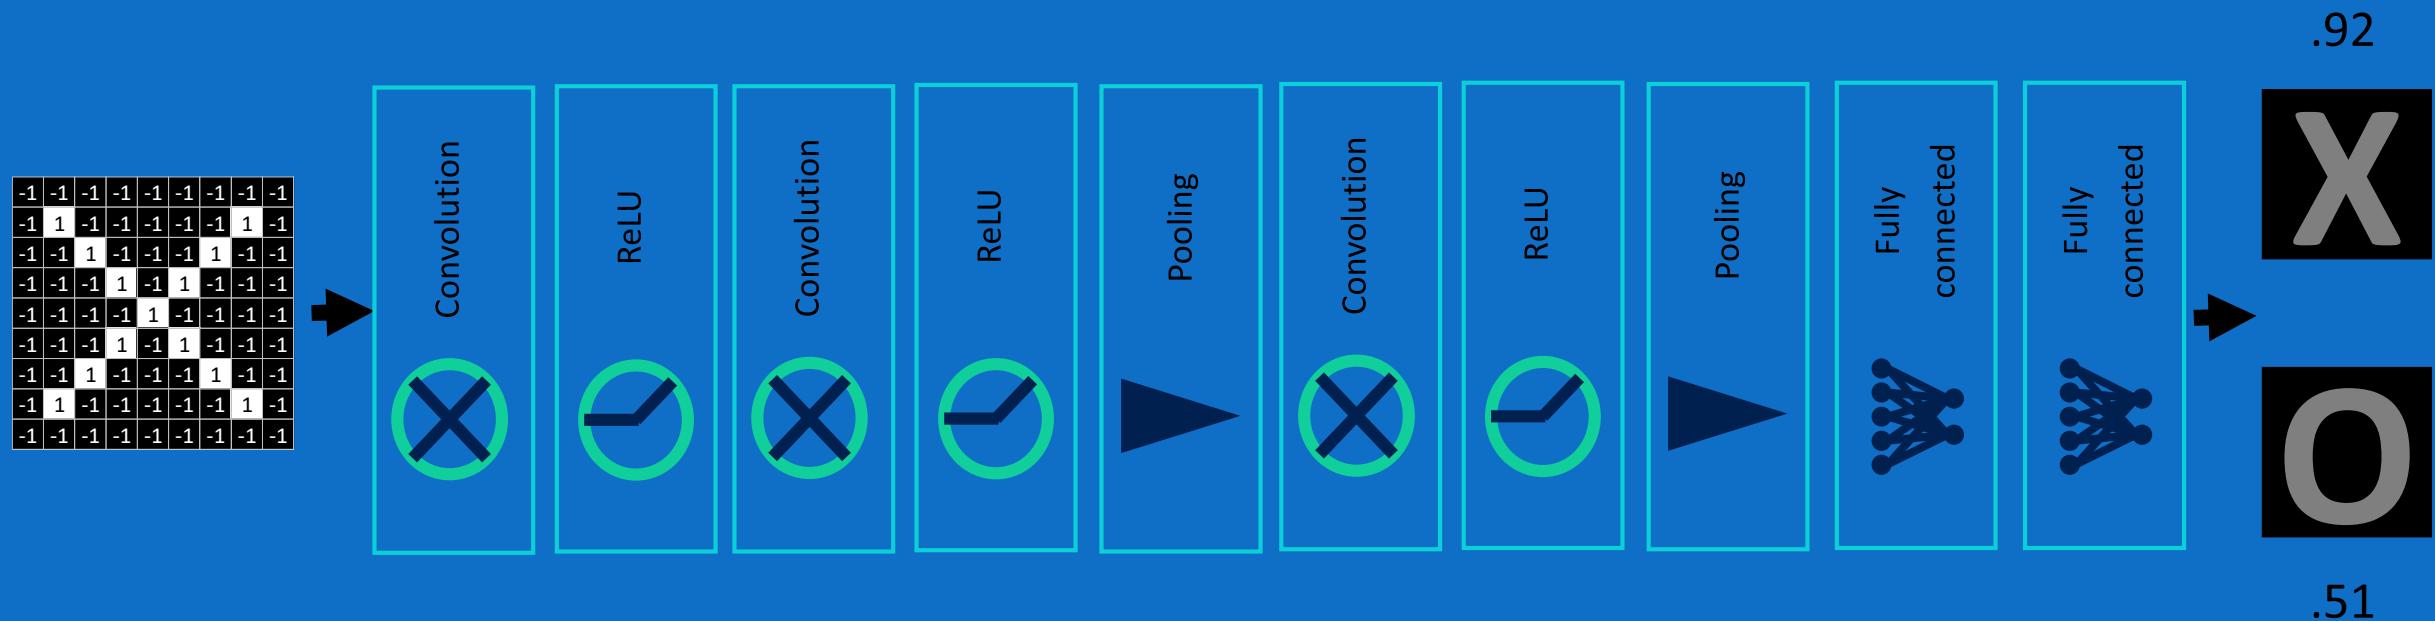
These can also be stacked.

0.9
0.65
0.45
0.87
0.96
0.73
0.23
0.63
0.44
0.89
0.94
0.53



# Putting it all together

A set of pixels becomes a set of votes.



# Learning

Q: Where do all the magic numbers come from?

Features in convolutional layers

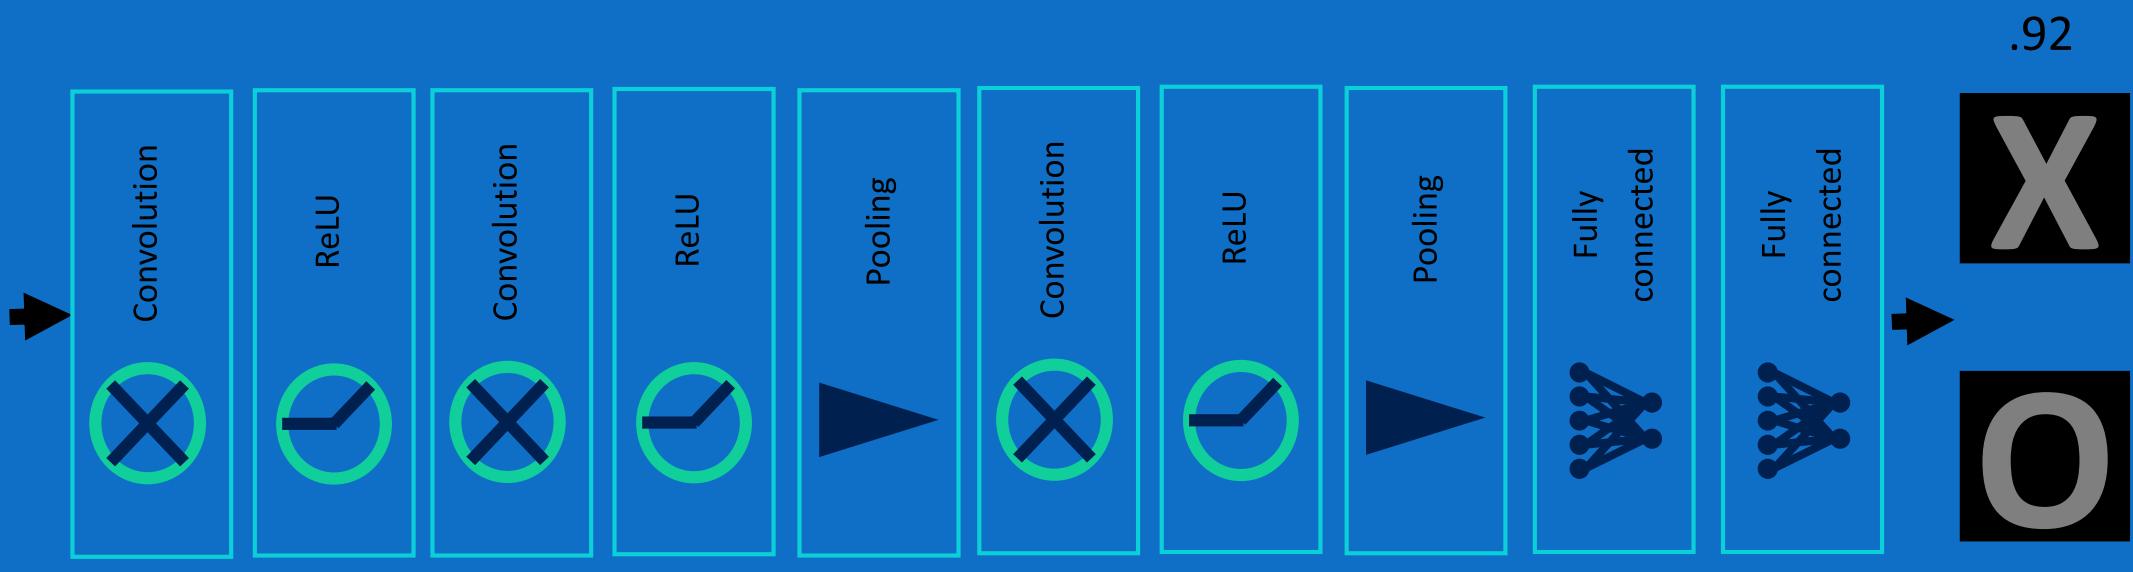
Voting weights in fully connected layers

A: Backpropagation

# Backprop

Error = right answer – actual answer

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	-1	1	1	-1	-1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	-1	1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	-1	1	1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

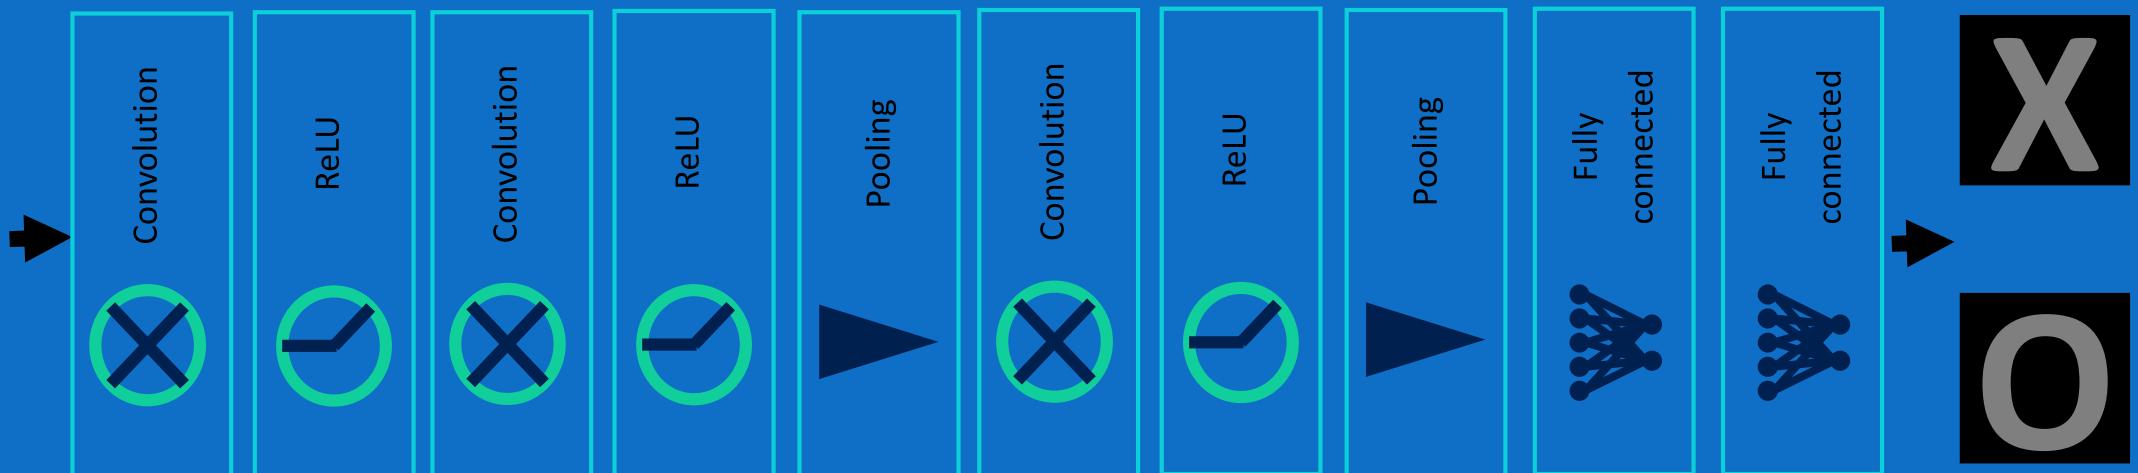


# Backprop

	Right answer	Actual answer	Error
X	1		
O			

.92

-1 -1 -1 -1 -1 -1 -1 -1 -1  
-1 1 -1 -1 -1 -1 -1 1 -1  
-1 -1 1 -1 -1 -1 1 -1 -1  
-1 -1 -1 1 -1 1 -1 -1 -1  
-1 -1 -1 -1 1 -1 -1 -1 -1  
-1 -1 -1 1 -1 -1 -1 -1 -1  
-1 -1 -1 1 -1 1 -1 -1 -1  
-1 -1 1 -1 -1 -1 1 -1 -1  
-1 1 -1 -1 -1 -1 1 1 -1  
-1 -1 -1 -1 -1 -1 -1 -1 -1

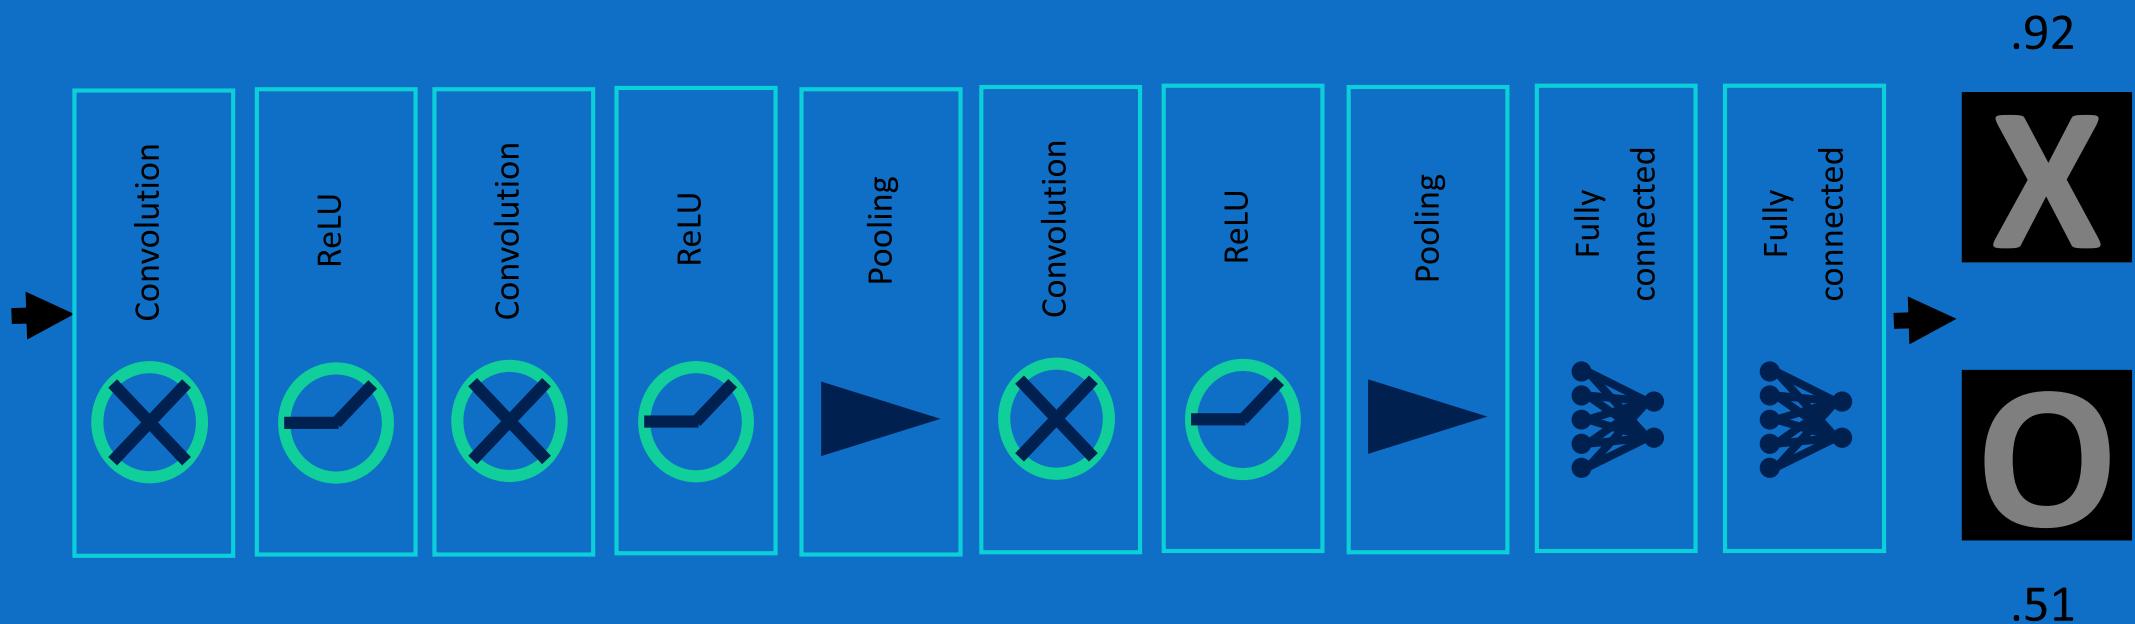


.51

# Backprop

	Right answer	Actual answer	Error
X	1	0.92	
O			

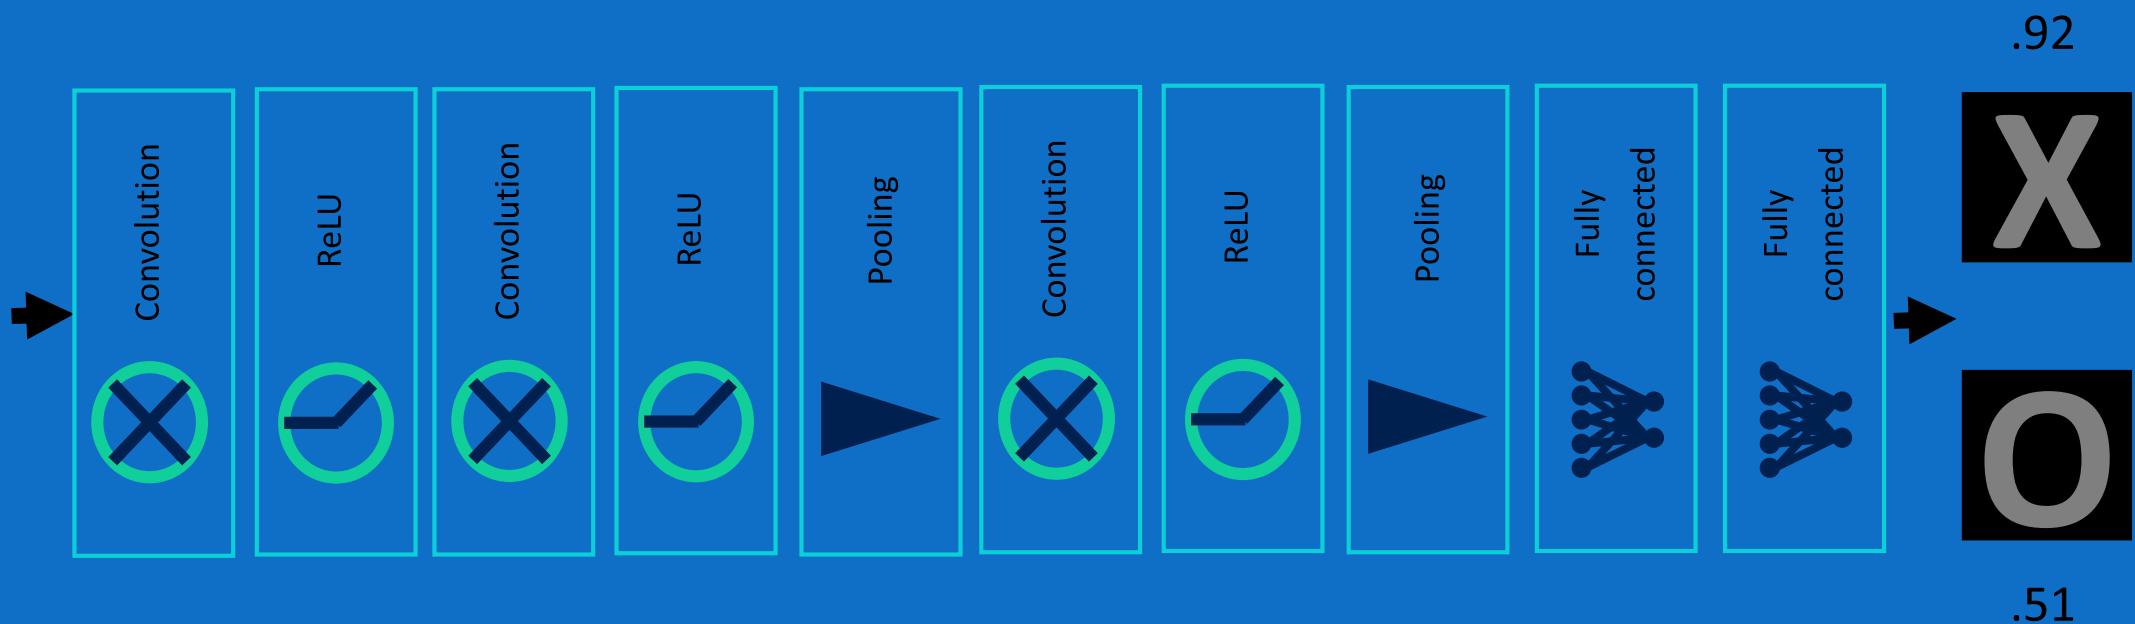
```
-1 -1 -1 -1 -1 -1 -1 -1 -1  
-1 1 -1 -1 -1 -1 -1 1 -1  
-1 -1 1 -1 -1 -1 1 -1 -1  
-1 -1 -1 1 -1 1 -1 -1 -1  
-1 -1 -1 -1 1 -1 -1 -1 -1  
-1 -1 -1 1 -1 -1 -1 -1 -1  
-1 -1 -1 1 -1 1 -1 -1 -1  
-1 -1 1 -1 -1 -1 1 -1 -1  
-1 1 -1 -1 -1 -1 1 1 -1  
-1 -1 -1 -1 -1 -1 -1 -1 -1
```



# Backprop

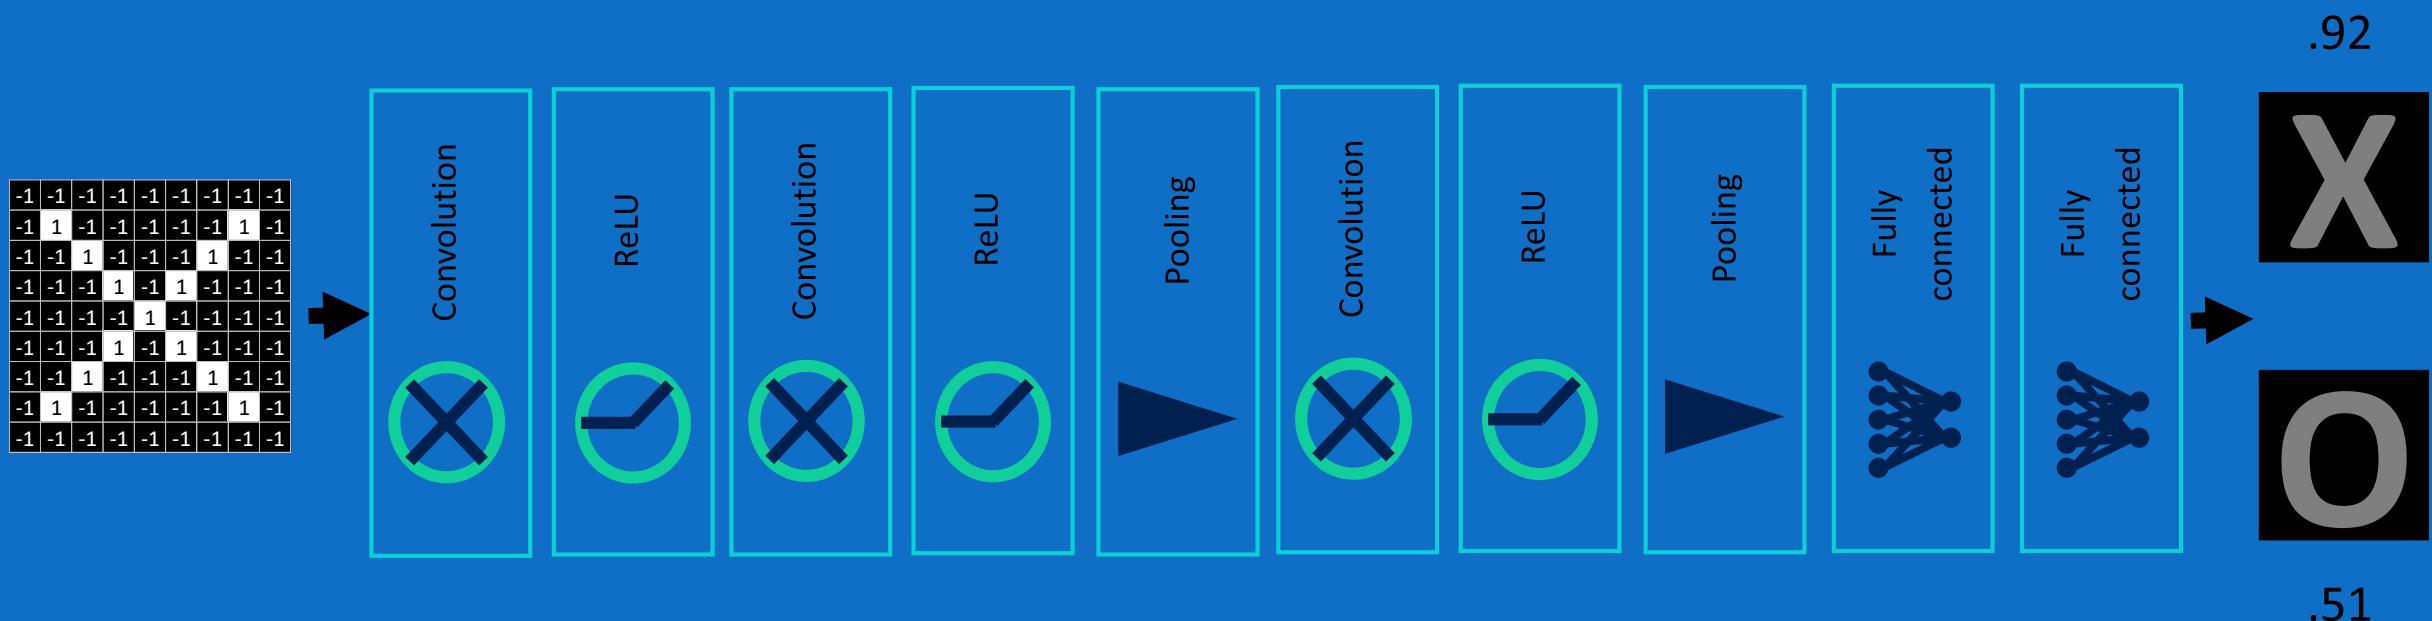
	Right answer	Actual answer	Error
X	1	0.92	0.08
O			

-1 -1 -1 -1 -1 -1 -1 -1 -1  
-1 1 -1 -1 -1 -1 -1 1 -1  
-1 -1 1 -1 -1 -1 1 -1 -1  
-1 -1 -1 1 -1 1 -1 -1 -1  
-1 -1 -1 -1 1 -1 -1 -1 -1  
-1 -1 -1 1 -1 -1 -1 -1 -1  
-1 -1 -1 1 -1 1 -1 -1 -1  
-1 -1 1 -1 -1 -1 1 -1 -1  
-1 1 -1 -1 -1 -1 1 -1 -1  
-1 -1 -1 -1 -1 -1 -1 -1 -1



# Backprop

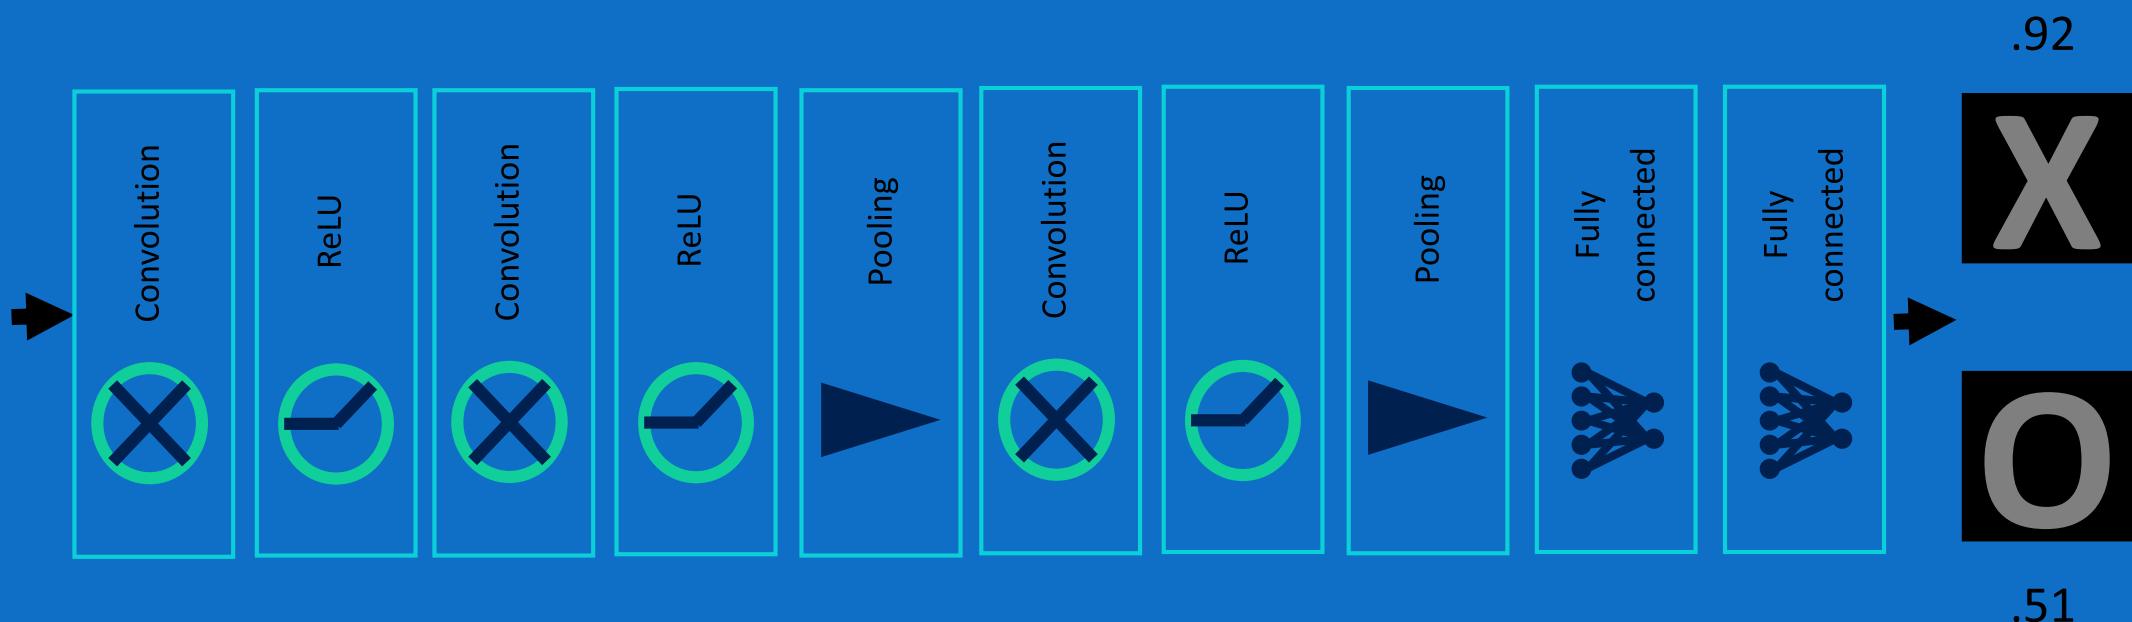
	Right answer	Actual answer	Error
X	1	0.92	0.08
O	0	0.51	0.49



# Backprop

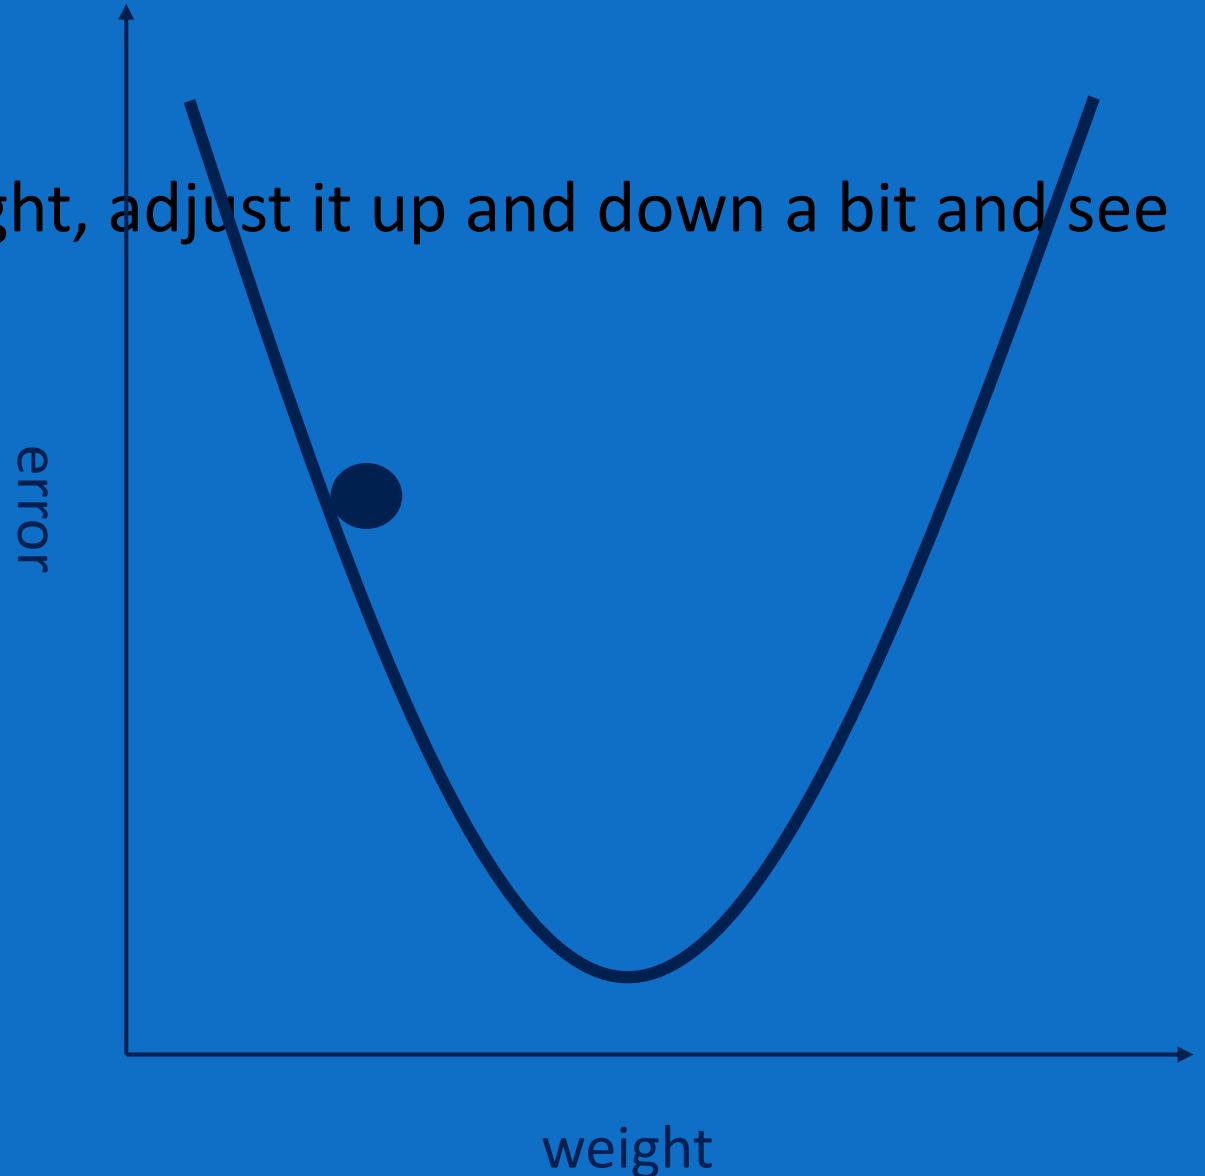
	Right answer	Actual answer	Error
X	1	0.92	0.08
O	0	0.51	0.49
Total			0.57

-1 -1 -1 -1 -1 -1 -1 -1 -1  
-1 1 -1 -1 -1 -1 -1 1 -1  
-1 -1 1 -1 -1 -1 1 -1 -1  
-1 -1 -1 1 -1 1 -1 -1 -1  
-1 -1 -1 -1 1 -1 -1 -1 -1  
-1 -1 -1 1 -1 -1 -1 -1 -1  
-1 -1 -1 1 -1 1 -1 -1 -1  
-1 -1 1 -1 -1 -1 1 -1 -1  
-1 1 -1 -1 -1 -1 1 -1 -1  
-1 -1 -1 -1 -1 -1 -1 -1 -1



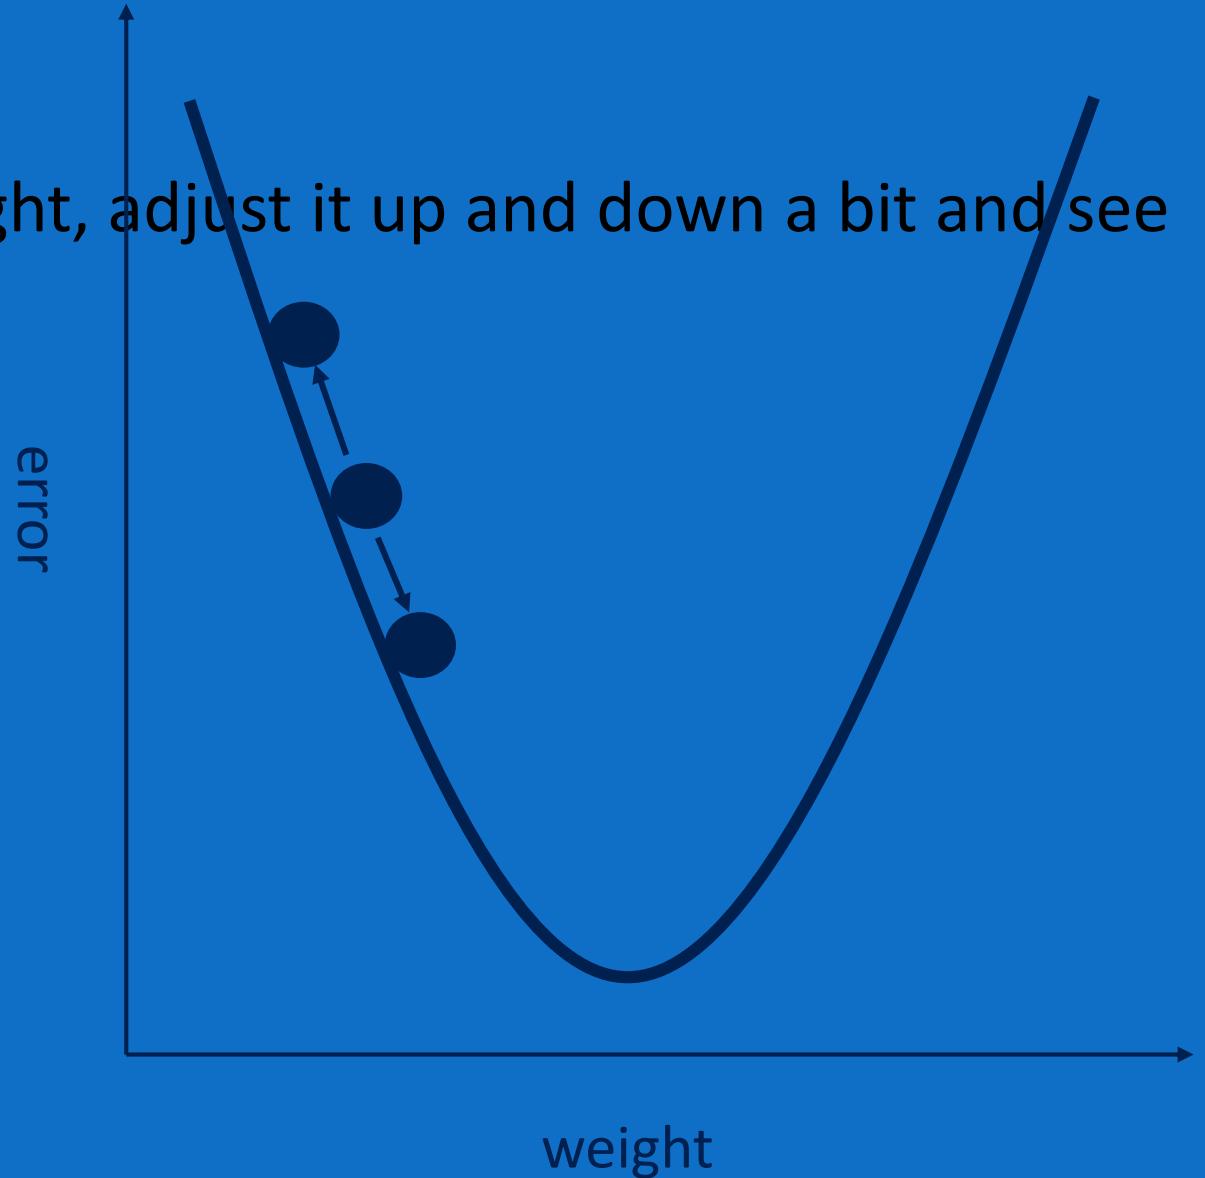
# Gradient descent

For each feature pixel and voting weight, adjust it up and down a bit and see how the error changes.



# Gradient descent

For each feature pixel and voting weight, adjust it up and down a bit and see how the error changes.



# Hyperparameters (knobs)

Convolution

- Number of features

- Size of features

Pooling

- Window size

- Window stride

Fully Connected

- Number of neurons

# Architecture

How many of each type of layer?

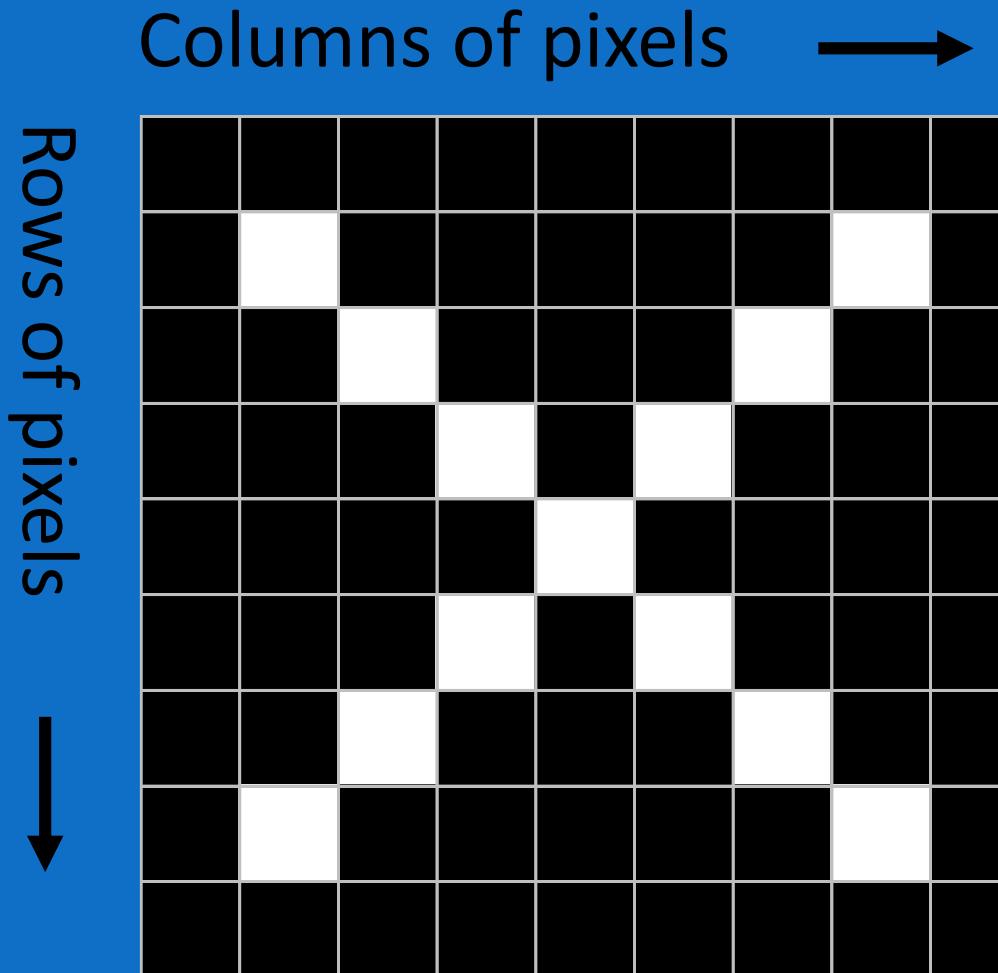
In what order?

Not just images

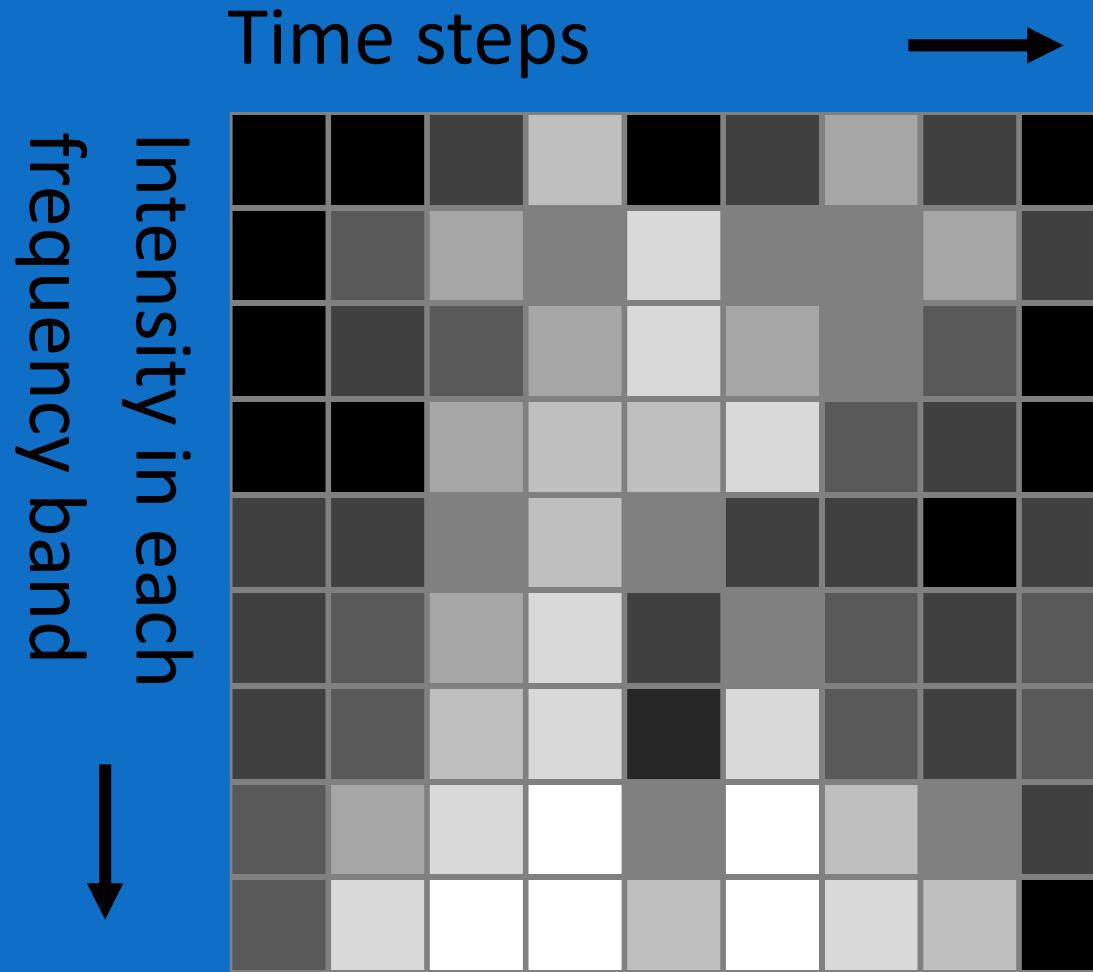
Any 2D (or 3D) data.

Things closer together are more closely related than things far away.

# Images



Sound

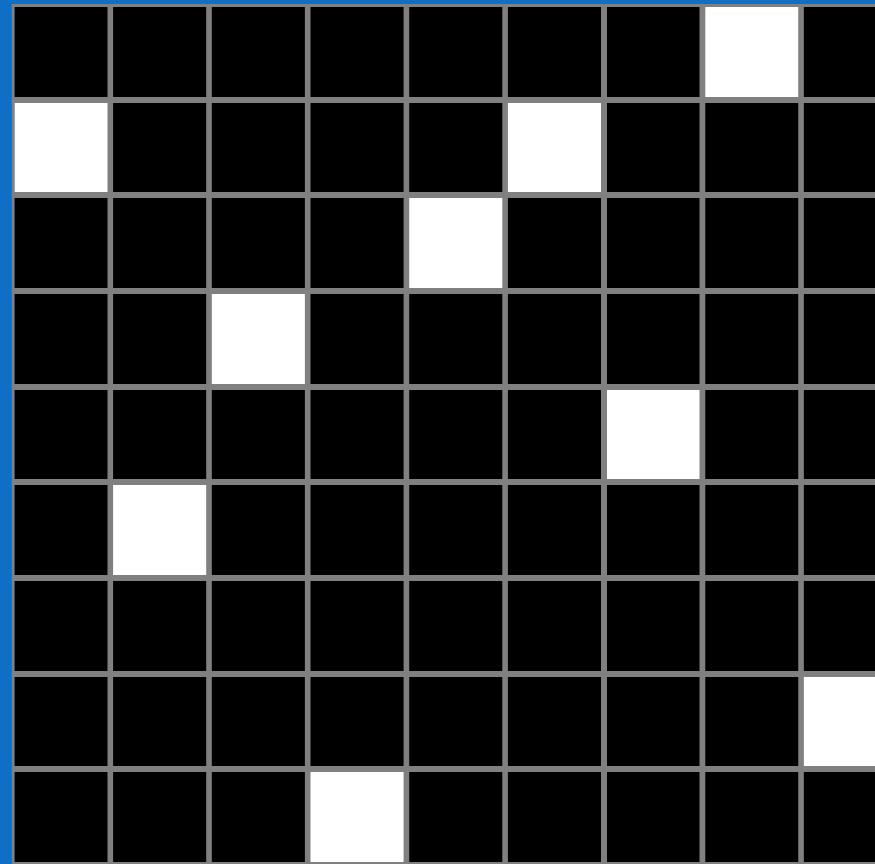


Text

Position in  
sentence



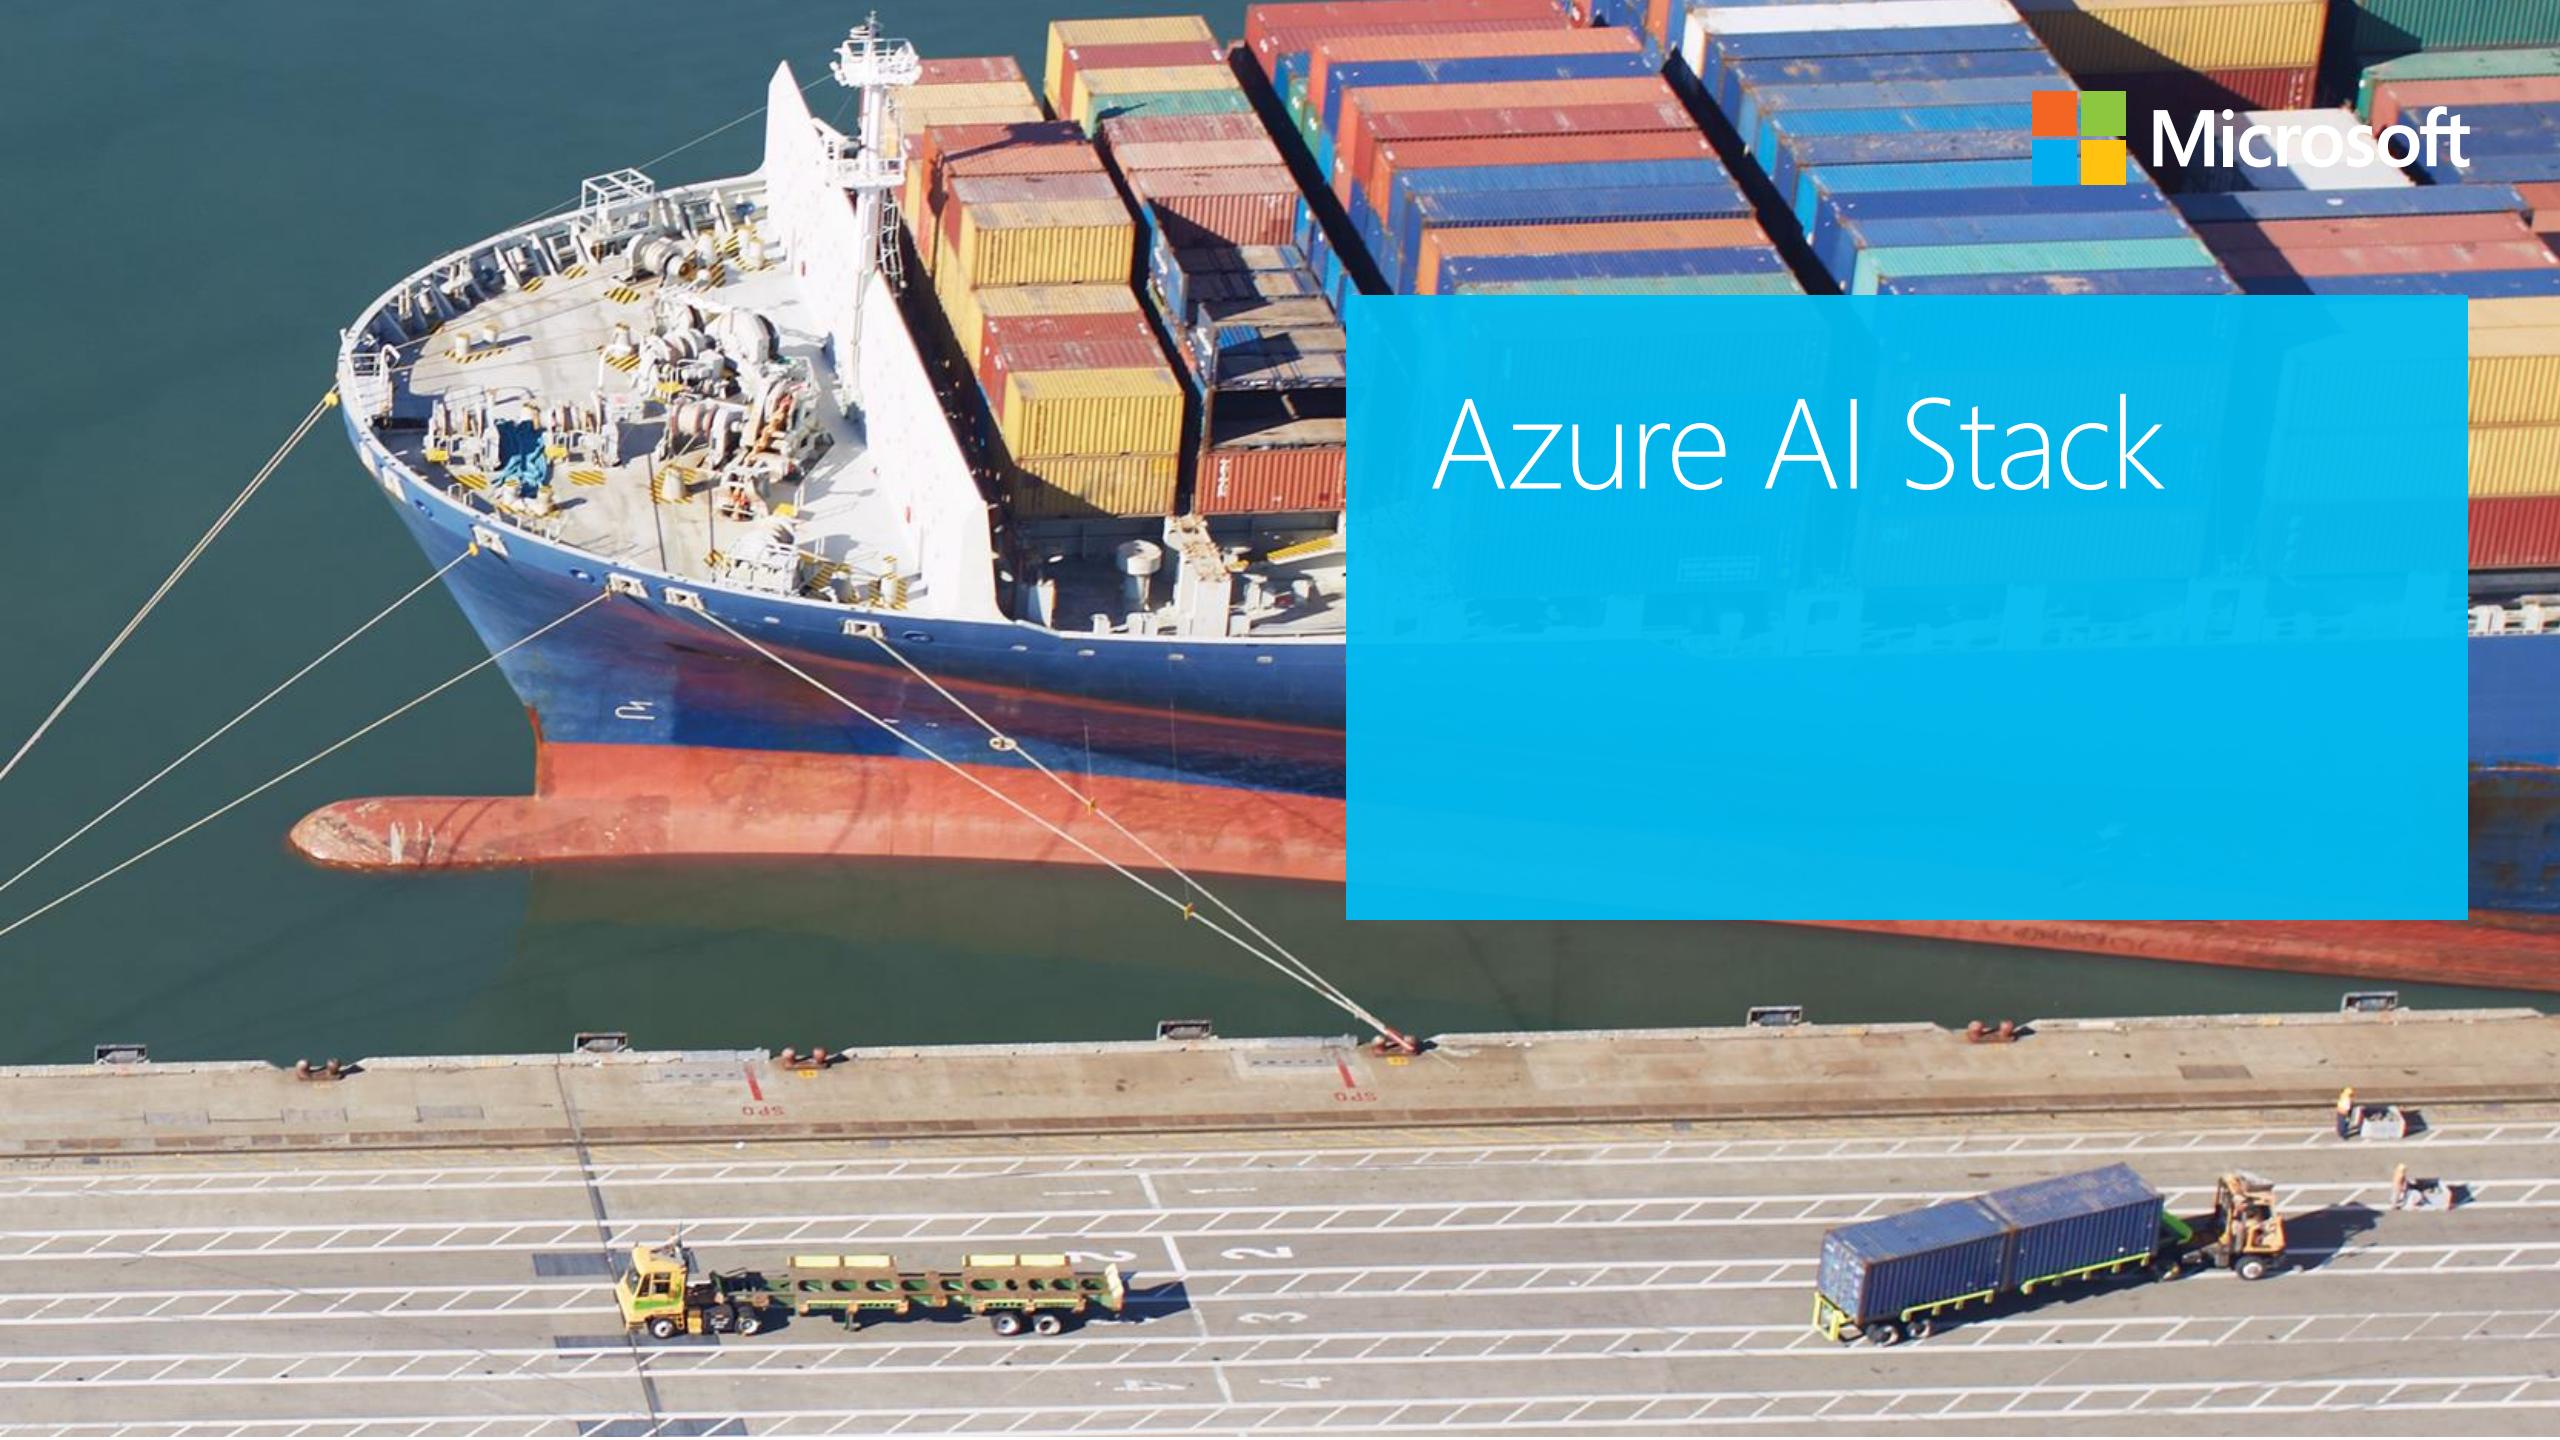
Words in  
dictionary





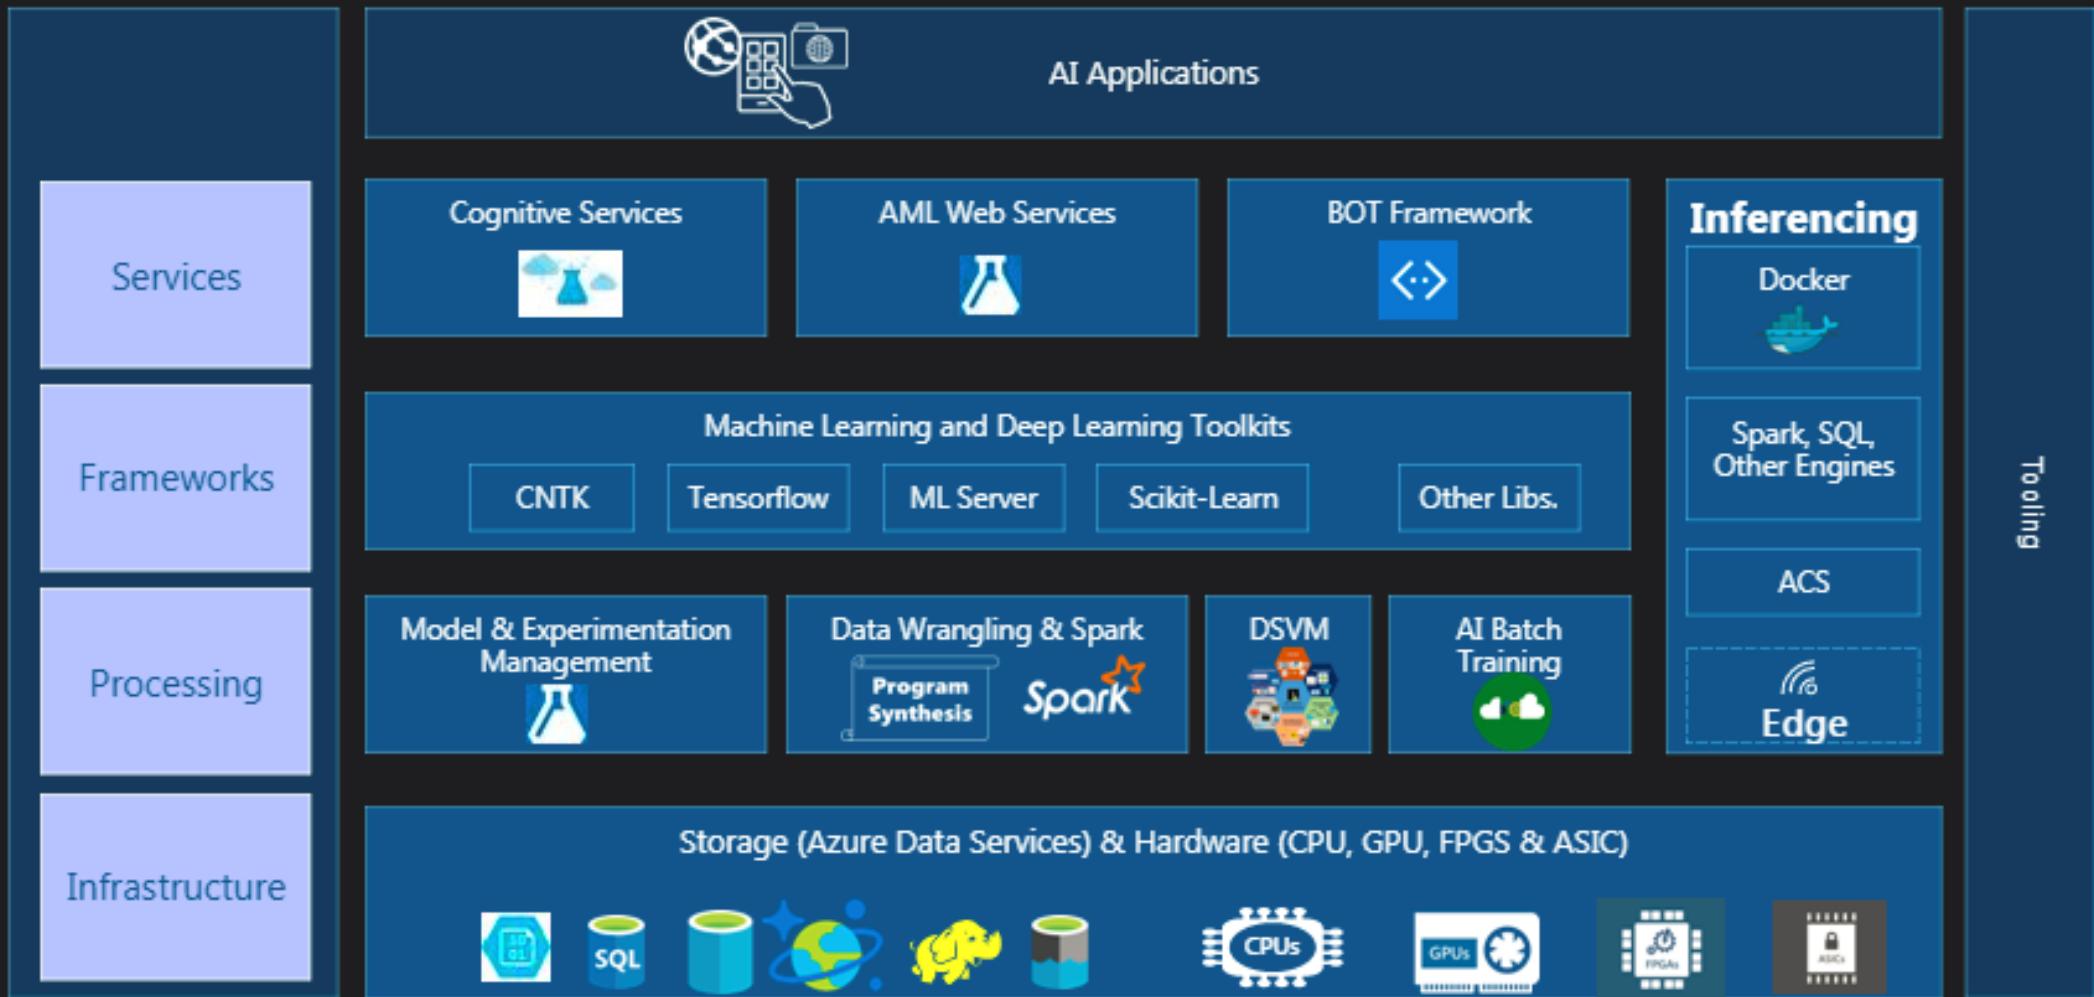
Microsoft

# Azure AI Stack



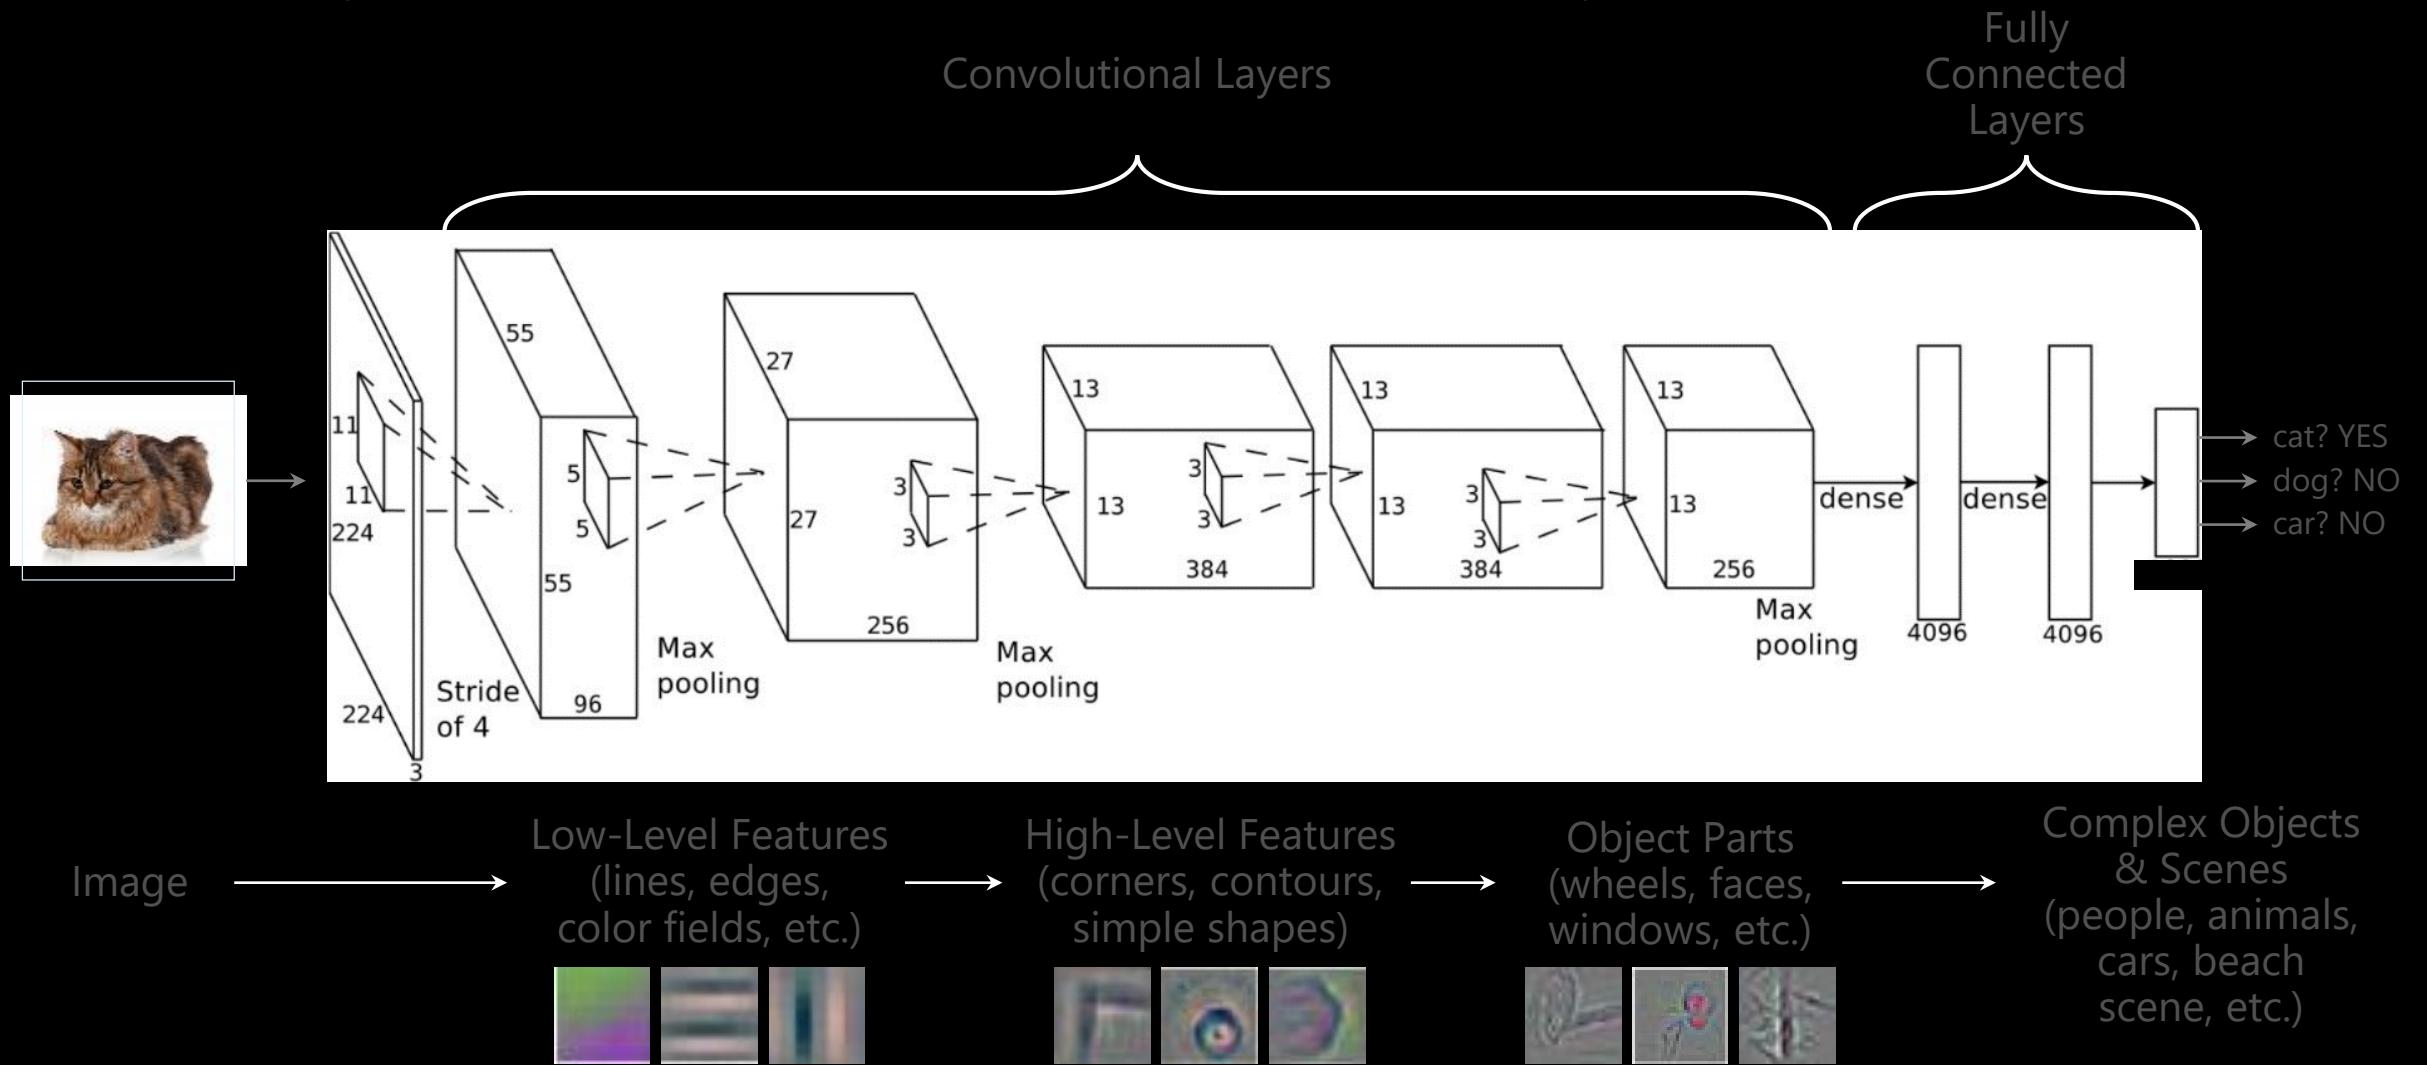


# Cloud AI Stack



# Deep Learning for Image Classification

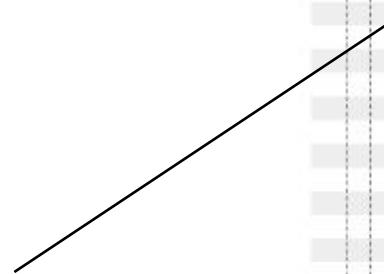
# Deep Neural Network for Computer Vision



# ImageNet dataset



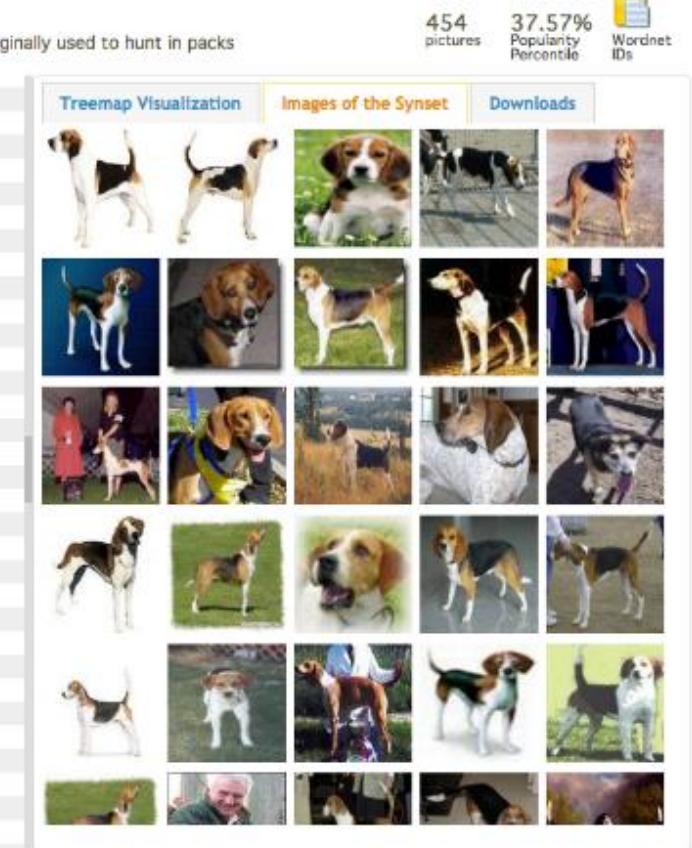
- Research dataset with >10 million images
- Image annotated with labels from an ontology (>22K labels)
- Generic images covering an extremely wide range of labels



## English foxhound

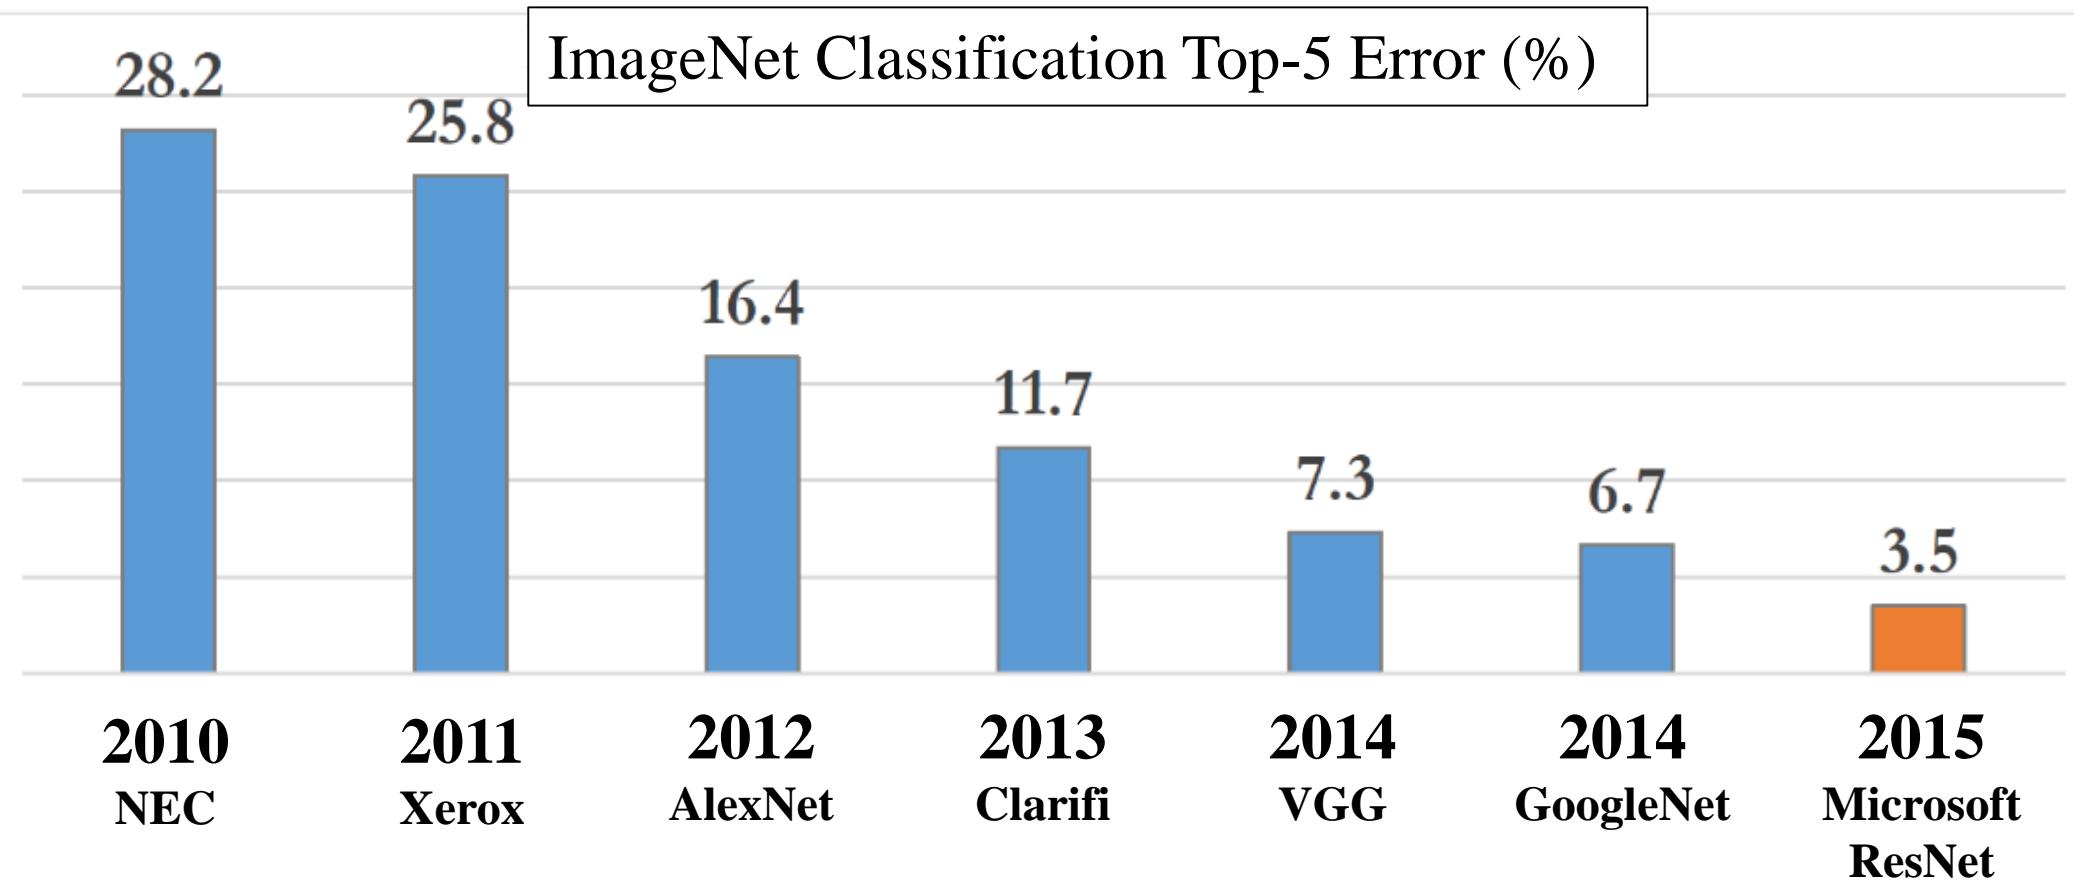
An English breed slightly larger than the American foxhounds originally used to hunt in packs

survivor (0)  
range animal (0)  
creepy-crawly (0)  
domestic animal, domesticated animal (213)  
- domestic cat, house cat, *Felis domesticus*, *Felis catus* (18)  
- dog, domestic dog, *Canis familiaris* (189)  
-- pooh, doggie, doggy, barker, bow-wow (0)  
- hunting dog (101)  
-- sporting dog, gun dog (28)  
-- dachshund, dachsie, badger dog (1)  
-- terrier (37)  
-- coursier (0)  
- hound, hound dog (29)  
-- Plott hound (0)  
-- wolfhound (2)  
-- Scottish deerhound, deerhound (0)  
-- coonhound (2)  
-- foxhound (3)  
--- Walker hound, Walker foxhound (0)  
--- American foxhound (0)  
--- English foxhound (0)  
--- Weimaraner (0)  
-- otterhound, otter hound (0)  
-- bloodhound, sleuthhound (0)  
-- Norwegian elkhound, elkhound (0)  
-- Saluki, gazelle hound (0)  
-- Afghan hound, Afghan (0)  
-- staghound (0)  
-- greyhound (2)  
-- beagle (0)  
-- harrier (0)  
-- basset, basset hound (0)  
-- bluetick (0)  
-- redbone (0)



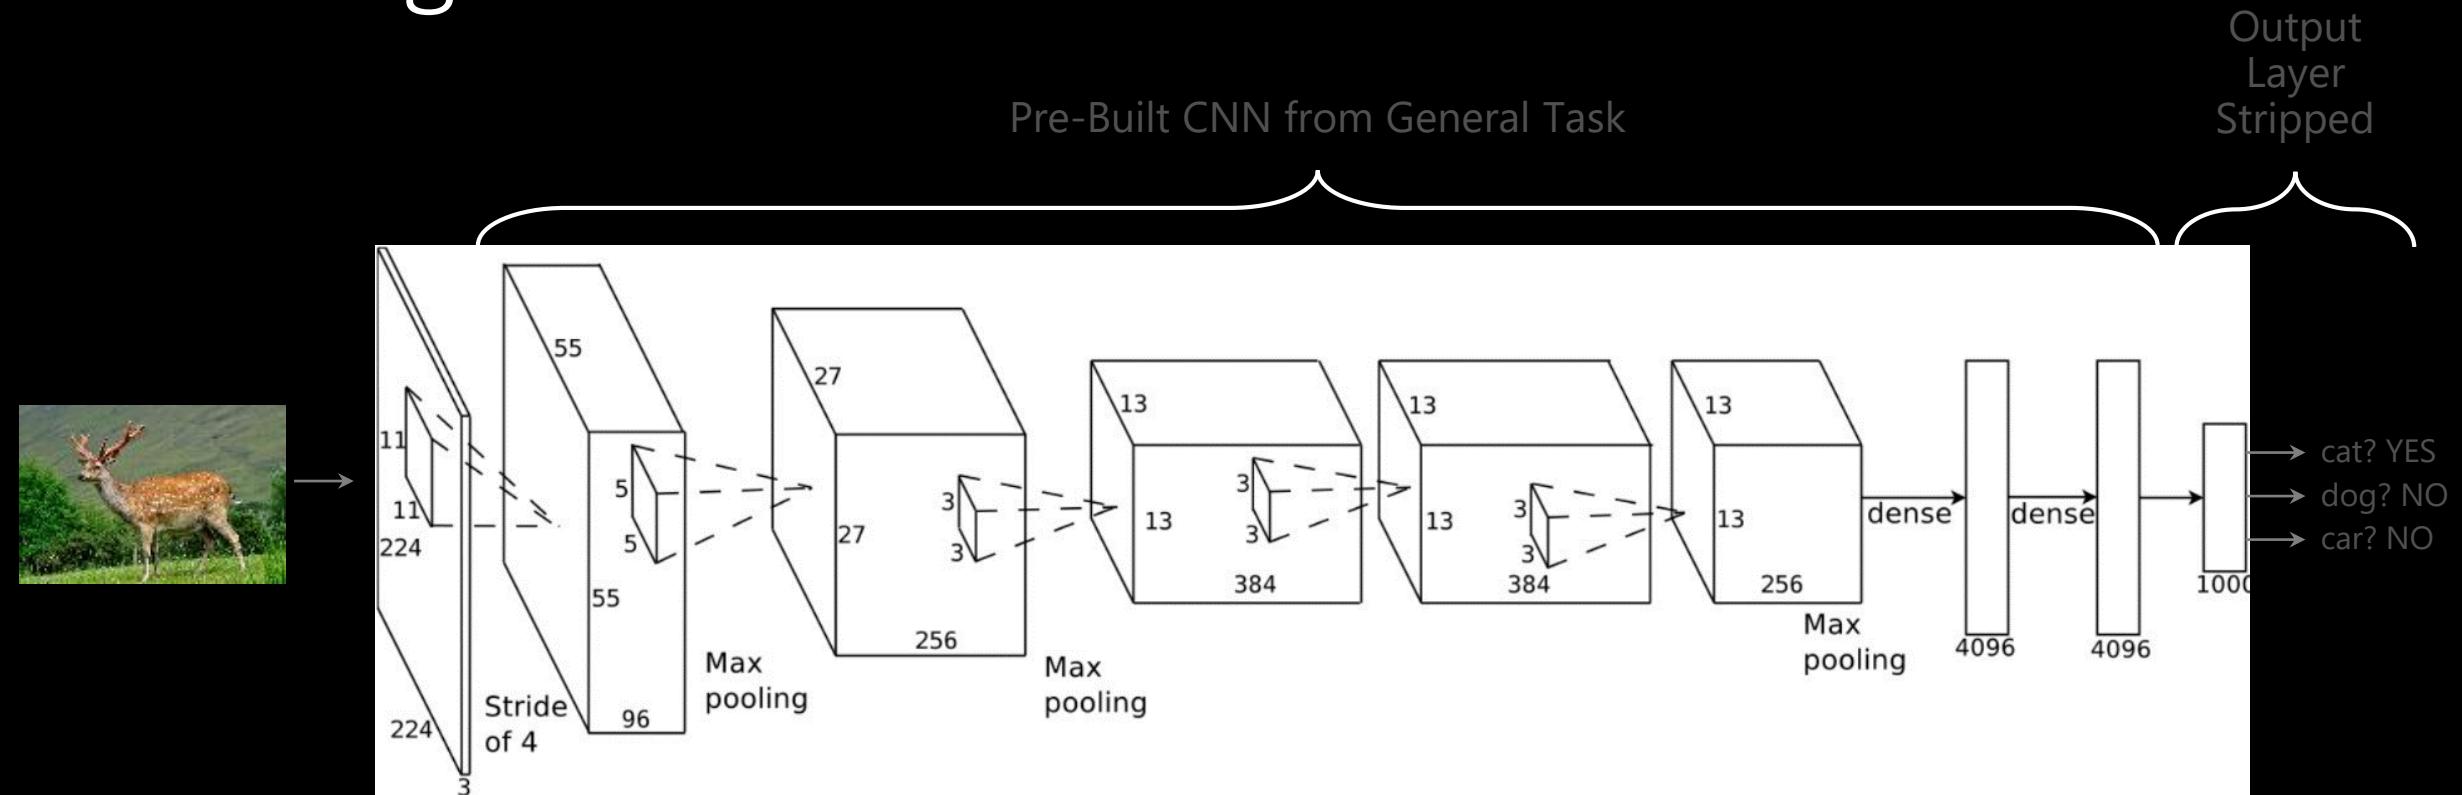
ImageNet taxonomy with example images for “English foxhound”

# Image Net ILSVRC Competition



# Customizing Deep Learning Model

# Using a Pre-Trained CNN as a Featurizer

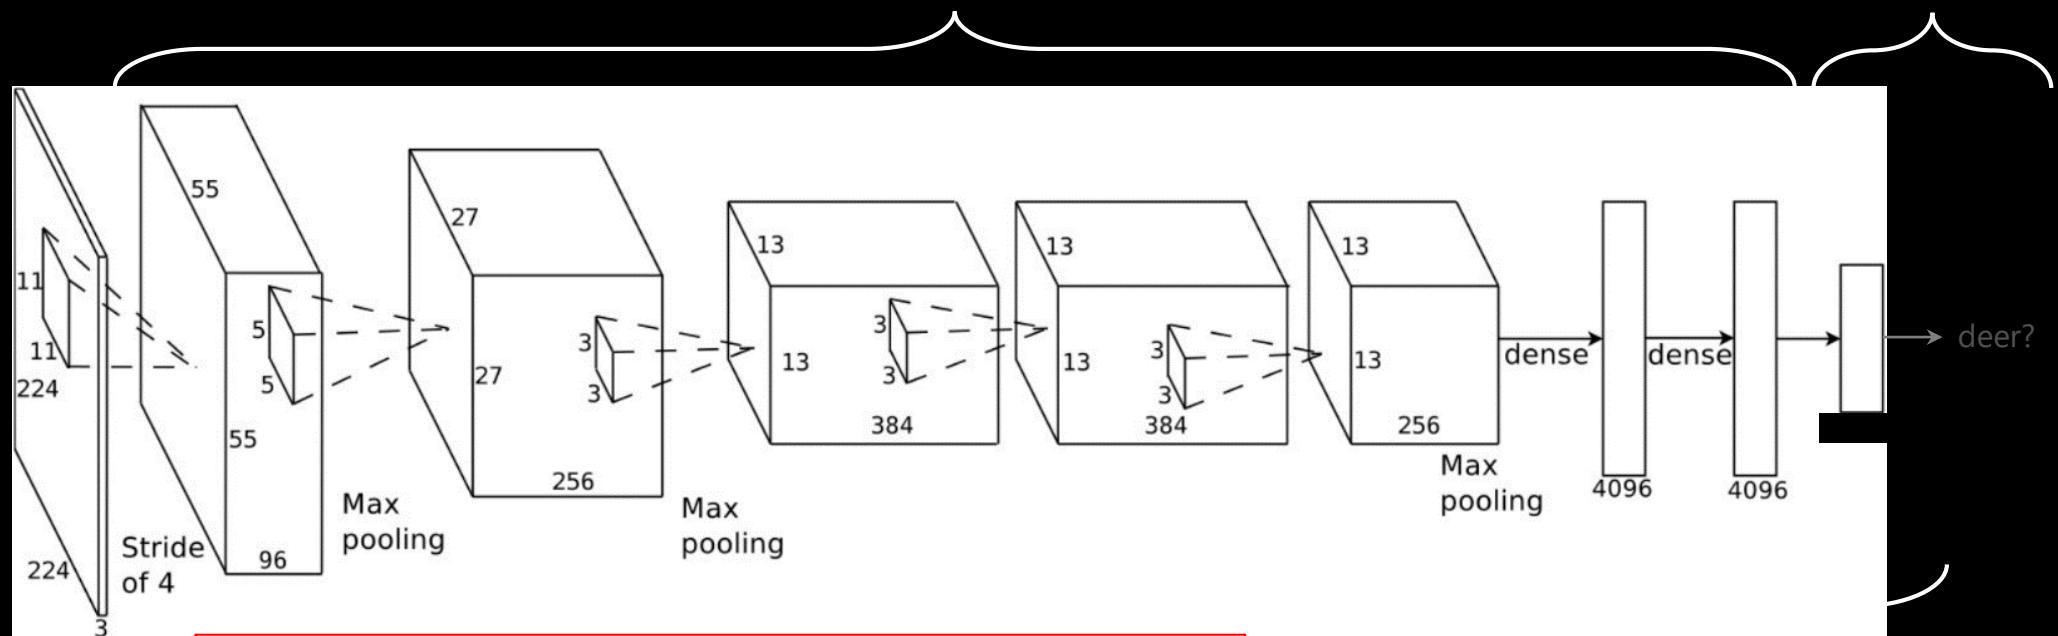


Outputs of penultimate layer of ImageNet Trained CNN provide excellent general purpose image features

# Transfer Learning

Pre-Built DNN from General Task

New  
Output  
Layer



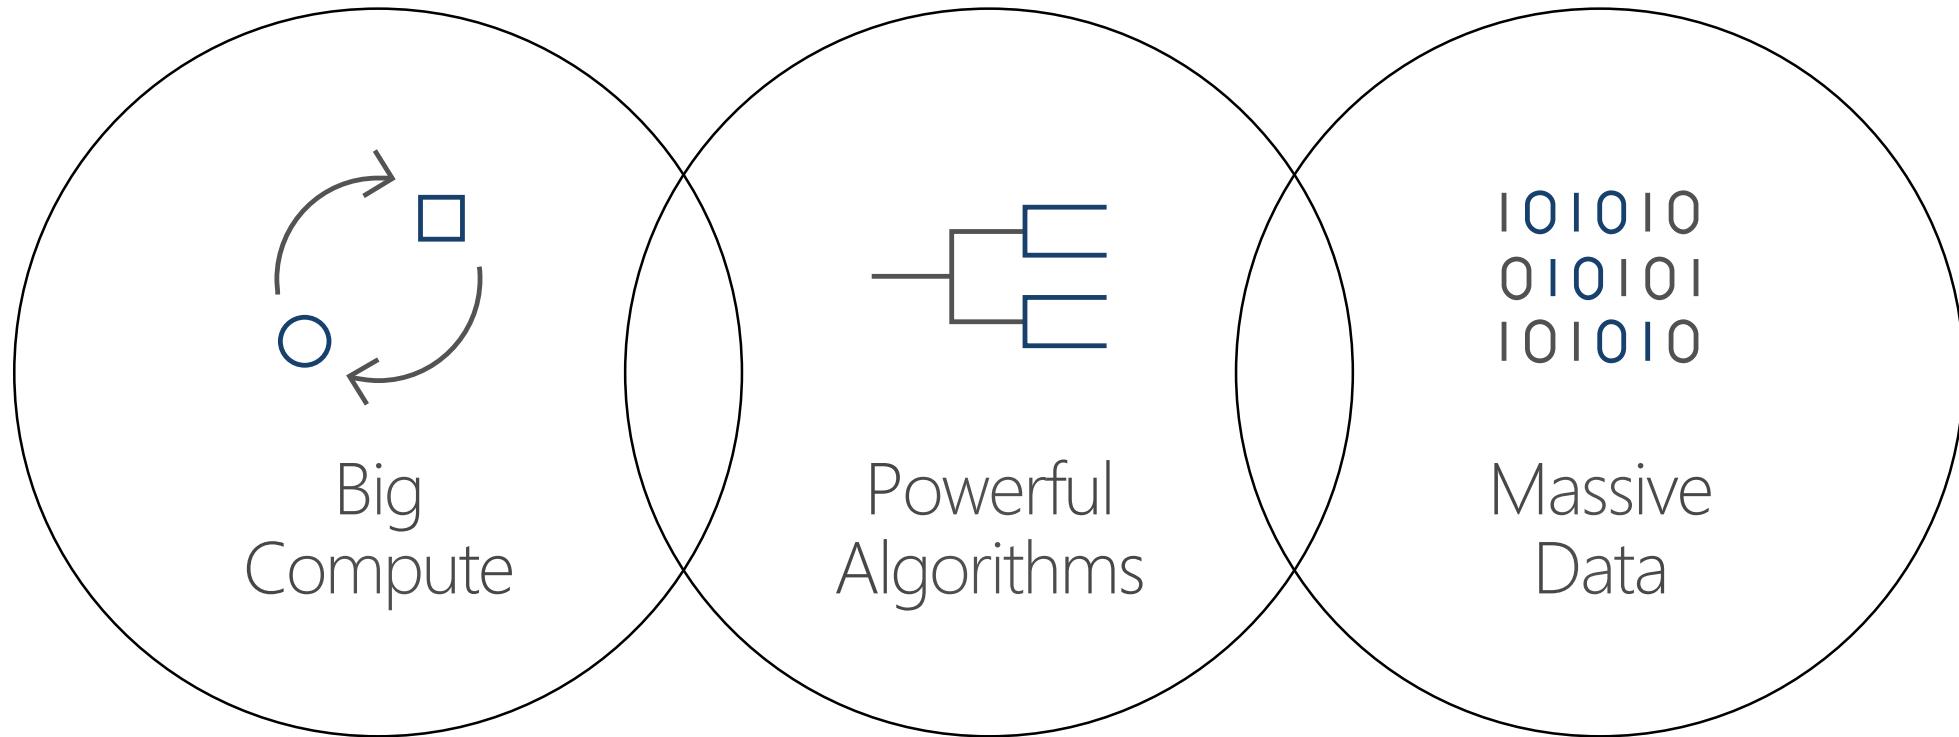
With transfer learning an accurate model can be achieved with thousands (or less) of labeled examples instead of millions

Train one or more layers in new network

# Microsoft AI Platform

Batch AI Training  
Azure Machine Learning  
Data Science Virtual Machine

# Recipe for AI Innovation



# Batch AI Training

## Easy Deployment and Flexibility

Managed service that enables data scientists to easily train and test deep learning and other AI models

## High Performance Training

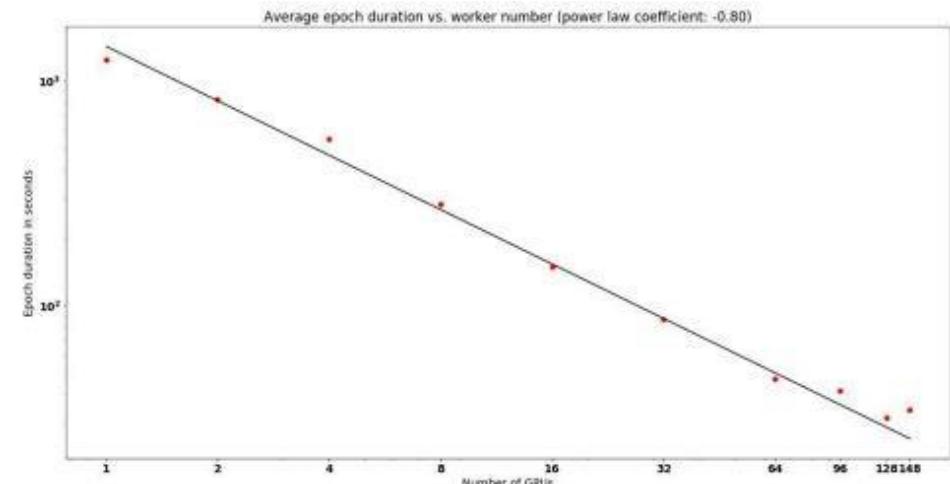
Work with clusters of GPUs to run experiments in parallel and at scale to reduce training time

## Any Framework and Tools

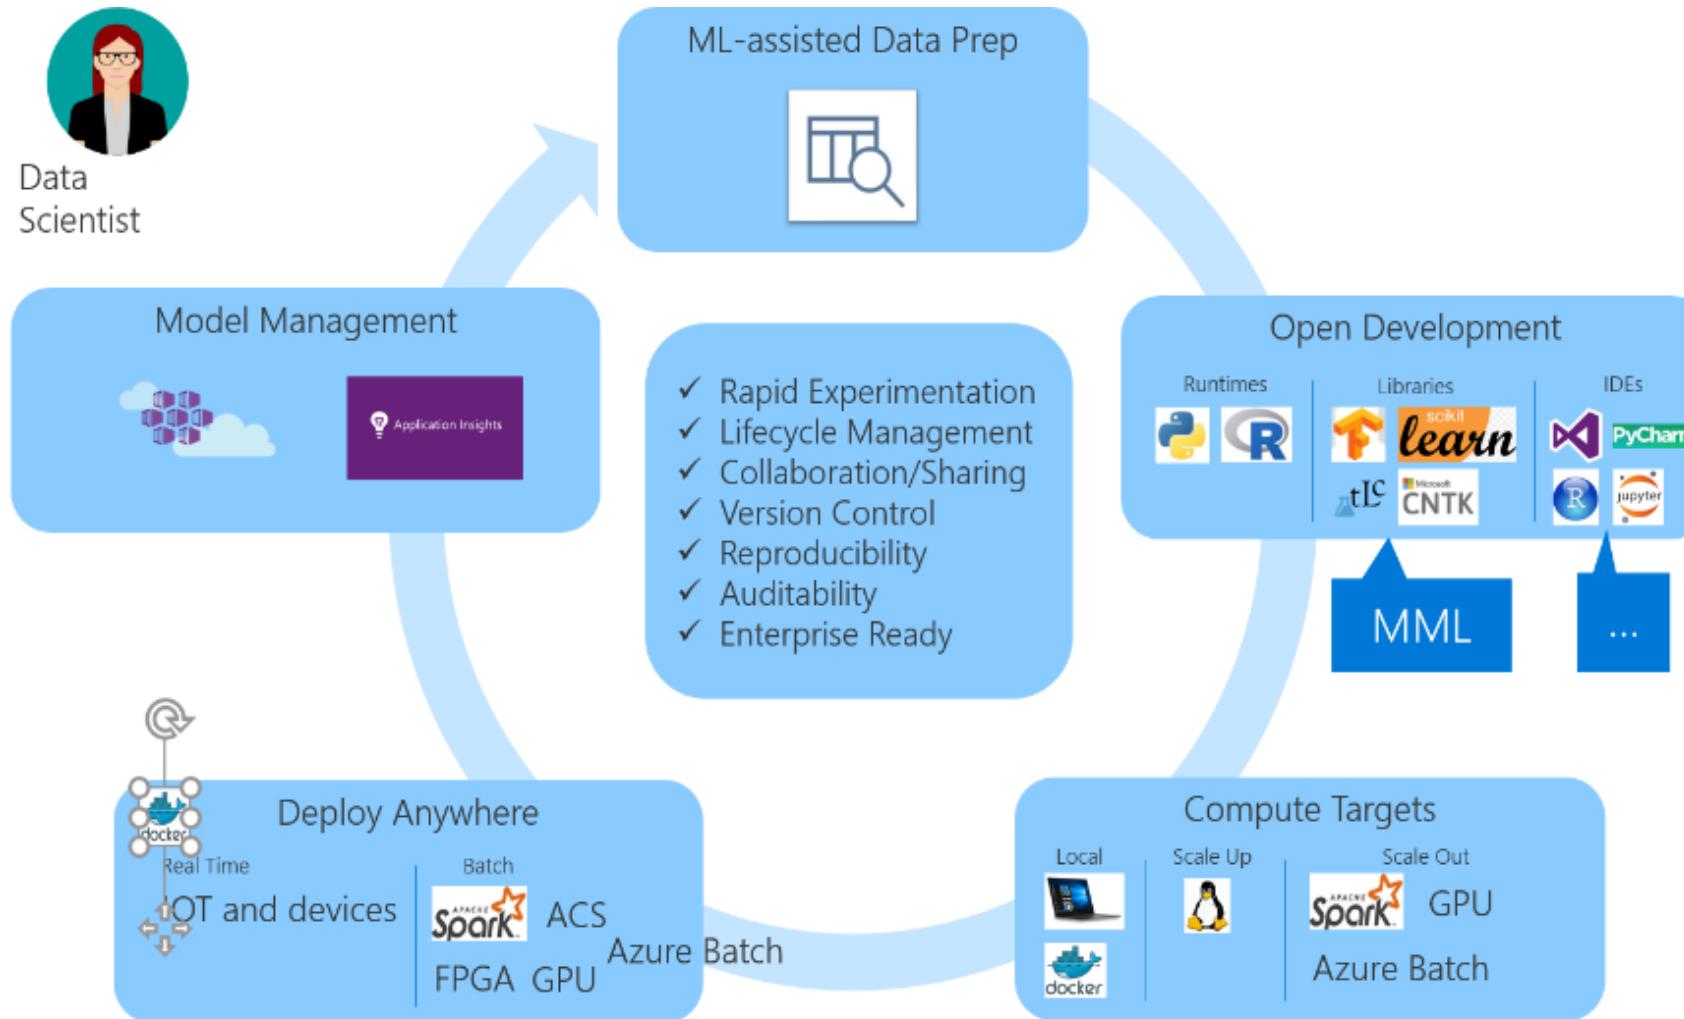
Do your AI with the scale and tools of Microsoft

## Near linear scale up on 1.6 TB Aerial Images

Training time decreases from  
1925 seconds/epoch to  
32 seconds/epoch.

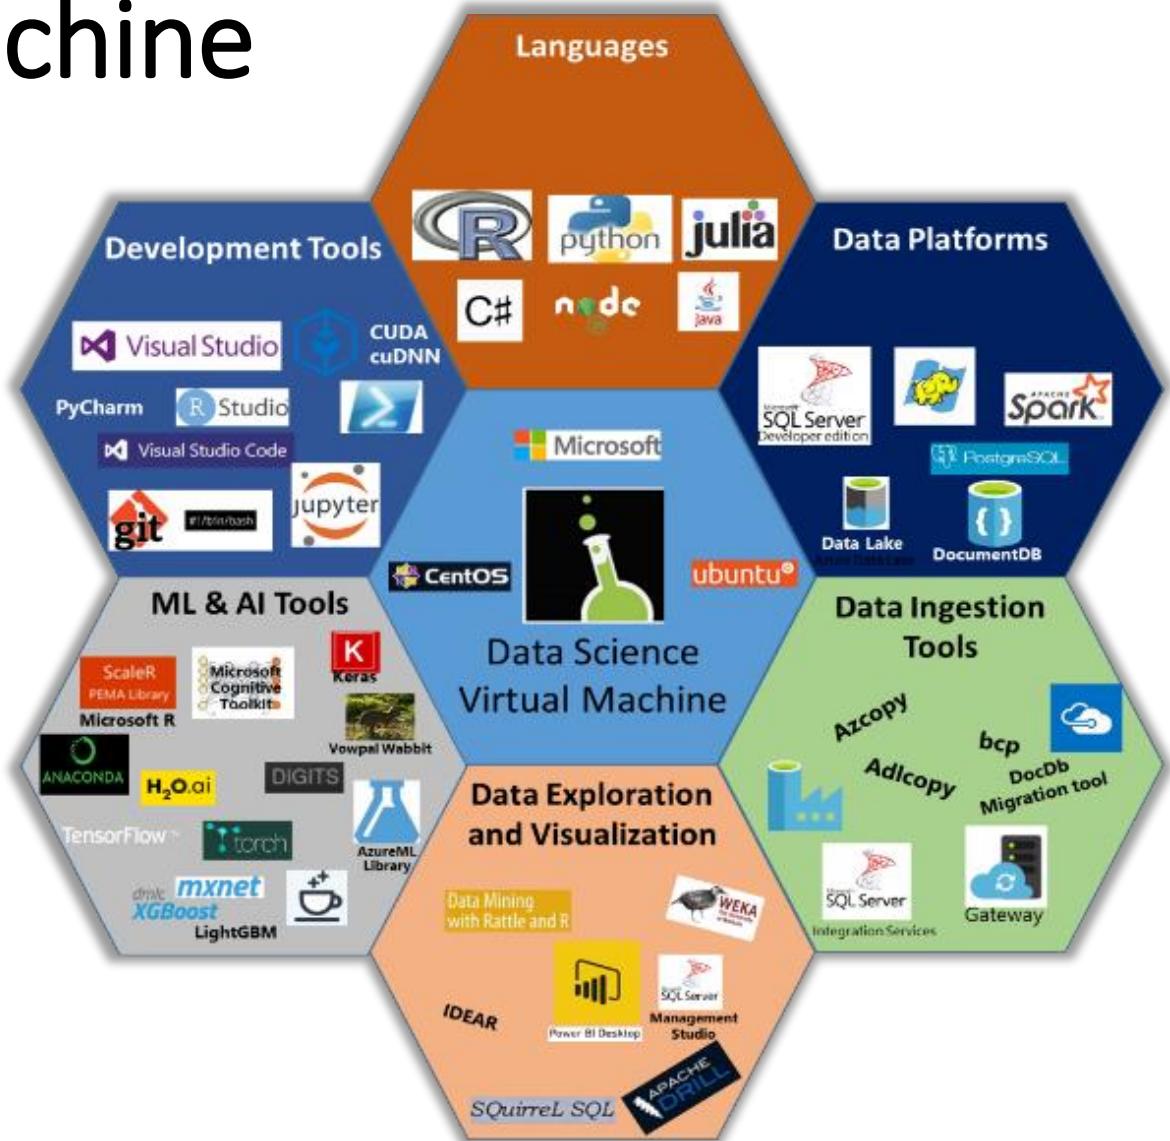


# Azure Machine Learning



# Data Science Virtual Machine

Running our Image Processing Pipeline in a Data Science Virtual Machine (DSVM) with Deep Learning products pre-stalled

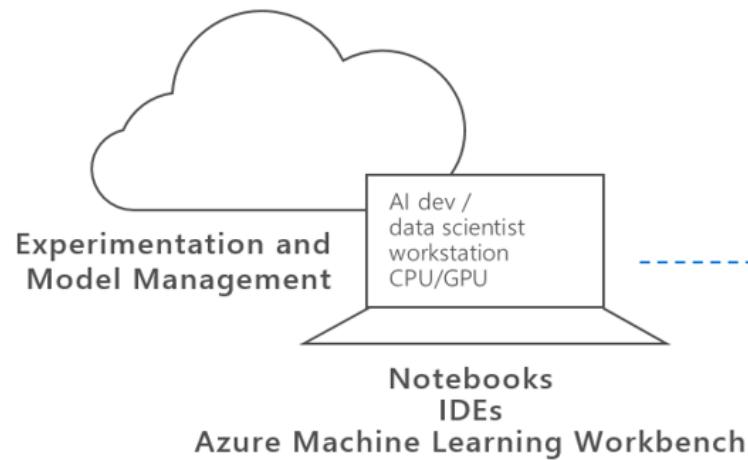


# Components of Azure Machine Learning

- Azure Machine Learning Workbench
- Azure Machine Learning Experimentation Service
- Azure Machine Learning Model Management Service
- Microsoft Machine Learning Libraries for Apache Spark (**MMLSpark Library**)
- Visual Studio Code Tools for AI

# AZURE MACHINE LEARNING

## AZURE MACHINE LEARNING SERVICES



## TRAIN & DEPLOY OPTIONS

AZURE



Spark  
SQL Server  
Virtual machines  
GPUs  
Container services

ON-PREMISES



SQL Server  
Machine Learning Server

EDGE COMPUTING



Azure IoT Edge

# Azure Machine Learning Workbench

- Azure Machine Learning Workbench is a desktop application plus command-line tools
- It allows you to manage machine learning solutions through the entire data science life cycle:
- Data ingestion and preparation
- Model development and experiment management
- Model deployment in various target environments

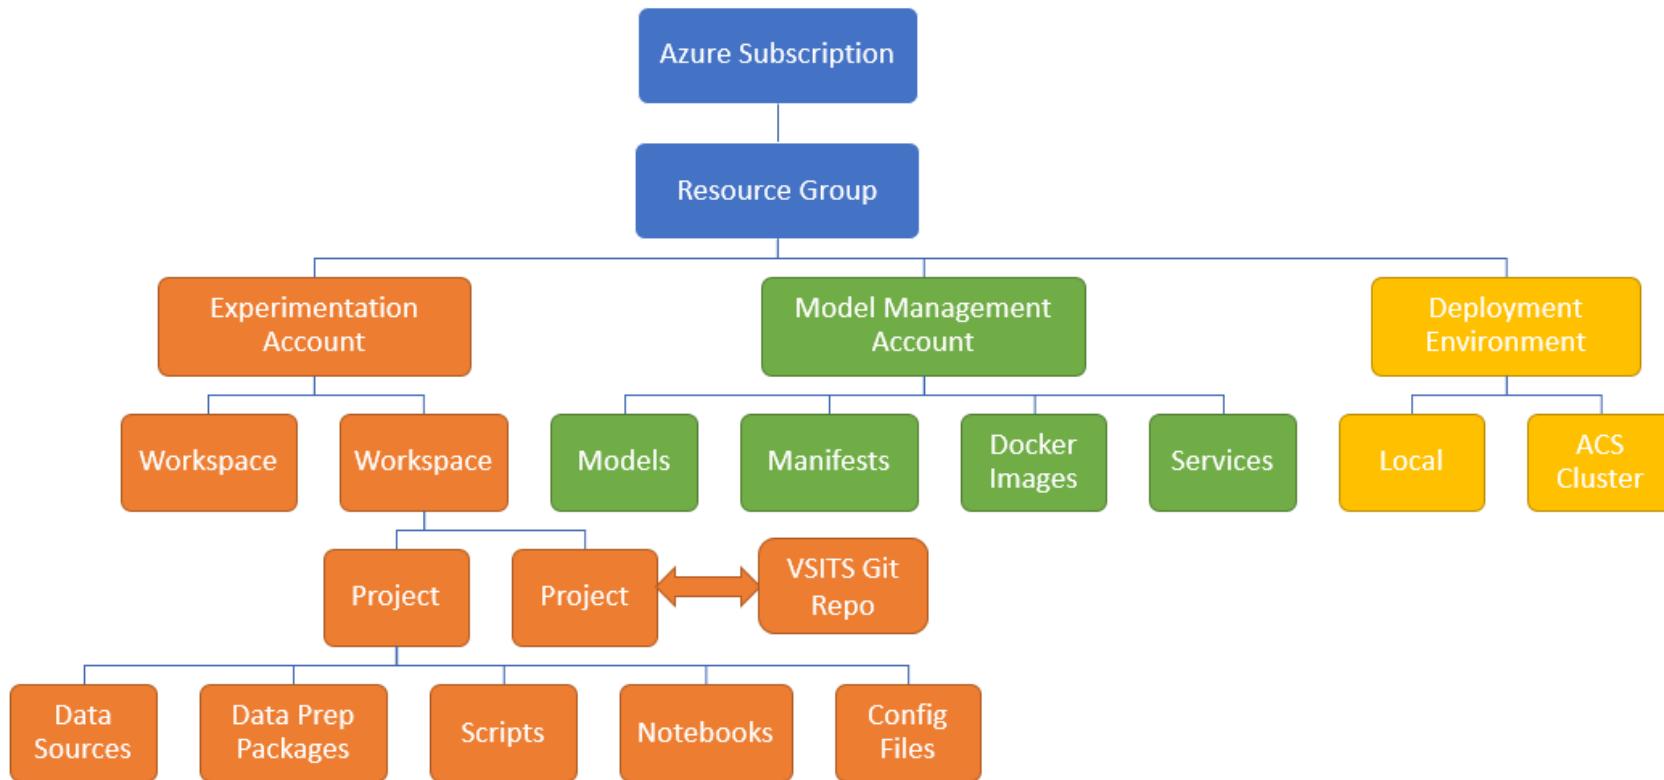
# Azure Machine Learning Experimentation Service

- The Experimentation Service handles the execution of machine learning experiments.
- It also supports the Workbench by providing project management, Git integration, access control, roaming, and sharing.
- Through easy configuration, you can execute your experiments across a range of compute environment options:
  - Local native
  - Local Docker container
  - Docker container on a remote VM
  - Scale out Spark cluster in Azure

# Azure Machine Learning Model Management Service

- Model Management Service allows data scientists and dev-ops teams to deploy predictive models into a wide variety of environments.
- Model versions and lineage are tracked from training runs to deployments. Models are stored, registered, and managed in the cloud
- Using simple CLI commands, you can containerize your model, scoring scripts and dependencies into Docker images. These images are registered in your own Docker registry hosted in Azure (Azure Container Registry). They can be reliably deployed to the following targets:
  - Local machines
  - On-premises servers
  - The cloud
  - IoT edge devices

# Azure Machine Learning Services



# Offering from Microsoft on AI

- Besides Azure Machine Learning, there are a wide variety of options in Azure to build, deploy, and manage machine learning models.
  - Microsoft Machine Learning Services in SQL Server
  - Microsoft Machine Learning Server
  - Data Science Virtual Machine
  - Spark MLlib in HDInsight
  - Batch AI Training Service
  - Microsoft Cognitive Toolkit
  - Microsoft Cognitive Services

# Image Classification Demo

- Training, evaluation and deployment to local or the cloud.
- All steps running within AML Workbench.
- Jupyter notebooks provided to visualize results, deploy as Rest API, and to perform load-testing.

Datasets:

Clothing texture images

# Azure Machine Learning Service Demo

1. Prepare data
2. Build Models
3. Deploy Model
4. Advanced Analytics

# Demo on Image Classification using CNTK

# Image Classification Demo

## Clothing texture dataset:

- Collected 1716 images from Bing, and annotated as one of 11 textures (striped, dotted, etc.)

Striped



Dotted

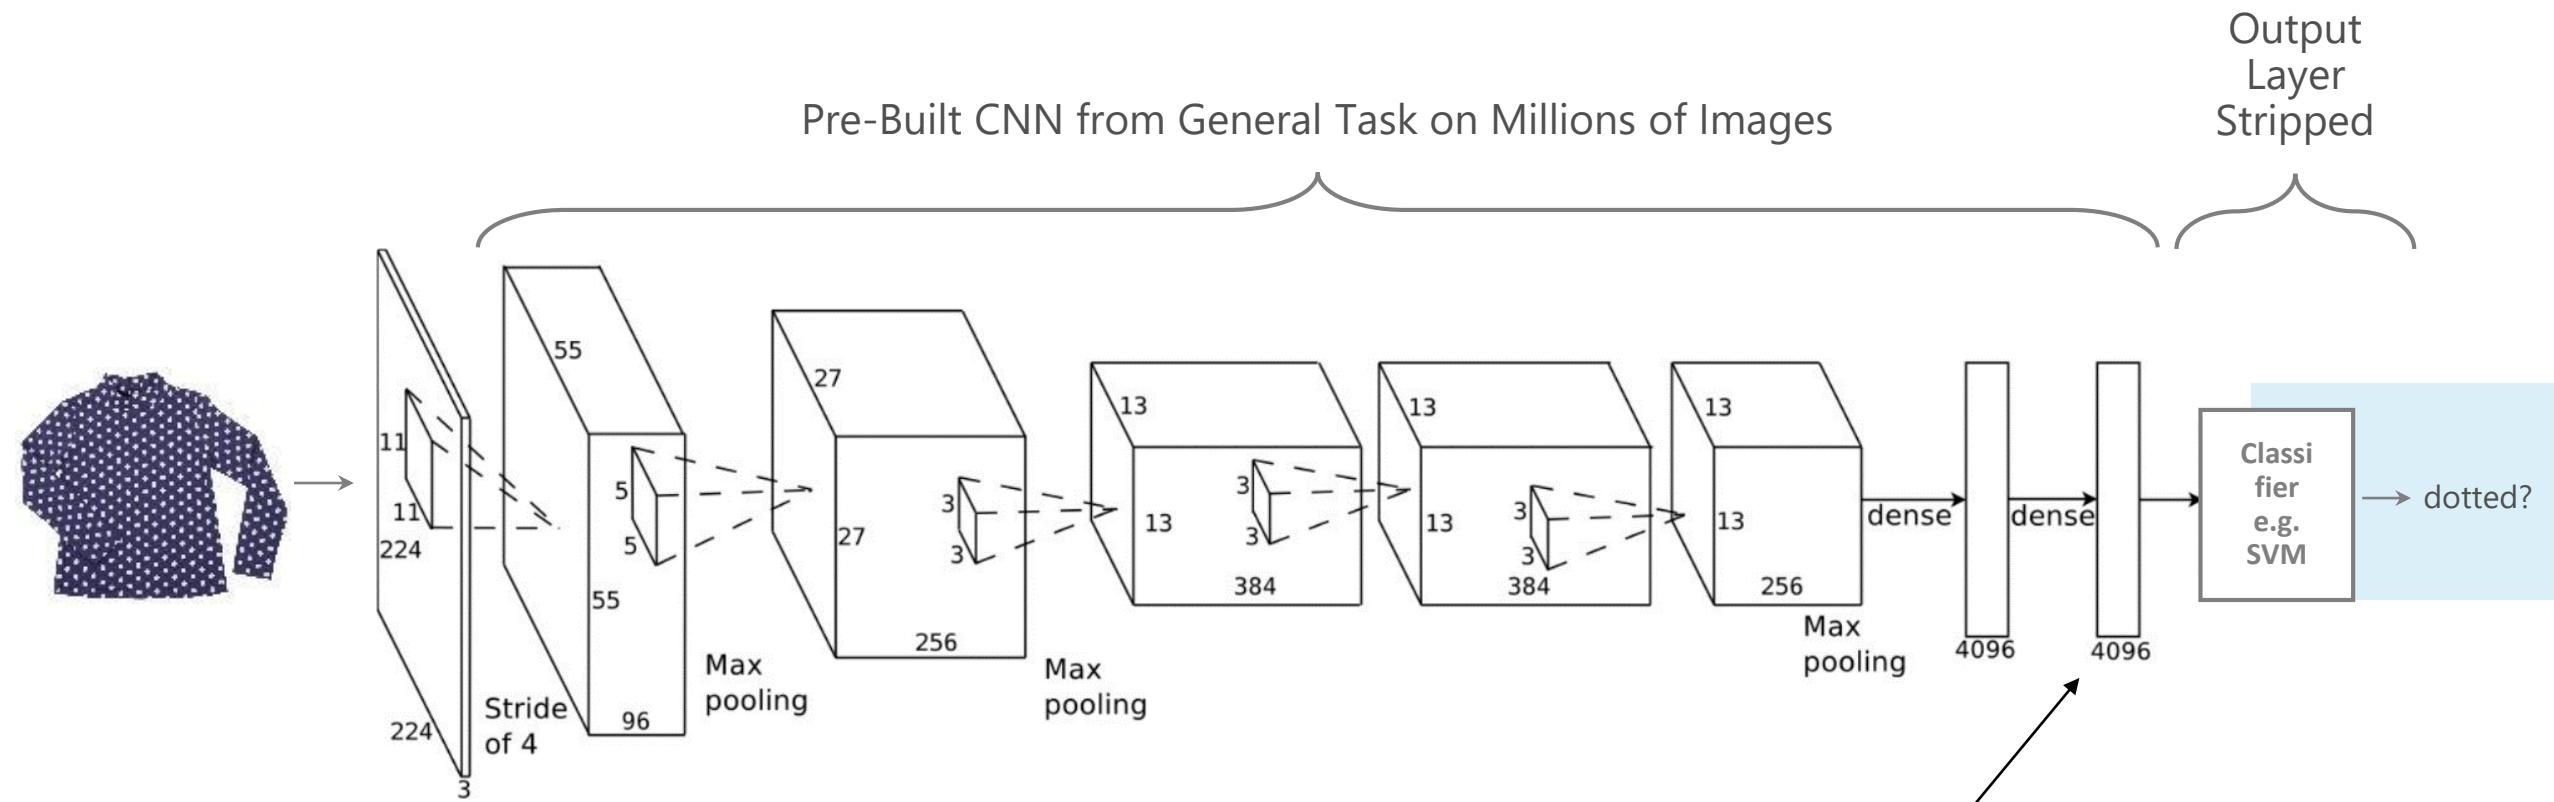


Argyle



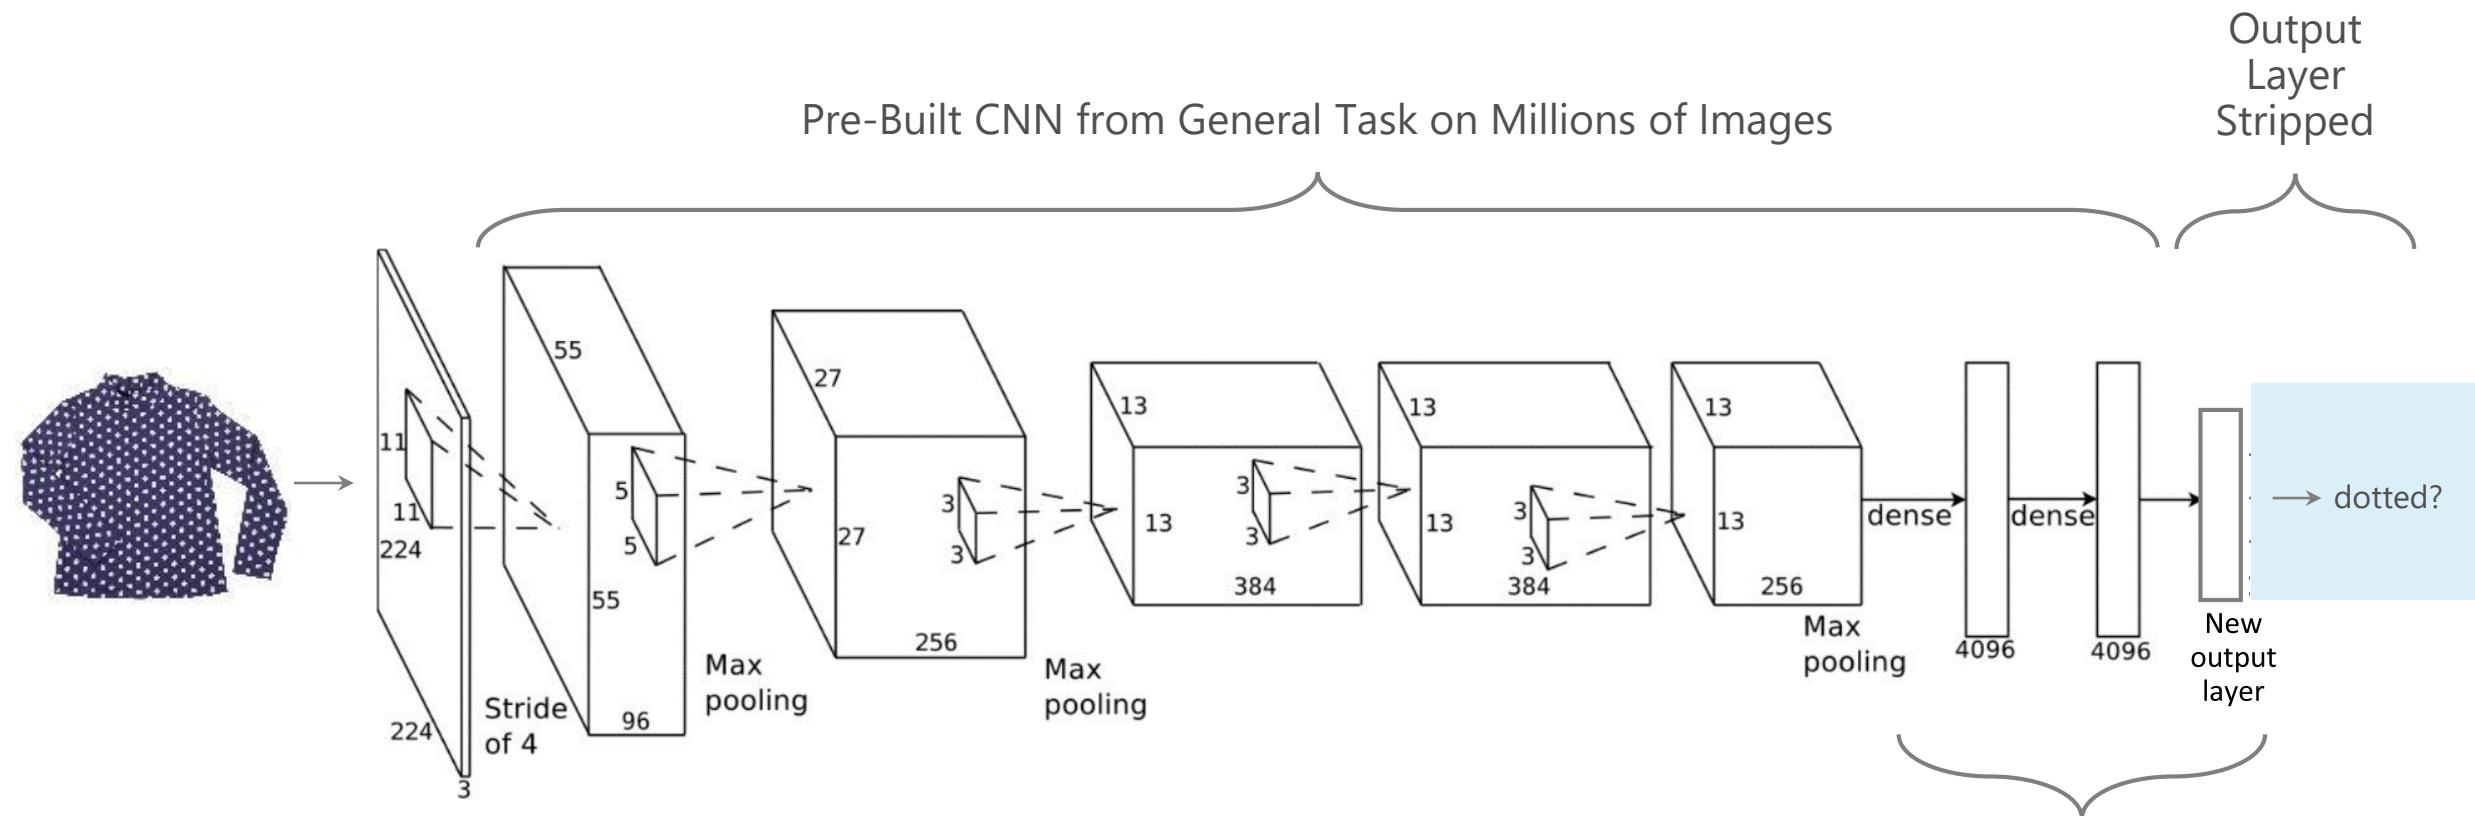
Can we apply transfer learning to accurately classify clothing texture?

# Using a Pre-Trained CNN as a Featurizer



Outputs of penultimate layer of  
ImageNet Trained CNN provide excellent  
general purpose image features

# Using a Pre-Trained CNN and Finetune



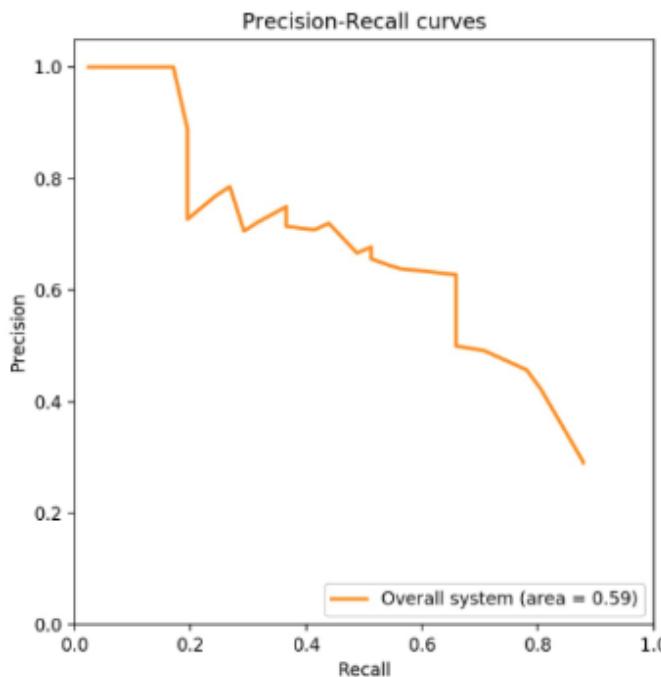
Using a pre-trained DNN, an accurate model can be achieved with thousands (or less) of labeled examples instead of millions

Train one or more layers in new network

# Transfer Learning Results - Texture Dataset

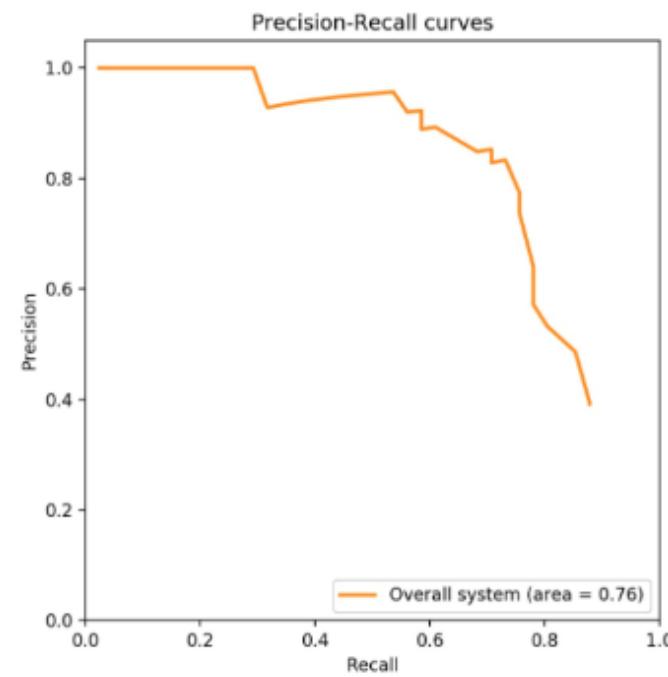
## DNN featurization

Input Image Size: 224x224 pixels  
Area Under Curve: 0.59  
Classification Accuracy: **69.0%**



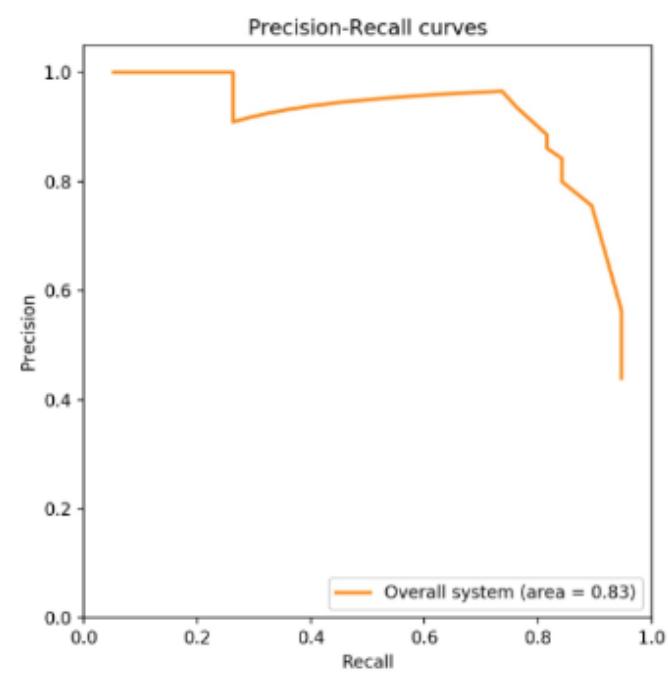
## Fine-tuning (full CNN)

Input Image Size: 224x224 pixels  
Area Under Curve: 0.76  
Classification Accuracy: **77.4%**



## Fine-tuning (full CNN)

Input Image Size: 896x886 pixels  
Area Under Curve: 0.83  
Classification Accuracy: **88.2%**



# Deployment as Rest API

Easy deployment using AML Workbench.

Cluster management using Kubernetes.

Notebook provided to do measure average response time

➤ Experiment: #requests: 100, #concurrent requests: 10.

Number of replicas	Average round trip response time
1	536.99ms
10	81.88ms

# Deep Learning using Azure Machine Learning Services.

- Bird Detection : [Link](#)
- Skin Cancer Detection : [Link](#)
- Shirts Classification : [Link](#)
- Aerial Image Classification : [Link](#)
- Document Collection Analysis : [Link](#)
- AI Real world Usecases : [Link](#)
- Energy Demand Forecasting : [Link](#)

Thank you